

```

mysql> SELECT PARTITION_NAME, TABLE_ROWS FROM INFORMATION_SCHEMA.PARTITIONS WHERE TABLE_NAME = 'e';
+-----+-----+
| PARTITION_NAME | TABLE_ROWS |
+-----+-----+
| p0            |      0 |
| p1            | 1000000 |
+-----+-----+
2 rows in set (0.00 sec)

# Exchange partition p0 of table e with the table e2 'WITH VALIDATION'

mysql> ALTER TABLE e EXCHANGE PARTITION p0 WITH TABLE e2 WITH VALIDATION;
Query OK, 0 rows affected (0.74 sec)

# Confirm that the partition was exchanged with table e2

mysql> SELECT PARTITION_NAME, TABLE_ROWS FROM INFORMATION_SCHEMA.PARTITIONS WHERE TABLE_NAME = 'e';
+-----+-----+
| PARTITION_NAME | TABLE_ROWS |
+-----+-----+
| p0            | 1000000 |
| p1            | 1000000 |
+-----+-----+
2 rows in set (0.00 sec)

# Once again, drop the rows from p0 of table e

mysql> DELETE FROM e WHERE id < 1000000;
Query OK, 1000000 rows affected (5.55 sec)

# Confirm that there are no rows in partition p0

mysql> SELECT PARTITION_NAME, TABLE_ROWS FROM INFORMATION_SCHEMA.PARTITIONS WHERE TABLE_NAME = 'e';
+-----+-----+
| PARTITION_NAME | TABLE_ROWS |
+-----+-----+
| p0            |      0 |
| p1            | 1000000 |
+-----+-----+
2 rows in set (0.00 sec)

# Exchange partition p0 of table e with the table e3 'WITHOUT VALIDATION'

mysql> ALTER TABLE e EXCHANGE PARTITION p0 WITH TABLE e3 WITHOUT VALIDATION;
Query OK, 0 rows affected (0.01 sec)

# Confirm that the partition was exchanged with table e3

mysql> SELECT PARTITION_NAME, TABLE_ROWS FROM INFORMATION_SCHEMA.PARTITIONS WHERE TABLE_NAME = 'e';
+-----+-----+
| PARTITION_NAME | TABLE_ROWS |
+-----+-----+
| p0            | 1000000 |
| p1            | 1000000 |
+-----+-----+
2 rows in set (0.00 sec)

```

If a partition is exchanged with a table that contains rows that do not match the partition definition, it is the responsibility of the database administrator to fix the non-matching rows, which can be performed using `REPAIR TABLE` or `ALTER TABLE ... REPAIR PARTITION`.

## Exchanging a Subpartition with a Nonpartitioned Table

You can also exchange a subpartition of a subpartitioned table (see [Section 24.2.6, “Subpartitioning”](#)) with a nonpartitioned table using an `ALTER TABLE ... EXCHANGE PARTITION` statement. In the following example, we first create a table `es` that is partitioned by `RANGE` and subpartitioned by `KEY`, populate this table as we did table `e`, and then create an empty, nonpartitioned copy `es2` of the table, as shown here:

```

mysql> CREATE TABLE es (
    ->     id INT NOT NULL,
    ->     fname VARCHAR(30),
    ->     lname VARCHAR(30)
    -> )
    -> PARTITION BY RANGE (id)
    -> SUBPARTITION BY KEY (lname)
    -> SUBPARTITIONS 2 (
        ->         PARTITION p0 VALUES LESS THAN (50),
        ->         PARTITION p1 VALUES LESS THAN (100),
        ->         PARTITION p2 VALUES LESS THAN (150),
        ->         PARTITION p3 VALUES LESS THAN (MAXVALUE)
    -> );
Query OK, 0 rows affected (2.76 sec)

mysql> INSERT INTO es VALUES
    ->     (1669, "Jim", "Smith"),
    ->     (337, "Mary", "Jones"),
    ->     (16, "Frank", "White"),
    ->     (2005, "Linda", "Black");
Query OK, 4 rows affected (0.04 sec)
Records: 4  Duplicates: 0  Warnings: 0

mysql> CREATE TABLE es2 LIKE es;
Query OK, 0 rows affected (1.27 sec)

mysql> ALTER TABLE es2 REMOVE PARTITIONING;
Query OK, 0 rows affected (0.70 sec)
Records: 0  Duplicates: 0  Warnings: 0

```

Although we did not explicitly name any of the subpartitions when creating table `es`, we can obtain generated names for these by including the `SUBPARTITION_NAME` column of the `PARTITIONS` table from `INFORMATION_SCHEMA` when selecting from that table, as shown here:

```

mysql> SELECT PARTITION_NAME, SUBPARTITION_NAME, TABLE_ROWS
    ->     FROM INFORMATION_SCHEMA.PARTITIONS
    ->     WHERE TABLE_NAME = 'es';
+-----+-----+-----+
| PARTITION_NAME | SUBPARTITION_NAME | TABLE_ROWS |
+-----+-----+-----+
| p0            | p0sp0          |      1 |
| p0            | p0sp1          |      0 |
| p1            | p1sp0          |      0 |
| p1            | p1sp1          |      0 |
| p2            | p2sp0          |      0 |
| p2            | p2sp1          |      0 |
| p3            | p3sp0          |      3 |
| p3            | p3sp1          |      0 |
+-----+-----+-----+
8 rows in set (0.00 sec)

```

The following `ALTER TABLE` statement exchanges subpartition `p3sp0` in table `es` with the nonpartitioned table `es2`:

```

mysql> ALTER TABLE es EXCHANGE PARTITION p3sp0 WITH TABLE es2;
Query OK, 0 rows affected (0.29 sec)

```

You can verify that the rows were exchanged by issuing the following queries:

```

mysql> SELECT PARTITION_NAME, SUBPARTITION_NAME, TABLE_ROWS
    ->     FROM INFORMATION_SCHEMA.PARTITIONS
    ->     WHERE TABLE_NAME = 'es';
+-----+-----+-----+
| PARTITION_NAME | SUBPARTITION_NAME | TABLE_ROWS |
+-----+-----+-----+
| p0            | p0sp0          |      1 |
| p0            | p0sp1          |      0 |
| p1            | p1sp0          |      0 |
| p1            | p1sp1          |      0 |
| p2            | p2sp0          |      0 |
| p2            | p2sp1          |      0 |
+-----+-----+-----+

```

```

| p3          | p3sp0          |          0 |
| p3          | p3sp1          |          0 |
+-----+-----+-----+
8 rows in set (0.00 sec)

mysql> SELECT * FROM es2;
+----+----+----+
| id | fname | lname |
+----+----+----+
| 1669 | Jim   | Smith |
| 337  | Mary  | Jones |
| 2005 | Linda | Black |
+----+----+----+
3 rows in set (0.00 sec)

```

If a table is subpartitioned, you can exchange only a subpartition of the table—not an entire partition—with an unpartitioned table, as shown here:

```

mysql> ALTER TABLE es EXCHANGE PARTITION p3 WITH TABLE es2;
ERROR 1704 (HY000): Subpartitioned table, use subpartition instead of partition

```

Table structures are compared in a strict fashion; the number, order, names, and types of columns and indexes of the partitioned table and the nonpartitioned table must match exactly. In addition, both tables must use the same storage engine:

```

mysql> CREATE TABLE es3 LIKE e;
Query OK, 0 rows affected (1.31 sec)

mysql> ALTER TABLE es3 REMOVE PARTITIONING;
Query OK, 0 rows affected (0.53 sec)
Records: 0  Duplicates: 0  Warnings: 0

mysql> SHOW CREATE TABLE es3\G
***** 1. row *****
    Table: es3
Create Table: CREATE TABLE `es3` (
  `id` int(11) NOT NULL,
  `fname` varchar(30) DEFAULT NULL,
  `lname` varchar(30) DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4
1 row in set (0.00 sec)

mysql> ALTER TABLE es3 ENGINE = MyISAM;
Query OK, 0 rows affected (0.15 sec)
Records: 0  Duplicates: 0  Warnings: 0

mysql> ALTER TABLE es EXCHANGE PARTITION p3sp0 WITH TABLE es3;
ERROR 1497 (HY000): The mix of handlers in the partitions is not allowed in this version of MySQL

```

## 24.3.4 Maintenance of Partitions

A number of table and partition maintenance tasks can be carried out on partitioned tables using SQL statements intended for such purposes.

Table maintenance of partitioned tables can be accomplished using the statements `CHECK TABLE`, `OPTIMIZE TABLE`, `ANALYZE TABLE`, and `REPAIR TABLE`, which are supported for partitioned tables.

You can use a number of extensions to `ALTER TABLE` for performing operations of this type on one or more partitions directly, as described in the following list:

- **Rebuilding partitions.** Rebuilds the partition; this has the same effect as dropping all records stored in the partition, then reinserting them. This can be useful for purposes of defragmentation.

Example:

```
ALTER TABLE t1 REBUILD PARTITION p0, p1;
```

- **Optimizing partitions.** If you have deleted a large number of rows from a partition or if you have made many changes to a partitioned table with variable-length rows (that is, having `VARCHAR`, `BLOB`,

or `TEXT` columns), you can use `ALTER TABLE ... OPTIMIZE PARTITION` to reclaim any unused space and to defragment the partition data file.

Example:

```
ALTER TABLE t1 OPTIMIZE PARTITION p0, p1;
```

Using `OPTIMIZE PARTITION` on a given partition is equivalent to running `CHECK PARTITION`, `ANALYZE PARTITION`, and `REPAIR PARTITION` on that partition.

Some MySQL storage engines, including `InnoDB`, do not support per-partition optimization; in these cases, `ALTER TABLE ... OPTIMIZE PARTITION` analyzes and rebuilds the entire table, and causes an appropriate warning to be issued. (Bug #11751825, Bug #42822) Use `ALTER TABLE ... REBUILD PARTITION` and `ALTER TABLE ... ANALYZE PARTITION` instead, to avoid this issue.

- **Analyzing partitions.** This reads and stores the key distributions for partitions.

Example:

```
ALTER TABLE t1 ANALYZE PARTITION p3;
```

- **Repairing partitions.** This repairs corrupted partitions.

Example:

```
ALTER TABLE t1 REPAIR PARTITION p0,p1;
```

Normally, `REPAIR PARTITION` fails when the partition contains duplicate key errors. You can use `ALTER IGNORE TABLE` with this option, in which case all rows that cannot be moved due to the presence of duplicate keys are removed from the partition (Bug #16900947).

- **Checking partitions.** You can check partitions for errors in much the same way that you can use `CHECK TABLE` with nonpartitioned tables.

Example:

```
ALTER TABLE trb3 CHECK PARTITION p1;
```

This statement tells you whether the data or indexes in partition `p1` of table `t1` are corrupted. If this is the case, use `ALTER TABLE ... REPAIR PARTITION` to repair the partition.

Normally, `CHECK PARTITION` fails when the partition contains duplicate key errors. You can use `ALTER IGNORE TABLE` with this option, in which case the statement returns the contents of each row in the partition where a duplicate key violation is found. Only the values for the columns in the partitioning expression for the table are reported. (Bug #16900947)

Each of the statements in the list just shown also supports the keyword `ALL` in place of the list of partition names. Using `ALL` causes the statement to act on all partitions in the table.

You can also truncate partitions using `ALTER TABLE ... TRUNCATE PARTITION`. This statement can be used to delete all rows from one or more partitions in much the same way that `TRUNCATE TABLE` deletes all rows from a table.

`ALTER TABLE ... TRUNCATE PARTITION ALL` truncates all partitions in the table.

## 24.3.5 Obtaining Information About Partitions

This section discusses obtaining information about existing partitions, which can be done in a number of ways. Methods of obtaining such information include the following:

- Using the `SHOW CREATE TABLE` statement to view the partitioning clauses used in creating a partitioned table.

- Using the `SHOW TABLE STATUS` statement to determine whether a table is partitioned.
- Querying the Information Schema `PARTITIONS` table.
- Using the statement `EXPLAIN SELECT` to see which partitions are used by a given `SELECT`.

From MySQL 8.0.16, when insertions, deletions, or updates are made to partitioned tables, the binary log records information about the partition and (if any) the subpartition in which the row event took place. A new row event is created for a modification that takes place in a different partition or subpartition, even if the table involved is the same. So if a transaction involves three partitions or subpartitions, three row events are generated. For an update event, the partition information is recorded for both the “before” image and the “after” image. The partition information is displayed if you specify the `-v` or `--verbose` option when viewing the binary log using `mysqlbinlog`. Partition information is only recorded when row-based logging is in use (`binlog_format=ROW`).

As discussed elsewhere in this chapter, `SHOW CREATE TABLE` includes in its output the `PARTITION BY` clause used to create a partitioned table. For example:

```
mysql> SHOW CREATE TABLE trb3\G
***** 1. row *****
      Table: trb3
Create Table: CREATE TABLE `trb3` (
  `id` int(11) DEFAULT NULL,
  `name` varchar(50) DEFAULT NULL,
  `purchased` date DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4
/*!50100 PARTITION BY RANGE (YEAR(purchased))
(PARTITION p0 VALUES LESS THAN (1990) ENGINE = InnoDB,
 PARTITION p1 VALUES LESS THAN (1995) ENGINE = InnoDB,
 PARTITION p2 VALUES LESS THAN (2000) ENGINE = InnoDB,
 PARTITION p3 VALUES LESS THAN (2005) ENGINE = InnoDB) */
0 row in set (0.00 sec)
```

The output from `SHOW TABLE STATUS` for partitioned tables is the same as that for nonpartitioned tables, except that the `Create_options` column contains the string `partitioned`. The `Engine` column contains the name of the storage engine used by all partitions of the table. (See Section 13.7.7.38, “`SHOW TABLE STATUS Statement`”, for more information about this statement.)

You can also obtain information about partitions from `INFORMATION_SCHEMA`, which contains a `PARTITIONS` table. See Section 26.3.21, “The `INFORMATION_SCHEMA PARTITIONS Table`”.

It is possible to determine which partitions of a partitioned table are involved in a given `SELECT` query using `EXPLAIN`. The `partitions` column in the `EXPLAIN` output lists the partitions from which records would be matched by the query.

Suppose that a table `trb1` is created and populated as follows:

```
CREATE TABLE trb1 (id INT, name VARCHAR(50), purchased DATE)
  PARTITION BY RANGE(id)
  (
    PARTITION p0 VALUES LESS THAN (3),
    PARTITION p1 VALUES LESS THAN (7),
    PARTITION p2 VALUES LESS THAN (9),
    PARTITION p3 VALUES LESS THAN (11)
  );
  
INSERT INTO trb1 VALUES
  (1, 'desk organiser', '2003-10-15'),
  (2, 'CD player', '1993-11-05'),
  (3, 'TV set', '1996-03-10'),
  (4, 'bookcase', '1982-01-10'),
  (5, 'exercise bike', '2004-05-09'),
  (6, 'sofa', '1987-06-05'),
  (7, 'popcorn maker', '2001-11-22'),
  (8, 'aquarium', '1992-08-04'),
  (9, 'study desk', '1984-09-16'),
```

```
(10, 'lava lamp', '1998-12-25');
```

You can see which partitions are used in a query such as `SELECT * FROM trbl;`, as shown here:

```
mysql> EXPLAIN SELECT * FROM trbl\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: trbl
    partitions: p0,p1,p2,p3
       type: ALL
possible_keys: NULL
         key: NULL
    key_len: NULL
       ref: NULL
      rows: 10
     Extra: Using filesort
```

In this case, all four partitions are searched. However, when a limiting condition making use of the partitioning key is added to the query, you can see that only those partitions containing matching values are scanned, as shown here:

```
mysql> EXPLAIN SELECT * FROM trbl WHERE id < 5\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: trbl
    partitions: p0,p1
       type: ALL
possible_keys: NULL
         key: NULL
    key_len: NULL
       ref: NULL
      rows: 10
     Extra: Using where
```

`EXPLAIN` also provides information about keys used and possible keys:

```
mysql> ALTER TABLE trbl ADD PRIMARY KEY (id);
Query OK, 10 rows affected (0.03 sec)
Records: 10  Duplicates: 0  Warnings: 0

mysql> EXPLAIN SELECT * FROM trbl WHERE id < 5\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: trbl
    partitions: p0,p1
       type: range
possible_keys: PRIMARY
         key: PRIMARY
    key_len: 4
       ref: NULL
      rows: 7
     Extra: Using where
```

If `EXPLAIN` is used to examine a query against a nonpartitioned table, no error is produced, but the value of the `partitions` column is always `NULL`.

The `rows` column of `EXPLAIN` output displays the total number of rows in the table.

See also [Section 13.8.2, “EXPLAIN Statement”](#).

## 24.4 Partition Pruning

The optimization known as *partition pruning* is based on a relatively simple concept which can be described as “Do not scan partitions where there can be no matching values”. Suppose a partitioned table `t1` is created by this statement:

```

CREATE TABLE t1 (
    fname VARCHAR(50) NOT NULL,
    lname VARCHAR(50) NOT NULL,
    region_code TINYINT UNSIGNED NOT NULL,
    dob DATE NOT NULL
)
PARTITION BY RANGE( region_code ) (
    PARTITION p0 VALUES LESS THAN (64),
    PARTITION p1 VALUES LESS THAN (128),
    PARTITION p2 VALUES LESS THAN (192),
    PARTITION p3 VALUES LESS THAN MAXVALUE
);

```

Suppose that you wish to obtain results from a `SELECT` statement such as this one:

```

SELECT fname, lname, region_code, dob
    FROM t1
   WHERE region_code > 125 AND region_code < 130;

```

It is easy to see that none of the rows which ought to be returned are in either of the partitions `p0` or `p3`; that is, we need search only in partitions `p1` and `p2` to find matching rows. By limiting the search, it is possible to expend much less time and effort in finding matching rows than by scanning all partitions in the table. This “cutting away” of unneeded partitions is known as *pruning*. When the optimizer can make use of partition pruning in performing this query, execution of the query can be an order of magnitude faster than the same query against a nonpartitioned table containing the same column definitions and data.

The optimizer can perform pruning whenever a `WHERE` condition can be reduced to either one of the following two cases:

- `partition_column = constant`
- `partition_column IN (constant1, constant2, ..., constantN)`

In the first case, the optimizer simply evaluates the partitioning expression for the value given, determines which partition contains that value, and scans only this partition. In many cases, the equal sign can be replaced with another arithmetic comparison, including `<`, `>`, `<=`, `>=`, and `<>`. Some queries using `BETWEEN` in the `WHERE` clause can also take advantage of partition pruning. See the examples later in this section.

In the second case, the optimizer evaluates the partitioning expression for each value in the list, creates a list of matching partitions, and then scans only the partitions in this partition list.

`SELECT`, `DELETE`, and `UPDATE` statements support partition pruning. An `INSERT` statement also accesses only one partition per inserted row; this is true even for a table that is partitioned by `HASH` or `KEY` although this is not currently shown in the output of `EXPLAIN`.

Pruning can also be applied to short ranges, which the optimizer can convert into equivalent lists of values. For instance, in the previous example, the `WHERE` clause can be converted to `WHERE region_code IN (126, 127, 128, 129)`. Then the optimizer can determine that the first two values in the list are found in partition `p1`, the remaining two values in partition `p2`, and that the other partitions contain no relevant values and so do not need to be searched for matching rows.

The optimizer can also perform pruning for `WHERE` conditions that involve comparisons of the preceding types on multiple columns for tables that use `RANGE COLUMNS` or `LIST COLUMNS` partitioning.

This type of optimization can be applied whenever the partitioning expression consists of an equality or a range which can be reduced to a set of equalities, or when the partitioning expression represents an increasing or decreasing relationship. Pruning can also be applied for tables partitioned on a `DATE` or `DATETIME` column when the partitioning expression uses the `YEAR()` or `TO_DAYS()` function. Pruning can also be applied for such tables when the partitioning expression uses the `TO_SECONDS()` function.

Suppose that table `t2`, partitioned on a `DATE` column, is created using the statement shown here:

```

CREATE TABLE t2 (
    fname VARCHAR(50) NOT NULL,
    lname VARCHAR(50) NOT NULL,
    region_code TINYINT UNSIGNED NOT NULL,
    dob DATE NOT NULL
)
PARTITION BY RANGE( YEAR(dob) ) (
    PARTITION d0 VALUES LESS THAN (1970),
    PARTITION d1 VALUES LESS THAN (1975),
    PARTITION d2 VALUES LESS THAN (1980),
    PARTITION d3 VALUES LESS THAN (1985),
    PARTITION d4 VALUES LESS THAN (1990),
    PARTITION d5 VALUES LESS THAN (2000),
    PARTITION d6 VALUES LESS THAN (2005),
    PARTITION d7 VALUES LESS THAN MAXVALUE
);

```

The following statements using `t2` can make use of partition pruning:

```

SELECT * FROM t2 WHERE dob = '1982-06-23';

UPDATE t2 SET region_code = 8 WHERE dob BETWEEN '1991-02-15' AND '1997-04-25';

DELETE FROM t2 WHERE dob >= '1984-06-21' AND dob <= '1999-06-21'

```

In the case of the last statement, the optimizer can also act as follows:

1. *Find the partition containing the low end of the range.*

`YEAR('1984-06-21')` yields the value `1984`, which is found in partition `d3`.

2. *Find the partition containing the high end of the range.*

`YEAR('1999-06-21')` evaluates to `1999`, which is found in partition `d5`.

3. *Scan only these two partitions and any partitions that may lie between them.*

In this case, this means that only partitions `d3`, `d4`, and `d5` are scanned. The remaining partitions may be safely ignored (and are ignored).



### Important

Invalid `DATE` and `DATETIME` values referenced in the `WHERE` condition of a statement against a partitioned table are treated as `NULL`. This means that a query such as `SELECT * FROM partitioned_table WHERE date_column < '2008-12-00'` does not return any values (see Bug #40972).

So far, we have looked only at examples using `RANGE` partitioning, but pruning can be applied with other partitioning types as well.

Consider a table that is partitioned by `LIST`, where the partitioning expression is increasing or decreasing, such as the table `t3` shown here. (In this example, we assume for the sake of brevity that the `region_code` column is limited to values between 1 and 10 inclusive.)

```

CREATE TABLE t3 (
    fname VARCHAR(50) NOT NULL,
    lname VARCHAR(50) NOT NULL,
    region_code TINYINT UNSIGNED NOT NULL,
    dob DATE NOT NULL
)
PARTITION BY LIST(region_code) (
    PARTITION r0 VALUES IN (1, 3),
    PARTITION r1 VALUES IN (2, 5, 8),
    PARTITION r2 VALUES IN (4, 9),
    PARTITION r3 VALUES IN (6, 7, 10)
);

```

For a statement such as `SELECT * FROM t3 WHERE region_code BETWEEN 1 AND 3`, the optimizer determines in which partitions the values 1, 2, and 3 are found (`r0` and `r1`) and skips the remaining ones (`r2` and `r3`).

For tables that are partitioned by `HASH` or `[LINEAR] KEY`, partition pruning is also possible in cases in which the `WHERE` clause uses a simple `=` relation against a column used in the partitioning expression. Consider a table created like this:

```
CREATE TABLE t4 (
    fname VARCHAR(50) NOT NULL,
    lname VARCHAR(50) NOT NULL,
    region_code TINYINT UNSIGNED NOT NULL,
    dob DATE NOT NULL
)
PARTITION BY KEY(region_code)
PARTITIONS 8;
```

A statement that compares a column value with a constant can be pruned:

```
UPDATE t4 WHERE region_code = 7;
```

Pruning can also be employed for short ranges, because the optimizer can turn such conditions into `IN` relations. For example, using the same table `t4` as defined previously, queries such as these can be pruned:

```
SELECT * FROM t4 WHERE region_code > 2 AND region_code < 6;
SELECT * FROM t4 WHERE region_code BETWEEN 3 AND 5;
```

In both these cases, the `WHERE` clause is transformed by the optimizer into `WHERE region_code IN (3, 4, 5)`.



### Important

This optimization is used only if the range size is smaller than the number of partitions. Consider this statement:

```
DELETE FROM t4 WHERE region_code BETWEEN 4 AND 12;
```

The range in the `WHERE` clause covers 9 values (4, 5, 6, 7, 8, 9, 10, 11, 12), but `t4` has only 8 partitions. This means that the `DELETE` cannot be pruned.

When a table is partitioned by `HASH` or `[LINEAR] KEY`, pruning can be used only on integer columns. For example, this statement cannot use pruning because `dob` is a `DATE` column:

```
SELECT * FROM t4 WHERE dob >= '2001-04-14' AND dob <= '2005-10-15';
```

However, if the table stores year values in an `INT` column, then a query having `WHERE year_col >= 2001 AND year_col <= 2005` can be pruned.

Tables using a storage engine that provides automatic partitioning, such as the `NDB` storage engine used by MySQL Cluster can be pruned if they are explicitly partitioned.

## 24.5 Partition Selection

Explicit selection of partitions and subpartitions for rows matching a given `WHERE` condition is supported. Partition selection is similar to partition pruning, in that only specific partitions are checked for matches, but differs in two key respects:

1. The partitions to be checked are specified by the issuer of the statement, unlike partition pruning, which is automatic.
2. Whereas partition pruning applies only to queries, explicit selection of partitions is supported for both queries and a number of DML statements.

SQL statements supporting explicit partition selection are listed here:

- `SELECT`
- `DELETE`
- `INSERT`
- `REPLACE`
- `UPDATE`
- `LOAD DATA`.
- `LOAD XML`.

The remainder of this section discusses explicit partition selection as it applies generally to the statements just listed, and provides some examples.

Explicit partition selection is implemented using a `PARTITION` option. For all supported statements, this option uses the syntax shown here:

```
PARTITION (partition_names)
partition_names:
  partition_name, ...
```

This option always follows the name of the table to which the partition or partitions belong. `partition_names` is a comma-separated list of partitions or subpartitions to be used. Each name in this list must be the name of an existing partition or subpartition of the specified table; if any of the partitions or subpartitions are not found, the statement fails with an error (`partition 'partition_name' doesn't exist`). Partitions and subpartitions named in `partition_names` may be listed in any order, and may overlap.

When the `PARTITION` option is used, only the partitions and subpartitions listed are checked for matching rows. This option can be used in a `SELECT` statement to determine which rows belong to a given partition. Consider a partitioned table named `employees`, created and populated using the statements shown here:

```
SET @@SQL_MODE = '';
CREATE TABLE employees (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    fname VARCHAR(25) NOT NULL,
    lname VARCHAR(25) NOT NULL,
    store_id INT NOT NULL,
    department_id INT NOT NULL
)
PARTITION BY RANGE(id) (
    PARTITION p0 VALUES LESS THAN (5),
    PARTITION p1 VALUES LESS THAN (10),
    PARTITION p2 VALUES LESS THAN (15),
    PARTITION p3 VALUES LESS THAN MAXVALUE
);
INSERT INTO employees VALUES
    ('', 'Bob', 'Taylor', 3, 2), ('', 'Frank', 'Williams', 1, 2),
    ('', 'Ellen', 'Johnson', 3, 4), ('', 'Jim', 'Smith', 2, 4),
    ('', 'Mary', 'Jones', 1, 1), ('', 'Linda', 'Black', 2, 3),
    ('', 'Ed', 'Jones', 2, 1), ('', 'June', 'Wilson', 3, 1),
    ('', 'Andy', 'Smith', 1, 3), ('', 'Lou', 'Waters', 2, 4),
    ('', 'Jill', 'Stone', 1, 4), ('', 'Roger', 'White', 3, 2),
    ('', 'Howard', 'Andrews', 1, 2), ('', 'Fred', 'Goldberg', 3, 3),
    ('', 'Barbara', 'Brown', 2, 3), ('', 'Alice', 'Rogers', 2, 2),
    ('', 'Mark', 'Morgan', 3, 3), ('', 'Karen', 'Cole', 3, 2);
```

You can see which rows are stored in partition `p1` like this:

```
mysql> SELECT * FROM employees PARTITION (p1);
+----+-----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+----+-----+-----+-----+-----+
| 5 | Mary | Jones | 1 | 1 |
| 6 | Linda | Black | 2 | 3 |
| 7 | Ed | Jones | 2 | 1 |
| 8 | June | Wilson | 3 | 1 |
| 9 | Andy | Smith | 1 | 3 |
+----+-----+-----+-----+
5 rows in set (0.00 sec)
```

The result is the same as obtained by the query `SELECT * FROM employees WHERE id BETWEEN 5 AND 9.`

To obtain rows from multiple partitions, supply their names as a comma-delimited list. For example, `SELECT * FROM employees PARTITION (p1, p2)` returns all rows from partitions `p1` and `p2` while excluding rows from the remaining partitions.

Any valid query against a partitioned table can be rewritten with a `PARTITION` option to restrict the result to one or more desired partitions. You can use `WHERE` conditions, `ORDER BY` and `LIMIT` options, and so on. You can also use aggregate functions with `HAVING` and `GROUP BY` options. Each of the following queries produces a valid result when run on the `employees` table as previously defined:

```
mysql> SELECT * FROM employees PARTITION (p0, p2)
    ->     WHERE lname LIKE 'S%';
+----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+----+-----+-----+-----+
| 4 | Jim | Smith | 2 | 4 |
| 11 | Jill | Stone | 1 | 4 |
+----+-----+-----+-----+
2 rows in set (0.00 sec)

mysql> SELECT id, CONCAT(fname, ' ', lname) AS name
    ->     FROM employees PARTITION (p0) ORDER BY lname;
+----+-----+
| id | name      |
+----+-----+
| 3 | Ellen Johnson |
| 4 | Jim Smith   |
| 1 | Bob Taylor   |
| 2 | Frank Williams |
+----+-----+
4 rows in set (0.06 sec)

mysql> SELECT store_id, COUNT(department_id) AS c
    ->     FROM employees PARTITION (p1,p2,p3)
    ->     GROUP BY store_id HAVING c > 4;
+----+-----+
| c | store_id |
+----+-----+
| 5 |          2 |
| 5 |          3 |
+----+-----+
2 rows in set (0.00 sec)
```

Statements using partition selection can be employed with tables using any of the supported partitioning types. When a table is created using `[LINEAR] HASH` or `[LINEAR] KEY` partitioning and the names of the partitions are not specified, MySQL automatically names the partitions `p0, p1, p2, ..., pN-1`, where `N` is the number of partitions. For subpartitions not explicitly named, MySQL assigns automatically to the subpartitions in each partition `pX` the names `pXsp0, pXsp1, pXsp2, ..., pXspM-1`, where `M` is the number of subpartitions. When executing against this table a `SELECT` (or other SQL statement for which explicit partition selection is allowed), you can use these generated names in a `PARTITION` option, as shown here:

```
mysql> CREATE TABLE employees_sub (
    ->     id INT NOT NULL AUTO_INCREMENT,
```

```

->      fname VARCHAR(25) NOT NULL,
->      lname VARCHAR(25) NOT NULL,
->      store_id INT NOT NULL,
->      department_id INT NOT NULL,
->      PRIMARY KEY pk (id, lname)
->  )
->      PARTITION BY RANGE(id)
->      SUBPARTITION BY KEY (lname)
->      SUBPARTITIONS 2 (
->          PARTITION p0 VALUES LESS THAN (5),
->          PARTITION p1 VALUES LESS THAN (10),
->          PARTITION p2 VALUES LESS THAN (15),
->          PARTITION p3 VALUES LESS THAN MAXVALUE
->      );
Query OK, 0 rows affected (1.14 sec)

mysql> INSERT INTO employees_sub    # reuse data in employees table
->      SELECT * FROM employees;
Query OK, 18 rows affected (0.09 sec)
Records: 18  Duplicates: 0  Warnings: 0

mysql> SELECT id, CONCAT(fname, ' ', lname) AS name
->      FROM employees_sub PARTITION (p2sp1);
+----+-----+
| id | name      |
+----+-----+
| 10 | Lou Waters |
| 14 | Fred Goldberg |
+----+-----+
2 rows in set (0.00 sec)

```

You may also use a `PARTITION` option in the `SELECT` portion of an `INSERT ... SELECT` statement, as shown here:

```

mysql> CREATE TABLE employees_copy LIKE employees;
Query OK, 0 rows affected (0.28 sec)

mysql> INSERT INTO employees_copy
->      SELECT * FROM employees PARTITION (p2);
Query OK, 5 rows affected (0.04 sec)
Records: 5  Duplicates: 0  Warnings: 0

mysql> SELECT * FROM employees_copy;
+----+-----+-----+-----+-----+
| id | fname | lname   | store_id | department_id |
+----+-----+-----+-----+-----+
| 10 | Lou   | Waters  |       2 |         4 |
| 11 | Jill  | Stone   |       1 |         4 |
| 12 | Roger | White   |       3 |         2 |
| 13 | Howard | Andrews |       1 |         2 |
| 14 | Fred   | Goldberg |       3 |         3 |
+----+-----+-----+-----+-----+
5 rows in set (0.00 sec)

```

Partition selection can also be used with joins. Suppose we create and populate two tables using the statements shown here:

```

CREATE TABLE stores (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    city VARCHAR(30) NOT NULL
)
PARTITION BY HASH(id)
PARTITIONS 2;

INSERT INTO stores VALUES
    ('', 'Nambucca'), ('', 'Uranga'),
    ('', 'Bellingen'), ('', 'Grafton');

CREATE TABLE departments (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(30) NOT NULL
)

```

```
PARTITION BY KEY(id)
PARTITIONS 2;

INSERT INTO departments VALUES
('Sales'), ('Customer Service'),
('Delivery'), ('Accounting');
```

You can explicitly select partitions (or subpartitions, or both) from any or all of the tables in a join. (The `PARTITION` option used to select partitions from a given table immediately follows the name of the table, before all other options, including any table alias.) For example, the following query gets the name, employee ID, department, and city of all employees who work in the Sales or Delivery department (partition `p1` of the `departments` table) at the stores in either of the cities of Nambucca and Bellingen (partition `p0` of the `stores` table):

```
mysql> SELECT
    ->     e.id AS 'Employee ID', CONCAT(e.fname, ' ', e.lname) AS Name,
    ->     s.city AS City, d.name AS department
    ->   FROM employees AS e
    ->   JOIN stores PARTITION (p1) AS s ON e.store_id=s.id
    ->   JOIN departments PARTITION (p0) AS d ON e.department_id=d.id
    -> ORDER BY e.lname;
+-----+-----+-----+-----+
| Employee ID | Name      | City     | department |
+-----+-----+-----+-----+
|      14 | Fred Goldberg | Bellingen | Delivery   |
|       5 | Mary Jones    | Nambucca | Sales      |
|      17 | Mark Morgan   | Bellingen | Delivery   |
|       9 | Andy Smith    | Nambucca | Delivery   |
|       8 | June Wilson   | Bellingen | Sales      |
+-----+-----+-----+-----+
5 rows in set (0.00 sec)
```

For general information about joins in MySQL, see [Section 13.2.13.2, “JOIN Clause”](#).

When the `PARTITION` option is used with `DELETE` statements, only those partitions (and subpartitions, if any) listed with the option are checked for rows to be deleted. Any other partitions are ignored, as shown here:

```
mysql> SELECT * FROM employees WHERE fname LIKE 'j%';
+-----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+-----+-----+-----+-----+
|  4 | Jim   | Smith |        2 |          4 |
|  8 | June  | Wilson |       3 |          1 |
| 11 | Jill  | Stone |        1 |          4 |
+-----+-----+-----+-----+
3 rows in set (0.00 sec)

mysql> DELETE FROM employees PARTITION (p0, p1)
    ->     WHERE fname LIKE 'j%';
Query OK, 2 rows affected (0.09 sec)

mysql> SELECT * FROM employees WHERE fname LIKE 'j%';
+-----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+-----+-----+-----+-----+
| 11 | Jill  | Stone |        1 |          4 |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Only the two rows in partitions `p0` and `p1` matching the `WHERE` condition were deleted. As you can see from the result when the `SELECT` is run a second time, there remains a row in the table matching the `WHERE` condition, but residing in a different partition (`p2`).

`UPDATE` statements using explicit partition selection behave in the same way; only rows in the partitions referenced by the `PARTITION` option are considered when determining the rows to be updated, as can be seen by executing the following statements:

```
mysql> UPDATE employees PARTITION (p0)
```

```

->      SET store_id = 2 WHERE fname = 'Jill';
Query OK, 0 rows affected (0.00 sec)
Rows matched: 0  Changed: 0  Warnings: 0

mysql> SELECT * FROM employees WHERE fname = 'Jill';
+----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+----+-----+-----+-----+
| 11 | Jill  | Stone |        1 |          4 |
+----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> UPDATE employees PARTITION (p2)
->      SET store_id = 2 WHERE fname = 'Jill';
Query OK, 1 row affected (0.09 sec)
Rows matched: 1  Changed: 1  Warnings: 0

mysql> SELECT * FROM employees WHERE fname = 'Jill';
+----+-----+-----+-----+
| id | fname | lname | store_id | department_id |
+----+-----+-----+-----+
| 11 | Jill  | Stone |        2 |          4 |
+----+-----+-----+-----+
1 row in set (0.00 sec)

```

In the same way, when `PARTITION` is used with `DELETE`, only rows in the partition or partitions named in the partition list are checked for deletion.

For statements that insert rows, the behavior differs in that failure to find a suitable partition causes the statement to fail. This is true for both `INSERT` and `REPLACE` statements, as shown here:

```

mysql> INSERT INTO employees PARTITION (p2) VALUES (20, 'Jan', 'Jones', 1, 3);
ERROR 1729 (HY000): Found a row not matching the given partition set
mysql> INSERT INTO employees PARTITION (p3) VALUES (20, 'Jan', 'Jones', 1, 3);
Query OK, 1 row affected (0.07 sec)

mysql> REPLACE INTO employees PARTITION (p0) VALUES (20, 'Jan', 'Jones', 3, 2);
ERROR 1729 (HY000): Found a row not matching the given partition set

mysql> REPLACE INTO employees PARTITION (p3) VALUES (20, 'Jan', 'Jones', 3, 2);
Query OK, 2 rows affected (0.09 sec)

```

For statements that write multiple rows to a partitioned table that using the `InnoDB` storage engine: If any row in the list following `VALUES` cannot be written to one of the partitions specified in the `partition_names` list, the entire statement fails and no rows are written. This is shown for `INSERT` statements in the following example, reusing the `employees` table created previously:

```

mysql> ALTER TABLE employees
->      REORGANIZE PARTITION p3 INTO (
->          PARTITION p3 VALUES LESS THAN (20),
->          PARTITION p4 VALUES LESS THAN (25),
->          PARTITION p5 VALUES LESS THAN MAXVALUE
->      );
Query OK, 6 rows affected (2.09 sec)
Records: 6  Duplicates: 0  Warnings: 0

mysql> SHOW CREATE TABLE employees\G
***** 1. row *****
      Table: employees
Create Table: CREATE TABLE `employees` (
  `id` int(11) NOT NULL AUTO_INCREMENT,
  `fname` varchar(25) NOT NULL,
  `lname` varchar(25) NOT NULL,
  `store_id` int(11) NOT NULL,
  `department_id` int(11) NOT NULL,
  PRIMARY KEY (`id`)
) ENGINE=InnoDB AUTO_INCREMENT=27 DEFAULT CHARSET=utf8mb4
/*!50100 PARTITION BY RANGE (id)
(PARTITION p0 VALUES LESS THAN (5) ENGINE = InnoDB,
 PARTITION p1 VALUES LESS THAN (10) ENGINE = InnoDB,
 PARTITION p2 VALUES LESS THAN (15) ENGINE = InnoDB,

```

```

PARTITION p3 VALUES LESS THAN (20) ENGINE = InnoDB,
PARTITION p4 VALUES LESS THAN (25) ENGINE = InnoDB,
PARTITION p5 VALUES LESS THAN MAXVALUE ENGINE = InnoDB) */
1 row in set (0.00 sec)

mysql> INSERT INTO employees PARTITION (p3, p4) VALUES
->      (24, 'Tim', 'Greene', 3, 1), (26, 'Linda', 'Mills', 2, 1);
ERROR 1729 (HY000): Found a row not matching the given partition set

mysql> INSERT INTO employees PARTITION (p3, p4, p5) VALUES
->      (24, 'Tim', 'Greene', 3, 1), (26, 'Linda', 'Mills', 2, 1);
Query OK, 2 rows affected (0.06 sec)
Records: 2  Duplicates: 0  Warnings: 0

```

The preceding is true for both `INSERT` statements and `REPLACE` statements that write multiple rows.

Partition selection is disabled for tables employing a storage engine that supplies automatic partitioning, such as `NDB`.

## 24.6 Restrictions and Limitations on Partitioning

This section discusses current restrictions and limitations on MySQL partitioning support.

**Prohibited constructs.** The following constructs are not permitted in partitioning expressions:

- Stored procedures, stored functions, loadable functions, or plugins.
- Declared variables or user variables.

For a list of SQL functions which are permitted in partitioning expressions, see [Section 24.6.3, “Partitioning Limitations Relating to Functions”](#).

**Arithmetic and logical operators.** Use of the arithmetic operators `+`, `-`, and `*` is permitted in partitioning expressions. However, the result must be an integer value or `NULL` (except in the case of `[LINEAR] KEY` partitioning, as discussed elsewhere in this chapter; see [Section 24.2, “Partitioning Types”](#), for more information).

The `DIV` operator is also supported; the `/` operator is not permitted.

The bit operators `|`, `&`, `^`, `<<`, `>>`, and `~` are not permitted in partitioning expressions.

**Server SQL mode.** Tables employing user-defined partitioning do not preserve the SQL mode in effect at the time that they were created. As discussed elsewhere in this Manual (see [Section 5.1.11, “Server SQL Modes”](#)), the results of many MySQL functions and operators may change according to the server SQL mode. Therefore, a change in the SQL mode at any time after the creation of partitioned tables may lead to major changes in the behavior of such tables, and could easily lead to corruption or loss of data. For these reasons, *it is strongly recommended that you never change the server SQL mode after creating partitioned tables*.

For one such change in the server SQL mode making a partitioned tables unusable, consider the following `CREATE TABLE` statement, which can be executed successfully only if the `NO_UNSIGNED_SUBTRACTION` mode is in effect:

```

mysql> SELECT @@sql_mode;
+-----+
| @@sql_mode |
+-----+
|          |
+-----+
1 row in set (0.00 sec)

mysql> CREATE TABLE tu (c1 BIGINT UNSIGNED)
->   PARTITION BY RANGE(c1 - 10) (
->     PARTITION p0 VALUES LESS THAN (-5),
->     PARTITION p1 VALUES LESS THAN (0),
->     PARTITION p2 VALUES LESS THAN (5),

```

```

->      PARTITION p3 VALUES LESS THAN (10),
->      PARTITION p4 VALUES LESS THAN (MAXVALUE)
->  );
ERROR 1563 (HY000): Partition constant is out of partition function domain

mysql> SET sql_mode='NO_UNSIGNED_SUBTRACTION';
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT @@sql_mode;
+-----+
| @@sql_mode           |
+-----+
| NO_UNSIGNED_SUBTRACTION |
+-----+
1 row in set (0.00 sec)

mysql> CREATE TABLE tu (c1 BIGINT UNSIGNED)
->      PARTITION BY RANGE(c1 - 10) (
->          PARTITION p0 VALUES LESS THAN (-5),
->          PARTITION p1 VALUES LESS THAN (0),
->          PARTITION p2 VALUES LESS THAN (5),
->          PARTITION p3 VALUES LESS THAN (10),
->          PARTITION p4 VALUES LESS THAN (MAXVALUE)
->      );
Query OK, 0 rows affected (0.05 sec)

```

If you remove the `NO_UNSIGNED_SUBTRACTION` server SQL mode after creating `tu`, you may no longer be able to access this table:

```

mysql> SET sql_mode='';
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT * FROM tu;
ERROR 1563 (HY000): Partition constant is out of partition function domain
mysql> INSERT INTO tu VALUES (20);
ERROR 1563 (HY000): Partition constant is out of partition function domain

```

See also [Section 5.1.11, “Server SQL Modes”](#).

Server SQL modes also impact replication of partitioned tables. Disparate SQL modes on source and replica can lead to partitioning expressions being evaluated differently; this can cause the distribution of data among partitions to be different in the source's and replica's copies of a given table, and may even cause inserts into partitioned tables that succeed on the source to fail on the replica. For best results, you should always use the same server SQL mode on the source and on the replica.

**Performance considerations.** Some effects of partitioning operations on performance are given in the following list:

- **File system operations.** Partitioning and repartitioning operations (such as `ALTER TABLE` with `PARTITION BY ...`, `REORGANIZE PARTITION`, or `REMOVE PARTITIONING`) depend on file system operations for their implementation. This means that the speed of these operations is affected by such factors as file system type and characteristics, disk speed, swap space, file handling efficiency of the operating system, and MySQL server options and variables that relate to file handling. In particular, you should make sure that `large_files_support` is enabled and that `open_files_limit` is set properly. Partitioning and repartitioning operations involving `InnoDB` tables may be made more efficient by enabling `innodb_file_per_table`.

See also [Maximum number of partitions](#).

- **Table locks.** Generally, the process executing a partitioning operation on a table takes a write lock on the table. Reads from such tables are relatively unaffected; pending `INSERT` and `UPDATE` operations are performed as soon as the partitioning operation has completed. For `InnoDB`-specific exceptions to this limitation, see [Partitioning Operations](#).
- **Indexes; partition pruning.** As with nonpartitioned tables, proper use of indexes can speed up queries on partitioned tables significantly. In addition, designing partitioned tables and queries

on these tables to take advantage of *partition pruning* can improve performance dramatically. See [Section 24.4, “Partition Pruning”](#), for more information.

Index condition pushdown is supported for partitioned tables. See [Section 8.2.1.6, “Index Condition Pushdown Optimization”](#).

- **Performance with LOAD DATA.** In MySQL 8.0, `LOAD DATA` uses buffering to improve performance. You should be aware that the buffer uses 130 KB memory per partition to achieve this.

#### Maximum number of partitions.

The maximum possible number of partitions for a given table not using the `NDB` storage engine is 8192. This number includes subpartitions.

The maximum possible number of user-defined partitions for a table using the `NDB` storage engine is determined according to the version of the NDB Cluster software being used, the number of data nodes, and other factors. See [NDB and user-defined partitioning](#), for more information.

If, when creating tables with a large number of partitions (but less than the maximum), you encounter an error message such as `Got error ... from storage engine: Out of resources when opening file`, you may be able to address the issue by increasing the value of the `open_files_limit` system variable. However, this is dependent on the operating system, and may not be possible or advisable on all platforms; see [Section B.3.2.16, “File Not Found and Similar Errors”](#), for more information. In some cases, using large numbers (hundreds) of partitions may also not be advisable due to other concerns, so using more partitions does not automatically lead to better results.

See also [File system operations](#).

#### Foreign keys not supported for partitioned InnoDB tables.

Partitioned tables using the `InnoDB` storage engine do not support foreign keys. More specifically, this means that the following two statements are true:

1. No definition of an `InnoDB` table employing user-defined partitioning may contain foreign key references; no `InnoDB` table whose definition contains foreign key references may be partitioned.
2. No `InnoDB` table definition may contain a foreign key reference to a user-partitioned table; no `InnoDB` table with user-defined partitioning may contain columns referenced by foreign keys.

The scope of the restrictions just listed includes all tables that use the `InnoDB` storage engine. `CREATE TABLE` and `ALTER TABLE` statements that would result in tables violating these restrictions are not allowed.

**ALTER TABLE ... ORDER BY.** An `ALTER TABLE ... ORDER BY column` statement run against a partitioned table causes ordering of rows only within each partition.

**ADD COLUMN ... ALGORITHM=INSTANT.** Once you perform `ALTER TABLE ... ADD COLUMN ... ALGORITHM=INSTANT` on a partitioned table, it is no longer possible to exchange partitions with this table.

**Effects on REPLACE statements by modification of primary keys.** It can be desirable in some cases (see [Section 24.6.1, “Partitioning Keys, Primary Keys, and Unique Keys”](#)) to modify a table's primary key. Be aware that, if your application uses `REPLACE` statements and you do this, the results of these statements can be drastically altered. See [Section 13.2.12, “REPLACE Statement”](#), for more information and an example.

#### FULLTEXT indexes.

Partitioned tables do not support `FULLTEXT` indexes or searches.

**Spatial columns.** Columns with spatial data types such as `POINT` or `GEOMETRY` cannot be used in partitioned tables.

#### Temporary tables.

Temporary tables cannot be partitioned.

**Log tables.** It is not possible to partition the log tables; an `ALTER TABLE ... PARTITION BY ...` statement on such a table fails with an error.

#### Data type of partitioning key.

A partitioning key must be either an integer column or an expression that resolves to an integer. Expressions employing `ENUM` columns cannot be used. The column or expression value may also be `NULL`; see [Section 24.2.7, “How MySQL Partitioning Handles NULL”](#).

There are two exceptions to this restriction:

- When partitioning by `[LINEAR] KEY`, it is possible to use columns of any valid MySQL data type other than `TEXT` or `BLOB` as partitioning keys, because the internal key-hashing functions produce the correct data type from these types. For example, the following two `CREATE TABLE` statements are valid:

```
CREATE TABLE tkc (c1 CHAR)
PARTITION BY KEY(c1)
PARTITIONS 4;

CREATE TABLE tke
  ( c1 ENUM('red', 'orange', 'yellow', 'green', 'blue', 'indigo', 'violet') )
PARTITION BY LINEAR KEY(c1)
PARTITIONS 6;
```

- When partitioning by `RANGE COLUMNS` or `LIST COLUMNS`, it is possible to use string, `DATE`, and `DATETIME` columns. For example, each of the following `CREATE TABLE` statements is valid:

```
CREATE TABLE rc (c1 INT, c2 DATE)
PARTITION BY RANGE COLUMNS(c2) (
    PARTITION p0 VALUES LESS THAN('1990-01-01'),
    PARTITION p1 VALUES LESS THAN('1995-01-01'),
    PARTITION p2 VALUES LESS THAN('2000-01-01'),
    PARTITION p3 VALUES LESS THAN('2005-01-01'),
    PARTITION p4 VALUES LESS THAN(MAXVALUE)
);

CREATE TABLE lc (c1 INT, c2 CHAR(1))
PARTITION BY LIST COLUMNS(c2) (
    PARTITION p0 VALUES IN('a', 'd', 'g', 'j', 'm', 'p', 's', 'v', 'y'),
    PARTITION p1 VALUES IN('b', 'e', 'h', 'k', 'n', 'q', 't', 'w', 'z'),
    PARTITION p2 VALUES IN('c', 'f', 'i', 'l', 'o', 'r', 'u', 'x', NULL)
);
```

Neither of the preceding exceptions applies to `BLOB` or `TEXT` column types.

#### Subqueries.

A partitioning key may not be a subquery, even if that subquery resolves to an integer value or `NULL`.

**Column index prefixes not supported for key partitioning.** When creating a table that is partitioned by key, any columns in the partitioning key which use column prefixes are not used in the table's partitioning function. Consider the following `CREATE TABLE` statement, which has three `VARCHAR` columns, and whose primary key uses all three columns and specifies prefixes for two of them:

```
CREATE TABLE t1 (
    a VARCHAR(10000),
    b VARCHAR(25),
    c VARCHAR(10),
    PRIMARY KEY (a(10), b, c(2))
) PARTITION BY KEY() PARTITIONS 2;
```

This statement is accepted, but the resulting table is actually created as if you had issued the following statement, using only the primary key column which does not include a prefix (column `b`) for the partitioning key:

```
CREATE TABLE t1 (
    a VARCHAR(10000),
    b VARCHAR(25),
    c VARCHAR(10),
    PRIMARY KEY (a(10), b, c(2))
) PARTITION BY KEY(b) PARTITIONS 2;
```

Prior to MySQL 8.0.21, no warning was issued or any other indication provided that this occurred, except in the event that all columns specified for the partitioning key used prefixes, in which case the statement failed, but with a misleading error message, as shown here:

```
mysql> CREATE TABLE t2 (
->     a VARCHAR(10000),
->     b VARCHAR(25),
->     c VARCHAR(10),
->     PRIMARY KEY (a(10), b(5), c(2))
-> ) PARTITION BY KEY() PARTITIONS 2;
ERROR 1503 (HY000): A PRIMARY KEY must include all columns in the
table's partitioning function
```

This also occurred when performing `ALTER TABLE` or when upgrading such tables.

This permissive behavior is deprecated as of MySQL 8.0.21 (and is subject to removal in a future version of MySQL). Beginning with MySQL 8.0.21, using one or more columns having a prefix in the partitioning key results in a warning for each such column, as shown here:

```
mysql> CREATE TABLE t1 (
->     a VARCHAR(10000),
->     b VARCHAR(25),
->     c VARCHAR(10),
->     PRIMARY KEY (a(10), b, c(2))
-> ) PARTITION BY KEY() PARTITIONS 2;
Query OK, 0 rows affected, 2 warnings (1.25 sec)

mysql> SHOW WARNINGS\G
***** 1. row *****
Level: Warning
Code: 1681
Message: Column 'test.t1.a' having prefix key part 'a(10)' is ignored by the
partitioning function. Use of prefixed columns in the PARTITION BY KEY() clause
is deprecated and will be removed in a future release.
***** 2. row *****
Level: Warning
Code: 1681
Message: Column 'test.t1.c' having prefix key part 'c(2)' is ignored by the
partitioning function. Use of prefixed columns in the PARTITION BY KEY() clause
is deprecated and will be removed in a future release.
2 rows in set (0.00 sec)
```

This includes cases in which the columns used in the partitioning function are defined implicitly as those in the table's primary key by employing an empty `PARTITION BY KEY()` clause.

In MySQL 8.0.21 and later, if all columns specified for the partitioning key employ prefixes, the `CREATE TABLE` statement used fails with an error message that identifies the issue correctly:

```
mysql> CREATE TABLE t1 (
->     a VARCHAR(10000),
->     b VARCHAR(25),
->     c VARCHAR(10),
->     PRIMARY KEY (a(10), b(5), c(2))
-> ) PARTITION BY KEY() PARTITIONS 2;
ERROR 1503 (HY000): A PRIMARY KEY must include all columns in the table's
partitioning function (prefixed columns are not considered).
```

For general information about partitioning tables by key, see [Section 24.2.5, “KEY Partitioning”](#).

### Issues with subpartitions.

Subpartitions must use `HASH` or `KEY` partitioning. Only `RANGE` and `LIST` partitions may be subpartitioned; `HASH` and `KEY` partitions cannot be subpartitioned.

`SUBPARTITION BY KEY` requires that the subpartitioning column or columns be specified explicitly, unlike the case with `PARTITION BY KEY`, where it can be omitted (in which case the table's primary key column is used by default). Consider the table created by this statement:

```
CREATE TABLE ts (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(30)
);
```

You can create a table having the same columns, partitioned by `KEY`, using a statement such as this one:

```
CREATE TABLE ts (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(30)
)
PARTITION BY KEY()
PARTITIONS 4;
```

The previous statement is treated as though it had been written like this, with the table's primary key column used as the partitioning column:

```
CREATE TABLE ts (
    id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    name VARCHAR(30)
)
PARTITION BY KEY(id)
PARTITIONS 4;
```

However, the following statement that attempts to create a subpartitioned table using the default column as the subpartitioning column fails, and the column must be specified for the statement to succeed, as shown here:

```
mysql> CREATE TABLE ts (
->     id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
->     name VARCHAR(30)
-> )
-> PARTITION BY RANGE(id)
-> SUBPARTITION BY KEY()
-> SUBPARTITIONS 4
-> (
->     PARTITION p0 VALUES LESS THAN (100),
->     PARTITION p1 VALUES LESS THAN (MAXVALUE)
-> );
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use near ')'

mysql> CREATE TABLE ts (
->     id INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
->     name VARCHAR(30)
-> )
-> PARTITION BY RANGE(id)
-> SUBPARTITION BY KEY(id)
-> SUBPARTITIONS 4
-> (
->     PARTITION p0 VALUES LESS THAN (100),
->     PARTITION p1 VALUES LESS THAN (MAXVALUE)
-> );
Query OK, 0 rows affected (0.07 sec)
```

This is a known issue (see Bug #51470).

**DATA DIRECTORY and INDEX DIRECTORY options.** Table-level `DATA DIRECTORY` and `INDEX DIRECTORY` options are ignored (see Bug #32091). You can employ these options for individual partitions or subpartitions of `InnoDB` tables. As of MySQL 8.0.21, the directory specified in a `DATA DIRECTORY` clause must be known to `InnoDB`. For more information, see [Using the DATA DIRECTORY Clause](#).

**Repairing and rebuilding partitioned tables.** The statements `CHECK TABLE`, `OPTIMIZE TABLE`, `ANALYZE TABLE`, and `REPAIR TABLE` are supported for partitioned tables.

In addition, you can use `ALTER TABLE ... REBUILD PARTITION` to rebuild one or more partitions of a partitioned table; `ALTER TABLE ... REORGANIZE PARTITION` also causes partitions to be rebuilt. See Section 13.1.9, “[ALTER TABLE Statement](#)”, for more information about these two statements.

`ANALYZE`, `CHECK`, `OPTIMIZE`, `REPAIR`, and `TRUNCATE` operations are supported with subpartitions. See Section 13.1.9.1, “[ALTER TABLE Partition Operations](#)”.

**File name delimiters for partitions and subpartitions.** Table partition and subpartition file names include generated delimiters such as `#P#` and `#SP#`. The lettercase of such delimiters can vary and should not be depended upon.

## 24.6.1 Partitioning Keys, Primary Keys, and Unique Keys

This section discusses the relationship of partitioning keys with primary keys and unique keys. The rule governing this relationship can be expressed as follows: All columns used in the partitioning expression for a partitioned table must be part of every unique key that the table may have.

In other words, *every unique key on the table must use every column in the table's partitioning expression.* (This also includes the table's primary key, since it is by definition a unique key. This particular case is discussed later in this section.) For example, each of the following table creation statements is invalid:

```
CREATE TABLE t1 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    UNIQUE KEY (col1, col2)
)
PARTITION BY HASH(col3)
PARTITIONS 4;

CREATE TABLE t2 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    UNIQUE KEY (col1),
    UNIQUE KEY (col3)
)
PARTITION BY HASH(col1 + col3)
PARTITIONS 4;
```

In each case, the proposed table would have at least one unique key that does not include all columns used in the partitioning expression.

Each of the following statements is valid, and represents one way in which the corresponding invalid table creation statement could be made to work:

```
CREATE TABLE t1 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    UNIQUE KEY (col1, col2, col3)
)
PARTITION BY HASH(col3)
PARTITIONS 4;

CREATE TABLE t2 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
```

```
        col4 INT NOT NULL,
        UNIQUE KEY (col1, col3)
    )
PARTITION BY HASH(col1 + col3)
PARTITIONS 4;
```

This example shows the error produced in such cases:

```
mysql> CREATE TABLE t3 (
    ->     col1 INT NOT NULL,
    ->     col2 DATE NOT NULL,
    ->     col3 INT NOT NULL,
    ->     col4 INT NOT NULL,
    ->     UNIQUE KEY (col1, col2),
    ->     UNIQUE KEY (col3)
    -> )
    -> PARTITION BY HASH(col1 + col3)
    -> PARTITIONS 4;
ERROR 1491 (HY000): A PRIMARY KEY must include all columns in the table's partitioning function
```

The `CREATE TABLE` statement fails because both `col1` and `col3` are included in the proposed partitioning key, but neither of these columns is part of both of unique keys on the table. This shows one possible fix for the invalid table definition:

```
mysql> CREATE TABLE t3 (
    ->     col1 INT NOT NULL,
    ->     col2 DATE NOT NULL,
    ->     col3 INT NOT NULL,
    ->     col4 INT NOT NULL,
    ->     UNIQUE KEY (col1, col2, col3),
    ->     UNIQUE KEY (col3)
    -> )
    -> PARTITION BY HASH(col3)
    -> PARTITIONS 4;
Query OK, 0 rows affected (0.05 sec)
```

In this case, the proposed partitioning key `col3` is part of both unique keys, and the table creation statement succeeds.

The following table cannot be partitioned at all, because there is no way to include in a partitioning key any columns that belong to both unique keys:

```
CREATE TABLE t4 (
    col1 INT NOT NULL,
    col2 INT NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    UNIQUE KEY (col1, col3),
    UNIQUE KEY (col2, col4)
);
```

Since every primary key is by definition a unique key, this restriction also includes the table's primary key, if it has one. For example, the next two statements are invalid:

```
CREATE TABLE t5 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    PRIMARY KEY(col1, col2)
)
PARTITION BY HASH(col3)
PARTITIONS 4;

CREATE TABLE t6 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    PRIMARY KEY(col1, col3),
```

```

        UNIQUE KEY(col2)
    )
PARTITION BY HASH( YEAR(col2) )
PARTITIONS 4;

```

In both cases, the primary key does not include all columns referenced in the partitioning expression. However, both of the next two statements are valid:

```

CREATE TABLE t7 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    PRIMARY KEY(col1, col2)
)
PARTITION BY HASH(col1 + YEAR(col2))
PARTITIONS 4;

CREATE TABLE t8 (
    col1 INT NOT NULL,
    col2 DATE NOT NULL,
    col3 INT NOT NULL,
    col4 INT NOT NULL,
    PRIMARY KEY(col1, col2, col4),
    UNIQUE KEY(col2, col1)
)
PARTITION BY HASH(col1 + YEAR(col2))
PARTITIONS 4;

```

If a table has no unique keys—this includes having no primary key—then this restriction does not apply, and you may use any column or columns in the partitioning expression as long as the column type is compatible with the partitioning type.

For the same reason, you cannot later add a unique key to a partitioned table unless the key includes all columns used by the table's partitioning expression. Consider the partitioned table created as shown here:

```

mysql> CREATE TABLE t_no_pk (c1 INT, c2 INT)
->     PARTITION BY RANGE(c1) (
->         PARTITION p0 VALUES LESS THAN (10),
->         PARTITION p1 VALUES LESS THAN (20),
->         PARTITION p2 VALUES LESS THAN (30),
->         PARTITION p3 VALUES LESS THAN (40)
->     );
Query OK, 0 rows affected (0.12 sec)

```

It is possible to add a primary key to `t_no_pk` using either of these `ALTER TABLE` statements:

```

# possible PK
mysql> ALTER TABLE t_no_pk ADD PRIMARY KEY(c1);
Query OK, 0 rows affected (0.13 sec)
Records: 0  Duplicates: 0  Warnings: 0

# drop this PK
mysql> ALTER TABLE t_no_pk DROP PRIMARY KEY;
Query OK, 0 rows affected (0.10 sec)
Records: 0  Duplicates: 0  Warnings: 0

# use another possible PK
mysql> ALTER TABLE t_no_pk ADD PRIMARY KEY(c1, c2);
Query OK, 0 rows affected (0.12 sec)
Records: 0  Duplicates: 0  Warnings: 0

# drop this PK
mysql> ALTER TABLE t_no_pk DROP PRIMARY KEY;
Query OK, 0 rows affected (0.09 sec)
Records: 0  Duplicates: 0  Warnings: 0

```

However, the next statement fails, because `c1` is part of the partitioning key, but is not part of the proposed primary key:

```
# fails with error 1503
mysql> ALTER TABLE t_no_pk ADD PRIMARY KEY(c2);
ERROR 1503 (HY000): A PRIMARY KEY must include all columns in the table's partitioning function
```

Since `t_no_pk` has only `c1` in its partitioning expression, attempting to add a unique key on `c2` alone fails. However, you can add a unique key that uses both `c1` and `c2`.

These rules also apply to existing nonpartitioned tables that you wish to partition using `ALTER TABLE ... PARTITION BY`. Consider a table `np_pk` created as shown here:

```
mysql> CREATE TABLE np_pk (
    ->     id INT NOT NULL AUTO_INCREMENT,
    ->     name VARCHAR(50),
    ->     added DATE,
    ->     PRIMARY KEY (id)
    -> );
Query OK, 0 rows affected (0.08 sec)
```

The following `ALTER TABLE` statement fails with an error, because the `added` column is not part of any unique key in the table:

```
mysql> ALTER TABLE np_pk
    ->     PARTITION BY HASH( TO_DAYS(added) )
    ->     PARTITIONS 4;
ERROR 1503 (HY000): A PRIMARY KEY must include all columns in the table's partitioning function
```

However, this statement using the `id` column for the partitioning column is valid, as shown here:

```
mysql> ALTER TABLE np_pk
    ->     PARTITION BY HASH(id)
    ->     PARTITIONS 4;
Query OK, 0 rows affected (0.11 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

In the case of `np_pk`, the only column that may be used as part of a partitioning expression is `id`; if you wish to partition this table using any other column or columns in the partitioning expression, you must first modify the table, either by adding the desired column or columns to the primary key, or by dropping the primary key altogether.

## 24.6.2 Partitioning Limitations Relating to Storage Engines

In MySQL 8.0, partitioning support is not actually provided by the MySQL Server, but rather by a table storage engine's own or native partitioning handler. In MySQL 8.0, only the `InnoDB` and `NDB` storage engines provide native partitioning handlers. This means that partitioned tables cannot be created using any other storage engine than these. (You must be using MySQL NDB Cluster with the `NDB` storage engine to create `NDB` tables.)

**InnoDB storage engine.** `InnoDB` foreign keys and MySQL partitioning are not compatible. Partitioned `InnoDB` tables cannot have foreign key references, nor can they have columns referenced by foreign keys. `InnoDB` tables which have or which are referenced by foreign keys cannot be partitioned.

`ALTER TABLE ... OPTIMIZE PARTITION` does not work correctly with partitioned tables that use `Innodb`. Use `ALTER TABLE ... REBUILD PARTITION` and `ALTER TABLE ... ANALYZE PARTITION`, instead, for such tables. For more information, see [Section 13.1.9.1, “ALTER TABLE Partition Operations”](#).

**User-defined partitioning and the NDB storage engine (NDB Cluster).** Partitioning by `KEY` (including `LINEAR KEY`) is the only type of partitioning supported for the `NDB` storage engine. It is not possible under normal circumstances in NDB Cluster to create an NDB Cluster table using any partitioning type other than `[LINEAR] KEY`, and attempting to do so fails with an error.

*Exception (not for production):* It is possible to override this restriction by setting the `new` system variable on NDB Cluster SQL nodes to `ON`. If you choose to do this, you should be aware that tables

using partitioning types other than `[LINEAR] KEY` are not supported in production. *In such cases, you can create and use tables with partitioning types other than KEY or LINEAR KEY, but you do this entirely at your own risk.*

The maximum number of partitions that can be defined for an `NDB` table depends on the number of data nodes and node groups in the cluster, the version of the `NDB` Cluster software in use, and other factors. See [NDB and user-defined partitioning](#), for more information.

The maximum amount of fixed-size data that can be stored per partition in an `NDB` table is 128 TB. Previously, this was 16 GB.

`CREATE TABLE` and `ALTER TABLE` statements that would cause a user-partitioned `NDB` table not to meet either or both of the following two requirements are not permitted, and fail with an error:

1. The table must have an explicit primary key.
2. All columns listed in the table's partitioning expression must be part of the primary key.

**Exception.** If a user-partitioned `NDB` table is created using an empty column-list (that is, using `PARTITION BY KEY()` or `PARTITION BY LINEAR KEY()`), then no explicit primary key is required.

**Partition selection.** Partition selection is not supported for `NDB` tables. See [Section 24.5, “Partition Selection”](#), for more information.

**Upgrading partitioned tables.** When performing an upgrade, tables which are partitioned by `KEY` must be dumped and reloaded. Partitioned tables using storage engines other than `InnoDB` cannot be upgraded from MySQL 5.7 or earlier to MySQL 8.0 or later; you must either drop the partitioning from such tables with `ALTER TABLE ... REMOVE PARTITIONING` or convert them to `InnoDB` using `ALTER TABLE ... ENGINE=INNODB` prior to the upgrade.

For information about converting `MyISAM` tables to `InnoDB`, see [Section 15.6.1.5, “Converting Tables from MyISAM to InnoDB”](#).

### 24.6.3 Partitioning Limitations Relating to Functions

This section discusses limitations in MySQL Partitioning relating specifically to functions used in partitioning expressions.

Only the MySQL functions shown in the following list are allowed in partitioning expressions:

- `ABS()`
- `CEILING()` (see [CEILING\(\)](#) and [FLOOR\(\)](#))
- `DATEDIFF()`
- `DAY()`
- `DAYOFMONTH()`
- `DAYOFWEEK()`
- `DAYOFYEAR()`
- `EXTRACT()` (see [EXTRACT\(\)](#) function with WEEK specifier)
- `FLOOR()` (see [CEILING\(\)](#) and [FLOOR\(\)](#))
- `HOUR()`
- `MICROSECOND()`
- `MINUTE()`

- `MOD()`
- `MONTH()`
- `QUARTER()`
- `SECOND()`
- `TIME_TO_SEC()`
- `TO_DAYS()`
- `TO_SECONDS()`
- `UNIX_TIMESTAMP()` (with `TIMESTAMP` columns)
- `WEEKDAY()`
- `YEAR()`
- `YEARWEEK()`

In MySQL 8.0, partition pruning is supported for the `TO_DAYS()`, `TO_SECONDS()`, `YEAR()`, and `UNIX_TIMESTAMP()` functions. See [Section 24.4, “Partition Pruning”](#), for more information.

**CEILING() and FLOOR().** Each of these functions returns an integer only if it is passed an argument of an exact numeric type, such as one of the `INT` types or `DECIMAL`. This means, for example, that the following `CREATE TABLE` statement fails with an error, as shown here:

```
mysql> CREATE TABLE t (c FLOAT) PARTITION BY LIST( FLOOR(c) )(
    ->     PARTITION p0 VALUES IN (1,3,5),
    ->     PARTITION p1 VALUES IN (2,4,6)
    -> );
ERROR 1490 (HY000): The PARTITION function returns the wrong type
```

**EXTRACT() function with WEEK specifier.** The value returned by the `EXTRACT()` function, when used as `EXTRACT(WEEK FROM col)`, depends on the value of the `default_week_format` system variable. For this reason, `EXTRACT()` is not permitted as a partitioning function when it specifies the unit as `WEEK`. (Bug #54483)

See [Section 12.6.2, “Mathematical Functions”](#), for more information about the return types of these functions, as well as [Section 11.1, “Numeric Data Types”](#).

---

# Chapter 25 Stored Objects

## Table of Contents

25.1 Defining Stored Programs .....	4800
25.2 Using Stored Routines .....	4801
25.2.1 Stored Routine Syntax .....	4802
25.2.2 Stored Routines and MySQL Privileges .....	4802
25.2.3 Stored Routine Metadata .....	4803
25.2.4 Stored Procedures, Functions, Triggers, and LAST_INSERT_ID() .....	4803
25.3 Using Triggers .....	4803
25.3.1 Trigger Syntax and Examples .....	4804
25.3.2 Trigger Metadata .....	4808
25.4 Using the Event Scheduler .....	4808
25.4.1 Event Scheduler Overview .....	4809
25.4.2 Event Scheduler Configuration .....	4809
25.4.3 Event Syntax .....	4812
25.4.4 Event Metadata .....	4812
25.4.5 Event Scheduler Status .....	4813
25.4.6 The Event Scheduler and MySQL Privileges .....	4813
25.5 Using Views .....	4816
25.5.1 View Syntax .....	4816
25.5.2 View Processing Algorithms .....	4816
25.5.3 Updatable and Insertable Views .....	4817
25.5.4 The View WITH CHECK OPTION Clause .....	4820
25.5.5 View Metadata .....	4821
25.6 Stored Object Access Control .....	4821
25.7 Stored Program Binary Logging .....	4825
25.8 Restrictions on Stored Programs .....	4831
25.9 Restrictions on Views .....	4834

This chapter discusses stored database objects that are defined in terms of SQL code that is stored on the server for later execution.

Stored objects include these object types:

- Stored procedure: An object created with `CREATE PROCEDURE` and invoked using the `CALL` statement. A procedure does not have a return value but can modify its parameters for later inspection by the caller. It can also generate result sets to be returned to the client program.
- Stored function: An object created with `CREATE FUNCTION` and used much like a built-in function. You invoke it in an expression and it returns a value during expression evaluation.
- Trigger: An object created with `CREATE TRIGGER` that is associated with a table. A trigger is activated when a particular event occurs for the table, such as an insert or update.
- Event: An object created with `CREATE EVENT` and invoked by the server according to schedule.
- View: An object created with `CREATE VIEW` that when referenced produces a result set. A view acts as a virtual table.

Terminology used in this document reflects the stored object hierarchy:

- Stored routines include stored procedures and functions.
- Stored programs include stored routines, triggers, and events.
- Stored objects include stored programs and views.

This chapter describes how to use stored objects. The following sections provide additional information about SQL syntax for statements related to these objects, and about object processing:

- For each object type, there are `CREATE`, `ALTER`, and `DROP` statements that control which objects exist and how they are defined. See [Section 13.1, “Data Definition Statements”](#).
- The `CALL` statement is used to invoke stored procedures. See [Section 13.2.1, “CALL Statement”](#).
- Stored program definitions include a body that may use compound statements, loops, conditionals, and declared variables. See [Section 13.6, “Compound Statement Syntax”](#).
- Metadata changes to objects referred to by stored programs are detected and cause automatic reparsing of the affected statements when the program is next executed. For more information, see [Section 8.10.3, “Caching of Prepared Statements and Stored Programs”](#).

## 25.1 Defining Stored Programs

Each stored program contains a body that consists of an SQL statement. This statement may be a compound statement made up of several statements separated by semicolon (`:`) characters. For example, the following stored procedure has a body made up of a `BEGIN ... END` block that contains a `SET` statement and a `REPEAT` loop that itself contains another `SET` statement:

```
CREATE PROCEDURE dorepeat(p1 INT)
BEGIN
    SET @x = 0;
    REPEAT SET @x = @x + 1; UNTIL @x > p1 END REPEAT;
END;
```

If you use the `mysql` client program to define a stored program containing semicolon characters, a problem arises. By default, `mysql` itself recognizes the semicolon as a statement delimiter, so you must redefine the delimiter temporarily to cause `mysql` to pass the entire stored program definition to the server.

To redefine the `mysql` delimiter, use the `delimiter` command. The following example shows how to do this for the `dorepeat()` procedure just shown. The delimiter is changed to `//` to enable the entire definition to be passed to the server as a single statement, and then restored to `:` before invoking the procedure. This enables the `:` delimiter used in the procedure body to be passed through to the server rather than being interpreted by `mysql` itself.

```
mysql> delimiter //
mysql> CREATE PROCEDURE dorepeat(p1 INT)
-> BEGIN
->     SET @x = 0;
->     REPEAT SET @x = @x + 1; UNTIL @x > p1 END REPEAT;
-> END
-> //
Query OK, 0 rows affected (0.00 sec)

mysql> delimiter ;
mysql> CALL dorepeat(1000);
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT @x;
+-----+
| @x   |
+-----+
| 1001 |
+-----+
1 row in set (0.00 sec)
```

You can redefine the delimiter to a string other than `//`, and the delimiter can consist of a single character or multiple characters. You should avoid the use of the backslash (`\`) character because that is the escape character for MySQL.

The following is an example of a function that takes a parameter, performs an operation using an SQL function, and returns the result. In this case, it is unnecessary to use `delimiter` because the function definition contains no internal `:` statement delimiters:

```
mysql> CREATE FUNCTION hello (s CHAR(20))
mysql> RETURNS CHAR(50) DETERMINISTIC
-> RETURN CONCAT('Hello, ',s,'!');
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT hello('world');
+-----+
| hello('world') |
+-----+
| Hello, world! |
+-----+
1 row in set (0.00 sec)
```

## 25.2 Using Stored Routines

MySQL supports stored routines (procedures and functions). A stored routine is a set of SQL statements that can be stored in the server. Once this has been done, clients don't need to keep reissuing the individual statements but can refer to the stored routine instead.

Stored routines can be particularly useful in certain situations:

- When multiple client applications are written in different languages or work on different platforms, but need to perform the same database operations.
- When security is paramount. Banks, for example, use stored procedures and functions for all common operations. This provides a consistent and secure environment, and routines can ensure that each operation is properly logged. In such a setup, applications and users would have no access to the database tables directly, but can only execute specific stored routines.

Stored routines can provide improved performance because less information needs to be sent between the server and the client. The tradeoff is that this does increase the load on the database server because more of the work is done on the server side and less is done on the client (application) side. Consider this if many client machines (such as Web servers) are serviced by only one or a few database servers.

Stored routines also enable you to have libraries of functions in the database server. This is a feature shared by modern application languages that enable such design internally (for example, by using classes). Using these client application language features is beneficial for the programmer even outside the scope of database use.

MySQL follows the SQL:2003 syntax for stored routines, which is also used by IBM's DB2. All syntax described here is supported and any limitations and extensions are documented where appropriate.

## Additional Resources

- You may find the [Stored Procedures User Forum](#) of use when working with stored procedures and functions.
- For answers to some commonly asked questions regarding stored routines in MySQL, see [Section A.4, “MySQL 8.0 FAQ: Stored Procedures and Functions”](#).
- There are some restrictions on the use of stored routines. See [Section 25.8, “Restrictions on Stored Programs”](#).
- Binary logging for stored routines takes place as described in [Section 25.7, “Stored Program Binary Logging”](#).

## 25.2.1 Stored Routine Syntax

A stored routine is either a procedure or a function. Stored routines are created with the [CREATE PROCEDURE](#) and [CREATE FUNCTION](#) statements (see [Section 13.1.17, “CREATE PROCEDURE and CREATE FUNCTION Statements”](#)). A procedure is invoked using a [CALL](#) statement (see [Section 13.2.1, “CALL Statement”](#)), and can only pass back values using output variables. A function can be called from inside a statement just like any other function (that is, by invoking the function's name), and can return a scalar value. The body of a stored routine can use compound statements (see [Section 13.6, “Compound Statement Syntax”](#)).

Stored routines can be dropped with the [DROP PROCEDURE](#) and [DROP FUNCTION](#) statements (see [Section 13.1.29, “DROP PROCEDURE and DROP FUNCTION Statements”](#)), and altered with the [ALTER PROCEDURE](#) and [ALTER FUNCTION](#) statements (see [Section 13.1.7, “ALTER PROCEDURE Statement”](#)).

A stored procedure or function is associated with a particular database. This has several implications:

- When the routine is invoked, an implicit [USE db\\_name](#) is performed (and undone when the routine terminates). [USE](#) statements within stored routines are not permitted.
- You can qualify routine names with the database name. This can be used to refer to a routine that is not in the current database. For example, to invoke a stored procedure `p` or function `f` that is associated with the `test` database, you can say `CALL test.p()` or `test.f()`.
- When a database is dropped, all stored routines associated with it are dropped as well.

Stored functions cannot be recursive.

Recursion in stored procedures is permitted but disabled by default. To enable recursion, set the [max\\_sp\\_recursion\\_depth](#) server system variable to a value greater than zero. Stored procedure recursion increases the demand on thread stack space. If you increase the value of [max\\_sp\\_recursion\\_depth](#), it may be necessary to increase thread stack size by increasing the value of [thread\\_stack](#) at server startup. See [Section 5.1.8, “Server System Variables”](#), for more information.

MySQL supports a very useful extension that enables the use of regular [SELECT](#) statements (that is, without using cursors or local variables) inside a stored procedure. The result set of such a query is simply sent directly to the client. Multiple [SELECT](#) statements generate multiple result sets, so the client must use a MySQL client library that supports multiple result sets. This means the client must use a client library from a version of MySQL at least as recent as 4.1. The client should also specify the [CLIENT\\_MULTI\\_RESULTS](#) option when it connects. For C programs, this can be done with the [mysql\\_real\\_connect\(\)](#) C API function. See [mysql\\_real\\_connect\(\)](#), and [Multiple Statement Execution Support](#).

In MySQL 8.0.22 and later, a user variable referenced by a statement in a stored procedure has its type determined the first time the procedure is invoked, and retains this type each time the procedure is invoked thereafter.

## 25.2.2 Stored Routines and MySQL Privileges

The MySQL grant system takes stored routines into account as follows:

- The [CREATE ROUTINE](#) privilege is needed to create stored routines.
- The [ALTER ROUTINE](#) privilege is needed to alter or drop stored routines. This privilege is granted automatically to the creator of a routine if necessary, and dropped from the creator when the routine is dropped.
- The [EXECUTE](#) privilege is required to execute stored routines. However, this privilege is granted automatically to the creator of a routine if necessary (and dropped from the creator when the routine is dropped). Also, the default [SQL SECURITY](#) characteristic for a routine is [DEFINER](#), which enables users who have access to the database with which the routine is associated to execute the routine.

- If the `automatic_sp_privileges` system variable is 0, the `EXECUTE` and `ALTER ROUTINE` privileges are not automatically granted to and dropped from the routine creator.
- The creator of a routine is the account used to execute the `CREATE` statement for it. This might not be the same as the account named as the `DEFINER` in the routine definition.
- The account named as a routine `DEFINER` can see all routine properties, including its definition. The account thus has full access to the routine output as produced by:
  - The contents of the Information Schema `ROUTINES` table.
  - The `SHOW CREATE FUNCTION` and `SHOW CREATE PROCEDURE` statements.
  - The `SHOW FUNCTION CODE` and `SHOW PROCEDURE CODE` statements.
  - The `SHOW FUNCTION STATUS` and `SHOW PROCEDURE STATUS` statements.
- For an account other than the account named as the routine `DEFINER`, access to routine properties depends on the privileges granted to the account:
  - With the `SHOW_ROUTINE` privilege or the global `SELECT` privilege, the account can see all routine properties, including its definition.
  - With the `CREATE ROUTINE`, `ALTER ROUTINE` or `EXECUTE` privilege granted at a scope that includes the routine, the account can see all routine properties except its definition.

### 25.2.3 Stored Routine Metadata

To obtain metadata about stored routines:

- Query the `ROUTINES` table of the `INFORMATION_SCHEMA` database. See [Section 26.3.30, “The INFORMATION\\_SCHEMA ROUTINES Table”](#).
- Use the `SHOW CREATE PROCEDURE` and `SHOW CREATE FUNCTION` statements to see routine definitions. See [Section 13.7.7.9, “SHOW CREATE PROCEDURE Statement”](#).
- Use the `SHOW PROCEDURE STATUS` and `SHOW FUNCTION STATUS` statements to see routine characteristics. See [Section 13.7.7.28, “SHOW PROCEDURE STATUS Statement”](#).
- Use the `SHOW PROCEDURE CODE` and `SHOW FUNCTION CODE` statements to see a representation of the internal implementation of the routine. See [Section 13.7.7.27, “SHOW PROCEDURE CODE Statement”](#).

### 25.2.4 Stored Procedures, Functions, Triggers, and LAST\_INSERT\_ID()

Within the body of a stored routine (procedure or function) or a trigger, the value of `LAST_INSERT_ID()` changes the same way as for statements executed outside the body of these kinds of objects (see [Section 12.16, “Information Functions”](#)). The effect of a stored routine or trigger upon the value of `LAST_INSERT_ID()` that is seen by following statements depends on the kind of routine:

- If a stored procedure executes statements that change the value of `LAST_INSERT_ID()`, the changed value is seen by statements that follow the procedure call.
- For stored functions and triggers that change the value, the value is restored when the function or trigger ends, so following statements do not see a changed value.

## 25.3 Using Triggers

A trigger is a named database object that is associated with a table, and that activates when a particular event occurs for the table. Some uses for triggers are to perform checks of values to be inserted into a table or to perform calculations on values involved in an update.

A trigger is defined to activate when a statement inserts, updates, or deletes rows in the associated table. These row operations are trigger events. For example, rows can be inserted by [INSERT](#) or [LOAD DATA](#) statements, and an insert trigger activates for each inserted row. A trigger can be set to activate either before or after the trigger event. For example, you can have a trigger activate before each row that is inserted into a table or after each row that is updated.



### Important

MySQL triggers activate only for changes made to tables by SQL statements. This includes changes to base tables that underlie updatable views. Triggers do not activate for changes to tables made by APIs that do not transmit SQL statements to the MySQL Server. This means that triggers are not activated by updates made using the [NDB API](#).

Triggers are not activated by changes in [INFORMATION\\_SCHEMA](#) or [performance\\_schema](#) tables. Those tables are actually views and triggers are not permitted on views.

The following sections describe the syntax for creating and dropping triggers, show some examples of how to use them, and indicate how to obtain trigger metadata.

## Additional Resources

- You may find the [MySQL User Forums](#) helpful when working with triggers.
- For answers to commonly asked questions regarding triggers in MySQL, see [Section A.5, “MySQL 8.0 FAQ: Triggers”](#).
- There are some restrictions on the use of triggers; see [Section 25.8, “Restrictions on Stored Programs”](#).
- Binary logging for triggers takes place as described in [Section 25.7, “Stored Program Binary Logging”](#).

### 25.3.1 Trigger Syntax and Examples

To create a trigger or drop a trigger, use the [CREATE TRIGGER](#) or [DROP TRIGGER](#) statement, described in [Section 13.1.22, “CREATE TRIGGER Statement”](#), and [Section 13.1.34, “DROP TRIGGER Statement”](#).

Here is a simple example that associates a trigger with a table, to activate for [INSERT](#) operations. The trigger acts as an accumulator, summing the values inserted into one of the columns of the table.

```
mysql> CREATE TABLE account (acct_num INT, amount DECIMAL(10,2));
Query OK, 0 rows affected (0.03 sec)

mysql> CREATE TRIGGER ins_sum BEFORE INSERT ON account
      FOR EACH ROW SET @sum = @sum + NEW.amount;
Query OK, 0 rows affected (0.01 sec)
```

The [CREATE TRIGGER](#) statement creates a trigger named [ins\\_sum](#) that is associated with the [account](#) table. It also includes clauses that specify the trigger action time, the triggering event, and what to do when the trigger activates:

- The keyword [BEFORE](#) indicates the trigger action time. In this case, the trigger activates before each row inserted into the table. The other permitted keyword here is [AFTER](#).
- The keyword [INSERT](#) indicates the trigger event; that is, the type of operation that activates the trigger. In the example, [INSERT](#) operations cause trigger activation. You can also create triggers for [DELETE](#) and [UPDATE](#) operations.
- The statement following [FOR EACH ROW](#) defines the trigger body; that is, the statement to execute each time the trigger activates, which occurs once for each row affected by the triggering event.

In the example, the trigger body is a simple `SET` that accumulates into a user variable the values inserted into the `amount` column. The statement refers to the column as `NEW.amount` which means "the value of the `amount` column to be inserted into the new row."

To use the trigger, set the accumulator variable to zero, execute an `INSERT` statement, and then see what value the variable has afterward:

```
mysql> SET @sum = 0;
mysql> INSERT INTO account VALUES(137,14.98),(141,1937.50),(97,-100.00);
mysql> SELECT @sum AS 'Total amount inserted';
+-----+
| Total amount inserted |
+-----+
| 1852.48 |
+-----+
```

In this case, the value of `@sum` after the `INSERT` statement has executed is `14.98 + 1937.50 - 100`, or `1852.48`.

To destroy the trigger, use a `DROP TRIGGER` statement. You must specify the schema name if the trigger is not in the default schema:

```
mysql> DROP TRIGGER test.ins_sum;
```

If you drop a table, any triggers for the table are also dropped.

Trigger names exist in the schema namespace, meaning that all triggers must have unique names within a schema. Triggers in different schemas can have the same name.

It is possible to define multiple triggers for a given table that have the same trigger event and action time. For example, you can have two `BEFORE UPDATE` triggers for a table. By default, triggers that have the same trigger event and action time activate in the order they were created. To affect trigger order, specify a clause after `FOR EACH ROW` that indicates `FOLLOWERS` or `PRECEDES` and the name of an existing trigger that also has the same trigger event and action time. With `FOLLOWERS`, the new trigger activates after the existing trigger. With `PRECEDES`, the new trigger activates before the existing trigger.

For example, the following trigger definition defines another `BEFORE INSERT` trigger for the `account` table:

```
mysql> CREATE TRIGGER ins_transaction BEFORE INSERT ON account
      FOR EACH ROW PRECEDES ins_sum
      SET
        @deposits = @deposits + IF(NEW.amount>0,NEW.amount,0),
        @withdrawals = @withdrawals + IF(NEW.amount<0,-NEW.amount,0);
Query OK, 0 rows affected (0.01 sec)
```

This trigger, `ins_transaction`, is similar to `ins_sum` but accumulates deposits and withdrawals separately. It has a `PRECEDES` clause that causes it to activate before `ins_sum`; without that clause, it would activate after `ins_sum` because it is created after `ins_sum`.

Within the trigger body, the `OLD` and `NEW` keywords enable you to access columns in the rows affected by a trigger. `OLD` and `NEW` are MySQL extensions to triggers; they are not case-sensitive.

In an `INSERT` trigger, only `NEW.col_name` can be used; there is no old row. In a `DELETE` trigger, only `OLD.col_name` can be used; there is no new row. In an `UPDATE` trigger, you can use `OLD.col_name` to refer to the columns of a row before it is updated and `NEW.col_name` to refer to the columns of the row after it is updated.

A column named with `OLD` is read only. You can refer to it (if you have the `SELECT` privilege), but not modify it. You can refer to a column named with `NEW` if you have the `SELECT` privilege for it. In a `BEFORE` trigger, you can also change its value with `SET NEW.col_name = value` if you have the `UPDATE` privilege for it. This means you can use a trigger to modify the values to be inserted into a new row or used to update a row. (Such a `SET` statement has no effect in an `AFTER` trigger because the row change has already occurred.)

In a `BEFORE` trigger, the `NEW` value for an `AUTO_INCREMENT` column is 0, not the sequence number that is generated automatically when the new row actually is inserted.

By using the `BEGIN ... END` construct, you can define a trigger that executes multiple statements. Within the `BEGIN` block, you also can use other syntax that is permitted within stored routines such as conditionals and loops. However, just as for stored routines, if you use the `mysql` program to define a trigger that executes multiple statements, it is necessary to redefine the `mysql` statement delimiter so that you can use the `;` statement delimiter within the trigger definition. The following example illustrates these points. It defines an `UPDATE` trigger that checks the new value to be used for updating each row, and modifies the value to be within the range from 0 to 100. This must be a `BEFORE` trigger because the value must be checked before it is used to update the row:

```
mysql> delimiter //
mysql> CREATE TRIGGER upd_check BEFORE UPDATE ON account
      FOR EACH ROW
      BEGIN
        IF NEW.amount < 0 THEN
          SET NEW.amount = 0;
        ELSEIF NEW.amount > 100 THEN
          SET NEW.amount = 100;
        END IF;
      END//;
mysql> delimiter ;
```

It can be easier to define a stored procedure separately and then invoke it from the trigger using a simple `CALL` statement. This is also advantageous if you want to execute the same code from within several triggers.

There are limitations on what can appear in statements that a trigger executes when activated:

- The trigger cannot use the `CALL` statement to invoke stored procedures that return data to the client or that use dynamic SQL. (Stored procedures are permitted to return data to the trigger through `OUT` or `INOUT` parameters.)
- The trigger cannot use statements that explicitly or implicitly begin or end a transaction, such as `START TRANSACTION`, `COMMIT`, or `ROLLBACK`. (`ROLLBACK` to `SAVEPOINT` is permitted because it does not end a transaction.).

See also [Section 25.8, “Restrictions on Stored Programs”](#).

MySQL handles errors during trigger execution as follows:

- If a `BEFORE` trigger fails, the operation on the corresponding row is not performed.
- A `BEFORE` trigger is activated by the *attempt* to insert or modify the row, regardless of whether the attempt subsequently succeeds.
- An `AFTER` trigger is executed only if any `BEFORE` triggers and the row operation execute successfully.
- An error during either a `BEFORE` or `AFTER` trigger results in failure of the entire statement that caused trigger invocation.
- For transactional tables, failure of a statement should cause rollback of all changes performed by the statement. Failure of a trigger causes the statement to fail, so trigger failure also causes rollback. For nontransactional tables, such rollback cannot be done, so although the statement fails, any changes performed prior to the point of the error remain in effect.

Triggers can contain direct references to tables by name, such as the trigger named `testref` shown in this example:

```
CREATE TABLE test1(a1 INT);
CREATE TABLE test2(a2 INT);
CREATE TABLE test3(a3 INT NOT NULL AUTO_INCREMENT PRIMARY KEY);
```

```

CREATE TABLE test4(
    a4 INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
    b4 INT DEFAULT 0
);

delimiter |

CREATE TRIGGER testref BEFORE INSERT ON test1
FOR EACH ROW
BEGIN
    INSERT INTO test2 SET a2 = NEW.a1;
    DELETE FROM test3 WHERE a3 = NEW.a1;
    UPDATE test4 SET b4 = b4 + 1 WHERE a4 = NEW.a1;
END;
|

delimiter ;

```

INSERT INTO test3 (a3) VALUES  
 (NULL), (NULL), (NULL), (NULL), (NULL),  
 (NULL), (NULL), (NULL), (NULL);

INSERT INTO test4 (a4) VALUES  
 (0), (0), (0), (0), (0), (0), (0), (0), (0);

Suppose that you insert the following values into table `test1` as shown here:

```

mysql> INSERT INTO test1 VALUES
    (1), (3), (1), (7), (1), (8), (4), (4);
Query OK, 8 rows affected (0.01 sec)
Records: 8  Duplicates: 0  Warnings: 0

```

As a result, the four tables contain the following data:

```

mysql> SELECT * FROM test1;
+---+
| a1 |
+---+
| 1 |
| 3 |
| 1 |
| 7 |
| 1 |
| 8 |
| 4 |
| 4 |
+---+
8 rows in set (0.00 sec)

mysql> SELECT * FROM test2;
+---+
| a2 |
+---+
| 1 |
| 3 |
| 1 |
| 7 |
| 1 |
| 8 |
| 4 |
| 4 |
+---+
8 rows in set (0.00 sec)

mysql> SELECT * FROM test3;
+---+
| a3 |
+---+
| 2 |
| 5 |
| 6 |
| 9 |

```

```
| 10 |
+---+
5 rows in set (0.00 sec)

mysql> SELECT * FROM test4;
+---+---+
| a4 | b4 |
+---+---+
| 1 | 3 |
| 2 | 0 |
| 3 | 1 |
| 4 | 2 |
| 5 | 0 |
| 6 | 0 |
| 7 | 1 |
| 8 | 1 |
| 9 | 0 |
| 10 | 0 |
+---+---+
10 rows in set (0.00 sec)
```

### 25.3.2 Trigger Metadata

To obtain metadata about triggers:

- Query the `TRIGGERS` table of the `INFORMATION_SCHEMA` database. See [Section 26.3.45, “The INFORMATION\\_SCHEMA TRIGGERS Table”](#).
- Use the `SHOW CREATE TRIGGER` statement. See [Section 13.7.7.11, “SHOW CREATE TRIGGER Statement”](#).
- Use the `SHOW TRIGGERS` statement. See [Section 13.7.7.40, “SHOW TRIGGERS Statement”](#).

## 25.4 Using the Event Scheduler

The *MySQL Event Scheduler* manages the scheduling and execution of events, that is, tasks that run according to a schedule. The following discussion covers the Event Scheduler and is divided into the following sections:

- [Section 25.4.1, “Event Scheduler Overview”](#), provides an introduction to and conceptual overview of MySQL Events.
- [Section 25.4.3, “Event Syntax”](#), discusses the SQL statements for creating, altering, and dropping MySQL Events.
- [Section 25.4.4, “Event Metadata”](#), shows how to obtain information about events and how this information is stored by the MySQL Server.
- [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#), discusses the privileges required to work with events and the ramifications that events have with regard to privileges when executing.

Stored routines require the `events` data dictionary table in the `mysql` system database. This table is created during the MySQL 8.0 installation procedure. If you are upgrading to MySQL 8.0 from an earlier version, be sure to perform the upgrade procedure to make sure that your system database is up to date. See [Section 2.10, “Upgrading MySQL”](#).

## Additional Resources

- There are some restrictions on the use of events; see [Section 25.8, “Restrictions on Stored Programs”](#).
- Binary logging for events takes place as described in [Section 25.7, “Stored Program Binary Logging”](#).
- You may also find the [MySQL User Forums](#) to be helpful.

### 25.4.1 Event Scheduler Overview

MySQL Events are tasks that run according to a schedule. Therefore, we sometimes refer to them as *scheduled events*. When you create an event, you are creating a named database object containing one or more SQL statements to be executed at one or more regular intervals, beginning and ending at a specific date and time. Conceptually, this is similar to the idea of the Unix `cron` (also known as a “cron job”) or the Windows Task Scheduler.

Scheduled tasks of this type are also sometimes known as “temporal triggers”, implying that these are objects that are triggered by the passage of time. While this is essentially correct, we prefer to use the term *events* to avoid confusion with triggers of the type discussed in [Section 25.3, “Using Triggers”](#). Events should more specifically not be confused with “temporary triggers”. Whereas a trigger is a database object whose statements are executed in response to a specific type of event that occurs on a given table, a (scheduled) event is an object whose statements are executed in response to the passage of a specified time interval.

While there is no provision in the SQL Standard for event scheduling, there are precedents in other database systems, and you may notice some similarities between these implementations and that found in the MySQL Server.

MySQL Events have the following major features and properties:

- In MySQL, an event is uniquely identified by its name and the schema to which it is assigned.
- An event performs a specific action according to a schedule. This action consists of an SQL statement, which can be a compound statement in a `BEGIN ... END` block if desired (see [Section 13.6, “Compound Statement Syntax”](#)). An event’s timing can be either *one-time* or *recurrent*. A one-time event executes one time only. A recurrent event repeats its action at a regular interval, and the schedule for a recurring event can be assigned a specific start day and time, end day and time, both, or neither. (By default, a recurring event’s schedule begins as soon as it is created, and continues indefinitely, until it is disabled or dropped.)

If a repeating event does not terminate within its scheduling interval, the result may be multiple instances of the event executing simultaneously. If this is undesirable, you should institute a mechanism to prevent simultaneous instances. For example, you could use the `GET_LOCK()` function, or row or table locking.

- Users can create, modify, and drop scheduled events using SQL statements intended for these purposes. Syntactically invalid event creation and modification statements fail with an appropriate error message. *A user may include statements in an event’s action which require privileges that the user does not actually have.* The event creation or modification statement succeeds but the event’s action fails. See [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#) for details.
- Many of the properties of an event can be set or modified using SQL statements. These properties include the event’s name, timing, persistence (that is, whether it is preserved following the expiration of its schedule), status (enabled or disabled), action to be performed, and the schema to which it is assigned. See [Section 13.1.3, “ALTER EVENT Statement”](#).

The default definer of an event is the user who created the event, unless the event has been altered, in which case the definer is the user who issued the last `ALTER EVENT` statement affecting that event. An event can be modified by any user having the `EVENT` privilege on the database for which the event is defined. See [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#).

- An event’s action statement may include most SQL statements permitted within stored routines. For restrictions, see [Section 25.8, “Restrictions on Stored Programs”](#).

### 25.4.2 Event Scheduler Configuration

Events are executed by a special *event scheduler thread*; when we refer to the Event Scheduler, we actually refer to this thread. When running, the event scheduler thread and its current state can be

seen by users having the `PROCESS` privilege in the output of `SHOW PROCESSLIST`, as shown in the discussion that follows.

The global `event_scheduler` system variable determines whether the Event Scheduler is enabled and running on the server. It has one of the following values, which affect event scheduling as described:

- `ON`: The Event Scheduler is started; the event scheduler thread runs and executes all scheduled events. `ON` is the default `event_scheduler` value.

When the Event Scheduler is `ON`, the event scheduler thread is listed in the output of `SHOW PROCESSLIST` as a daemon process, and its state is represented as shown here:

```
mysql> SHOW PROCESSLIST\G
***** 1. row *****
    Id: 1
  User: root
  Host: localhost
    db: NULL
Command: Query
  Time: 0
  State: NULL
  Info: show processlist
***** 2. row *****
    Id: 2
  User: event_scheduler
  Host: localhost
    db: NULL
Command: Daemon
  Time: 3
  State: Waiting for next activation
  Info: NULL
2 rows in set (0.00 sec)
```

Event scheduling can be stopped by setting the value of `event_scheduler` to `OFF`.

- `OFF`: The Event Scheduler is stopped. The event scheduler thread does not run, is not shown in the output of `SHOW PROCESSLIST`, and no scheduled events execute.

When the Event Scheduler is stopped (`event_scheduler` is `OFF`), it can be started by setting the value of `event_scheduler` to `ON`. (See next item.)

- `DISABLED`: This value renders the Event Scheduler nonoperational. When the Event Scheduler is `DISABLED`, the event scheduler thread does not run (and so does not appear in the output of `SHOW PROCESSLIST`). In addition, the Event Scheduler state cannot be changed at runtime.

If the Event Scheduler status has not been set to `DISABLED`, `event_scheduler` can be toggled between `ON` and `OFF` (using `SET`). It is also possible to use `0` for `OFF`, and `1` for `ON` when setting this variable. Thus, any of the following 4 statements can be used in the `mysql` client to turn on the Event Scheduler:

```
SET GLOBAL event_scheduler = ON;
SET @@GLOBAL.event_scheduler = ON;
SET GLOBAL event_scheduler = 1;
SET @@GLOBAL.event_scheduler = 1;
```

Similarly, any of these 4 statements can be used to turn off the Event Scheduler:

```
SET GLOBAL event_scheduler = OFF;
SET @@GLOBAL.event_scheduler = OFF;
SET GLOBAL event_scheduler = 0;
SET @@GLOBAL.event_scheduler = 0;
```



#### Note

If the Event Scheduler is enabled, enabling the `super_read_only` system variable prevents it from updating event “last executed” timestamps in the

`events` data dictionary table. This causes the Event Scheduler to stop the next time it tries to execute a scheduled event, after writing a message to the server error log. (In this situation the `event_scheduler` system variable does not change from `ON` to `OFF`. An implication is that this variable rejects the DBA *intent* that the Event Scheduler be enabled or disabled, where its actual status of started or stopped may be distinct.). If `super_read_only` is subsequently disabled after being enabled, the server automatically restarts the Event Scheduler as needed, as of MySQL 8.0.26. Prior to MySQL 8.0.26, it is necessary to manually restart the Event Scheduler by enabling it again.

Although `ON` and `OFF` have numeric equivalents, the value displayed for `event_scheduler` by `SELECT` or `SHOW VARIABLES` is always one of `OFF`, `ON`, or `DISABLED`. `DISABLED` has no numeric equivalent. For this reason, `ON` and `OFF` are usually preferred over `1` and `0` when setting this variable.

Note that attempting to set `event_scheduler` without specifying it as a global variable causes an error:

```
mysql> SET @@event_scheduler = OFF;
ERROR 1229 (HY000): Variable 'event_scheduler' is a GLOBAL
variable and should be set with SET GLOBAL
```



### Important

It is possible to set the Event Scheduler to `DISABLED` only at server startup. If `event_scheduler` is `ON` or `OFF`, you cannot set it to `DISABLED` at runtime. Also, if the Event Scheduler is set to `DISABLED` at startup, you cannot change the value of `event_scheduler` at runtime.

To disable the event scheduler, use one of the following two methods:

- As a command-line option when starting the server:

```
--event-scheduler=DISABLED
```

- In the server configuration file (`my.cnf`, or `my.ini` on Windows systems), include the line where it can be read by the server (for example, in a `[mysqld]` section):

```
event_scheduler=DISABLED
```

To enable the Event Scheduler, restart the server without the `--event-scheduler=DISABLED` command-line option, or after removing or commenting out the line containing `event-scheduler=DISABLED` in the server configuration file, as appropriate. Alternatively, you can use `ON` (or `1`) or `OFF` (or `0`) in place of the `DISABLED` value when starting the server.



### Note

You can issue event-manipulation statements when `event_scheduler` is set to `DISABLED`. No warnings or errors are generated in such cases (provided that the statements are themselves valid). However, scheduled events cannot execute until this variable is set to `ON` (or `1`). Once this has been done, the event scheduler thread executes all events whose scheduling conditions are satisfied.

Starting the MySQL server with the `--skip-grant-tables` option causes `event_scheduler` to be set to `DISABLED`, overriding any other value set either on the command line or in the `my.cnf` or `my.ini` file (Bug #26807).

For SQL statements used to create, alter, and drop events, see [Section 25.4.3, “Event Syntax”](#).

MySQL provides an `EVENTS` table in the `INFORMATION_SCHEMA` database. This table can be queried to obtain information about scheduled events which have been defined on the server. See [Section 25.4.4, “Event Metadata”](#), and [Section 26.3.14, “The INFORMATION\\_SCHEMA EVENTS Table”](#), for more information.

For information regarding event scheduling and the MySQL privilege system, see [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#).

### 25.4.3 Event Syntax

MySQL provides several SQL statements for working with scheduled events:

- New events are defined using the `CREATE EVENT` statement. See [Section 13.1.13, “CREATE EVENT Statement”](#).
- The definition of an existing event can be changed by means of the `ALTER EVENT` statement. See [Section 13.1.3, “ALTER EVENT Statement”](#).
- When a scheduled event is no longer wanted or needed, it can be deleted from the server by its definer using the `DROP EVENT` statement. See [Section 13.1.25, “DROP EVENT Statement”](#). Whether an event persists past the end of its schedule also depends on its `ON COMPLETION` clause, if it has one. See [Section 13.1.13, “CREATE EVENT Statement”](#).

An event can be dropped by any user having the `EVENT` privilege for the database on which the event is defined. See [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#).

### 25.4.4 Event Metadata

To obtain metadata about events:

- Query the `EVENTS` table of the `INFORMATION_SCHEMA` database. See [Section 26.3.14, “The INFORMATION\\_SCHEMA EVENTS Table”](#).
- Use the `SHOW CREATE EVENT` statement. See [Section 13.7.7.7, “SHOW CREATE EVENT Statement”](#).
- Use the `SHOW EVENTS` statement. See [Section 13.7.7.18, “SHOW EVENTS Statement”](#).

#### Event Scheduler Time Representation

Each session in MySQL has a session time zone (STZ). This is the session `time_zone` value that is initialized from the server's global `time_zone` value when the session begins but may be changed during the session.

The session time zone that is current when a `CREATE EVENT` or `ALTER EVENT` statement executes is used to interpret times specified in the event definition. This becomes the event time zone (ETZ); that is, the time zone that is used for event scheduling and is in effect within the event as it executes.

For representation of event information in the data dictionary, the `execute_at`, `starts`, and `ends` times are converted to UTC and stored along with the event time zone. This enables event execution to proceed as defined regardless of any subsequent changes to the server time zone or daylight saving time effects. The `last_executed` time is also stored in UTC.

Event times can be obtained by selecting from the Information Schema `EVENTS` table or from `SHOW EVENTS`, but they are reported as ETZ or STZ values. The following table summarizes representation of event times.

Value	EVENTS Table	SHOW EVENTS
Execute at	ETZ	ETZ
Starts	ETZ	ETZ
Ends	ETZ	ETZ
Last executed	ETZ	n/a
Created	STZ	n/a
Last altered	STZ	n/a

## 25.4.5 Event Scheduler Status

The Event Scheduler writes information about event execution that terminates with an error or warning to the MySQL Server's error log. See [Section 25.4.6, “The Event Scheduler and MySQL Privileges”](#) for an example.

To obtain information about the state of the Event Scheduler for debugging and troubleshooting purposes, run `mysqladmin debug` (see [Section 4.5.2, “mysqladmin — A MySQL Server Administration Program”](#)); after running this command, the server's error log contains output relating to the Event Scheduler, similar to what is shown here:

```
Events status:
LLA = Last Locked At LUA = Last Unlocked At
WOC = Waiting On Condition DL = Data Locked

Event scheduler status:
State      : INITIALIZED
Thread id  : 0
LLA        : n/a:0
LUA        : n/a:0
WOC        : NO
Workers    : 0
Executed   : 0
Data locked: NO

Event queue status:
Element count   : 0
Data locked     : NO
Attempting lock : NO
LLA            : init_queue:95
LUA            : init_queue:103
WOC            : NO
Next activation : never
```

In statements that occur as part of events executed by the Event Scheduler, diagnostics messages (not only errors, but also warnings) are written to the error log, and, on Windows, to the application event log. For frequently executed events, it is possible for this to result in many logged messages. For example, for `SELECT ... INTO var_list` statements, if the query returns no rows, a warning with error code 1329 occurs ([No data](#)), and the variable values remain unchanged. If the query returns multiple rows, error 1172 occurs ([Result consisted of more than one row](#)). For either condition, you can avoid having the warnings be logged by declaring a condition handler; see [Section 13.6.7.2, “DECLARE ... HANDLER Statement”](#). For statements that may retrieve multiple rows, another strategy is to use `LIMIT 1` to limit the result set to a single row.

## 25.4.6 The Event Scheduler and MySQL Privileges

To enable or disable the execution of scheduled events, it is necessary to set the value of the global `event_scheduler` system variable. This requires privileges sufficient to set global system variables. See [Section 5.1.9.1, “System Variable Privileges”](#).

The `EVENT` privilege governs the creation, modification, and deletion of events. This privilege can be bestowed using `GRANT`. For example, this `GRANT` statement confers the `EVENT` privilege for the schema named `myschema` on the user `jon@ghidora`:

```
GRANT EVENT ON myschema.* TO jon@ghidora;
```

(We assume that this user account already exists, and that we wish for it to remain unchanged otherwise.)

To grant this same user the `EVENT` privilege on all schemas, use the following statement:

```
GRANT EVENT ON *.* TO jon@ghidora;
```

The `EVENT` privilege has global or schema-level scope. Therefore, trying to grant it on a single table results in an error as shown:

```
mysql> GRANT EVENT ON myschema.mytable TO jon@ghidora;
ERROR 1144 (42000): Illegal GRANT/REVOKE command; please
consult the manual to see which privileges can be used
```

It is important to understand that an event is executed with the privileges of its definer, and that it cannot perform any actions for which its definer does not have the requisite privileges. For example, suppose that `jon@ghidora` has the `EVENT` privilege for `myschema`. Suppose also that this user has the `SELECT` privilege for `myschema`, but no other privileges for this schema. It is possible for `jon@ghidora` to create a new event such as this one:

```
CREATE EVENT e_store_ts
  ON SCHEDULE
    EVERY 10 SECOND
  DO
    INSERT INTO myschema.mytable VALUES (UNIX_TIMESTAMP());
```

The user waits for a minute or so, and then performs a `SELECT * FROM mytable;` query, expecting to see several new rows in the table. Instead, the table is empty. Since the user does not have the `INSERT` privilege for the table in question, the event has no effect.

If you inspect the MySQL error log (`hostname.err`), you can see that the event is executing, but the action it is attempting to perform fails:

```
2013-09-24T12:41:31.261992Z 25 [ERROR] Event Scheduler:
[jon@ghidora][cookbook.e_store_ts] INSERT command denied to user
'jon'@'ghidora' for table 'mytable'
2013-09-24T12:41:31.262022Z 25 [Note] Event Scheduler:
[jon@ghidora].[myschema.e_store_ts] event execution failed.
2013-09-24T12:41:41.271796Z 26 [ERROR] Event Scheduler:
[jon@ghidora][cookbook.e_store_ts] INSERT command denied to user
'jon'@'ghidora' for table 'mytable'
2013-09-24T12:41:41.272761Z 26 [Note] Event Scheduler:
[jon@ghidora].[myschema.e_store_ts] event execution failed.
```

Since this user very likely does not have access to the error log, it is possible to verify whether the event's action statement is valid by executing it directly:

```
mysql> INSERT INTO myschema.mytable VALUES (UNIX_TIMESTAMP());
ERROR 1142 (42000): INSERT command denied to user
'jon'@'ghidora' for table 'mytable'
```

Inspection of the Information Schema `EVENTS` table shows that `e_store_ts` exists and is enabled, but its `LAST_EXECUTED` column is `NULL`:

```
mysql> SELECT * FROM INFORMATION_SCHEMA.EVENTS
  >   WHERE EVENT_NAME='e_store_ts'
  >     AND EVENT_SCHEMA='myschema'\G
***** 1. row *****
EVENT_CATALOG: NULL
EVENT_SCHEMA: myschema
EVENT_NAME: e_store_ts
DEFINER: jon@ghidora
EVENT_BODY: SQL
EVENT_DEFINITION: INSERT INTO myschema.mytable VALUES (UNIX_TIMESTAMP())
EVENT_TYPE: RECURRING
EXECUTE_AT: NULL
INTERVAL_VALUE: 5
INTERVAL_FIELD: SECOND
SQL_MODE: NULL
STARTS: 0000-00-00 00:00:00
ENDS: 0000-00-00 00:00:00
STATUS: ENABLED
ON_COMPLETION: NOT PRESERVE
CREATED: 2006-02-09 22:36:06
LAST_ALTERED: 2006-02-09 22:36:06
LAST_EXECUTED: NULL
EVENT_COMMENT:
1 row in set (0.00 sec)
```

To rescind the `EVENT` privilege, use the `REVOKE` statement. In this example, the `EVENT` privilege on the schema `myschema` is removed from the `jon@ghidora` user account:

```
REVOKE EVENT ON myschema.* FROM jon@ghidora;
```



### Important

Revoking the `EVENT` privilege from a user does not delete or disable any events that may have been created by that user.

An event is not migrated or dropped as a result of renaming or dropping the user who created it.

Suppose that the user `jon@ghidora` has been granted the `EVENT` and `INSERT` privileges on the `myschema` schema. This user then creates the following event:

```
CREATE EVENT e_insert
  ON SCHEDULE
    EVERY 7 SECOND
  DO
    INSERT INTO myschema.mytable;
```

After this event has been created, `root` revokes the `EVENT` privilege for `jon@ghidora`. However, `e_insert` continues to execute, inserting a new row into `mytable` each seven seconds. The same would be true if `root` had issued either of these statements:

- `DROP USER jon@ghidora;`
- `RENAME USER jon@ghidora TO someotherguy@ghidora;`

You can verify that this is true by examining the Information Schema `EVENTS` table before and after issuing a `DROP USER` or `RENAME USER` statement.

Event definitions are stored in the data dictionary. To drop an event created by another user account, you must be the MySQL `root` user or another user with the necessary privileges.

Users' `EVENT` privileges are stored in the `Event_priv` columns of the `mysql.user` and `mysql.db` tables. In both cases, this column holds one of the values '`Y`' or '`N`'. '`N`' is the default. `mysql.user.Event_priv` is set to '`Y`' for a given user only if that user has the global `EVENT` privilege (that is, if the privilege was bestowed using `GRANT EVENT ON *.*`). For a schema-level `EVENT` privilege, `GRANT` creates a row in `mysql.db` and sets that row's `Db` column to the name of the schema, the `User` column to the name of the user, and the `Event_priv` column to '`Y`'. There should never be any need to manipulate these tables directly, since the `GRANT EVENT` and `REVOKE EVENT` statements perform the required operations on them.

Five status variables provide counts of event-related operations (but *not* of statements executed by events; see [Section 25.8, “Restrictions on Stored Programs”](#)). These are:

- `Com_create_event`: The number of `CREATE EVENT` statements executed since the last server restart.
- `Com_alter_event`: The number of `ALTER EVENT` statements executed since the last server restart.
- `Com_drop_event`: The number of `DROP EVENT` statements executed since the last server restart.
- `Com_show_create_event`: The number of `SHOW CREATE EVENT` statements executed since the last server restart.
- `Com_show_events`: The number of `SHOW EVENTS` statements executed since the last server restart.

You can view current values for all of these at one time by running the statement `SHOW STATUS LIKE '%event%';`.

## 25.5 Using Views

MySQL supports views, including updatable views. Views are stored queries that when invoked produce a result set. A view acts as a virtual table.

The following discussion describes the syntax for creating and dropping views, and shows some examples of how to use them.

### Additional Resources

- You may find the [MySQL User Forums](#) helpful when working with views.
- For answers to some commonly asked questions regarding views in MySQL, see [Section A.6, “MySQL 8.0 FAQ: Views”](#).
- There are some restrictions on the use of views; see [Section 25.9, “Restrictions on Views”](#).

#### 25.5.1 View Syntax

The `CREATE VIEW` statement creates a new view (see [Section 13.1.23, “CREATE VIEW Statement”](#)). To alter the definition of a view or drop a view, use `ALTER VIEW` (see [Section 13.1.11, “ALTER VIEW Statement”](#)), or `DROP VIEW` (see [Section 13.1.35, “DROP VIEW Statement”](#)).

A view can be created from many kinds of `SELECT` statements. It can refer to base tables or other views. It can use joins, `UNION`, and subqueries. The `SELECT` need not even refer to any tables. The following example defines a view that selects two columns from another table, as well as an expression calculated from those columns:

```
mysql> CREATE TABLE t (qty INT, price INT);
mysql> INSERT INTO t VALUES(3, 50), (5, 60);
mysql> CREATE VIEW v AS SELECT qty, price, qty*price AS value FROM t;
mysql> SELECT * FROM v;
+-----+-----+-----+
| qty | price | value |
+-----+-----+-----+
|    3 |     50 |    150 |
|    5 |     60 |    300 |
+-----+-----+-----+
mysql> SELECT * FROM v WHERE qty = 5;
+-----+-----+-----+
| qty | price | value |
+-----+-----+-----+
|    5 |     60 |    300 |
+-----+-----+-----+
```

#### 25.5.2 View Processing Algorithms

The optional `ALGORITHM` clause for `CREATE VIEW` or `ALTER VIEW` is a MySQL extension to standard SQL. It affects how MySQL processes the view. `ALGORITHM` takes three values: `MERGE`, `TEMPTABLE`, or `UNDEFINED`.

- For `MERGE`, the text of a statement that refers to the view and the view definition are merged such that parts of the view definition replace corresponding parts of the statement.
- For `TEMPTABLE`, the results from the view are retrieved into a temporary table, which then is used to execute the statement.
- For `UNDEFINED`, MySQL chooses which algorithm to use. It prefers `MERGE` over `TEMPTABLE` if possible, because `MERGE` is usually more efficient and because a view cannot be updatable if a temporary table is used.
- If no `ALGORITHM` clause is present, the default algorithm is determined by the value of the `derived_merge` flag of the `optimizer_switch` system variable. For additional discussion, see

#### [Section 8.2.2.4, “Optimizing Derived Tables, View References, and Common Table Expressions with Merging or Materialization”.](#)

A reason to specify `TEMPTABLE` explicitly is that locks can be released on underlying tables after the temporary table has been created and before it is used to finish processing the statement. This might result in quicker lock release than the `MERGE` algorithm so that other clients that use the view are not blocked as long.

A view algorithm can be `UNDEFINED` for three reasons:

- No `ALGORITHM` clause is present in the `CREATE VIEW` statement.
- The `CREATE VIEW` statement has an explicit `ALGORITHM = UNDEFINED` clause.
- `ALGORITHM = MERGE` is specified for a view that can be processed only with a temporary table. In this case, MySQL generates a warning and sets the algorithm to `UNDEFINED`.

As mentioned earlier, `MERGE` is handled by merging corresponding parts of a view definition into the statement that refers to the view. The following examples briefly illustrate how the `MERGE` algorithm works. The examples assume that there is a view `v_merge` that has this definition:

```
CREATE ALGORITHM = MERGE VIEW v_merge (vc1, vc2) AS
SELECT c1, c2 FROM t WHERE c3 > 100;
```

Example 1: Suppose that we issue this statement:

```
SELECT * FROM v_merge;
```

MySQL handles the statement as follows:

- `v_merge` becomes `t`
- `*` becomes `vc1, vc2`, which corresponds to `c1, c2`
- The view `WHERE` clause is added

The resulting statement to be executed becomes:

```
SELECT c1, c2 FROM t WHERE c3 > 100;
```

Example 2: Suppose that we issue this statement:

```
SELECT * FROM v_merge WHERE vc1 < 100;
```

This statement is handled similarly to the previous one, except that `vc1 < 100` becomes `c1 < 100` and the view `WHERE` clause is added to the statement `WHERE` clause using an `AND` connective (and parentheses are added to make sure the parts of the clause are executed with correct precedence). The resulting statement to be executed becomes:

```
SELECT c1, c2 FROM t WHERE (c3 > 100) AND (c1 < 100);
```

Effectively, the statement to be executed has a `WHERE` clause of this form:

```
WHERE (select WHERE) AND (view WHERE)
```

If the `MERGE` algorithm cannot be used, a temporary table must be used instead. Constructs that prevent merging are the same as those that prevent merging in derived tables and common table expressions. Examples are `SELECT DISTINCT` or `LIMIT` in the subquery. For details, see [Section 8.2.2.4, “Optimizing Derived Tables, View References, and Common Table Expressions with Merging or Materialization”](#).

### 25.5.3 Updatable and Insertable Views

Some views are updatable and references to them can be used to specify tables to be updated in data change statements. That is, you can use them in statements such as `UPDATE`, `DELETE`, or `INSERT` to update the contents of the underlying table. Derived tables and common table expressions can also be

specified in multiple-table `UPDATE` and `DELETE` statements, but can only be used for reading data to specify rows to be updated or deleted. Generally, the view references must be updatable, meaning that they may be merged and not materialized. Composite views have more complex rules.

For a view to be updatable, there must be a one-to-one relationship between the rows in the view and the rows in the underlying table. There are also certain other constructs that make a view nonupdatable. To be more specific, a view is not updatable if it contains any of the following:

- Aggregate functions or window functions (`SUM()`, `MIN()`, `MAX()`, `COUNT()`, and so forth)
- `DISTINCT`
- `GROUP BY`
- `HAVING`
- `UNION` or `UNION ALL`
- Subquery in the select list

Nondependent subqueries in the select list fail for `INSERT`, but are okay for `UPDATE`, `DELETE`. For dependent subqueries in the select list, no data change statements are permitted.

- Certain joins (see additional join discussion later in this section)
- Reference to nonupdatable view in the `FROM` clause
- Subquery in the `WHERE` clause that refers to a table in the `FROM` clause
- Refers only to literal values (in this case, there is no underlying table to update)
- `ALGORITHM = TEMPTABLE` (use of a temporary table always makes a view nonupdatable)
- Multiple references to any column of a base table (fails for `INSERT`, okay for `UPDATE`, `DELETE`)

A generated column in a view is considered updatable because it is possible to assign to it. However, if such a column is updated explicitly, the only permitted value is `DEFAULT`. For information about generated columns, see [Section 13.1.20.8, “CREATE TABLE and Generated Columns”](#).

It is sometimes possible for a multiple-table view to be updatable, assuming that it can be processed with the `MERGE` algorithm. For this to work, the view must use an inner join (not an outer join or a `UNION`). Also, only a single table in the view definition can be updated, so the `SET` clause must name only columns from one of the tables in the view. Views that use `UNION ALL` are not permitted even though they might be theoretically updatable.

With respect to insertability (being updatable with `INSERT` statements), an updatable view is insertable if it also satisfies these additional requirements for the view columns:

- There must be no duplicate view column names.
- The view must contain all columns in the base table that do not have a default value.
- The view columns must be simple column references. They must not be expressions, such as these:

```
3.14159
col1 + 3
UPPER(col2)
col3 / col4
(subquery)
```

MySQL sets a flag, called the view updatability flag, at `CREATE VIEW` time. The flag is set to `YES` (true) if `UPDATE` and `DELETE` (and similar operations) are legal for the view. Otherwise, the flag is set to `NO` (false). The `IS_UPDATABLE` column in the Information Schema `VIEWS` table displays the status of this flag. It means that the server always knows whether a view is updatable.

If a view is not updatable, statements such `UPDATE`, `DELETE`, and `INSERT` are illegal and are rejected. (Even if a view is updatable, it might not be possible to insert into it, as described elsewhere in this section.)

The updatability of views may be affected by the value of the `updatable_views_with_limit` system variable. See [Section 5.1.8, “Server System Variables”](#).

For the following discussion, suppose that these tables and views exist:

```
CREATE TABLE t1 (x INTEGER);
CREATE TABLE t2 (c INTEGER);
CREATE VIEW vmat AS SELECT SUM(x) AS s FROM t1;
CREATE VIEW vup AS SELECT * FROM t2;
CREATE VIEW vjoin AS SELECT * FROM vmat JOIN vup ON vmat.s=vup.c;
```

`INSERT`, `UPDATE`, and `DELETE` statements are permitted as follows:

- `INSERT`: The insert table of an `INSERT` statement may be a view reference that is merged. If the view is a join view, all components of the view must be updatable (not materialized). For a multiple-table updatable view, `INSERT` can work if it inserts into a single table.

This statement is invalid because one component of the join view is nonupdatable:

```
INSERT INTO vjoin (c) VALUES (1);
```

This statement is valid; the view contains no materialized components:

```
INSERT INTO vup (c) VALUES (1);
```

- `UPDATE`: The table or tables to be updated in an `UPDATE` statement may be view references that are merged. If a view is a join view, at least one component of the view must be updatable (this differs from `INSERT`).

In a multiple-table `UPDATE` statement, the updated table references of the statement must be base tables or updatable view references. Nonupdated table references may be materialized views or derived tables.

This statement is valid; column `c` is from the updatable part of the join view:

```
UPDATE vjoin SET c=c+1;
```

This statement is invalid; column `x` is from the nonupdatable part:

```
UPDATE vjoin SET x=x+1;
```

This statement is valid; the updated table reference of the multiple-table `UPDATE` is an updatable view (`vup`):

```
UPDATE vup JOIN (SELECT SUM(x) AS s FROM t1) AS dt ON ...
SET c=c+1;
```

This statement is invalid; it tries to update a materialized derived table:

```
UPDATE vup JOIN (SELECT SUM(x) AS s FROM t1) AS dt ON ...
SET s=s+1;
```

- `DELETE`: The table or tables to be deleted from in a `DELETE` statement must be merged views. Join views are not allowed (this differs from `INSERT` and `UPDATE`).

This statement is invalid because the view is a join view:

```
DELETE vjoin WHERE ...;
```

This statement is valid because the view is a merged (updatable) view:

```
DELETE vup WHERE ...;
```

This statement is valid because it deletes from a merged (updatable) view:

```
DELETE vup FROM vup JOIN (SELECT SUM(x) AS s FROM t1) AS dt ON ...;
```

Additional discussion and examples follow.

Earlier discussion in this section pointed out that a view is not insertable if not all columns are simple column references (for example, if it contains columns that are expressions or composite expressions). Although such a view is not insertable, it can be updatable if you update only columns that are not expressions. Consider this view:

```
CREATE VIEW v AS SELECT col1, 1 AS col2 FROM t;
```

This view is not insertable because `col2` is an expression. But it is updatable if the update does not try to update `col2`. This update is permissible:

```
UPDATE v SET col1 = 0;
```

This update is not permissible because it attempts to update an expression column:

```
UPDATE v SET col2 = 0;
```

If a table contains an `AUTO_INCREMENT` column, inserting into an insertable view on the table that does not include the `AUTO_INCREMENT` column does not change the value of `LAST_INSERT_ID()`, because the side effects of inserting default values into columns not part of the view should not be visible.

## 25.5.4 The View WITH CHECK OPTION Clause

The `WITH CHECK OPTION` clause can be given for an updatable view to prevent inserts to rows for which the `WHERE` clause in the `select_statement` is not true. It also prevents updates to rows for which the `WHERE` clause is true but the update would cause it to be not true (in other words, it prevents visible rows from being updated to nonvisible rows).

In a `WITH CHECK OPTION` clause for an updatable view, the `LOCAL` and `CASCDED` keywords determine the scope of check testing when the view is defined in terms of another view. When neither keyword is given, the default is `CASCDED`.

`WITH CHECK OPTION` testing is standard-compliant:

- With `LOCAL`, the view `WHERE` clause is checked, then checking recurses to underlying views and applies the same rules.
- With `CASCDED`, the view `WHERE` clause is checked, then checking recurses to underlying views, adds `WITH CASCDED CHECK OPTION` to them (for purposes of the check; their definitions remain unchanged), and applies the same rules.
- With no check option, the view `WHERE` clause is not checked, then checking recurses to underlying views, and applies the same rules.

Consider the definitions for the following table and set of views:

```
CREATE TABLE t1 (a INT);
CREATE VIEW v1 AS SELECT * FROM t1 WHERE a < 2
WITH CHECK OPTION;
CREATE VIEW v2 AS SELECT * FROM v1 WHERE a > 0
WITH LOCAL CHECK OPTION;
CREATE VIEW v3 AS SELECT * FROM v1 WHERE a > 0
WITH CASCDED CHECK OPTION;
```

Here the `v2` and `v3` views are defined in terms of another view, `v1`.

Inserts for `v2` are checked against its `LOCAL` check option, then the check recurses to `v1` and the rules are applied again. The rules for `v1` cause a check failure. The check for `v3` also fails:

```
mysql> INSERT INTO v2 VALUES (2);
ERROR 1369 (HY000): CHECK OPTION failed 'test.v2'
mysql> INSERT INTO v3 VALUES (2);
ERROR 1369 (HY000): CHECK OPTION failed 'test.v3'
```

## 25.5.5 View Metadata

To obtain metadata about views:

- Query the `VIEWS` table of the `INFORMATION_SCHEMA` database. See [Section 26.3.48, “The INFORMATION\\_SCHEMA VIEWS Table”](#).
- Use the `SHOW CREATE VIEW` statement. See [Section 13.7.7.13, “SHOW CREATE VIEW Statement”](#).

## 25.6 Stored Object Access Control

Stored programs (procedures, functions, triggers, and events) and views are defined prior to use and, when referenced, execute within a security context that determines their privileges. The privileges applicable to execution of a stored object are controlled by its `DEFINER` attribute and `SQL SECURITY` characteristic.

- [The DEFINER Attribute](#)
- [The SQL SECURITY Characteristic](#)
- [Examples](#)
- [Orphan Stored Objects](#)
- [Risk-Minimization Guidelines](#)

### The DEFINER Attribute

A stored object definition can include a `DEFINER` attribute that names a MySQL account. If a definition omits the `DEFINER` attribute, the default object definer is the user who creates it.

The following rules determine which accounts you can specify as the `DEFINER` attribute for a stored object:

- If you have the `SET_USER_ID` privilege (or the deprecated `SUPER` privilege), you can specify any account as the `DEFINER` attribute. If the account does not exist, a warning is generated. Additionally, to set a stored object `DEFINER` attribute to an account that has the `SYSTEM_USER` privilege, you must have the `SYSTEM_USER` privilege.
- Otherwise, the only permitted account is your own, specified either literally or as `CURRENT_USER` or `CURRENT_USER()`. You cannot set the definer to any other account.

Creating a stored object with a nonexistent `DEFINER` account creates an orphan object, which may have negative consequences; see [Orphan Stored Objects](#).

### The SQL SECURITY Characteristic

For stored routines (procedures and functions) and views, the object definition can include an `SQL SECURITY` characteristic with a value of `DEFINER` or `INVOKER` to specify whether the object executes in definer or invoker context. If the definition omits the `SQL SECURITY` characteristic, the default is definer context.

Triggers and events have no `SQL SECURITY` characteristic and always execute in definer context. The server invokes these objects automatically as necessary, so there is no invoking user.

Definer and invoker security contexts differ as follows:

- A stored object that executes in definer security context executes with the privileges of the account named by its `DEFINER` attribute. These privileges may be entirely different from those of the invoking user. The invoker must have appropriate privileges to reference the object (for example, `EXECUTE` to call a stored procedure or `SELECT` to select from a view), but during object execution, the invoker's privileges are ignored and only the `DEFINER` account privileges matter. If the `DEFINER` account has few privileges, the object is correspondingly limited in the operations it can perform. If the `DEFINER` account is highly privileged (such as an administrative account), the object can perform powerful operations *no matter who invokes it*.
- A stored routine or view that executes in invoker security context can perform only operations for which the invoker has privileges. The `DEFINER` attribute has no effect on object execution.

## Examples

Consider the following stored procedure, which is declared with `SQL SECURITY DEFINER` to execute in definer security context:

```
CREATE DEFINER = 'admin'@'localhost' PROCEDURE p1()
SQL SECURITY DEFINER
BEGIN
    UPDATE t1 SET counter = counter + 1;
END;
```

Any user who has the `EXECUTE` privilege for `p1` can invoke it with a `CALL` statement. However, when `p1` executes, it does so in definer security context and thus executes with the privileges of `'admin'@'localhost'`, the account named as its `DEFINER` attribute. This account must have the `EXECUTE` privilege for `p1` as well as the `UPDATE` privilege for the table `t1` referenced within the object body. Otherwise, the procedure fails.

Now consider this stored procedure, which is identical to `p1` except that its `SQL SECURITY` characteristic is `INVOKER`:

```
CREATE DEFINER = 'admin'@'localhost' PROCEDURE p2()
SQL SECURITY INVOKER
BEGIN
    UPDATE t1 SET counter = counter + 1;
END;
```

Unlike `p1`, `p2` executes in invoker security context and thus with the privileges of the invoking user regardless of the `DEFINER` attribute value. `p2` fails if the invoker lacks the `EXECUTE` privilege for `p2` or the `UPDATE` privilege for the table `t1`.

## Orphan Stored Objects

An orphan stored object is one for which its `DEFINER` attribute names a nonexistent account:

- An orphan stored object can be created by specifying a nonexistent `DEFINER` account at object-creation time.
- An existing stored object can become orphaned through execution of a `DROP USER` statement that drops the object `DEFINER` account, or a `RENAME USER` statement that renames the object `DEFINER` account.

An orphan stored object may be problematic in these ways:

- Because the `DEFINER` account does not exist, the object may not work as expected if it executes in definer security context:
  - For a stored routine, an error occurs at routine execution time if the `SQL SECURITY` value is `DEFINER` but the definer account does not exist.
  - For a trigger, it is not a good idea for trigger activation to occur until the account actually does exist. Otherwise, the behavior with respect to privilege checking is undefined.

- For an event, an error occurs at event execution time if the account does not exist.
- For a view, an error occurs when the view is referenced if the `SQL SECURITY` value is `DEFINER` but the definer account does not exist.
- The object may present a security risk if the nonexistent `DEFINER` account is subsequently re-created for a purpose unrelated to the object. In this case, the account “adopts” the object and, with the appropriate privileges, is able to execute it even if that is not intended.

As of MySQL 8.0.22, the server imposes additional account-management security checks designed to prevent operations that (perhaps inadvertently) cause stored objects to become orphaned or that cause adoption of stored objects that are currently orphaned:

- `DROP USER` fails with an error if any account to be dropped is named as the `DEFINER` attribute for any stored object. (That is, the statement fails if dropping an account would cause a stored object to become orphaned.)
- `RENAME USER` fails with an error if any account to be renamed is named as the `DEFINER` attribute for any stored object. (That is, the statement fails if renaming an account would cause a stored object to become orphaned.)
- `CREATE USER` fails with an error if any account to be created is named as the `DEFINER` attribute for any stored object. (That is, the statement fails if creating an account would cause the account to adopt a currently orphaned stored object.)

In certain situations, it may be necessary to deliberately execute those account-management statements even when they would otherwise fail. To make this possible, if a user has the `SET_USER_ID` privilege, that privilege overrides the orphan object security checks and the statements succeed with a warning rather than failing with an error.

To obtain information about the accounts used as stored object definers in a MySQL installation, query the `INFORMATION_SCHEMA`.

This query identifies which `INFORMATION_SCHEMA` tables describe objects that have a `DEFINER` attribute:

```
mysql> SELECT TABLE_SCHEMA, TABLE_NAME FROM INFORMATION_SCHEMA.COLUMNS
      WHERE COLUMN_NAME = 'DEFINER';
+-----+-----+
| TABLE_SCHEMA | TABLE_NAME |
+-----+-----+
| information_schema | EVENTS |
| information_schema | ROUTINES |
| information_schema | TRIGGERS |
| information_schema | VIEWS |
+-----+-----+
```

The result tells you which tables to query to discover which stored object `DEFINER` values exist and which objects have a particular `DEFINER` value:

- To identify which `DEFINER` values exist in each table, use these queries:

```
SELECT DISTINCT DEFINER FROM INFORMATION_SCHEMA.EVENTS;
SELECT DISTINCT DEFINER FROM INFORMATION_SCHEMA.ROUTINES;
SELECT DISTINCT DEFINER FROM INFORMATION_SCHEMA.TRIGGERS;
SELECT DISTINCT DEFINER FROM INFORMATION_SCHEMA.VIEWS;
```

The query results are significant for any account displayed as follows:

- If the account exists, dropping or renaming it causes stored objects to become orphaned. If you plan to drop or rename the account, consider first dropping its associated stored objects or redefining them to have a different definer.

- If the account does not exist, creating it causes it to adopt currently orphaned stored objects. If you plan to create the account, consider whether the orphaned objects should be associated with it. If not, redefine them to have a different definer.

To redefine an object with a different definer, you can use `ALTER EVENT` or `ALTER VIEW` to directly modify the `DEFINER` account of events and views. For stored procedures and functions and for triggers, you must drop the object and re-create it to assign a different `DEFINER` account

- To identify which objects have a given `DEFINER` account, use these queries, substituting the account of interest for `user_name@host_name`:

```
SELECT EVENT_SCHEMA, EVENT_NAME FROM INFORMATION_SCHEMA.EVENTS
WHERE DEFINER = 'user_name@host_name';
SELECT ROUTINE_SCHEMA, ROUTINE_NAME, ROUTINE_TYPE
FROM INFORMATION_SCHEMA.ROUTINES
WHERE DEFINER = 'user_name@host_name';
SELECT TRIGGER_SCHEMA, TRIGGER_NAME FROM INFORMATION_SCHEMA.TRIGGERS
WHERE DEFINER = 'user_name@host_name';
SELECT TABLE_SCHEMA, TABLE_NAME FROM INFORMATION_SCHEMA.VIEWS
WHERE DEFINER = 'user_name@host_name';
```

For the `ROUTINES` table, the query includes the `ROUTINE_TYPE` column so that output rows distinguish whether the `DEFINER` is for a stored procedure or stored function.

If the account you are searching for does not exist, any objects displayed by those queries are orphan objects.

## Risk-Minimization Guidelines

To minimize the risk potential for stored object creation and use, follow these guidelines:

- Do not create orphan stored objects; that is, objects for which the `DEFINER` attribute names a nonexistent account. Do not cause stored objects to become orphaned by dropping or renaming an account named by the `DEFINER` attribute of any existing object.
- For a stored routine or view, use `SQL SECURITY INVOKER` in the object definition when possible so that it can be used only by users with permissions appropriate for the operations performed by the object.
- If you create definer-context stored objects while using an account that has the `SET_USER_ID` privilege (or the deprecated `SUPER` privilege), specify an explicit `DEFINER` attribute that names an account possessing only the privileges required for the operations performed by the object. Specify a highly privileged `DEFINER` account only when absolutely necessary.
- Administrators can prevent users from creating stored objects that specify highly privileged `DEFINER` accounts by not granting them the `SET_USER_ID` privilege (or the deprecated `SUPER` privilege).
- Definer-context objects should be written keeping in mind that they may be able to access data for which the invoking user has no privileges. In some cases, you can prevent references to these objects by not granting unauthorized users particular privileges:
  - A stored routine cannot be referenced by a user who does not have the `EXECUTE` privilege for it.
  - A view cannot be referenced by a user who does not have the appropriate privilege for it (`SELECT` to select from it, `INSERT` to insert into it, and so forth).

However, no such control exists for triggers and events because they always execute in definer context. The server invokes these objects automatically as necessary, and users do not reference them directly:

- A trigger is activated by access to the table with which it is associated, even ordinary table accesses by users with no special privileges.

- An event is executed by the server on a scheduled basis.

In both cases, if the `DEFINER` account is highly privileged, the object may be able to perform sensitive or dangerous operations. This remains true if the privileges needed to create the object are revoked from the account of the user who created it. Administrators should be especially careful about granting users object-creation privileges.

- By default, when a routine with the `SQL SECURITY DEFINER` characteristic is executed, MySQL Server does not set any active roles for the MySQL account named in the `DEFINER` clause, only the default roles. The exception is if the `activate_all_roles_on_login` system variable is enabled, in which case MySQL Server sets all roles granted to the `DEFINER` user, including mandatory roles. Any privileges granted through roles are therefore not checked by default when the `CREATE PROCEDURE` or `CREATE FUNCTION` statement is issued. For stored programs, if execution should occur with roles different from the default, the program body can execute `SET ROLE` to activate the required roles. This must be done with caution since the privileges assigned to roles can be changed.

## 25.7 Stored Program Binary Logging

The binary log contains information about SQL statements that modify database contents. This information is stored in the form of “events” that describe the modifications. (Binary log events differ from scheduled event stored objects.) The binary log has two important purposes:

- For replication, the binary log is used on source replication servers as a record of the statements to be sent to replica servers. The source sends the events contained in its binary log to its replicas, which execute those events to make the same data changes that were made on the source. See [Section 17.2, “Replication Implementation”](#).
- Certain data recovery operations require use of the binary log. After a backup file has been restored, the events in the binary log that were recorded after the backup was made are re-executed. These events bring databases up to date from the point of the backup. See [Section 7.3.2, “Using Backups for Recovery”](#).

However, if logging occurs at the statement level, there are certain binary logging issues with respect to stored programs (stored procedures and functions, triggers, and events):

- In some cases, a statement might affect different sets of rows on source and replica.
- Replicated statements executed on a replica are processed by the replica's applier thread. Unless you implement replication privilege checks, which are available from MySQL 8.0.18 (see [Section 17.3.3, “Replication Privilege Checks”](#)), the applier thread has full privileges. In this situation, it is possible for a procedure to follow different execution paths on source and replica servers, so a user could write a routine containing a dangerous statement that executes only on the replica.
- If a stored program that modifies data is nondeterministic, it is not repeatable. This can result in different data on source and replica, or cause restored data to differ from the original data.

This section describes how MySQL handles binary logging for stored programs. It states the current conditions that the implementation places on the use of stored programs, and what you can do to avoid logging problems. It also provides additional information about the reasons for these conditions.

Unless noted otherwise, the remarks here assume that binary logging is enabled on the server (see [Section 5.4.4, “The Binary Log”](#).) If the binary log is not enabled, replication is not possible, nor is the binary log available for data recovery. From MySQL 8.0, binary logging is enabled by default, and is only disabled if you specify the `--skip-log-bin` or `--disable-log-bin` option at startup.

In general, the issues described here result when binary logging occurs at the SQL statement level (statement-based binary logging). If you use row-based binary logging, the log contains changes made to individual rows as a result of executing SQL statements. When routines or triggers execute, row changes are logged, not the statements that make the changes. For stored procedures, this means

that the `CALL` statement is not logged. For stored functions, row changes made within the function are logged, not the function invocation. For triggers, row changes made by the trigger are logged. On the replica side, only the row changes are seen, not the stored program invocation.

Mixed format binary logging (`binlog_format=MIXED`) uses statement-based binary logging, except for cases where only row-based binary logging is guaranteed to lead to proper results. With mixed format, when a stored function, stored procedure, trigger, event, or prepared statement contains anything that is not safe for statement-based binary logging, the entire statement is marked as unsafe and logged in row format. The statements used to create and drop procedures, functions, triggers, and events are always safe, and are logged in statement format. For more information about row-based, mixed, and statement-based logging, and how safe and unsafe statements are determined, see [Section 17.2.1, “Replication Formats”](#).

The conditions on the use of stored functions in MySQL can be summarized as follows. These conditions do not apply to stored procedures or Event Scheduler events and they do not apply unless binary logging is enabled.

- To create or alter a stored function, you must have the `SET_USER_ID` privilege (or the deprecated `SUPER` privilege), in addition to the `CREATE ROUTINE` or `ALTER ROUTINE` privilege that is normally required. (Depending on the `DEFINER` value in the function definition, `SET_USER_ID` or `SUPER` might be required regardless of whether binary logging is enabled. See [Section 13.1.17, “CREATE PROCEDURE and CREATE FUNCTION Statements”](#).)
- When you create a stored function, you must declare either that it is deterministic or that it does not modify data. Otherwise, it may be unsafe for data recovery or replication.

By default, for a `CREATE FUNCTION` statement to be accepted, at least one of `DETERMINISTIC`, `NO SQL`, or `READS SQL DATA` must be specified explicitly. Otherwise an error occurs:

```
ERROR 1418 (HY000): This function has none of DETERMINISTIC, NO SQL,
or READS SQL DATA in its declaration and binary logging is enabled
(you *might* want to use the less safe log_bin_trust_function_creators
variable)
```

This function is deterministic (and does not modify data), so it is safe:

```
CREATE FUNCTION f1(i INT)
RETURNS INT
DETERMINISTIC
READS SQL DATA
BEGIN
    RETURN i;
END;
```

This function uses `UUID()`, which is not deterministic, so the function also is not deterministic and is not safe:

```
CREATE FUNCTION f2()
RETURNS CHAR(36) CHARACTER SET utf8mb4
BEGIN
    RETURN UUID();
END;
```

This function modifies data, so it may not be safe:

```
CREATE FUNCTION f3(p_id INT)
RETURNS INT
BEGIN
    UPDATE t SET modtime = NOW() WHERE id = p_id;
    RETURN ROW_COUNT();
END;
```

Assessment of the nature of a function is based on the “honesty” of the creator. MySQL does not check that a function declared `DETERMINISTIC` is free of statements that produce nondeterministic results.

- When you attempt to execute a stored function, if `binlog_format=STATEMENT` is set, the `DETERMINISTIC` keyword must be specified in the function definition. If this is not the case, an error is generated and the function does not run, unless `log_bin_trust_function_creators=1` is specified to override this check (see below). For recursive function calls, the `DETERMINISTIC` keyword is required on the outermost call only. If row-based or mixed binary logging is in use, the statement is accepted and replicated even if the function was defined without the `DETERMINISTIC` keyword.
- Because MySQL does not check if a function really is deterministic at creation time, the invocation of a stored function with the `DETERMINISTIC` keyword might carry out an action that is unsafe for statement-based logging, or invoke a function or procedure containing unsafe statements. If this occurs when `binlog_format=STATEMENT` is set, a warning message is issued. If row-based or mixed binary logging is in use, no warning is issued, and the statement is replicated in row-based format.
- To relax the preceding conditions on function creation (that you must have the `SUPER` privilege and that a function must be declared deterministic or to not modify data), set the global `log_bin_trust_function_creators` system variable to 1. By default, this variable has a value of 0, but you can change it like this:

```
mysql> SET GLOBAL log_bin_trust_function_creators = 1;
```

You can also set this variable at server startup.

If binary logging is not enabled, `log_bin_trust_function_creators` does not apply. `SUPER` is not required for function creation unless, as described previously, the `DEFINER` value in the function definition requires it.

- For information about built-in functions that may be unsafe for replication (and thus cause stored functions that use them to be unsafe as well), see [Section 17.5.1, “Replication Features and Issues”](#).

Triggers are similar to stored functions, so the preceding remarks regarding functions also apply to triggers with the following exception: `CREATE TRIGGER` does not have an optional `DETERMINISTIC` characteristic, so triggers are assumed to be always deterministic. However, this assumption might be invalid in some cases. For example, the `UUID()` function is nondeterministic (and does not replicate). Be careful about using such functions in triggers.

Triggers can update tables, so error messages similar to those for stored functions occur with `CREATE TRIGGER` if you do not have the required privileges. On the replica side, the replica uses the trigger `DEFINER` attribute to determine which user is considered to be the creator of the trigger.

The rest of this section provides additional detail about the logging implementation and its implications. You need not read it unless you are interested in the background on the rationale for the current logging-related conditions on stored routine use. This discussion applies only for statement-based logging, and not for row-based logging, with the exception of the first item: `CREATE` and `DROP` statements are logged as statements regardless of the logging mode.

- The server writes `CREATE EVENT`, `CREATE PROCEDURE`, `CREATE FUNCTION`, `ALTER EVENT`, `ALTER PROCEDURE`, `ALTER FUNCTION`, `DROP EVENT`, `DROP PROCEDURE`, and `DROP FUNCTION` statements to the binary log.
- A stored function invocation is logged as a `SELECT` statement if the function changes data and occurs within a statement that would not otherwise be logged. This prevents nonreplication of data changes that result from use of stored functions in nonlogged statements. For example, `SELECT` statements are not written to the binary log, but a `SELECT` might invoke a stored function that makes changes. To handle this, a `SELECT func_name()` statement is written to the binary log when the given function makes a change. Suppose that the following statements are executed on the source server:

```
CREATE FUNCTION f1(a INT) RETURNS INT
BEGIN
```

```

IF (a < 3) THEN
    INSERT INTO t2 VALUES (a);
END IF;
RETURN 0;
END;

CREATE TABLE t1 (a INT);
INSERT INTO t1 VALUES (1),(2),(3);

SELECT f1(a) FROM t1;

```

When the `SELECT` statement executes, the function `f1()` is invoked three times. Two of those invocations insert a row, and MySQL logs a `SELECT` statement for each of them. That is, MySQL writes the following statements to the binary log:

```

SELECT f1(1);
SELECT f1(2);

```

The server also logs a `SELECT` statement for a stored function invocation when the function invokes a stored procedure that causes an error. In this case, the server writes the `SELECT` statement to the log along with the expected error code. On the replica, if the same error occurs, that is the expected result and replication continues. Otherwise, replication stops.

- Logging stored function invocations rather than the statements executed by a function has a security implication for replication, which arises from two factors:
  - It is possible for a function to follow different execution paths on source and replica servers.
  - Statements executed on a replica are processed by the replica's applier thread. Unless you implement replication privilege checks, which are available from MySQL 8.0.18 (see [Section 17.3.3, “Replication Privilege Checks”](#)), the applier thread has full privileges.

The implication is that although a user must have the `CREATE ROUTINE` privilege to create a function, the user can write a function containing a dangerous statement that executes only on the replica where it is processed by a thread that has full privileges. For example, if the source and replica servers have server ID values of 1 and 2, respectively, a user on the source server could create and invoke an unsafe function `unsafe_func()` as follows:

```

mysql> delimiter //
mysql> CREATE FUNCTION unsafe_func () RETURNS INT
    -> BEGIN
    ->     IF @@server_id=2 THEN dangerous_statement; END IF;
    ->     RETURN 1;
    -> END;
    -> //
mysql> delimiter ;
mysql> INSERT INTO t VALUES(unsafe_func());

```

The `CREATE FUNCTION` and `INSERT` statements are written to the binary log, so the replica executes them. Because the replica's applier thread has full privileges, it executes the dangerous statement. Thus, the function invocation has different effects on the source and replica and is not replication-safe.

To guard against this danger for servers that have binary logging enabled, stored function creators must have the `SUPER` privilege, in addition to the usual `CREATE ROUTINE` privilege that is required. Similarly, to use `ALTER FUNCTION`, you must have the `SUPER` privilege in addition to the `ALTER ROUTINE` privilege. Without the `SUPER` privilege, an error occurs:

```

ERROR 1419 (HY000): You do not have the SUPER privilege and
binary logging is enabled (you *might* want to use the less safe
log_bin_trust_function_creators variable)

```

If you do not want to require function creators to have the `SUPER` privilege (for example, if all users with the `CREATE ROUTINE` privilege on your system are experienced application developers), set the global `log_bin_trust_function_creators` system variable to 1. You can also set this

variable at server startup. If binary logging is not enabled, `log_bin_trust_function_creators` does not apply. `SUPER` is not required for function creation unless, as described previously, the `DEFINER` value in the function definition requires it.

- The use of replication privilege checks where available (from MySQL 8.0.18) is recommended whatever choice you make about privileges for function creators. Replication privilege checks can be set up to ensure that only expected and relevant operations are authorized for the replication channel. For instructions to do this, see [Section 17.3.3, “Replication Privilege Checks”](#).
- If a function that performs updates is nondeterministic, it is not repeatable. This can have two undesirable effects:
  - It causes a replica to differ from the source.
  - Restored data does not match the original data.

To deal with these problems, MySQL enforces the following requirement: On a source server, creation and alteration of a function is refused unless you declare the function to be deterministic or to not modify data. Two sets of function characteristics apply here:

- The `DETERMINISTIC` and `NOT DETERMINISTIC` characteristics indicate whether a function always produces the same result for given inputs. The default is `NOT DETERMINISTIC` if neither characteristic is given. To declare that a function is deterministic, you must specify `DETERMINISTIC` explicitly.
- The `CONTAINS SQL`, `NO SQL`, `READS SQL DATA`, and `MODIFIES SQL DATA` characteristics provide information about whether the function reads or writes data. Either `NO SQL` or `READS SQL DATA` indicates that a function does not change data, but you must specify one of these explicitly because the default is `CONTAINS SQL` if no characteristic is given.

By default, for a `CREATE FUNCTION` statement to be accepted, at least one of `DETERMINISTIC`, `NO SQL`, or `READS SQL DATA` must be specified explicitly. Otherwise an error occurs:

```
ERROR 1418 (HY000): This function has none of DETERMINISTIC, NO SQL,
or READS SQL DATA in its declaration and binary logging is enabled
(you *might* want to use the less safe log_bin_trust_function_creators
variable)
```

If you set `log_bin_trust_function_creators` to 1, the requirement that functions be deterministic or not modify data is dropped.

- Stored procedure calls are logged at the statement level rather than at the `CALL` level. That is, the server does not log the `CALL` statement, it logs those statements within the procedure that actually execute. As a result, the same changes that occur on the source server also occur on replicas. This prevents problems that could result from a procedure having different execution paths on different machines.

In general, statements executed within a stored procedure are written to the binary log using the same rules that would apply were the statements to be executed in standalone fashion. Some special care is taken when logging procedure statements because statement execution within procedures is not quite the same as in nonprocedure context:

- A statement to be logged might contain references to local procedure variables. These variables do not exist outside of stored procedure context, so a statement that refers to such a variable cannot be logged literally. Instead, each reference to a local variable is replaced by this construct for logging purposes:

```
NAME_CONST(var_name, var_value)
```

`var_name` is the local variable name, and `var_value` is a constant indicating the value that the variable has at the time the statement is logged. `NAME_CONST()` has a value of `var_value`, and a “name” of `var_name`. Thus, if you invoke this function directly, you get a result like this:

```
mysql> SELECT NAME_CONST('myname', 14);
+-----+
| myname |
+-----+
|     14 |
+-----+
```

`NAME_CONST()` enables a logged standalone statement to be executed on a replica with the same effect as the original statement that was executed on the source within a stored procedure.

The use of `NAME_CONST()` can result in a problem for `CREATE TABLE ... SELECT` statements when the source column expressions refer to local variables. Converting these references to `NAME_CONST()` expressions can result in column names that are different on the source and replica servers, or names that are too long to be legal column identifiers. A workaround is to supply aliases for columns that refer to local variables. Consider this statement when `myvar` has a value of 1:

```
CREATE TABLE t1 SELECT myvar;
```

This is rewritten as follows:

```
CREATE TABLE t1 SELECT NAME_CONST(myvar, 1);
```

To ensure that the source and replica tables have the same column names, write the statement like this:

```
CREATE TABLE t1 SELECT myvar AS myvar;
```

The rewritten statement becomes:

```
CREATE TABLE t1 SELECT NAME_CONST(myvar, 1) AS myvar;
```

- A statement to be logged might contain references to user-defined variables. To handle this, MySQL writes a `SET` statement to the binary log to make sure that the variable exists on the replica with the same value as on the source. For example, if a statement refers to a variable `@my_var`, that statement is preceded in the binary log by the following statement, where `value` is the value of `@my_var` on the source:

```
SET @my_var = value;
```

- Procedure calls can occur within a committed or rolled-back transaction. Transactional context is accounted for so that the transactional aspects of procedure execution are replicated correctly. That is, the server logs those statements within the procedure that actually execute and modify data, and also logs `BEGIN`, `COMMIT`, and `ROLLBACK` statements as necessary. For example, if a procedure updates only transactional tables and is executed within a transaction that is rolled back, those updates are not logged. If the procedure occurs within a committed transaction, `BEGIN` and `COMMIT` statements are logged with the updates. For a procedure that executes within a rolled-back transaction, its statements are logged using the same rules that would apply if the statements were executed in standalone fashion:
  - Updates to transactional tables are not logged.
  - Updates to nontransactional tables are logged because rollback does not cancel them.
  - Updates to a mix of transactional and nontransactional tables are logged surrounded by `BEGIN` and `ROLLBACK` so that replicas make the same changes and rollbacks as on the source.
- A stored procedure call is *not* written to the binary log at the statement level if the procedure is invoked from within a stored function. In that case, the only thing logged is the statement that invokes the function (if it occurs within a statement that is logged) or a `DO` statement (if it occurs within a statement that is not logged). For this reason, care should be exercised in the use of stored functions that invoke a procedure, even if the procedure is otherwise safe in itself.

## 25.8 Restrictions on Stored Programs

- SQL Statements Not Permitted in Stored Routines
- Restrictions for Stored Functions
- Restrictions for Triggers
- Name Conflicts within Stored Routines
- Replication Considerations
- Debugging Considerations
- Unsupported Syntax from the SQL:2003 Standard
- Stored Routine Concurrency Considerations
- Event Scheduler Restrictions
- Stored routines and triggers in NDB Cluster

These restrictions apply to the features described in [Chapter 25, \*Stored Objects\*](#).

Some of the restrictions noted here apply to all stored routines; that is, both to stored procedures and stored functions. There are also some [restrictions specific to stored functions](#) but not to stored procedures.

The restrictions for stored functions also apply to triggers. There are also some [restrictions specific to triggers](#).

The restrictions for stored procedures also apply to the [DO](#) clause of Event Scheduler event definitions. There are also some [restrictions specific to events](#).

### SQL Statements Not Permitted in Stored Routines

Stored routines cannot contain arbitrary SQL statements. The following statements are not permitted:

- The locking statements [LOCK TABLES](#) and [UNLOCK TABLES](#).
- [ALTER VIEW](#).
- [LOAD DATA](#) and [LOAD XML](#).
- SQL prepared statements ([PREPARE](#), [EXECUTE](#), [DEALLOCATE PREPARE](#)) can be used in stored procedures, but not stored functions or triggers. Thus, stored functions and triggers cannot use dynamic SQL (where you construct statements as strings and then execute them).
- Generally, statements not permitted in SQL prepared statements are also not permitted in stored programs. For a list of statements supported as prepared statements, see [Section 13.5, “Prepared Statements”](#). Exceptions are [SIGNAL](#), [RESIGNAL](#), and [GET DIAGNOSTICS](#), which are not permissible as prepared statements but are permitted in stored programs.
- Because local variables are in scope only during stored program execution, references to them are not permitted in prepared statements created within a stored program. Prepared statement scope is the current session, not the stored program, so the statement could be executed after the program ends, at which point the variables would no longer be in scope. For example, [SELECT ... INTO local\\_var](#) cannot be used as a prepared statement. This restriction also applies to stored procedure and function parameters. See [Section 13.5.1, “PREPARE Statement”](#).
- Within all stored programs (stored procedures and functions, triggers, and events), the parser treats [BEGIN \[WORK\]](#) as the beginning of a [BEGIN ... END](#) block. To begin a transaction in this context, use [START TRANSACTION](#) instead.

## Restrictions for Stored Functions

The following additional statements or operations are not permitted within stored functions. They are permitted within stored procedures, except stored procedures that are invoked from within a stored function or trigger. For example, if you use `FLUSH` in a stored procedure, that stored procedure cannot be called from a stored function or trigger.

- Statements that perform explicit or implicit commit or rollback. Support for these statements is not required by the SQL standard, which states that each DBMS vendor may decide whether to permit them.
- Statements that return a result set. This includes `SELECT` statements that do not have an `INTO var_list` clause and other statements such as `SHOW`, `EXPLAIN`, and `CHECK TABLE`. A function can process a result set either with `SELECT ... INTO var_list` or by using a cursor and `FETCH` statements. See [Section 13.2.13.1, “SELECT ... INTO Statement”](#), and [Section 13.6.6, “Cursors”](#).
- `FLUSH` statements.
- Stored functions cannot be used recursively.
- A stored function or trigger cannot modify a table that is already being used (for reading or writing) by the statement that invoked the function or trigger.
- If you refer to a temporary table multiple times in a stored function under different aliases, a `Can't reopen table: 'tbl_name'` error occurs, even if the references occur in different statements within the function.
- `HANDLER ... READ` statements that invoke stored functions can cause replication errors and are disallowed.

## Restrictions for Triggers

For triggers, the following additional restrictions apply:

- Triggers are not activated by foreign key actions.
- When using row-based replication, triggers on the replica are not activated by statements originating on the source. The triggers on the replica are activated when using statement-based replication. For more information, see [Section 17.5.1.36, “Replication and Triggers”](#).
- The `RETURN` statement is not permitted in triggers, which cannot return a value. To exit a trigger immediately, use the `LEAVE` statement.
- Triggers are not permitted on tables in the `mysql` database. Nor are they permitted on `INFORMATION_SCHEMA` or `performance_schema` tables. Those tables are actually views and triggers are not permitted on views.
- The trigger cache does not detect when metadata of the underlying objects has changed. If a trigger uses a table and the table has changed since the trigger was loaded into the cache, the trigger operates using the outdated metadata.

## Name Conflicts within Stored Routines

The same identifier might be used for a routine parameter, a local variable, and a table column. Also, the same local variable name can be used in nested blocks. For example:

```
CREATE PROCEDURE p (i INT)
BEGIN
    DECLARE i INT DEFAULT 0;
    SELECT i FROM t;
    BEGIN
```

```
DECLARE i INT DEFAULT 1;
    SELECT i FROM t;
END;
END;
```

In such cases, the identifier is ambiguous and the following precedence rules apply:

- A local variable takes precedence over a routine parameter or table column.
- A routine parameter takes precedence over a table column.
- A local variable in an inner block takes precedence over a local variable in an outer block.

The behavior that variables take precedence over table columns is nonstandard.

## Replication Considerations

Use of stored routines can cause replication problems. This issue is discussed further in [Section 25.7, “Stored Program Binary Logging”](#).

The `--replicate-wild-do-table=db_name.tbl_name` option applies to tables, views, and triggers. It does not apply to stored procedures and functions, or events. To filter statements operating on the latter objects, use one or more of the `--replicate-*--db` options.

## Debugging Considerations

There are no stored routine debugging facilities.

## Unsupported Syntax from the SQL:2003 Standard

The MySQL stored routine syntax is based on the SQL:2003 standard. The following items from that standard are not currently supported:

- `UNDO` handlers
- `FOR` loops

## Stored Routine Concurrency Considerations

To prevent problems of interaction between sessions, when a client issues a statement, the server uses a snapshot of routines and triggers available for execution of the statement. That is, the server calculates a list of procedures, functions, and triggers that may be used during execution of the statement, loads them, and then proceeds to execute the statement. While the statement executes, it does not see changes to routines performed by other sessions.

For maximum concurrency, stored functions should minimize their side-effects; in particular, updating a table within a stored function can reduce concurrent operations on that table. A stored function acquires table locks before executing, to avoid inconsistency in the binary log due to mismatch of the order in which statements execute and when they appear in the log. When statement-based binary logging is used, statements that invoke a function are recorded rather than the statements executed within the function. Consequently, stored functions that update the same underlying tables do not execute in parallel. In contrast, stored procedures do not acquire table-level locks. All statements executed within stored procedures are written to the binary log, even for statement-based binary logging. See [Section 25.7, “Stored Program Binary Logging”](#).

## Event Scheduler Restrictions

The following limitations are specific to the Event Scheduler:

- Event names are handled in case-insensitive fashion. For example, you cannot have two events in the same database with the names `anEvent` and `AnEvent`.

- An event may not be created from within a stored program. An event may not be altered, or dropped from within a stored program, if the event name is specified by means of a variable. An event also may not create, alter, or drop stored routines or triggers.
- DDL statements on events are prohibited while a `LOCK TABLES` statement is in effect.
- Event timings using the intervals `YEAR`, `QUARTER`, `MONTH`, and `YEAR_MONTH` are resolved in months; those using any other interval are resolved in seconds. There is no way to cause events scheduled to occur at the same second to execute in a given order. In addition—due to rounding, the nature of threaded applications, and the fact that a nonzero length of time is required to create events and to signal their execution—events may be delayed by as much as 1 or 2 seconds. However, the time shown in the Information Schema `EVENTS` table's `LAST_EXECUTED` column is always accurate to within one second of the actual event execution time. (See also Bug #16522.)
- Each execution of the statements contained in the body of an event takes place in a new connection; thus, these statements have no effect in a given user session on the server's statement counts such as `Com_select` and `Com_insert` that are displayed by `SHOW STATUS`. However, such counts are updated in the global scope. (Bug #16422)
- Events do not support times later than the end of the Unix Epoch; this is approximately the beginning of the year 2038. Such dates are specifically not permitted by the Event Scheduler. (Bug #16396)
- References to stored functions, loadable functions, and tables in the `ON SCHEDULE` clauses of `CREATE EVENT` and `ALTER EVENT` statements are not supported. These sorts of references are not permitted. (See Bug #22830 for more information.)

## Stored routines and triggers in NDB Cluster

While stored procedures, stored functions, triggers, and scheduled events are all supported by tables using the `NDB` storage engine, you must keep in mind that these do *not* propagate automatically between MySQL Servers acting as Cluster SQL nodes. This is because stored routine and trigger definitions are stored in tables in the `mysql` system database using `InnoDB` tables, which are not copied between Cluster nodes.

Any stored routine or trigger that interacts with MySQL Cluster tables must be re-created by running the appropriate `CREATE PROCEDURE`, `CREATE FUNCTION`, or `CREATE TRIGGER` statements on each MySQL Server that participates in the cluster where you wish to use the stored routine or trigger. Similarly, any changes to existing stored routines or triggers must be carried out explicitly on all Cluster SQL nodes, using the appropriate `ALTER` or `DROP` statements on each MySQL Server accessing the cluster.



### Warning

Do *not* attempt to work around the issue just described by converting any `mysql` database tables to use the `NDB` storage engine. *Altering the system tables in the mysql database is not supported* and is very likely to produce undesirable results.

## 25.9 Restrictions on Views

The maximum number of tables that can be referenced in the definition of a view is 61.

View processing is not optimized:

- It is not possible to create an index on a view.
- Indexes can be used for views processed using the merge algorithm. However, a view that is processed with the temptable algorithm is unable to take advantage of indexes on its underlying tables (although indexes can be used during generation of the temporary tables).

There is a general principle that you cannot modify a table and select from the same table in a subquery. See [Section 13.2.15.12, “Restrictions on Subqueries”](#).

The same principle also applies if you select from a view that selects from the table, if the view selects from the table in a subquery and the view is evaluated using the merge algorithm. Example:

```
CREATE VIEW v1 AS
SELECT * FROM t2 WHERE EXISTS (SELECT 1 FROM t1 WHERE t1.a = t2.a);

UPDATE t1, v2 SET t1.a = 1 WHERE t1.b = v2.b;
```

If the view is evaluated using a temporary table, you *can* select from the table in the view subquery and still modify that table in the outer query. In this case, the view is stored in a temporary table and thus you are not really selecting from the table in a subquery and modifying it at the same time. (This is another reason you might wish to force MySQL to use the temptable algorithm by specifying `ALGORITHM = TEMPTABLE` in the view definition.)

You can use `DROP TABLE` or `ALTER TABLE` to drop or alter a table that is used in a view definition. No warning results from the `DROP` or `ALTER` operation, even though this invalidates the view. Instead, an error occurs later, when the view is used. `CHECK TABLE` can be used to check for views that have been invalidated by `DROP` or `ALTER` operations.

With regard to view updatability, the overall goal for views is that if any view is theoretically updatable, it should be updatable in practice. Many theoretically updatable views can be updated now, but limitations still exist. For details, see [Section 25.5.3, “Updatable and Insertable Views”](#).

There exists a shortcoming with the current implementation of views. If a user is granted the basic privileges necessary to create a view (the `CREATE VIEW` and `SELECT` privileges), that user cannot call `SHOW CREATE VIEW` on that object unless the user is also granted the `SHOW VIEW` privilege.

That shortcoming can lead to problems backing up a database with `mysqldump`, which may fail due to insufficient privileges. This problem is described in Bug #22062.

The workaround to the problem is for the administrator to manually grant the `SHOW VIEW` privilege to users who are granted `CREATE VIEW`, since MySQL doesn't grant it implicitly when views are created.

Views do not have indexes, so index hints do not apply. Use of index hints when selecting from a view is not permitted.

`SHOW CREATE VIEW` displays view definitions using an `AS alias_name` clause for each column. If a column is created from an expression, the default alias is the expression text, which can be quite long. Aliases for column names in `CREATE VIEW` statements are checked against the maximum column length of 64 characters (not the maximum alias length of 256 characters). As a result, views created from the output of `SHOW CREATE VIEW` fail if any column alias exceeds 64 characters. This can cause problems in the following circumstances for views with too-long aliases:

- View definitions fail to replicate to newer replicas that enforce the column-length restriction.
- Dump files created with `mysqldump` cannot be loaded into servers that enforce the column-length restriction.

A workaround for either problem is to modify each problematic view definition to use aliases that provide shorter column names. Then the view replicates properly, and can be dumped and reloaded without causing an error. To modify the definition, drop and create the view again with `DROP VIEW` and `CREATE VIEW`, or replace the definition with `CREATE OR REPLACE VIEW`.

For problems that occur when reloading view definitions in dump files, another workaround is to edit the dump file to modify its `CREATE VIEW` statements. However, this does not change the original view definitions, which may cause problems for subsequent dump operations.



---

# Chapter 26 INFORMATION\_SCHEMA Tables

## Table of Contents

26.1 Introduction .....	4838
26.2 INFORMATION_SCHEMA Table Reference .....	4841
26.3 INFORMATION_SCHEMA General Tables .....	4846
26.3.1 INFORMATION_SCHEMA General Table Reference .....	4846
26.3.2 The INFORMATION_SCHEMA ADMINISTRABLE_ROLE_AUTHORIZATIONS Table	4848
26.3.3 The INFORMATION_SCHEMA APPLICABLE_ROLES Table .....	4848
26.3.4 The INFORMATION_SCHEMA CHARACTER_SETS Table .....	4849
26.3.5 The INFORMATION_SCHEMA CHECK_CONSTRAINTS Table .....	4849
26.3.6 The INFORMATION_SCHEMA COLLATIONS Table .....	4850
26.3.7 The INFORMATION_SCHEMA COLLATION_CHARACTER_SET_APPLICABILITY Table .....	4851
26.3.8 The INFORMATION_SCHEMA COLUMNS Table .....	4851
26.3.9 The INFORMATION_SCHEMA COLUMNS_EXTENSIONS Table .....	4854
26.3.10 The INFORMATION_SCHEMA COLUMN_PRIVILEGES Table .....	4854
26.3.11 The INFORMATION_SCHEMA COLUMN_STATISTICS Table .....	4855
26.3.12 The INFORMATION_SCHEMA ENABLED_ROLES Table .....	4855
26.3.13 The INFORMATION_SCHEMA ENGINES Table .....	4856
26.3.14 The INFORMATION_SCHEMA EVENTS Table .....	4857
26.3.15 The INFORMATION_SCHEMA FILES Table .....	4860
26.3.16 The INFORMATION_SCHEMA KEY_COLUMN_USAGE Table .....	4868
26.3.17 The INFORMATION_SCHEMA KEYWORDS Table .....	4869
26.3.18 The INFORMATION_SCHEMA ndb_transid_mysql_connection_map Table .....	4869
26.3.19 The INFORMATION_SCHEMA OPTIMIZER_TRACE Table .....	4871
26.3.20 The INFORMATION_SCHEMA PARAMETERS Table .....	4871
26.3.21 The INFORMATION_SCHEMA PARTITIONS Table .....	4872
26.3.22 The INFORMATION_SCHEMA PLUGINS Table .....	4875
26.3.23 The INFORMATION_SCHEMA PROCESSLIST Table .....	4877
26.3.24 The INFORMATION_SCHEMA PROFILING Table .....	4878
26.3.25 The INFORMATION_SCHEMA REFERENTIAL_CONSTRAINTS Table .....	4879
26.3.26 The INFORMATION_SCHEMA RESOURCE_GROUPS Table .....	4880
26.3.27 The INFORMATION_SCHEMA ROLE_COLUMN_GRANTS Table .....	4880
26.3.28 The INFORMATION_SCHEMA ROLE_ROUTINE_GRANTS Table .....	4881
26.3.29 The INFORMATION_SCHEMA ROLE_TABLE_GRANTS Table .....	4882
26.3.30 The INFORMATION_SCHEMA ROUTINES Table .....	4883
26.3.31 The INFORMATION_SCHEMA SCHEMATA Table .....	4885
26.3.32 The INFORMATION_SCHEMA SCHEMATA_EXTENSIONS Table .....	4886
26.3.33 The INFORMATION_SCHEMA SCHEMA_PRIVILEGES Table .....	4887
26.3.34 The INFORMATION_SCHEMA STATISTICS Table .....	4887
26.3.35 The INFORMATION_SCHEMA ST_GeOMETRY_COLUMNS Table .....	4890
26.3.36 The INFORMATION_SCHEMA ST_SPATIAL_REFERENCE_SYSTEMS Table .....	4890
26.3.37 The INFORMATION_SCHEMA ST_UNITS_OF_MEASURE Table .....	4892
26.3.38 The INFORMATION_SCHEMA TABLES Table .....	4892
26.3.39 The INFORMATION_SCHEMA TABLES_EXTENSIONS Table .....	4896
26.3.40 The INFORMATION_SCHEMA TABLESPACES Table .....	4896
26.3.41 The INFORMATION_SCHEMA TABLESPACES_EXTENSIONS Table .....	4897
26.3.42 The INFORMATION_SCHEMA TABLE_CONSTRAINTS Table .....	4897
26.3.43 The INFORMATION_SCHEMA TABLE_CONSTRAINTS_EXTENSIONS Table .....	4898
26.3.44 The INFORMATION_SCHEMA TABLE_PRIVILEGES Table .....	4898
26.3.45 The INFORMATION_SCHEMA TRIGGERS Table .....	4899
26.3.46 The INFORMATION_SCHEMA USER_ATTRIBUTES Table .....	4901
26.3.47 The INFORMATION_SCHEMA USER_PRIVILEGES Table .....	4902
26.3.48 The INFORMATION_SCHEMA VIEWS Table .....	4902

26.3.49 The INFORMATION_SCHEMA VIEW_ROUTINE_USAGE Table .....	4904
26.3.50 The INFORMATION_SCHEMA VIEW_TABLE_USAGE Table .....	4904
26.4 INFORMATION_SCHEMA InnoDB Tables .....	4905
26.4.1 INFORMATION_SCHEMA InnoDB Table Reference .....	4905
26.4.2 The INFORMATION_SCHEMA INNODB_BUFFER_PAGE Table .....	4906
26.4.3 The INFORMATION_SCHEMA INNODB_BUFFER_PAGE_LRU Table .....	4910
26.4.4 The INFORMATION_SCHEMA INNODB_BUFFER_POOL_STATS Table .....	4913
26.4.5 The INFORMATION_SCHEMA INNODB_CACHED_INDEXES Table .....	4916
26.4.6 The INFORMATION_SCHEMA INNODB_CMP and INNODB_CMP_RESET Tables ..	4917
26.4.7 The INFORMATION_SCHEMA INNODB_CMPMEM and INNODB_CMPMEM_RESET Tables .....	4918
26.4.8 The INFORMATION_SCHEMA INNODB_CMP_PER_INDEX and INNODB_CMP_PER_INDEX_RESET Tables .....	4920
26.4.9 The INFORMATION_SCHEMA INNODB_COLUMNS Table .....	4921
26.4.10 The INFORMATION_SCHEMA INNODB_DATAFILES Table .....	4923
26.4.11 The INFORMATION_SCHEMA INNODB_FIELDS Table .....	4923
26.4.12 The INFORMATION_SCHEMA INNODB_FOREIGN Table .....	4924
26.4.13 The INFORMATION_SCHEMA INNODB_FOREIGN_COLS Table .....	4925
26.4.14 The INFORMATION_SCHEMA INNODB_FT_BEING_DELETED Table .....	4925
26.4.15 The INFORMATION_SCHEMA INNODB_FT_CONFIG Table .....	4926
26.4.16 The INFORMATION_SCHEMA INNODB_FT_DEFAULT_STOPWORD Table .....	4927
26.4.17 The INFORMATION_SCHEMA INNODB_FT_DELETED Table .....	4928
26.4.18 The INFORMATION_SCHEMA INNODB_FT_INDEX_CACHE Table .....	4929
26.4.19 The INFORMATION_SCHEMA INNODB_FT_INDEX_TABLE Table .....	4930
26.4.20 The INFORMATION_SCHEMA INNODB_INDEXES Table .....	4932
26.4.21 The INFORMATION_SCHEMA INNODB_METRICS Table .....	4933
26.4.22 The INFORMATION_SCHEMA INNODB_SESSION_TEMP_TABLESPACES Table	4935
26.4.23 The INFORMATION_SCHEMA INNODB_TABLES Table .....	4936
26.4.24 The INFORMATION_SCHEMA INNODB_TABLESPACES Table .....	4937
26.4.25 The INFORMATION_SCHEMA INNODB_TABLESPACES_BRIEF Table .....	4940
26.4.26 The INFORMATION_SCHEMA INNODB_TABLESTATS View .....	4940
26.4.27 The INFORMATION_SCHEMA INNODB_TEMP_TABLE_INFO Table .....	4942
26.4.28 The INFORMATION_SCHEMA INNODB_TRX Table .....	4943
26.4.29 The INFORMATION_SCHEMA INNODB_VIRTUAL Table .....	4945
26.5 INFORMATION_SCHEMA Thread Pool Tables .....	4947
26.5.1 INFORMATION_SCHEMA Thread Pool Table Reference .....	4947
26.5.2 The INFORMATION_SCHEMA TP_THREAD_GROUP_STATE Table .....	4947
26.5.3 The INFORMATION_SCHEMA TP_THREAD_GROUP_STATS Table .....	4948
26.5.4 The INFORMATION_SCHEMA TP_THREAD_STATE Table .....	4948
26.6 INFORMATION_SCHEMA Connection-Control Tables .....	4949
26.6.1 INFORMATION_SCHEMA Connection-Control Table Reference .....	4949
26.6.2 The INFORMATION_SCHEMA CONNECTION_FAILED_LOGIN_ATTEMPTS Table .....	4949
26.7 INFORMATION_SCHEMA MySQL Enterprise Firewall Tables .....	4949
26.7.1 INFORMATION_SCHEMA Firewall Table Reference .....	4949
26.7.2 The INFORMATION_SCHEMA MYSQL_FIREWALL_USERS Table .....	4950
26.7.3 The INFORMATION_SCHEMA MYSQL_FIREWALL_WHITELIST Table .....	4950
26.8 Extensions to SHOW Statements .....	4950

[INFORMATION\\_SCHEMA](#) provides access to database *metadata*, information about the MySQL server such as the name of a database or table, the data type of a column, or access privileges. Other terms that are sometimes used for this information are *data dictionary* and *system catalog*.

## 26.1 Introduction

[INFORMATION\\_SCHEMA](#) provides access to database *metadata*, information about the MySQL server such as the name of a database or table, the data type of a column, or access privileges. Other terms that are sometimes used for this information are *data dictionary* and *system catalog*.

- [INFORMATION\\_SCHEMA Usage Notes](#)
- [Character Set Considerations](#)
- [INFORMATION\\_SCHEMA as Alternative to SHOW Statements](#)
- [INFORMATION\\_SCHEMA and Privileges](#)
- [Performance Considerations](#)
- [Standards Considerations](#)
- [Conventions in the INFORMATION\\_SCHEMA Reference Sections](#)
- [Related Information](#)

## INFORMATION\_SCHEMA Usage Notes

[INFORMATION\\_SCHEMA](#) is a database within each MySQL instance, the place that stores information about all the other databases that the MySQL server maintains. The [INFORMATION\\_SCHEMA](#) database contains several read-only tables. They are actually views, not base tables, so there are no files associated with them, and you cannot set triggers on them. Also, there is no database directory with that name.

Although you can select [INFORMATION\\_SCHEMA](#) as the default database with a [USE](#) statement, you can only read the contents of tables, not perform [INSERT](#), [UPDATE](#), or [DELETE](#) operations on them.

Here is an example of a statement that retrieves information from [INFORMATION\\_SCHEMA](#):

```
mysql> SELECT table_name, table_type, engine
      FROM information_schema.tables
     WHERE table_schema = 'db5'
       ORDER BY table_name;
+-----+-----+-----+
| table_name | table_type | engine |
+-----+-----+-----+
| fk          | BASE TABLE | InnoDB |
| fk2         | BASE TABLE | InnoDB |
| goto        | BASE TABLE | MyISAM |
| into        | BASE TABLE | MyISAM |
| k           | BASE TABLE | MyISAM |
| kurs        | BASE TABLE | MyISAM |
| loop        | BASE TABLE | MyISAM |
| pk          | BASE TABLE | InnoDB |
| t           | BASE TABLE | MyISAM |
| t2          | BASE TABLE | MyISAM |
| t3          | BASE TABLE | MyISAM |
| t7          | BASE TABLE | MyISAM |
| tables      | BASE TABLE | MyISAM |
| v           | VIEW        | NULL   |
| v2          | VIEW        | NULL   |
| v3          | VIEW        | NULL   |
| v56         | VIEW        | NULL   |
+-----+-----+-----+
17 rows in set (0.01 sec)
```

Explanation: The statement requests a list of all the tables in database `db5`, showing just three pieces of information: the name of the table, its type, and its storage engine.

Beginning with MySQL 8.0.30, information about generated invisible primary keys is visible by default in all [INFORMATION\\_SCHEMA](#) tables describing table columns, keys, or both, such as the [COLUMNS](#) and [STATISTICS](#) tables. If you wish to make such information hidden from queries that select from these tables, you can do so by setting the value of the [show\\_gipk\\_in\\_create\\_table\\_and\\_information\\_schema](#) server system variable to [OFF](#). For more information, see [Section 13.1.20.11, “Generated Invisible Primary Keys”](#).

## Character Set Considerations

The definition for character columns (for example, `TABLES . TABLE_NAME`) is generally `VARCHAR(N) CHARACTER SET utf8mb3` where `N` is at least 64. MySQL uses the default collation for this character set (`utf8mb3_general_ci`) for all searches, sorts, comparisons, and other string operations on such columns.

Because some MySQL objects are represented as files, searches in `INFORMATION_SCHEMA` string columns can be affected by file system case sensitivity. For more information, see [Section 10.8.7, “Using Collation in INFORMATION\\_SCHEMA Searches”](#).

## INFORMATION\_SCHEMA as Alternative to SHOW Statements

The `SELECT ... FROM INFORMATION_SCHEMA` statement is intended as a more consistent way to provide access to the information provided by the various `SHOW` statements that MySQL supports (`SHOW DATABASES`, `SHOW TABLES`, and so forth). Using `SELECT` has these advantages, compared to `SHOW`:

- It conforms to Codd's rules, because all access is done on tables.
- You can use the familiar syntax of the `SELECT` statement, and only need to learn some table and column names.
- The implementor need not worry about adding keywords.
- You can filter, sort, concatenate, and transform the results from `INFORMATION_SCHEMA` queries into whatever format your application needs, such as a data structure or a text representation to parse.
- This technique is more interoperable with other database systems. For example, Oracle Database users are familiar with querying tables in the Oracle data dictionary.

Because `SHOW` is familiar and widely used, the `SHOW` statements remain as an alternative. In fact, along with the implementation of `INFORMATION_SCHEMA`, there are enhancements to `SHOW` as described in [Section 26.8, “Extensions to SHOW Statements”](#).

## INFORMATION\_SCHEMA and Privileges

For most `INFORMATION_SCHEMA` tables, each MySQL user has the right to access them, but can see only the rows in the tables that correspond to objects for which the user has the proper access privileges. In some cases (for example, the `ROUTINE_DEFINITION` column in the `INFORMATION_SCHEMA ROUTINES` table), users who have insufficient privileges see `NULL`. Some tables have different privilege requirements; for these, the requirements are mentioned in the applicable table descriptions. For example, `InnoDB` tables (tables with names that begin with `INNODB_`) require the `PROCESS` privilege.

The same privileges apply to selecting information from `INFORMATION_SCHEMA` and viewing the same information through `SHOW` statements. In either case, you must have some privilege on an object to see information about it.

## Performance Considerations

`INFORMATION_SCHEMA` queries that search for information from more than one database might take a long time and impact performance. To check the efficiency of a query, you can use `EXPLAIN`. For information about using `EXPLAIN` output to tune `INFORMATION_SCHEMA` queries, see [Section 8.2.3, “Optimizing INFORMATION\\_SCHEMA Queries”](#).

## Standards Considerations

The implementation for the `INFORMATION_SCHEMA` table structures in MySQL follows the ANSI/ISO SQL:2003 standard Part 11 *Schemata*. Our intent is approximate compliance with SQL:2003 core feature F021 *Basic information schema*.

Users of SQL Server 2000 (which also follows the standard) may notice a strong similarity. However, MySQL has omitted many columns that are not relevant for our implementation, and added columns that are MySQL-specific. One such added column is the `ENGINE` column in the `INFORMATION_SCHEMA TABLES` table.

Although other DBMSs use a variety of names, like `syscat` or `system`, the standard name is `INFORMATION_SCHEMA`.

To avoid using any name that is reserved in the standard or in DB2, SQL Server, or Oracle, we changed the names of some columns marked “MySQL extension”. (For example, we changed `COLLATION` to `TABLE_COLLATION` in the `TABLES` table.) See the list of reserved words near the end of this article: <https://web.archive.org/web/20070428032454/http://www.dbazine.com/db2/db2-disarticles/gulutzan5>.

## Conventions in the INFORMATION\_SCHEMA Reference Sections

The following sections describe each of the tables and columns in `INFORMATION_SCHEMA`. For each column, there are three pieces of information:

- “`INFORMATION_SCHEMA Name`” indicates the name for the column in the `INFORMATION_SCHEMA` table. This corresponds to the standard SQL name unless the “Remarks” field says “MySQL extension.”
- “`SHOW Name`” indicates the equivalent field name in the closest `SHOW` statement, if there is one.
- “Remarks” provides additional information where applicable. If this field is `NULL`, it means that the value of the column is always `NULL`. If this field says “MySQL extension,” the column is a MySQL extension to standard SQL.

Many sections indicate what `SHOW` statement is equivalent to a `SELECT` that retrieves information from `INFORMATION_SCHEMA`. For `SHOW` statements that display information for the default database if you omit a `FROM db_name` clause, you can often select information for the default database by adding an `AND TABLE_SCHEMA = SCHEMA()` condition to the `WHERE` clause of a query that retrieves information from an `INFORMATION_SCHEMA` table.

## Related Information

These sections discuss additional `INFORMATION_SCHEMA`-related topics:

- information about `INFORMATION_SCHEMA` tables specific to the `InnoDB` storage engine: [Section 26.4, “INFORMATION\\_SCHEMA InnoDB Tables”](#)
- information about `INFORMATION_SCHEMA` tables specific to the thread pool plugin: [Section 26.5, “INFORMATION\\_SCHEMA Thread Pool Tables”](#)
- information about `INFORMATION_SCHEMA` tables specific to the `CONNECTION_CONTROL` plugin: [Section 26.6, “INFORMATION\\_SCHEMA Connection-Control Tables”](#)
- Answers to questions that are often asked concerning the `INFORMATION_SCHEMA` database: [Section A.7, “MySQL 8.0 FAQ: INFORMATION\\_SCHEMA”](#)
- `INFORMATION_SCHEMA` queries and the optimizer: [Section 8.2.3, “Optimizing INFORMATION\\_SCHEMA Queries”](#)
- The effect of collation on `INFORMATION_SCHEMA` comparisons: [Section 10.8.7, “Using Collation in INFORMATION\\_SCHEMA Searches”](#)

## 26.2 INFORMATION\_SCHEMA Table Reference

The following table summarizes all available `INFORMATION_SCHEMA` tables. For greater detail, see the individual table descriptions.

**Table 26.1 INFORMATION\_SCHEMA Tables**

Table Name	Description	Introduced	Deprecated
<code>ADMINISTRABLE_ROLE</code>	Grantable user or roles for current user or role	8.0.19	
<code>APPLICABLE_ROLES</code>	Applicable roles for current user	8.0.19	
<code>CHARACTER_SETS</code>	Available character sets		
<code>CHECK_CONSTRAINTS</code>	Table and column CHECK constraints	8.0.16	
<code>COLLATION_CHARACTER_SET_APPLICABILITY</code>	Character set applicable to each collation		
<code>COLLATIONS</code>	Collations for each character set		
<code>COLUMN_PRIVILEGES</code>	Privileges defined on columns		
<code>COLUMN_STATISTICS</code>	Histogram statistics for column values		
<code>COLUMNS</code>	Columns in each table		
<code>COLUMNS_EXTENSIONS</code>	Column attributes for primary and secondary storage engines	8.0.21	
<code>CONNECTION_CONTROL</code>	Current number of consecutive failed connection attempts per account		
<code>ENABLED_ROLES</code>	Roles enabled within current session	8.0.19	
<code>ENGINES</code>	Storage engine properties		
<code>EVENTS</code>	Event Manager events		
<code>FILES</code>	Files that store tablespace data		
<code>INNODB_BUFFER_PAGE</code>	Pages in InnoDB buffer pool		
<code>INNODB_BUFFER_PAGE_LIST</code>	LRU ordering of pages in InnoDB buffer pool		
<code>INNODB_BUFFER_POOL_STATS</code>	InnoDB buffer pool statistics		
<code>INNODB_CACHED_INDEXES</code>	Number of index pages cached per index in InnoDB buffer pool		
<code>INNODB_CMP</code>	Status for operations related to compressed InnoDB tables		
<code>INNODB_CMP_PER_INDEX_INFO</code>	Status for operations related to compressed InnoDB tables and indexes		

Table Name	Description	Introduced	Deprecated
<code>INNODB_CMP_PER_INDEX</code>	Status for operations related to compressed InnoDB tables and indexes		
<code>INNODB_CMP_RESET</code>	Status for operations related to compressed InnoDB tables		
<code>INNODB_CMPMEM</code>	Status for compressed pages within InnoDB buffer pool		
<code>INNODB_CMPMEM_RESET</code>	Status for compressed pages within InnoDB buffer pool		
<code>INNODB_COLUMNS</code>	Columns in each InnoDB table		
<code>INNODB_DATAFILES</code>	Data file path information for InnoDB file-per-table and general tablespaces		
<code>INNODB_FIELDS</code>	Key columns of InnoDB indexes		
<code>INNODB_FOREIGN</code>	InnoDB foreign-key metadata		
<code>INNODB_FOREIGN_COLS</code>	InnoDB foreign-key column status information		
<code>INNODB_FT_BEING_DELETED</code>	Snapshot of INNODB_FT_DELETED table		
<code>INNODB_FT_CONFIG</code>	Metadata for InnoDB table FULLTEXT index and associated processing		
<code>INNODB_FT_DEFAULT_STOPWORD</code>	Default list of stopwords for InnoDB FULLTEXT indexes		
<code>INNODB_FT_DELETED</code>	Rows deleted from InnoDB table FULLTEXT index		
<code>INNODB_FT_INDEX_CACHE</code>	Token information for newly inserted rows in InnoDB FULLTEXT index		
<code>INNODB_FT_INDEX_TABLE</code>	Inverted index information for processing text searches against InnoDB table FULLTEXT index		
<code>INNODB_INDEXES</code>	InnoDB index metadata		

## INFORMATION\_SCHEMA Table Reference

Table Name	Description	Introduced	Deprecated
INNODB_METRICS	InnoDB performance information		
INNODB_SESSION_TEMPSPACE	Session temporary-tablespace metadata	8.0.13	
INNODB_TABLES	InnoDB table metadata		
INNODB_TABLESPACES	InnoDB file-per-table, general, and undo tablespace metadata		
INNODB_TABLESPACES_BRIEF	Brief file-per-table, general, undo, and system tablespace metadata		
INNODB_TABLESTATS	InnoDB table low-level status information		
INNODB_TEMP_TABLE_INFO	Information about active user-created InnoDB temporary tables		
INNODB_TRX	Active InnoDB transaction information		
INNODB_VIRTUAL	InnoDB virtual generated column metadata		
KEY_COLUMN_USAGE	Which key columns have constraints		
KEYWORDS	MySQL keywords		
MYSQL_FIREWALL_USER	Firewall in-memory data for account profiles		8.0.26
MYSQL_FIREWALL_WHITELIST	Firewall in-memory data for account profile allowlists		8.0.26
ndb_transid_mysql_cdc	NDB transaction log information		
OPTIMIZER_TRACE	Information produced by optimizer trace activity		
PARAMETERS	Stored routine parameters and stored function return values		
PARTITIONS	Table partition information		
PLUGINS	Plugin information		
PROCESSLIST	Information about currently executing threads		
PROFILING	Statement profiling information		
REFERENTIAL_CONSTRAINTS	Foreign key information		
RESOURCE_GROUPS	Resource group information		

Table Name	Description	Introduced	Deprecated
ROLE_COLUMN_GRANTS	Column privileges for roles available to or granted by currently enabled roles	8.0.19	
ROLE_ROUTINE_GRANTS	Routine privileges for roles available to or granted by currently enabled roles	8.0.19	
ROLE_TABLE_GRANTS	Table privileges for roles available to or granted by currently enabled roles	8.0.19	
ROUTINES	Stored routine information		
SCHEMA_PRIVILEGES	Privileges defined on schemas		
SCHEMATA	Schema information		
SCHEMATA_EXTENSIONS	Schema options	8.0.22	
ST_Geometry_Columns	Columns in each table that store spatial data		
ST_Spatial_Reference_Systems	Available spatial reference systems		
ST_Units_of_Measure	Acceptable units for ST_Distance()	8.0.14	
STATISTICS	Table index statistics		
TABLE_CONSTRAINTS	Which tables have constraints		
TABLE_CONSTRAINTS_EXTENSIONS	Table constraint attributes for primary and secondary storage engines	8.0.21	
TABLE_PRIVILEGES	Privileges defined on tables		
TABLES	Table information		
TABLES_EXTENSIONS	Table attributes for primary and secondary storage engines	8.0.21	
TABLESPACES	Tablespace information		
TABLESPACES_EXTENSIONS	Tablespace attributes for primary storage engines	8.0.21	
TP_THREAD_GROUP_STATUS	Thread pool thread group states		
TP_THREAD_GROUP_STATISTICS	Thread pool thread group statistics		
TP_THREAD_STATE	Thread pool thread information		
TRIGGERS	Trigger information		

Table Name	Description	Introduced	Deprecated
USER_ATTRIBUTES	User comments and attributes	8.0.21	
USER_PRIVILEGES	Privileges defined globally per user		
VIEW_ROUTINE_USAGE	Stored functions used in views	8.0.13	
VIEW_TABLE_USAGE	Tables and views used in views	8.0.13	
VIEWS	View information		

## 26.3 INFORMATION\_SCHEMA General Tables

The following sections describe what may be denoted as the “general” set of [INFORMATION\\_SCHEMA](#) tables. These are the tables not associated with particular storage engines, components, or plugins.

### 26.3.1 INFORMATION\_SCHEMA General Table Reference

The following table summarizes [INFORMATION\\_SCHEMA](#) general tables. For greater detail, see the individual table descriptions.

**Table 26.2 INFORMATION\_SCHEMA General Tables**

Table Name	Description	Introduced
ADMINISTRABLE_ROLE_AUTHORITIES	Grantable users or roles for current user or role	8.0.19
APPLICABLE_ROLES	Applicable roles for current user	8.0.19
CHARACTER_SETS	Available character sets	
CHECK_CONSTRAINTS	Table and column CHECK constraints	8.0.16
COLLATION_CHARACTER_SET_APPLICABILITY	Character set applicable to each collation	
COLLATIONS	Collations for each character set	
COLUMN_PRIVILEGES	Privileges defined on columns	
COLUMN_STATISTICS	Histogram statistics for column values	
COLUMNS	Columns in each table	
COLUMNS_EXTENSIONS	Column attributes for primary and secondary storage engines	8.0.21
ENABLED_ROLES	Roles enabled within current session	8.0.19
ENGINES	Storage engine properties	
EVENTS	Event Manager events	
FILES	Files that store tablespace data	
KEY_COLUMN_USAGE	Which key columns have constraints	
KEYWORDS	MySQL keywords	
ndb_transid_mysql_connect	NDB transaction information	
OPTIMIZER_TRACE	Information produced by optimizer trace activity	

Table Name	Description	Introduced
PARAMETERS	Stored routine parameters and stored function return values	
PARTITIONS	Table partition information	
PLUGINS	Plugin information	
PROCESSLIST	Information about currently executing threads	
PROFILING	Statement profiling information	
REFERENTIAL_CONSTRAINTS	Foreign key information	
RESOURCE_GROUPS	Resource group information	
ROLE_COLUMN_GRANTS	Column privileges for roles available to or granted by currently enabled roles	8.0.19
ROLE_ROUTINE_GRANTS	Routine privileges for roles available to or granted by currently enabled roles	8.0.19
ROLE_TABLE_GRANTS	Table privileges for roles available to or granted by currently enabled roles	8.0.19
ROUTINES	Stored routine information	
SCHEMA_PRIVILEGES	Privileges defined on schemas	
SCHEMATA	Schema information	
SCHEMATA_EXTENSIONS	Schema options	8.0.22
ST_GEOMETRY_COLUMNS	Columns in each table that store spatial data	
ST_SPATIAL_REFERENCE_SYSTEMS	Available spatial reference systems	
ST_UNITS_OF_MEASURE	Acceptable units for ST_Distance()	8.0.14
STATISTICS	Table index statistics	
TABLE_CONSTRAINTS	Which tables have constraints	
TABLE_CONSTRAINTS_EXTENSIONS	Table constraint attributes for primary and secondary storage engines	8.0.21
TABLE_PRIVILEGES	Privileges defined on tables	
TABLES	Table information	
TABLES_EXTENSIONS	Table attributes for primary and secondary storage engines	8.0.21
TABLESPACES	Tablespace information	
TABLESPACES_EXTENSIONS	Tablespace attributes for primary storage engines	8.0.21
TRIGGERS	Trigger information	
USER_ATTRIBUTES	User comments and attributes	8.0.21
USER_PRIVILEGES	Privileges defined globally per user	
VIEW_ROUTINE_USAGE	Stored functions used in views	8.0.13

Table Name	Description	Introduced
VIEW_TABLE_USAGE	Tables and views used in views	8.0.13
VIEWS	View information	

## 26.3.2 The INFORMATION\_SCHEMA ADMINISTRABLE\_ROLE\_AUTHORIZATIONS Table

The [ADMINISTRABLE\\_ROLE\\_AUTHORIZATIONS](#) table (available as of MySQL 8.0.19) provides information about which roles applicable for the current user or role can be granted to other users or roles.

The [ADMINISTRABLE\\_ROLE\\_AUTHORIZATIONS](#) table has these columns:

- [USER](#)

The user name part of the current user account.

- [HOST](#)

The host name part of the current user account.

- [GRANTEE](#)

The user name part of the account to which the role is granted.

- [GRANTEE\\_HOST](#)

The host name part of the account to which the role is granted.

- [ROLE\\_NAME](#)

The user name part of the granted role.

- [ROLE\\_HOST](#)

The host name part of the granted role.

- [IS\\_GRANTABLE](#)

[YES](#) or [NO](#), depending on whether the role is grantable to other accounts.

- [IS\\_DEFAULT](#)

[YES](#) or [NO](#), depending on whether the role is a default role.

- [IS\\_MANDATORY](#)

[YES](#) or [NO](#), depending on whether the role is mandatory.

## 26.3.3 The INFORMATION\_SCHEMA APPLICABLE\_ROLES Table

The [APPLICABLE\\_ROLES](#) table (available as of MySQL 8.0.19) provides information about the roles that are applicable for the current user.

The [APPLICABLE\\_ROLES](#) table has these columns:

- [USER](#)

The user name part of the current user account.

- [HOST](#)

The host name part of the current user account.

- `GRANTEE`

The user name part of the account to which the role is granted.

- `GRANTEE_HOST`

The host name part of the account to which the role is granted.

- `ROLE_NAME`

The user name part of the granted role.

- `ROLE_HOST`

The host name part of the granted role.

- `IS_GRANTABLE`

`YES` or `NO`, depending on whether the role is grantable to other accounts.

- `IS_DEFAULT`

`YES` or `NO`, depending on whether the role is a default role.

- `IS_MANDATORY`

`YES` or `NO`, depending on whether the role is mandatory.

#### 26.3.4 The INFORMATION\_SCHEMA CHARACTER\_SETS Table

The `CHARACTER_SETS` table provides information about available character sets.

The `CHARACTER_SETS` table has these columns:

- `CHARACTER_SET_NAME`

The character set name.

- `DEFAULT_COLLATE_NAME`

The default collation for the character set.

- `DESCRIPTION`

A description of the character set.

- `MAXLEN`

The maximum number of bytes required to store one character.

#### Notes

Character set information is also available from the `SHOW CHARACTER SET` statement. See Section 13.7.7.3, “[SHOW CHARACTER SET Statement](#)”. The following statements are equivalent:

```
SELECT * FROM INFORMATION_SCHEMA.CHARACTER_SETS
[WHERE CHARACTER_SET_NAME LIKE 'wild']

SHOW CHARACTER SET
[LIKE 'wild']
```

#### 26.3.5 The INFORMATION\_SCHEMA CHECK\_CONSTRAINTS Table

As of MySQL 8.0.16, `CREATE TABLE` permits the core features of table and column `CHECK` constraints, and the `CHECK_CONSTRAINTS` table provides information about these constraints.

The `CHECK_CONSTRAINTS` table has these columns:

- `CONSTRAINT_CATALOG`

The name of the catalog to which the constraint belongs. This value is always `def`.

- `CONSTRAINT_SCHEMA`

The name of the schema (database) to which the constraint belongs.

- `CONSTRAINT_NAME`

The name of the constraint.

- `CHECK_CLAUSE`

The expression that specifies the constraint condition.

## 26.3.6 The INFORMATION\_SCHEMA COLLATIONS Table

The `COLLATIONS` table provides information about collations for each character set.

The `COLLATIONS` table has these columns:

- `COLLATION_NAME`

The collation name.

- `CHARACTER_SET_NAME`

The name of the character set with which the collation is associated.

- `ID`

The collation ID.

- `IS_DEFAULT`

Whether the collation is the default for its character set.

- `IS_COMPILED`

Whether the character set is compiled into the server.

- `SORTLEN`

This is related to the amount of memory required to sort strings expressed in the character set.

- `PAD_ATTRIBUTE`

The collation pad attribute, either `NO PAD` or `PAD SPACE`. This attribute affects whether trailing spaces are significant in string comparisons; see [Trailing Space Handling in Comparisons](#).

### Notes

Collation information is also available from the `SHOW COLLATION` statement. See [Section 13.7.7.4, “SHOW COLLATION Statement”](#). The following statements are equivalent:

```
SELECT COLLATION_NAME FROM INFORMATION_SCHEMA.COLLATIONS
```

```
[WHERE COLLATION_NAME LIKE 'wild']
```

```
SHOW COLLATION
```

[LIKE '*wild*']

## 26.3.7 The INFORMATION\_SCHEMA COLLATION\_CHARACTER\_SET\_APPLICABILITY Table

The `COLLATION_CHARACTER_SET_APPLICABILITY` table indicates what character set is applicable for what collation.

The `COLLATION_CHARACTER_SET_APPLICABILITY` table has these columns:

- `COLLATION_NAME`

The collation name.

- `CHARACTER_SET_NAME`

The name of the character set with which the collation is associated.

### Notes

The `COLLATION_CHARACTER_SET_APPLICABILITY` columns are equivalent to the first two columns displayed by the `SHOW COLLATION` statement.

## 26.3.8 The INFORMATION\_SCHEMA COLUMNS Table

The `COLUMNS` table provides information about columns in tables. The related `ST_GeOMETRY_COLUMNS` table provides information about table columns that store spatial data. See Section 26.3.35, “The INFORMATION\_SCHEMA ST\_GeOMETRY\_COLUMNS Table”.

The `COLUMNS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table containing the column belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table containing the column belongs.

- `TABLE_NAME`

The name of the table containing the column.

- `COLUMN_NAME`

The name of the column.

- `ORDINAL_POSITION`

The position of the column within the table. `ORDINAL_POSITION` is necessary because you might want to say `ORDER BY ORDINAL_POSITION`. Unlike `SHOW COLUMNS`, `SELECT` from the `COLUMNS` table does not have automatic ordering.

- `COLUMN_DEFAULT`

The default value for the column. This is `NULL` if the column has an explicit default of `NULL`, or if the column definition includes no `DEFAULT` clause.

- `IS_NULLABLE`

The column nullability. The value is `YES` if `NULL` values can be stored in the column, `NO` if not.

- `DATA_TYPE`

The column data type.

The `DATA_TYPE` value is the type name only with no other information. The `COLUMN_TYPE` value contains the type name and possibly other information such as the precision or length.

- `CHARACTER_MAXIMUM_LENGTH`

For string columns, the maximum length in characters.

- `CHARACTER_OCTET_LENGTH`

For string columns, the maximum length in bytes.

- `NUMERIC_PRECISION`

For numeric columns, the numeric precision.

- `NUMERIC_SCALE`

For numeric columns, the numeric scale.

- `DATETIME_PRECISION`

For temporal columns, the fractional seconds precision.

- `CHARACTER_SET_NAME`

For character string columns, the character set name.

- `COLLATION_NAME`

For character string columns, the collation name.

- `COLUMN_TYPE`

The column data type.

The `DATA_TYPE` value is the type name only with no other information. The `COLUMN_TYPE` value contains the type name and possibly other information such as the precision or length.

- `COLUMN_KEY`

Whether the column is indexed:

- If `COLUMN_KEY` is empty, the column either is not indexed or is indexed only as a secondary column in a multiple-column, nonunique index.
- If `COLUMN_KEY` is `PRI`, the column is a `PRIMARY KEY` or is one of the columns in a multiple-column `PRIMARY KEY`.
- If `COLUMN_KEY` is `UNI`, the column is the first column of a `UNIQUE` index. (A `UNIQUE` index permits multiple `NULL` values, but you can tell whether the column permits `NULL` by checking the `Null` column.)
- If `COLUMN_KEY` is `MUL`, the column is the first column of a nonunique index in which multiple occurrences of a given value are permitted within the column.

If more than one of the `COLUMN_KEY` values applies to a given column of a table, `COLUMN_KEY` displays the one with the highest priority, in the order `PRI`, `UNI`, `MUL`.

A `UNIQUE` index may be displayed as `PRI` if it cannot contain `NULL` values and there is no `PRIMARY KEY` in the table. A `UNIQUE` index may display as `MUL` if several columns form a composite `UNIQUE`

index; although the combination of the columns is unique, each column can still hold multiple occurrences of a given value.

- **EXTRA**

Any additional information that is available about a given column. The value is nonempty in these cases:

- `auto_increment` for columns that have the `AUTO_INCREMENT` attribute.
- `on update CURRENT_TIMESTAMP` for `TIMESTAMP` or `DATETIME` columns that have the `ON UPDATE CURRENT_TIMESTAMP` attribute.
- `STORED GENERATED` or `VIRTUAL GENERATED` for generated columns.
- `DEFAULT_GENERATED` for columns that have an expression default value.

- **PRIVILEGES**

The privileges you have for the column.

- **COLUMN\_COMMENT**

Any comment included in the column definition.

- **GENERATION\_EXPRESSION**

For generated columns, displays the expression used to compute column values. Empty for nongenerated columns. For information about generated columns, see [Section 13.1.20.8, “CREATE TABLE and Generated Columns”](#).

- **SRS\_ID**

This value applies to spatial columns. It contains the column `SRID` value that indicates the spatial reference system for values stored in the column. See [Section 11.4.1, “Spatial Data Types”](#), and [Section 11.4.5, “Spatial Reference System Support”](#). The value is `NULL` for nonspatial columns and spatial columns with no `SRID` attribute.

## Notes

- In `SHOW COLUMNS`, the `Type` display includes values from several different `COLUMNS` columns.
- `CHARACTER_OCTET_LENGTH` should be the same as `CHARACTER_MAXIMUM_LENGTH`, except for multibyte character sets.
- `CHARACTER_SET_NAME` can be derived from `COLLATION_NAME`. For example, if you say `SHOW FULL COLUMNS FROM t`, and you see in the `COLLATION_NAME` column a value of `utf8mb4_swedish_ci`, the character set is what appears before the first underscore: `utf8mb4`.

Column information is also available from the `SHOW COLUMNS` statement. See [Section 13.7.7.5, “SHOW COLUMNS Statement”](#). The following statements are nearly equivalent:

```
SELECT COLUMN_NAME, DATA_TYPE, IS_NULLABLE, COLUMN_DEFAULT
  FROM INFORMATION_SCHEMA.COLUMNS
 WHERE table_name = 'tbl_name'
   [AND table_schema = 'db_name']
   [AND column_name LIKE 'wild']

SHOW COLUMNS
  FROM tbl_name
  [FROM db_name]
  [LIKE 'wild']
```

In MySQL 8.0.30 and later, information about generated invisible primary key columns is visible in this table by default. You can cause such information to be hidden by setting

`show_gipk_in_create_table_and_information_schema = OFF`. For more information, see Section 13.1.20.11, “Generated Invisible Primary Keys”.

## 26.3.9 The INFORMATION\_SCHEMA COLUMNS\_EXTENSIONS Table

The `COLUMNS_EXTENSIONS` table (available as of MySQL 8.0.21) provides information about column attributes defined for primary and secondary storage engines.



### Note

The `COLUMNS_EXTENSIONS` table is reserved for future use.

The `COLUMNS_EXTENSIONS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table belongs.

- `TABLE_NAME`

The name of the table.

- `COLUMN_NAME`

The name of the column.

- `ENGINE_ATTRIBUTE`

Column attributes defined for the primary storage engine. Reserved for future use.

- `SECONDARY_ENGINE_ATTRIBUTE`

Column attributes defined for the secondary storage engine. Reserved for future use.

## 26.3.10 The INFORMATION\_SCHEMA COLUMN\_PRIVILEGES Table

The `COLUMN_PRIVILEGES` table provides information about column privileges. It takes its values from the `mysql.columns_priv` system table.

The `COLUMN_PRIVILEGES` table has these columns:

- `GRANTEE`

The name of the account to which the privilege is granted, in '`user_name`'@'`host_name`' format.

- `TABLE_CATALOG`

The name of the catalog to which the table containing the column belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table containing the column belongs.

- `TABLE_NAME`

The name of the table containing the column.

- `COLUMN_NAME`

The name of the column.

- `PRIVILEGE_TYPE`

The privilege granted. The value can be any privilege that can be granted at the column level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per column privilege held by the grantee.

In the output from `SHOW FULL COLUMNS`, the privileges are all in one column and in lowercase, for example, `select,insert,update,references`. In `COLUMN_PRIVILEGES`, there is one privilege per row, in uppercase.

- `IS_GRANTABLE`

`YES` if the user has the `GRANT OPTION` privilege, `NO` otherwise. The output does not list `GRANT OPTION` as a separate row with `PRIVILEGE_TYPE='GRANT OPTION'`.

## Notes

- `COLUMN_PRIVILEGES` is a nonstandard `INFORMATION_SCHEMA` table.

The following statements are *not* equivalent:

```
SELECT ... FROM INFORMATION_SCHEMA.COLUMN_PRIVILEGES  
SHOW GRANTS ...
```

### 26.3.11 The INFORMATION\_SCHEMA COLUMN\_STATISTICS Table

The `COLUMN_STATISTICS` table provides access to histogram statistics for column values.

For information about histogram statistics, see [Section 8.9.6, “Optimizer Statistics”](#), and [Section 13.7.3.1, “ANALYZE TABLE Statement”](#).

You can see information only for columns for which you have some privilege.

The `COLUMN_STATISTICS` table has these columns:

- `SCHEMA_NAME`

The names of the schema for which the statistics apply.

- `TABLE_NAME`

The names of the column for which the statistics apply.

- `COLUMN_NAME`

The names of the column for which the statistics apply.

- `HISTOGRAM`

A `JSON` object describing the column statistics, stored as a histogram.

### 26.3.12 The INFORMATION\_SCHEMA ENABLED\_ROLES Table

The `ENABLED_ROLES` table (available as of MySQL 8.0.19) provides information about the roles that are enabled within the current session.

The `ENABLED_ROLES` table has these columns:

- `ROLE_NAME`

The user name part of the granted role.

- `ROLE_HOST`

The host name part of the granted role.

- [IS\\_DEFAULT](#)

[YES](#) or [NO](#), depending on whether the role is a default role.

- [IS\\_MANDATORY](#)

[YES](#) or [NO](#), depending on whether the role is mandatory.

### 26.3.13 The INFORMATION\_SCHEMA ENGINES Table

The [ENGINES](#) table provides information about storage engines. This is particularly useful for checking whether a storage engine is supported, or to see what the default engine is.

The [ENGINES](#) table has these columns:

- [ENGINE](#)

The name of the storage engine.

- [SUPPORT](#)

The server's level of support for the storage engine, as shown in the following table.

Value	Meaning
<a href="#">YES</a>	The engine is supported and is active
<a href="#">DEFAULT</a>	Like <a href="#">YES</a> , plus this is the default engine
<a href="#">NO</a>	The engine is not supported
<a href="#">DISABLED</a>	The engine is supported but has been disabled

A value of [NO](#) means that the server was compiled without support for the engine, so it cannot be enabled at runtime.

A value of [DISABLED](#) occurs either because the server was started with an option that disables the engine, or because not all options required to enable it were given. In the latter case, the error log should contain a reason indicating why the option is disabled. See [Section 5.4.2, “The Error Log”](#).

You might also see [DISABLED](#) for a storage engine if the server was compiled to support it, but was started with a `--skip-engine_name` option. For the [NDB](#) storage engine, [DISABLED](#) means the server was compiled with support for NDB Cluster, but was not started with the `--ndbcluster` option.

All MySQL servers support [MyISAM](#) tables. It is not possible to disable [MyISAM](#).

- [COMMENT](#)

A brief description of the storage engine.

- [TRANSACTIONS](#)

Whether the storage engine supports transactions.

- [XA](#)

Whether the storage engine supports XA transactions.

- [SAVEPOINTS](#)

Whether the storage engine supports savepoints.

## Notes

- `ENGINES` is a nonstandard `INFORMATION_SCHEMA` table.

Storage engine information is also available from the `SHOW ENGINES` statement. See [Section 13.7.7.16, “SHOW ENGINES Statement”](#). The following statements are equivalent:

```
SELECT * FROM INFORMATION_SCHEMA.ENGINES  
SHOW ENGINES
```

## 26.3.14 The INFORMATION\_SCHEMA EVENTS Table

The `EVENTS` table provides information about Event Manager events, which are discussed in [Section 25.4, “Using the Event Scheduler”](#).

The `EVENTS` table has these columns:

- `EVENT_CATALOG`

The name of the catalog to which the event belongs. This value is always `def`.

- `EVENT_SCHEMA`

The name of the schema (database) to which the event belongs.

- `EVENT_NAME`

The name of the event.

- `DEFINER`

The account named in the `DEFINER` clause (often the user who created the event), in `'user_name'@'host_name'` format.

- `TIME_ZONE`

The event time zone, which is the time zone used for scheduling the event and that is in effect within the event as it executes. The default value is `SYSTEM`.

- `EVENT_BODY`

The language used for the statements in the event's `DO` clause. The value is always `SQL`.

- `EVENT_DEFINITION`

The text of the SQL statement making up the event's `DO` clause; in other words, the statement executed by this event.

- `EVENT_TYPE`

The event repetition type, either `ONE TIME` (transient) or `RECURRING` (repeating).

- `EXECUTE_AT`

For a one-time event, this is the `DATETIME` value specified in the `AT` clause of the `CREATE EVENT` statement used to create the event, or of the last `ALTER EVENT` statement that modified the event. The value shown in this column reflects the addition or subtraction of any `INTERVAL` value included in the event's `AT` clause. For example, if an event is created using `ON SCHEDELE AT CURRENT_TIMESTAMP + '1:6' DAY_HOUR`, and the event was created at 2018-02-09 14:05:30, the value shown in this column would be `'2018-02-10 20:05:30'`. If the event's timing is determined by an `EVERY` clause instead of an `AT` clause (that is, if the event is recurring), the value of this column is `NULL`.

- [INTERVAL\\_VALUE](#)

For a recurring event, the number of intervals to wait between event executions. For a transient event, the value is always [NULL](#).

- [INTERVAL\\_FIELD](#)

The time units used for the interval which a recurring event waits before repeating. For a transient event, the value is always [NULL](#).

- [SQL\\_MODE](#)

The SQL mode in effect when the event was created or altered, and under which the event executes. For the permitted values, see [Section 5.1.11, “Server SQL Modes”](#).

- [STARTS](#)

The start date and time for a recurring event. This is displayed as a [DATETIME](#) value, and is [NULL](#) if no start date and time are defined for the event. For a transient event, this column is always [NULL](#). For a recurring event whose definition includes a [STARTS](#) clause, this column contains the corresponding [DATETIME](#) value. As with the [EXECUTE\\_AT](#) column, this value resolves any expressions used. If there is no [STARTS](#) clause affecting the timing of the event, this column is [NULL](#).

- [ENDS](#)

For a recurring event whose definition includes a [ENDS](#) clause, this column contains the corresponding [DATETIME](#) value. As with the [EXECUTE\\_AT](#) column, this value resolves any expressions used. If there is no [ENDS](#) clause affecting the timing of the event, this column is [NULL](#).

- [STATUS](#)

The event status. One of [ENABLED](#), [DISABLED](#), or [SLAVESIDE\\_DISABLED](#). [SLAVESIDE\\_DISABLED](#) indicates that the creation of the event occurred on another MySQL server acting as a replication source and replicated to the current MySQL server which is acting as a replica, but the event is not presently being executed on the replica. For more information, see [Section 17.5.1.16, “Replication of Invoked Features”](#). information.

- [ON\\_COMPLETION](#)

One of the two values [PRESERVE](#) or [NOT PRESERVE](#).

- [CREATED](#)

The date and time when the event was created. This is a [TIMESTAMP](#) value.

- [LAST\\_ALTERED](#)

The date and time when the event was last modified. This is a [TIMESTAMP](#) value. If the event has not been modified since its creation, this value is the same as the [CREATED](#) value.

- [LAST\\_EXECUTED](#)

The date and time when the event last executed. This is a [DATETIME](#) value. If the event has never executed, this column is [NULL](#).

[LAST\\_EXECUTED](#) indicates when the event started. As a result, the [ENDS](#) column is never less than [LAST\\_EXECUTED](#).

- [EVENT\\_COMMENT](#)

The text of the comment, if the event has one. If not, this value is empty.

- [ORIGINATOR](#)

The server ID of the MySQL server on which the event was created; used in replication. This value may be updated by `ALTER EVENT` to the server ID of the server on which that statement occurs, if executed on a replication source. The default value is 0.

- `CHARACTER_SET_CLIENT`

The session value of the `character_set_client` system variable when the event was created.

- `COLLATION_CONNECTION`

The session value of the `collation_connection` system variable when the event was created.

- `DATABASE_COLLATION`

The collation of the database with which the event is associated.

## Notes

- `EVENTS` is a nonstandard `INFORMATION_SCHEMA` table.
- Times in the `EVENTS` table are displayed using the event time zone, the current session time zone, or UTC, as described in [Section 25.4.4, “Event Metadata”](#).
- For more information about `SLAVESIDE_DISABLED` and the `ORIGINATOR` column, see [Section 17.5.1.16, “Replication of Invoked Features”](#).

## Example

Suppose that the user '`jon`'@'`ghidora`' creates an event named `e_daily`, and then modifies it a few minutes later using an `ALTER EVENT` statement, as shown here:

```
DELIMITER |

CREATE EVENT e_daily
    ON SCHEDULE
        EVERY 1 DAY
    COMMENT 'Saves total number of sessions then clears the table each day'
    DO
        BEGIN
            INSERT INTO site_activity.totals (time, total)
            SELECT CURRENT_TIMESTAMP, COUNT(*)
            FROM site_activity.sessions;
            DELETE FROM site_activity.sessions;
        END |

DELIMITER ;
ALTER EVENT e_daily
    ENABLE;
```

(Note that comments can span multiple lines.)

This user can then run the following `SELECT` statement, and obtain the output shown:

```
mysql> SELECT * FROM INFORMATION_SCHEMA.EVENTS
    WHERE EVENT_NAME = 'e_daily'
    AND EVENT_SCHEMA = 'myschema'\G
***** 1. row *****
    EVENT_CATALOG: def
    EVENT_SCHEMA: myschema
    EVENT_NAME: e_daily
    DEFINER: jon@ghidora
    TIME_ZONE: SYSTEM
    EVENT_BODY: SQL
    EVENT_DEFINITION: BEGIN
        INSERT INTO site_activity.totals (time, total)
        SELECT CURRENT_TIMESTAMP, COUNT(*)
```

```

        FROM site_activity.sessions;
    DELETE FROM site_activity.sessions;
END
    EVENT_TYPE: RECURRING
    EXECUTE_AT: NULL
    INTERVAL_VALUE: 1
    INTERVAL_FIELD: DAY
    SQL_MODE: ONLY_FULL_GROUP_BY, STRICT_TRANS_TABLES,
               NO_ZERO_IN_DATE, NO_ZERO_DATE,
               ERROR_FOR_DIVISION_BY_ZERO,
               NO_ENGINE_SUBSTITUTION
    STARTS: 2018-08-08 11:06:34
    ENDS: NULL
    STATUS: ENABLED
ON_COMPLETION: NOT PRESERVE
    CREATED: 2018-08-08 11:06:34
    LAST_ALTERED: 2018-08-08 11:06:34
    LAST_EXECUTED: 2018-08-08 16:06:34
EVENT_COMMENT: Saves total number of sessions then clears the
               table each day
    ORIGINATOR: 1
CHARACTER_SET_CLIENT: utf8mb4
COLLATION_CONNECTION: utf8mb4_0900_ai_ci
DATABASE_COLLATION: utf8mb4_0900_ai_ci

```

Event information is also available from the `SHOW EVENTS` statement. See [Section 13.7.7.18, “SHOW EVENTS Statement”](#). The following statements are equivalent:

```

SELECT
    EVENT_SCHEMA, EVENT_NAME, DEFINER, TIME_ZONE, EVENT_TYPE, EXECUTE_AT,
    INTERVAL_VALUE, INTERVAL_FIELD, STARTS, ENDS, STATUS, ORIGINATOR,
    CHARACTER_SET_CLIENT, COLLATION_CONNECTION, DATABASE_COLLATION
FROM INFORMATION_SCHEMA.EVENTS
WHERE table_schema = 'db_name'
[AND column_name LIKE 'wild']

SHOW EVENTS
[FROM db_name]
[LIKE 'wild']

```

### 26.3.15 The INFORMATION\_SCHEMA FILES Table

The `FILES` table provides information about the files in which MySQL tablespace data is stored.

The `FILES` table provides information about `InnoDB` data files. In NDB Cluster, this table also provides information about the files in which NDB Cluster Disk Data tables are stored. For additional information specific to `InnoDB`, see [InnoDB Notes](#), later in this section; for additional information specific to NDB Cluster, see [NDB Notes](#).

The `FILES` table has these columns:

- `FILE_ID`

For `InnoDB`: The tablespace ID, also referred to as the `space_id` or `fil_space_t::id`.

For `NDB`: A file identifier. `FILE_ID` column values are auto-generated.

- `FILE_NAME`

For `InnoDB`: The name of the data file. File-per-table and general tablespaces have an `.ibd` file name extension. Undo tablespaces are prefixed by `undo`. The system tablespace is prefixed by `ibdata`. The global temporary tablespace is prefixed by `ibtmp`. The file name includes the file path, which may be relative to the MySQL data directory (the value of the `datadir` system variable).

For `NDB`: The name of an undo log file created by `CREATE LOGFILE GROUP` or `ALTER LOGFILE GROUP`, or of a data file created by `CREATE TABLESPACE` or `ALTER TABLESPACE`. In NDB 8.0, the file name is shown with a relative path; for an undo log file, this path is relative to the

directory `DataDir/ndb_NodeId_fs/LG`; for a data file, it is relative to the directory `DataDir/ndb_NodeId_fs/TS`. This means, for example, that the name of a data file created with `ALTER TABLESPACE ts ADD DATAFILE 'data_2.dat' INITIAL SIZE 256M` is shown as `./data_2.dat`.

- **FILE\_TYPE**

For `InnoDB`: The tablespace file type. There are three possible file types for `InnoDB` files. `TABLESPACE` is the file type for any system, general, or file-per-table tablespace file that holds tables, indexes, or other forms of user data. `TEMPORARY` is the file type for temporary tablespaces. `UNDO LOG` is the file type for undo tablespaces, which hold undo records.

For `NDB`: One of the values `UNDO LOG` or `DATAFILE`. Prior to NDB 8.0.13, `TABLESPACE` was also a possible value.

- **TABLESPACE\_NAME**

The name of the tablespace with which the file is associated.

For `InnoDB`: General tablespace names are as specified when created. File-per-table tablespace names are shown in the following format: `schema_name/table_name`. The `InnoDB` system tablespace name is `innodb_system`. The global temporary tablespace name is `innodb_temporary`. Default undo tablespace names are `innodb_undo_001` and `innodb_undo_002`. User-created undo tablespace names are as specified when created.

- **TABLE\_CATALOG**

This value is always empty.

- **TABLE\_SCHEMA**

This is always `NULL`.

- **TABLE\_NAME**

This is always `NULL`.

- **LOGFILE\_GROUP\_NAME**

For `InnoDB`: This is always `NULL`.

For `NDB`: The name of the log file group to which the log file or data file belongs.

- **LOGFILE\_GROUP\_NUMBER**

For `InnoDB`: This is always `NULL`.

For `NDB`: For a Disk Data undo log file, the auto-generated ID number of the log file group to which the log file belongs. This is the same as the value shown for the `id` column in the `ndbinfo.dict_obj_info` table and the `log_id` column in the `ndbinfo.logspaces` and `ndbinfo.logs` tables for this undo log file.

- **ENGINE**

For `InnoDB`: This value is always `InnoDB`.

For `NDB`: This value is always `ndbcluster`.

- **FULLTEXT\_KEYS**

This is always `NULL`.

- **DELETED\_ROWS**

This is always `NULL`.

- `UPDATE_COUNT`

This is always `NULL`.

- `FREE_EXTENTS`

For `InnoDB`: The number of fully free extents in the current data file.

For `NDB`: The number of extents which have not yet been used by the file.

- `TOTAL_EXTENTS`

For `InnoDB`: The number of full extents used in the current data file. Any partial extent at the end of the file is not counted.

For `NDB`: The total number of extents allocated to the file.

- `EXTENT_SIZE`

For `InnoDB`: Extent size is 1048576 (1MB) for files with a 4KB, 8KB, or 16KB page size. Extent size is 2097152 bytes (2MB) for files with a 32KB page size, and 4194304 (4MB) for files with a 64KB page size. `FILES` does not report `InnoDB` page size. Page size is defined by the `innodb_page_size` system variable. Extent size information can also be retrieved from the `INNODB_TABLESPACES` table where `FILES.FILE_ID = INNODB_TABLESPACES.SPACE`.

For `NDB`: The size of an extent for the file in bytes.

- `INITIAL_SIZE`

For `InnoDB`: The initial size of the file in bytes.

For `NDB`: The size of the file in bytes. This is the same value that was used in the `INITIAL_SIZE` clause of the `CREATE LOGFILE GROUP`, `ALTER LOGFILE GROUP`, `CREATE TABLESPACE`, or `ALTER TABLESPACE` statement used to create the file.

- `MAXIMUM_SIZE`

For `InnoDB`: The maximum number of bytes permitted in the file. The value is `NULL` for all data files except for predefined system tablespace data files. Maximum system tablespace file size is defined by `innodb_data_file_path`. Maximum global temporary tablespace file size is defined by `innodb_temp_data_file_path`. A `NULL` value for a predefined system tablespace data file indicates that a file size limit was not defined explicitly.

For `NDB`: This value is always the same as the `INITIAL_SIZE` value.

- `AUTOEXTEND_SIZE`

The auto-extend size of the tablespace. For `NDB`, `AUTOEXTEND_SIZE` is always `NULL`.

- `CREATION_TIME`

This is always `NULL`.

- `LAST_UPDATE_TIME`

This is always `NULL`.

- `LAST_ACCESS_TIME`

This is always `NULL`.

- **RECOVER\_TIME**

This is always **NULL**.

- **TRANSACTION\_COUNTER**

This is always **NULL**.

- **VERSION**

For **InnoDB**: This is always **NULL**.

For **NDB**: The version number of the file.

- **ROW\_FORMAT**

For **InnoDB**: This is always **NULL**.

For **NDB**: One of **FIXED** or **DYNAMIC**.

- **TABLE\_ROWS**

This is always **NULL**.

- **AVG\_ROW\_LENGTH**

This is always **NULL**.

- **DATA\_LENGTH**

This is always **NULL**.

- **MAX\_DATA\_LENGTH**

This is always **NULL**.

- **INDEX\_LENGTH**

This is always **NULL**.

- **DATA\_FREE**

For **InnoDB**: The total amount of free space (in bytes) for the entire tablespace. Predefined system tablespaces, which include the system tablespace and temporary table tablespaces, may have one or more data files.

For **NDB**: This is always **NULL**.

- **CREATE\_TIME**

This is always **NULL**.

- **UPDATE\_TIME**

This is always **NULL**.

- **CHECK\_TIME**

This is always **NULL**.

- **CHECKSUM**

This is always **NULL**.

- **STATUS**

For [InnoDB](#): This value is [NORMAL](#) by default. [InnoDB](#) file-per-table tablespaces may report [IMPORTING](#), which indicates that the tablespace is not yet available.

For [NDB](#): For NDB Cluster Disk Data files, this value is always [NORMAL](#).

- [EXTRA](#)

For [InnoDB](#): This is always [NULL](#).

For [NDB](#): (*NDB 8.0.15 and later*) For undo log files, this column shows the undo log buffer size; for data files, it is always [NULL](#). A more detailed explanation is provided in the next few paragraphs.

[NDBCLUSTER](#) stores a copy of each data file and each undo log file on each data node in the cluster. In NDB 8.0.13 and later, the [FILES](#) table contains only one row for each such file. Suppose that you run the following two statements on an NDB Cluster with four data nodes:

```
CREATE LOGFILE GROUP mygroup
    ADD UNDOFILE 'new_undo.dat'
    INITIAL_SIZE 2G
    ENGINE NDBCLUSTER;

CREATE TABLESPACE myts
    ADD DATAFILE 'data_1.dat'
    USE LOGFILE GROUP mygroup
    INITIAL_SIZE 256M
    ENGINE NDBCLUSTER;
```

After running these two statements successfully, you should see a result similar to the one shown here for this query against the [FILES](#) table:

```
mysql> SELECT LOGFILE_GROUP_NAME, FILE_TYPE, EXTRA
    ->     FROM INFORMATION_SCHEMA.FILES
    ->     WHERE ENGINE = 'ndbcluster';

+-----+-----+-----+
| LOGFILE_GROUP_NAME | FILE_TYPE | EXTRA          |
+-----+-----+-----+
| mygroup           | UNDO LOG  | UNDO_BUFFER_SIZE=8388608 |
| mygroup           | DATAFILE   | NULL            |
+-----+-----+-----+
```

The undo log buffer size information was inadvertently removed in NDB 8.0.13, but was restored in NDB 8.0.15. (Bug #92796, Bug #28800252)

Prior to NDB 8.0.13, the [FILES](#) table contained a row for each of these files on each data node the file belonged to, as well as the size of its undo buffer. In these versions, the result of the same query contains one row per data node, as shown here:

```
+-----+-----+-----+
| LOGFILE_GROUP_NAME | FILE_TYPE | EXTRA          |
+-----+-----+-----+
| mygroup           | UNDO LOG  | CLUSTER_NODE=5 ; UNDO_BUFFER_SIZE=8388608 |
| mygroup           | UNDO LOG  | CLUSTER_NODE=6 ; UNDO_BUFFER_SIZE=8388608 |
| mygroup           | UNDO LOG  | CLUSTER_NODE=7 ; UNDO_BUFFER_SIZE=8388608 |
| mygroup           | UNDO LOG  | CLUSTER_NODE=8 ; UNDO_BUFFER_SIZE=8388608 |
| mygroup           | DATAFILE   | CLUSTER_NODE=5 |
| mygroup           | DATAFILE   | CLUSTER_NODE=6 |
| mygroup           | DATAFILE   | CLUSTER_NODE=7 |
| mygroup           | DATAFILE   | CLUSTER_NODE=8 |
+-----+-----+-----+
```

## Notes

- [FILES](#) is a nonstandard [INFORMATION\\_SCHEMA](#) table.
- As of MySQL 8.0.21, you must have the [PROCESS](#) privilege to query this table.

## InnoDB Notes

The following notes apply to InnoDB data files.

- Information reported by `FILES` is obtained from the InnoDB in-memory cache for open files, whereas `INNODB_DATAFILES` gets its data from the `INNODB_SYS_DATAFILES` internal data dictionary table.
- The information provided by `FILES` includes global temporary tablespace information which is not available in the `INNODB_SYS_DATAFILES` internal data dictionary table, and is therefore not included in `INNODB_DATAFILES`.
- Undo tablespace information is shown in `FILES` when separate undo tablespaces are present, as they are by default in MySQL 8.0.
- The following query returns all `FILES` table information relating to InnoDB tablespaces.

```
SELECT
    FILE_ID, FILE_NAME, FILE_TYPE, TABLESPACE_NAME, FREE_EXTENTS,
    TOTAL_EXTENTS, EXTENT_SIZE, INITIAL_SIZE, MAXIMUM_SIZE,
    AUTOEXTEND_SIZE, DATA_FREE, STATUS
FROM INFORMATION_SCHEMA.FILES
WHERE ENGINE='InnoDB' \G
```

## NDB Notes

- The `FILES` table provides information about Disk Data *files* only; you cannot use it for determining disk space allocation or availability for individual NDB tables. However, it is possible to see how much space is allocated for each NDB table having data stored on disk—as well as how much remains available for storage of data on disk for that table—using `ndb_desc`.
- Beginning with NDB 8.0.29 much of the information in the `FILES` table can also be found in the `ndbinfo.files` table.
- The `CREATION_TIME`, `LAST_UPDATE_TIME`, and `LAST_ACCESSED` values are as reported by the operating system, and are not supplied by the NDB storage engine. Where no value is provided by the operating system, these columns display `NULL`.
- The difference between the `TOTAL_EXTENTS` and `FREE_EXTENTS` columns is the number of extents currently in use by the file:

```
SELECT TOTAL_EXTENTS - FREE_EXTENTS AS extents_used
    FROM INFORMATION_SCHEMA.FILES
    WHERE FILE_NAME = './myfile.dat';
```

To approximate the amount of disk space in use by the file, multiply that difference by the value of the `EXTENT_SIZE` column, which gives the size of an extent for the file in bytes:

```
SELECT (TOTAL_EXTENTS - FREE_EXTENTS) * EXTENT_SIZE AS bytes_used
    FROM INFORMATION_SCHEMA.FILES
    WHERE FILE_NAME = './myfile.dat';
```

Similarly, you can estimate the amount of space that remains available in a given file by multiplying `FREE_EXTENTS` by `EXTENT_SIZE`:

```
SELECT FREE_EXTENTS * EXTENT_SIZE AS bytes_free
    FROM INFORMATION_SCHEMA.FILES
    WHERE FILE_NAME = './myfile.dat';
```



### Important

The byte values produced by the preceding queries are approximations only, and their precision is inversely proportional to the value of `EXTENT_SIZE`. That is, the larger `EXTENT_SIZE` becomes, the less accurate the approximations are.

It is also important to remember that once an extent is used, it cannot be freed again without dropping the data file of which it is a part. This means that deletes from a Disk Data table do *not* release disk space.

The extent size can be set in a [CREATE TABLESPACE](#) statement. For more information, see [Section 13.1.21, “CREATE TABLESPACE Statement”](#).

- Prior to NDB 8.0.13, an additional row was present in the `FILES` table following the creation of a logfile group, having `NULL` in the `FILE_NAME` column. In NDB 8.0.13 and later, this row—which did not correspond to any file—is no longer shown, and it is necessary to query the `ndbinfo.logspaces` table to obtain undo log file usage information. See the description of this table as well as [Section 23.6.11.1, “NDB Cluster Disk Data Objects”](#), for more information.

The remainder of the discussion in this item applies only to NDB 8.0.12 and earlier. For the row having `NULL` in the `FILE_NAME` column, the value of the `FILE_ID` column is always `0`, that of the `FILE_TYPE` column is always `UNDO LOG`, and that of the `STATUS` column is always `NORMAL`. The value of the `ENGINE` column is always `ndbcluster`.

The `FREE_EXTENTS` column in this row shows the total number of free extents available to all undo files belonging to a given log file group whose name and number are shown in the `LOGFILE_GROUP_NAME` and `LOGFILE_GROUP_NUMBER` columns, respectively.

Suppose there are no existing log file groups on your NDB Cluster, and you create one using the following statement:

```
mysql> CREATE LOGFILE GROUP lg1
      ADD UNDOFILE 'undofile.dat'
      INITIAL_SIZE = 16M
      UNDO_BUFFER_SIZE = 1M
      ENGINE = NDB;
```

You can now see this `NULL` row when you query the `FILES` table:

```
mysql> SELECT DISTINCT
      FILE_NAME AS File,
      FREE_EXTENTS AS Free,
      TOTAL_EXTENTS AS Total,
      EXTENT_SIZE AS Size,
      INITIAL_SIZE AS Initial
      FROM INFORMATION_SCHEMA.FILES;
+-----+-----+-----+-----+
| File | Free | Total | Size | Initial |
+-----+-----+-----+-----+
| undofile.dat | NULL | 4194304 | 4 | 16777216 |
| NULL | 4184068 | NULL | 4 | NULL |
+-----+-----+-----+-----+
```

The total number of free extents available for undo logging is always somewhat less than the sum of the `TOTAL_EXTENTS` column values for all undo files in the log file group due to overhead required for maintaining the undo files. This can be seen by adding a second undo file to the log file group, then repeating the previous query against the `FILES` table:

```
mysql> ALTER LOGFILE GROUP lg1
      ADD UNDOFILE 'undofile02.dat'
      INITIAL_SIZE = 4M
      ENGINE = NDB;

mysql> SELECT DISTINCT
      FILE_NAME AS File,
      FREE_EXTENTS AS Free,
      TOTAL_EXTENTS AS Total,
      EXTENT_SIZE AS Size,
      INITIAL_SIZE AS Initial
      FROM INFORMATION_SCHEMA.FILES;
+-----+-----+-----+-----+
```

## The INFORMATION\_SCHEMA.FILES Table

File	Free	Total	Size	Initial
undofile.dat	NULL	4194304	4	16777216
undofile02.dat	NULL	1048576	4	4194304
NULL	5223944	NULL	4	NULL

The amount of free space in bytes which is available for undo logging by Disk Data tables using this log file group can be approximated by multiplying the number of free extents by the initial size:

```
mysql> SELECT
    FREE_EXTENTS AS 'Free Extents',
    FREE_EXTENTS * EXTENT_SIZE AS 'Free Bytes'
    FROM INFORMATION_SCHEMA.FILES
    WHERE LOGFILE_GROUP_NAME = 'lg1'
    AND FILE_NAME IS NULL;
+-----+-----+
| Free Extents | Free Bytes |
+-----+-----+
|      5223944 |     20895776 |
+-----+-----+
```

If you create an NDB Cluster Disk Data table and then insert some rows into it, you can see approximately how much space remains for undo logging afterward, for example:

```
mysql> CREATE TABLESPACE ts1
    ADD DATAFILE 'data1.dat'
    USE LOGFILE GROUP lg1
    INITIAL_SIZE 512M
    ENGINE = NDB;

mysql> CREATE TABLE dd (
    c1 INT NOT NULL PRIMARY KEY,
    c2 INT,
    c3 DATE
)
TABLESPACE ts1 STORAGE DISK
ENGINE = NDB;

mysql> INSERT INTO dd VALUES
    (NULL, 1234567890, '2007-02-02'),
    (NULL, 1126789005, '2007-02-03'),
    (NULL, 1357924680, '2007-02-04'),
    (NULL, 1642097531, '2007-02-05');

mysql> SELECT
    FREE_EXTENTS AS 'Free Extents',
    FREE_EXTENTS * EXTENT_SIZE AS 'Free Bytes'
    FROM INFORMATION_SCHEMA.FILES
    WHERE LOGFILE_GROUP_NAME = 'lg1'
    AND FILE_NAME IS NULL;
+-----+-----+
| Free Extents | Free Bytes |
+-----+-----+
|      5207565 |     20830260 |
+-----+-----+
```

- Prior to NDB 8.0.13, an additional row was present in the `FILES` table for each NDB Cluster Disk Data tablespace. Because it did not correspond to an actual file, it was removed in NDB 8.0.13. This row had `NULL` for the value of the `FILE_NAME` column, the value of the `FILE_ID` column was always `0`, that of the `FILE_TYPE` column was always `TABLESPACE`, that of the `STATUS` column was always `NORMAL`, and the value of the `ENGINE` column is always `NDBCLUSTER`.

In NDB 8.0.13 and later, you can obtain information about Disk Data tablespaces using the `ndb_desc` utility. For more information, see [Section 23.6.11.1, “NDB Cluster Disk Data Objects”](#), as well as the description of `ndb_desc`.

- For additional information, and examples of creating, dropping, and obtaining information about NDB Cluster Disk Data objects, see [Section 23.6.11, “NDB Cluster Disk Data Tables”](#).

### 26.3.16 The INFORMATION\_SCHEMA KEY\_COLUMN\_USAGE Table

The `KEY_COLUMN_USAGE` table describes which key columns have constraints. This table provides no information about functional key parts because they are expressions and the table provides information only about columns.

The `KEY_COLUMN_USAGE` table has these columns:

- `CONSTRAINT_CATALOG`

The name of the catalog to which the constraint belongs. This value is always `def`.

- `CONSTRAINT_SCHEMA`

The name of the schema (database) to which the constraint belongs.

- `CONSTRAINT_NAME`

The name of the constraint.

- `TABLE_CATALOG`

The name of the catalog to which the table belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table belongs.

- `TABLE_NAME`

The name of the table that has the constraint.

- `COLUMN_NAME`

The name of the column that has the constraint.

If the constraint is a foreign key, then this is the column of the foreign key, not the column that the foreign key references.

- `ORDINAL_POSITION`

The column's position within the constraint, not the column's position within the table. Column positions are numbered beginning with 1.

- `POSITION_IN_UNIQUE_CONSTRAINT`

`NULL` for unique and primary-key constraints. For foreign-key constraints, this column is the ordinal position in key of the table that is being referenced.

- `REFERENCED_TABLE_SCHEMA`

The name of the schema referenced by the constraint.

- `REFERENCED_TABLE_NAME`

The name of the table referenced by the constraint.

- `REFERENCED_COLUMN_NAME`

The name of the column referenced by the constraint.

Suppose that there are two tables name `t1` and `t3` that have the following definitions:

```
CREATE TABLE t1
```

```

(
    s1 INT,
    s2 INT,
    s3 INT,
    PRIMARY KEY(s3)
) ENGINE=InnoDB;

CREATE TABLE t3
(
    s1 INT,
    s2 INT,
    s3 INT,
    KEY(s1),
    CONSTRAINT CO FOREIGN KEY (s2) REFERENCES t1(s3)
) ENGINE=InnoDB;

```

For those two tables, the `KEY_COLUMN_USAGE` table has two rows:

- One row with `CONSTRAINT_NAME = 'PRIMARY'`, `TABLE_NAME = 't1'`, `COLUMN_NAME = 's3'`, `ORDINAL_POSITION = 1`, `POSITION_IN_UNIQUE_CONSTRAINT = NULL`.  
For NDB: This value is always `NULL`.
- One row with `CONSTRAINT_NAME = 'CO'`, `TABLE_NAME = 't3'`, `COLUMN_NAME = 's2'`, `ORDINAL_POSITION = 1`, `POSITION_IN_UNIQUE_CONSTRAINT = 1`.

### 26.3.17 The INFORMATION\_SCHEMA KEYWORDS Table

The `KEYWORDS` table lists the words considered keywords by MySQL and, for each one, indicates whether it is reserved. Reserved keywords may require special treatment in some contexts, such as special quoting when used as identifiers (see [Section 9.3, “Keywords and Reserved Words”](#)). This table provides applications a runtime source of MySQL keyword information.

Prior to MySQL 8.0.13, selecting from the `KEYWORDS` table with no default database selected produced an error. (Bug #90160, Bug #27729859)

The `KEYWORDS` table has these columns:

- `WORD`

The keyword.

- `RESERVED`

An integer indicating whether the keyword is reserved (1) or nonreserved (0).

These queries lists all keywords, all reserved keywords, and all nonreserved keywords, respectively:

```

SELECT * FROM INFORMATION_SCHEMA.KEYWORDS;
SELECT WORD FROM INFORMATION_SCHEMA.KEYWORDS WHERE RESERVED = 1;
SELECT WORD FROM INFORMATION_SCHEMA.KEYWORDS WHERE RESERVED = 0;

```

The latter two queries are equivalent to:

```

SELECT WORD FROM INFORMATION_SCHEMA.KEYWORDS WHERE RESERVED;
SELECT WORD FROM INFORMATION_SCHEMA.KEYWORDS WHERE NOT RESERVED;

```

If you build MySQL from source, the build process generates a `keyword_list.h` header file containing an array of keywords and their reserved status. This file can be found in the `sql` directory under the build directory. This file may be useful for applications that require a static source for the keyword list.

### 26.3.18 The INFORMATION\_SCHEMA ndb\_transid\_mysql\_connection\_map Table

## The INFORMATION\_SCHEMA ndb\_transid\_mysql\_connection\_map Table

The `ndb_transid_mysql_connection_map` table provides a mapping between NDB transactions, NDB transaction coordinators, and MySQL Servers attached to an NDB Cluster as API nodes. This information is used when populating the `server_operations` and `server_transactions` tables of the `ndbinfo` NDB Cluster information database.

INFORMATION_SCHEMA Name	SHOW Name	Remarks
<code>mysql_connection_id</code>		MySQL Server connection ID
<code>node_id</code>		Transaction coordinator node ID
<code>ndb_transid</code>		NDB transaction ID

The `mysql_connection_id` is the same as the connection or session ID shown in the output of `SHOW PROCESSLIST`.

There are no `SHOW` statements associated with this table.

This is a nonstandard table, specific to NDB Cluster. It is implemented as an `INFORMATION_SCHEMA` plugin; you can verify that it is supported by checking the output of `SHOW PLUGINS`. If `ndb_transid_mysql_connection_map` support is enabled, the output from this statement includes a plugin having this name, of type `INFORMATION SCHEMA`, and having status `ACTIVE`, as shown here (using emphasized text):

```
mysql> SHOW PLUGINS;
+-----+-----+-----+-----+-----+
| Name           | Status | Type      | Library | License |
+-----+-----+-----+-----+-----+
| binlog          | ACTIVE | STORAGE ENGINE | NULL    | GPL     |
| mysql_native_password | ACTIVE | AUTHENTICATION | NULL    | GPL     |
| sha256_password | ACTIVE | AUTHENTICATION | NULL    | GPL     |
| caching_sha2_password | ACTIVE | AUTHENTICATION | NULL    | GPL     |
| sha2_cache_cleaner | ACTIVE | AUDIT      | NULL    | GPL     |
| daemon_keyring_proxy_plugin | ACTIVE | DAEMON    | NULL    | GPL     |
| CSV             | ACTIVE | STORAGE ENGINE | NULL    | GPL     |
| MEMORY          | ACTIVE | STORAGE ENGINE | NULL    | GPL     |
| InnoDB           | ACTIVE | STORAGE ENGINE | NULL    | GPL     |
| INNODB_TRX       | ACTIVE | INFORMATION SCHEMA | NULL    | GPL     |
| INNODB_CMP       | ACTIVE | INFORMATION SCHEMA | NULL    | GPL     |
|
...
| INNODB_SESSION_TEMP_TABLESPACES | ACTIVE | INFORMATION SCHEMA | NULL    | GPL     |
| MyISAM            | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| MRG_MYISAM        | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| PERFORMANCE_SCHEMA | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| TempTable         | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| ARCHIVE           | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| BLACKHOLE          | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| ndbcluster         | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| ndbinfo            | ACTIVE | STORAGE ENGINE   | NULL    | GPL     |
| /ndb_transid_mysql_connection_map | ACTIVE | INFORMATION SCHEMA | NULL    | GPL     |
| ngram              | ACTIVE | FTPPARSER      | NULL    | GPL     |
| mysqlx_cache_cleaner | ACTIVE | AUDIT      | NULL    | GPL     |
| mysqlx              | ACTIVE | DAEMON    | NULL    | GPL     |
+-----+-----+-----+-----+-----+
47 rows in set (0.01 sec)
```

The plugin is enabled by default. You can disable it (or force the server not to run unless the plugin starts) by starting the server with the `--ndb-transid-mysql-connection-map` option. If the plugin is disabled, the status is shown by `SHOW PLUGINS` as `DISABLED`. The plugin cannot be enabled or disabled at runtime.

Although the names of this table and its columns are displayed using lowercase, you can use uppercase or lowercase when referring to them in SQL statements.

For this table to be created, the MySQL Server must be a binary supplied with the NDB Cluster distribution, or one built from the NDB Cluster sources with NDB storage engine support enabled. It is not available in the standard MySQL 8.0 Server.

### 26.3.19 The INFORMATION\_SCHEMA OPTIMIZER\_TRACE Table

The `OPTIMIZER_TRACE` table provides information produced by the optimizer tracing capability for traced statements. To enable tracking, use the `optimizer_trace` system variable. For details, see [MySQL Internals: Tracing the Optimizer](#).

The `OPTIMIZER_TRACE` table has these columns:

- `QUERY`

The text of the traced statement.

- `TRACE`

The trace, in `JSON` format.

- `MISSING_BYTES_BEYOND_MAX_MEM_SIZE`

Each remembered trace is a string that is extended as optimization progresses and appends data to it. The `optimizer_trace_max_mem_size` variable sets a limit on the total amount of memory used by all currently remembered traces. If this limit is reached, the current trace is not extended (and thus is incomplete), and the `MISSING_BYTES_BEYOND_MAX_MEM_SIZE` column shows the number of bytes missing from the trace.

- `INSUFFICIENT_PRIVILEGES`

If a traced query uses views or stored routines that have `SQL SECURITY` with a value of `DEFINER`, it may be that a user other than the definer is denied from seeing the trace of the query. In that case, the trace is shown as empty and `INSUFFICIENT_PRIVILEGES` has a value of 1. Otherwise, the value is 0.

### 26.3.20 The INFORMATION\_SCHEMA PARAMETERS Table

The `PARAMETERS` table provides information about parameters for stored routines (stored procedures and stored functions), and about return values for stored functions. The `PARAMETERS` table does not include built-in (native) functions or loadable functions.

The `PARAMETERS` table has these columns:

- `SPECIFIC_CATALOG`

The name of the catalog to which the routine containing the parameter belongs. This value is always `def`.

- `SPECIFIC_SCHEMA`

The name of the schema (database) to which the routine containing the parameter belongs.

- `SPECIFIC_NAME`

The name of the routine containing the parameter.

- `ORDINAL_POSITION`

For successive parameters of a stored procedure or function, the `ORDINAL_POSITION` values are 1, 2, 3, and so forth. For a stored function, there is also a row that applies to the function return value (as described by the `RETURNS` clause). The return value is not a true parameter, so the row that describes it has these unique characteristics:

- The `ORDINAL_POSITION` value is 0.
- The `PARAMETER_NAME` and `PARAMETER_MODE` values are `NULL` because the return value has no name and the mode does not apply.

- [PARAMETER\\_MODE](#)

The mode of the parameter. This value is one of [IN](#), [OUT](#), or [INOUT](#). For a stored function return value, this value is [NULL](#).

- [PARAMETER\\_NAME](#)

The name of the parameter. For a stored function return value, this value is [NULL](#).

- [DATA\\_TYPE](#)

The parameter data type.

The [DATA\\_TYPE](#) value is the type name only with no other information. The [DTD\\_IDENTIFIER](#) value contains the type name and possibly other information such as the precision or length.

- [CHARACTER\\_MAXIMUM\\_LENGTH](#)

For string parameters, the maximum length in characters.

- [CHARACTER\\_OCTET\\_LENGTH](#)

For string parameters, the maximum length in bytes.

- [NUMERIC\\_PRECISION](#)

For numeric parameters, the numeric precision.

- [NUMERIC\\_SCALE](#)

For numeric parameters, the numeric scale.

- [DATETIME\\_PRECISION](#)

For temporal parameters, the fractional seconds precision.

- [CHARACTER\\_SET\\_NAME](#)

For character string parameters, the character set name.

- [COLLATION\\_NAME](#)

For character string parameters, the collation name.

- [DTD\\_IDENTIFIER](#)

The parameter data type.

The [DATA\\_TYPE](#) value is the type name only with no other information. The [DTD\\_IDENTIFIER](#) value contains the type name and possibly other information such as the precision or length.

- [ROUTINE\\_TYPE](#)

[PROCEDURE](#) for stored procedures, [FUNCTION](#) for stored functions.

### 26.3.21 The INFORMATION\_SCHEMA PARTITIONS Table

The [PARTITIONS](#) table provides information about table partitions. Each row in this table corresponds to an individual partition or subpartition of a partitioned table. For more information about partitioning tables, see [Chapter 24, Partitioning](#).

The [PARTITIONS](#) table has these columns:

- [TABLE\\_CATALOG](#)

The name of the catalog to which the table belongs. This value is always `def`.

- **TABLE\_SCHEMA**

The name of the schema (database) to which the table belongs.

- **TABLE\_NAME**

The name of the table containing the partition.

- **PARTITION\_NAME**

The name of the partition.

- **SUBPARTITION\_NAME**

If the `PARTITIONS` table row represents a subpartition, the name of subpartition; otherwise `NULL`.

For `NDB`: This value is always `NULL`.

- **PARTITION\_ORDINAL\_POSITION**

All partitions are indexed in the same order as they are defined, with `1` being the number assigned to the first partition. The indexing can change as partitions are added, dropped, and reorganized; the number shown is this column reflects the current order, taking into account any indexing changes.

- **SUBPARTITION\_ORDINAL\_POSITION**

Subpartitions within a given partition are also indexed and reindexed in the same manner as partitions are indexed within a table.

- **PARTITION\_METHOD**

One of the values `RANGE`, `LIST`, `HASH`, `LINEAR HASH`, `KEY`, or `LINEAR KEY`; that is, one of the available partitioning types as discussed in [Section 24.2, “Partitioning Types”](#).

- **SUBPARTITION\_METHOD**

One of the values `HASH`, `LINEAR HASH`, `KEY`, or `LINEAR KEY`; that is, one of the available subpartitioning types as discussed in [Section 24.2.6, “Subpartitioning”](#).

- **PARTITION\_EXPRESSION**

The expression for the partitioning function used in the `CREATE TABLE` or `ALTER TABLE` statement that created the table's current partitioning scheme.

For example, consider a partitioned table created in the `test` database using this statement:

```
CREATE TABLE tp (
    c1 INT,
    c2 INT,
    c3 VARCHAR(25)
)
PARTITION BY HASH(c1 + c2)
PARTITIONS 4;
```

The `PARTITION_EXPRESSION` column in a `PARTITIONS` table row for a partition from this table displays `c1 + c2`, as shown here:

```
mysql> SELECT DISTINCT PARTITION_EXPRESSION
      FROM INFORMATION_SCHEMA.PARTITIONS
     WHERE TABLE_NAME='tp' AND TABLE_SCHEMA='test';
+-----+
| PARTITION_EXPRESSION |
+-----+
```

c1 + c2

For a table that is not explicitly partitioned, this column is always `NULL`, regardless of storage engine.

- [SUBPARTITION\\_EXPRESSION](#)

This works in the same fashion for the subpartitioning expression that defines the subpartitioning for a table as [PARTITION\\_EXPRESSION](#) does for the partitioning expression used to define a table's partitioning.

If the table has no subpartitions, this column is `NULL`.

- [PARTITION\\_DESCRIPTION](#)

This column is used for RANGE and LIST partitions. For a [RANGE](#) partition, it contains the value set in the partition's `VALUES LESS THAN` clause, which can be either an integer or `MAXVALUE`. For a [LIST](#) partition, this column contains the values defined in the partition's `VALUES IN` clause, which is a list of comma-separated integer values.

For partitions whose `PARTITION_METHOD` is other than [RANGE](#) or [LIST](#), this column is always `NULL`.

- [TABLE\\_ROWS](#)

The number of table rows in the partition.

For partitioned [InnoDB](#) tables, the row count given in the [TABLE\\_ROWS](#) column is only an estimated value used in SQL optimization, and may not always be exact.

For [NDB](#) tables, you can also obtain this information using the `ndb_desc` utility.

- [AVG\\_ROW\\_LENGTH](#)

The average length of the rows stored in this partition or subpartition, in bytes. This is the same as [DATA\\_LENGTH](#) divided by [TABLE\\_ROWS](#).

For [NDB](#) tables, you can also obtain this information using the `ndb_desc` utility.

- [DATA\\_LENGTH](#)

The total length of all rows stored in this partition or subpartition, in bytes; that is, the total number of bytes stored in the partition or subpartition.

For [NDB](#) tables, you can also obtain this information using the `ndb_desc` utility.

- [MAX\\_DATA\\_LENGTH](#)

The maximum number of bytes that can be stored in this partition or subpartition.

For [NDB](#) tables, you can also obtain this information using the `ndb_desc` utility.

- [INDEX\\_LENGTH](#)

The length of the index file for this partition or subpartition, in bytes.

For partitions of [NDB](#) tables, whether the tables use implicit or explicit partitioning, the [INDEX\\_LENGTH](#) column value is always 0. However, you can obtain equivalent information using the `ndb_desc` utility.

- [DATA\\_FREE](#)

The number of bytes allocated to the partition or subpartition but not used.

For [NDB](#) tables, you can also obtain this information using the `ndb_desc` utility.

- [CREATE\\_TIME](#)  
The time that the partition or subpartition was created.
- [UPDATE\\_TIME](#)  
The time that the partition or subpartition was last modified.
- [CHECK\\_TIME](#)  
The last time that the table to which this partition or subpartition belongs was checked.  
For partitioned [InnoDB](#) tables, the value is always [NULL](#).
- [CHECKSUM](#)  
The checksum value, if any; otherwise [NULL](#).
- [PARTITION\\_COMMENT](#)  
The text of the comment, if the partition has one. If not, this value is empty.  
The maximum length for a partition comment is defined as 1024 characters, and the display width of the [PARTITION\\_COMMENT](#) column is also 1024, characters to match this limit.
- [NODEGROUP](#)  
This is the nodegroup to which the partition belongs. For NDB Cluster tables, this is always [default](#). For partitioned tables using storage engines other than [NDB](#), the value is also [default](#). Otherwise, this column is empty.
- [TABLESPACE\\_NAME](#)  
The name of the tablespace to which the partition belongs. The value is always [DEFAULT](#), unless the table uses the [NDB](#) storage engine (see the *Notes* at the end of this section).

## Notes

- [PARTITIONS](#) is a nonstandard [INFORMATION\\_SCHEMA](#) table.
- A table using any storage engine other than [NDB](#) and which is not partitioned has one row in the [PARTITIONS](#) table. However, the values of the [PARTITION\\_NAME](#), [SUBPARTITION\\_NAME](#), [PARTITION\\_ORDINAL\\_POSITION](#), [SUBPARTITION\\_ORDINAL\\_POSITION](#), [PARTITION\\_METHOD](#), [SUBPARTITION\\_METHOD](#), [PARTITION\\_EXPRESSION](#), [SUBPARTITION\\_EXPRESSION](#), and [PARTITION\\_DESCRIPTION](#) columns are all [NULL](#). Also, the [PARTITION\\_COMMENT](#) column in this case is blank.
- An [NDB](#) table which is not explicitly partitioned has one row in the [PARTITIONS](#) table for each data node in the [NDB](#) cluster. For each such row:
  - The [SUBPARTITION\\_NAME](#), [SUBPARTITION\\_ORDINAL\\_POSITION](#), [SUBPARTITION\\_METHOD](#), [PARTITION\\_EXPRESSION](#), [SUBPARTITION\\_EXPRESSION](#), [CREATE\\_TIME](#), [UPDATE\\_TIME](#), [CHECK\\_TIME](#), [CHECKSUM](#), and [TABLESPACE\\_NAME](#) columns are all [NULL](#).
  - The [PARTITION\\_METHOD](#) is always [AUTO](#).
  - The [NODEGROUP](#) column is [default](#).
  - The [PARTITION\\_COMMENT](#) column is empty.

## 26.3.22 The INFORMATION\_SCHEMA PLUGINS Table

The [PLUGINS](#) table provides information about server plugins.

The `PLUGINS` table has these columns:

- `PLUGIN_NAME`

The name used to refer to the plugin in statements such as `INSTALL PLUGIN` and `UNINSTALL PLUGIN`.

- `PLUGIN_VERSION`

The version from the plugin's general type descriptor.

- `PLUGIN_STATUS`

The plugin status, one of `ACTIVE`, `INACTIVE`, `DISABLED`, `DELETING`, or `DELETED`.

- `PLUGIN_TYPE`

The type of plugin, such as `STORAGE ENGINE`, `INFORMATION_SCHEMA`, or `AUTHENTICATION`.

- `PLUGIN_TYPE_VERSION`

The version from the plugin's type-specific descriptor.

- `PLUGIN_LIBRARY`

The name of the plugin shared library file. This is the name used to refer to the plugin file in statements such as `INSTALL PLUGIN` and `UNINSTALL PLUGIN`. This file is located in the directory named by the `plugin_dir` system variable. If the library name is `NULL`, the plugin is compiled in and cannot be uninstalled with `UNINSTALL PLUGIN`.

- `PLUGIN_LIBRARY_VERSION`

The plugin API interface version.

- `PLUGIN_AUTHOR`

The plugin author.

- `PLUGIN_DESCRIPTION`

A short description of the plugin.

- `PLUGIN_LICENSE`

How the plugin is licensed (for example, [GPL](#)).

- `LOAD_OPTION`

How the plugin was loaded. The value is `OFF`, `ON`, `FORCE`, or `FORCE_PLUS_PERMANENT`. See [Section 5.6.1, “Installing and Uninstalling Plugins”](#).

## Notes

- `PLUGINS` is a nonstandard `INFORMATION_SCHEMA` table.
- For plugins installed with `INSTALL PLUGIN`, the `PLUGIN_NAME` and `PLUGIN_LIBRARY` values are also registered in the `mysql.plugin` table.
- For information about plugin data structures that form the basis of the information in the `PLUGINS` table, see [The MySQL Plugin API](#).

Plugin information is also available from the `SHOW PLUGINS` statement. See [Section 13.7.7.25, “SHOW PLUGINS Statement”](#). These statements are equivalent:

```
SELECT
```

```

PLUGIN_NAME, PLUGIN_STATUS, PLUGIN_TYPE,
PLUGIN_LIBRARY, PLUGIN_LICENSE
FROM INFORMATION_SCHEMA.PLUGINS;

SHOW PLUGINS;

```

### 26.3.23 The INFORMATION\_SCHEMA PROCESSLIST Table

The MySQL process list indicates the operations currently being performed by the set of threads executing within the server. The [PROCESSLIST](#) table is one source of process information. For a comparison of this table with other sources, see [Sources of Process Information](#).

The [PROCESSLIST](#) table has these columns:

- [ID](#)

The connection identifier. This is the same value displayed in the [Id](#) column of the [SHOW PROCESSLIST](#) statement, displayed in the [PROCESSLIST\\_ID](#) column of the Performance Schema [threads](#) table, and returned by the [CONNECTION\\_ID\(\)](#) function within the thread.

- [USER](#)

The MySQL user who issued the statement. A value of [system user](#) refers to a nonclient thread spawned by the server to handle tasks internally, for example, a delayed-row handler thread or an I/O or SQL thread used on replica hosts. For [system user](#), there is no host specified in the [Host](#) column. [unauthenticated user](#) refers to a thread that has become associated with a client connection but for which authentication of the client user has not yet occurred. [event\\_scheduler](#) refers to the thread that monitors scheduled events (see [Section 25.4, “Using the Event Scheduler”](#)).



#### Note

A [USER](#) value of [system user](#) is distinct from the [SYSTEM\\_USER](#) privilege. The former designates internal threads. The latter distinguishes the system user and regular user account categories (see [Section 6.2.11, “Account Categories”](#)).

- [HOST](#)

The host name of the client issuing the statement (except for [system user](#), for which there is no host). The host name for TCP/IP connections is reported in [host\\_name:client\\_port](#) format to make it easier to determine which client is doing what.

- [DB](#)

The default database for the thread, or [NULL](#) if none has been selected.

- [COMMAND](#)

The type of command the thread is executing on behalf of the client, or [Sleep](#) if the session is idle. For descriptions of thread commands, see [Section 8.14, “Examining Server Thread \(Process Information”](#). The value of this column corresponds to the [COM\\_xxx](#) commands of the client/server protocol and [Com\\_xxx](#) status variables. See [Section 5.1.10, “Server Status Variables”](#).

- [TIME](#)

The time in seconds that the thread has been in its current state. For a replica SQL thread, the value is the number of seconds between the timestamp of the last replicated event and the real time of the replica host. See [Section 17.2.3, “Replication Threads”](#).

- [STATE](#)

An action, event, or state that indicates what the thread is doing. For descriptions of [STATE](#) values, see [Section 8.14, “Examining Server Thread \(Process\) Information”](#).

Most states correspond to very quick operations. If a thread stays in a given state for many seconds, there might be a problem that needs to be investigated.

- [INFO](#)

The statement the thread is executing, or `NULL` if it is executing no statement. The statement might be the one sent to the server, or an innermost statement if the statement executes other statements. For example, if a `CALL` statement executes a stored procedure that is executing a `SELECT` statement, the `INFO` value shows the `SELECT` statement.

## Notes

- `PROCESSLIST` is a nonstandard `INFORMATION_SCHEMA` table.
- Like the output from the `SHOW PROCESSLIST` statement, the `PROCESSLIST` table provides information about all threads, even those belonging to other users, if you have the `PROCESS` privilege. Otherwise (without the `PROCESS` privilege), nonanonymous users have access to information about their own threads but not threads for other users, and anonymous users have no access to thread information.
- If an SQL statement refers to the `PROCESSLIST` table, MySQL populates the entire table once, when statement execution begins, so there is read consistency during the statement. There is no read consistency for a multi-statement transaction.

The following statements are equivalent:

```
SELECT * FROM INFORMATION_SCHEMA.PROCESSLIST  
SHOW FULL PROCESSLIST
```

### 26.3.24 The INFORMATION\_SCHEMA PROFILING Table

The `PROFILING` table provides statement profiling information. Its contents correspond to the information produced by the `SHOW PROFILE` and `SHOW PROFILES` statements (see [Section 13.7.7.30, “SHOW PROFILE Statement”](#)). The table is empty unless the `profiling` session variable is set to 1.



#### Note

This table is deprecated; expect it to be removed in a future MySQL release. Use the `Performance Schema` instead; see [Section 27.19.1, “Query Profiling Using Performance Schema”](#).

The `PROFILING` table has these columns:

- [QUERY\\_ID](#)

A numeric statement identifier.

- [SEQ](#)

A sequence number indicating the display order for rows with the same `QUERY_ID` value.

- [STATE](#)

The profiling state to which the row measurements apply.

- [DURATION](#)

How long statement execution remained in the given state, in seconds.

- [CPU\\_USER, CPU\\_SYSTEM](#)

User and system CPU use, in seconds.

- [CONTEXT\\_VOLUNTARY](#), [CONTEXT\\_INVOLUNTARY](#)

How many voluntary and involuntary context switches occurred.

- [BLOCK\\_OPS\\_IN](#), [BLOCK\\_OPS\\_OUT](#)

The number of block input and output operations.

- [MESSAGES\\_SENT](#), [MESSAGES\\_RECEIVED](#)

The number of communication messages sent and received.

- [PAGEFAULTS\\_MAJOR](#), [PAGEFAULTS\\_MINOR](#)

The number of major and minor page faults.

- [SWAPS](#)

How many swaps occurred.

- [SOURCE\\_FUNCTION](#), [SOURCE\\_FILE](#), and [SOURCE\\_LINE](#)

Information indicating where in the source code the profiled state executes.

## Notes

- [PROFILING](#) is a nonstandard [INFORMATION\\_SCHEMA](#) table.

Profiling information is also available from the [SHOW PROFILE](#) and [SHOW PROFILES](#) statements. See [Section 13.7.7.30, “SHOW PROFILE Statement”](#). For example, the following queries are equivalent:

```
SHOW PROFILE FOR QUERY 2;

SELECT STATE, FORMAT(DURATION, 6) AS DURATION
FROM INFORMATION_SCHEMA.PROFILING
WHERE QUERY_ID = 2 ORDER BY SEQ;
```

## 26.3.25 The INFORMATION\_SCHEMA REFERENTIAL\_CONSTRAINTS Table

The [REFERENTIAL\\_CONSTRAINTS](#) table provides information about foreign keys.

The [REFERENTIAL\\_CONSTRAINTS](#) table has these columns:

- [CONSTRAINT\\_CATALOG](#)

The name of the catalog to which the constraint belongs. This value is always [def](#).

- [CONSTRAINT\\_SCHEMA](#)

The name of the schema (database) to which the constraint belongs.

- [CONSTRAINT\\_NAME](#)

The name of the constraint.

- [UNIQUE\\_CONSTRAINT\\_CATALOG](#)

The name of the catalog containing the unique constraint that the constraint references. This value is always [def](#).

- [UNIQUE\\_CONSTRAINT\\_SCHEMA](#)

The name of the schema containing the unique constraint that the constraint references.

- [UNIQUE\\_CONSTRAINT\\_NAME](#)

The name of the unique constraint that the constraint references.

- [MATCH\\_OPTION](#)

The value of the constraint [MATCH](#) attribute. The only valid value at this time is [NONE](#).

- [UPDATE\\_RULE](#)

The value of the constraint [ON UPDATE](#) attribute. The possible values are [CASCADE](#), [SET NULL](#), [SET DEFAULT](#), [RESTRICT](#), [NO ACTION](#).

- [DELETE\\_RULE](#)

The value of the constraint [ON DELETE](#) attribute. The possible values are [CASCADE](#), [SET NULL](#), [SET DEFAULT](#), [RESTRICT](#), [NO ACTION](#).

- [TABLE\\_NAME](#)

The name of the table. This value is the same as in the [TABLE\\_CONSTRAINTS](#) table.

- [REFERENCED\\_TABLE\\_NAME](#)

The name of the table referenced by the constraint.

### 26.3.26 The INFORMATION\_SCHEMA RESOURCE\_GROUPS Table

The [RESOURCE\\_GROUPS](#) table provides access to information about resource groups. For general discussion of the resource group capability, see [Section 5.1.16, “Resource Groups”](#).

You can see information only for columns for which you have some privilege.

The [RESOURCE\\_GROUPS](#) table has these columns:

- [RESOURCE\\_GROUP\\_NAME](#)

The name of the resource group.

- [RESOURCE\\_GROUP\\_TYPE](#)

The resource group type, either [SYSTEM](#) or [USER](#).

- [RESOURCE\\_GROUP\\_ENABLED](#)

Whether the resource group is enabled (1) or disabled (0);

- [VCPU\\_IDS](#)

The CPU affinity; that is, the set of virtual CPUs that the resource group can use. The value is a list of comma-separated CPU numbers or ranges.

- [THREAD\\_PRIORITY](#)

The priority for threads assigned to the resource group. The priority ranges from -20 (highest priority) to 19 (lowest priority). System resource groups have a priority that ranges from -20 to 0. User resource groups have a priority that ranges from 0 to 19.

### 26.3.27 The INFORMATION\_SCHEMA ROLE\_COLUMN\_GRANTS Table

The [ROLE\\_COLUMN\\_GRANTS](#) table (available as of MySQL 8.0.19) provides information about the column privileges for roles that are available to or granted by the currently enabled roles.

The [ROLE\\_COLUMN\\_GRANTS](#) table has these columns:

- [GRANTOR](#)

The user name part of the account that granted the role.

- [GRANTOR](#)

The host name part of the account that granted the role.

- [GRANTEE](#)

The user name part of the account to which the role is granted.

- [GRANTEE\\_HOST](#)

The host name part of the account to which the role is granted.

- [TABLE\\_CATALOG](#)

The name of the catalog to which the role applies. This value is always [def](#).

- [TABLE\\_SCHEMA](#)

The name of the schema (database) to which the role applies.

- [TABLE\\_NAME](#)

The name of the table to which the role applies.

- [COLUMN\\_NAME](#)

The name of the column to which the role applies.

- [PRIVILEGE\\_TYPE](#)

The privilege granted. The value can be any privilege that can be granted at the column level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per column privilege held by the grantee.

- [IS\\_GRANTABLE](#)

[YES](#) or [NO](#), depending on whether the role is grantable to other accounts.

### 26.3.28 The INFORMATION\_SCHEMA ROLE\_ROUTINE\_GRANTS Table

The [ROLE\\_ROUTINE\\_GRANTS](#) table (available as of MySQL 8.0.19) provides information about the routine privileges for roles that are available to or granted by the currently enabled roles.

The [ROLE\\_ROUTINE\\_GRANTS](#) table has these columns:

- [GRANTOR](#)

The user name part of the account that granted the role.

- [GRANTOR\\_HOST](#)

The host name part of the account that granted the role.

- [GRANTEE](#)

The user name part of the account to which the role is granted.

- [GRANTEE\\_HOST](#)

The host name part of the account to which the role is granted.

- [SPECIFIC\\_CATALOG](#)

The name of the catalog to which the routine belongs. This value is always `def`.

- `SPECIFIC_SCHEMA`

The name of the schema (database) to which the routine belongs.

- `SPECIFIC_NAME`

The name of the routine.

- `ROUTINE_CATALOG`

The name of the catalog to which the routine belongs. This value is always `def`.

- `ROUTINE_SCHEMA`

The name of the schema (database) to which the routine belongs.

- `ROUTINE_NAME`

The name of the routine.

- `PRIVILEGE_TYPE`

The privilege granted. The value can be any privilege that can be granted at the routine level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per column privilege held by the grantee.

- `IS_GRANTABLE`

`YES` or `NO`, depending on whether the role is grantable to other accounts.

### 26.3.29 The INFORMATION\_SCHEMA ROLE\_TABLE\_GRANTS Table

The `ROLE_TABLE_GRANTS` table (available as of MySQL 8.0.19) provides information about the table privileges for roles that are available to or granted by the currently enabled roles.

The `ROLE_TABLE_GRANTS` table has these columns:

- `GRANTOR`

The user name part of the account that granted the role.

- `GRANTOR_HOST`

The host name part of the account that granted the role.

- `GRANTEE`

The user name part of the account to which the role is granted.

- `GRANTEE_HOST`

The host name part of the account to which the role is granted.

- `TABLE_CATALOG`

The name of the catalog to which the role applies. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the role applies.

- `TABLE_NAME`

The name of the table to which the role applies.

- [PRIVILEGE\\_TYPE](#)

The privilege granted. The value can be any privilege that can be granted at the table level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per column privilege held by the grantee.

- [IS\\_GRANTABLE](#)

[YES](#) or [NO](#), depending on whether the role is grantable to other accounts.

### 26.3.30 The INFORMATION\_SCHEMA ROUTINES Table

The [ROUTINES](#) table provides information about stored routines (stored procedures and stored functions). The [ROUTINES](#) table does not include built-in (native) functions or loadable functions.

The [ROUTINES](#) table has these columns:

- [SPECIFIC\\_NAME](#)

The name of the routine.

- [ROUTINE\\_CATALOG](#)

The name of the catalog to which the routine belongs. This value is always [def](#).

- [ROUTINE\\_SCHEMA](#)

The name of the schema (database) to which the routine belongs.

- [ROUTINE\\_NAME](#)

The name of the routine.

- [ROUTINE\\_TYPE](#)

[PROCEDURE](#) for stored procedures, [FUNCTION](#) for stored functions.

- [DATA\\_TYPE](#)

If the routine is a stored function, the return value data type. If the routine is a stored procedure, this value is empty.

The [DATA\\_TYPE](#) value is the type name only with no other information. The [DTD\\_IDENTIFIER](#) value contains the type name and possibly other information such as the precision or length.

- [CHARACTER\\_MAXIMUM\\_LENGTH](#)

For stored function string return values, the maximum length in characters. If the routine is a stored procedure, this value is [NULL](#).

- [CHARACTER\\_OCTET\\_LENGTH](#)

For stored function string return values, the maximum length in bytes. If the routine is a stored procedure, this value is [NULL](#).

- [NUMERIC\\_PRECISION](#)

For stored function numeric return values, the numeric precision. If the routine is a stored procedure, this value is [NULL](#).

- [NUMERIC\\_SCALE](#)

For stored function numeric return values, the numeric scale. If the routine is a stored procedure, this value is [NULL](#).

- [DATETIME\\_PRECISION](#)

For stored function temporal return values, the fractional seconds precision. If the routine is a stored procedure, this value is [NULL](#).

- [CHARACTER\\_SET\\_NAME](#)

For stored function character string return values, the character set name. If the routine is a stored procedure, this value is [NULL](#).

- [COLLATION\\_NAME](#)

For stored function character string return values, the collation name. If the routine is a stored procedure, this value is [NULL](#).

- [DTD\\_IDENTIFIER](#)

If the routine is a stored function, the return value data type. If the routine is a stored procedure, this value is empty.

The [DATA\\_TYPE](#) value is the type name only with no other information. The [DTD\\_IDENTIFIER](#) value contains the type name and possibly other information such as the precision or length.

- [ROUTINE\\_BODY](#)

The language used for the routine definition. This value is always [SQL](#).

- [ROUTINE\\_DEFINITION](#)

The text of the SQL statement executed by the routine.

- [EXTERNAL\\_NAME](#)

This value is always [NULL](#).

- [EXTERNAL\\_LANGUAGE](#)

The language of the stored routine. The value is read from the [external\\_language](#) column of the [mysql.routines](#) data dictionary table.

- [PARAMETER\\_STYLE](#)

This value is always [SQL](#).

- [IS\\_DETERMINISTIC](#)

[YES](#) or [NO](#), depending on whether the routine is defined with the [DETERMINISTIC](#) characteristic.

- [SQL\\_DATA\\_ACCESS](#)

The data access characteristic for the routine. The value is one of [CONTAINS\\_SQL](#), [NO\\_SQL](#), [READS\\_SQL\\_DATA](#), or [MODIFIES\\_SQL\\_DATA](#).

- [SQL\\_PATH](#)

This value is always [NULL](#).

- [SECURITY\\_TYPE](#)

The routine [SQL\\_SECURITY](#) characteristic. The value is one of [DEFINER](#) or [INVOKER](#).

- [CREATED](#)

The date and time when the routine was created. This is a [TIMESTAMP](#) value.

- [LAST\\_ALTERED](#)

The date and time when the routine was last modified. This is a [TIMESTAMP](#) value. If the routine has not been modified since its creation, this value is the same as the [CREATED](#) value.

- [SQL\\_MODE](#)

The SQL mode in effect when the routine was created or altered, and under which the routine executes. For the permitted values, see [Section 5.1.11, “Server SQL Modes”](#).

- [ROUTINE\\_COMMENT](#)

The text of the comment, if the routine has one. If not, this value is empty.

- [DEFINER](#)

The account named in the [DEFINER](#) clause (often the user who created the routine), in '*user\_name*' '@' '*host\_name*' format.

- [CHARACTER\\_SET\\_CLIENT](#)

The session value of the [character\\_set\\_client](#) system variable when the routine was created.

- [COLLATION\\_CONNECTION](#)

The session value of the [collation\\_connection](#) system variable when the routine was created.

- [DATABASE\\_COLLATION](#)

The collation of the database with which the routine is associated.

## Notes

- To see information about a routine, you must be the user named as the routine [DEFINER](#), have the [SHOW\\_ROUTINE](#) privilege, have the [SELECT](#) privilege at the global level, or have the [CREATE ROUTINE](#), [ALTER ROUTINE](#), or [EXECUTE](#) privilege granted at a scope that includes the routine. The [ROUTINE\\_DEFINITION](#) column is [NULL](#) if you have only [CREATE ROUTINE](#), [ALTER ROUTINE](#), or [EXECUTE](#).
- Information about stored function return values is also available in the [PARAMETERS](#) table. The return value row for a stored function can be identified as the row that has an [ORDINAL\\_POSITION](#) value of 0.

### 26.3.31 The INFORMATION\_SCHEMA SCHEMATA Table

A schema is a database, so the [SCHEMATA](#) table provides information about databases.

The [SCHEMATA](#) table has these columns:

- [CATALOG\\_NAME](#)

The name of the catalog to which the schema belongs. This value is always [def](#).

- [SCHEMA\\_NAME](#)

The name of the schema.

- [DEFAULT\\_CHARACTER\\_SET\\_NAME](#)

The schema default character set.

- `DEFAULT_COLLATION_NAME`

The schema default collation.

- `SQL_PATH`

This value is always `NULL`.

- `DEFAULT_ENCRYPTION`

The schema default encryption. This column was added in MySQL 8.0.16.

Schema names are also available from the `SHOW DATABASES` statement. See [Section 13.7.7.14, “SHOW DATABASES Statement”](#). The following statements are equivalent:

```
SELECT SCHEMA_NAME AS `Database`
  FROM INFORMATION_SCHEMA.SCHEMATA
 [WHERE SCHEMA_NAME LIKE 'wild']

SHOW DATABASES
 [LIKE 'wild']
```

You see only those databases for which you have some kind of privilege, unless you have the global `SHOW DATABASES` privilege.



### Caution

Because any static global privilege is considered a privilege for all databases, any static global privilege enables a user to see all database names with `SHOW DATABASES` or by examining the `SCHEMATA` table of `INFORMATION_SCHEMA`, except databases that have been restricted at the database level by partial revokes.

## Notes

- The `SCHEMATA_EXTENSIONS` table augments the `SCHEMATA` table with information about schema options.

### 26.3.32 The INFORMATION\_SCHEMA SCHEMATA\_EXTENSIONS Table

The `SCHEMATA_EXTENSIONS` table (available as of MySQL 8.0.22) augments the `SCHEMATA` table with information about schema options.

The `SCHEMATA_EXTENSIONS` table has these columns:

- `CATALOG_NAME`

The name of the catalog to which the schema belongs. This value is always `def`.

- `SCHEMA_NAME`

The name of the schema.

- `OPTIONS`

The options for the schema. If the schema is read only, the value contains `READ ONLY=1`. If the schema is not read only, no `READ ONLY` option appears.

## Example

```
mysql> ALTER SCHEMA mydb READ ONLY = 1;
mysql> SELECT * FROM INFORMATION_SCHEMA.SCHEMATA_EXTENSIONS
      WHERE SCHEMA_NAME = 'mydb';
+-----+-----+-----+
| CATALOG_NAME | SCHEMA_NAME | OPTIONS |
+-----+-----+-----+
```

```
+-----+-----+-----+
| def      | mydb      | READ ONLY=1 |
+-----+-----+-----+
mysql> ALTER SCHEMA mydb READ ONLY = 0;
mysql> SELECT * FROM INFORMATION_SCHEMA.SCHEMATA_EXTENSIONS
      WHERE SCHEMA_NAME = 'mydb';
+-----+-----+-----+
| CATALOG_NAME | SCHEMA_NAME | OPTIONS   |
+-----+-----+-----+
| def          | mydb        |           |
+-----+-----+-----+
```

## Notes

- `SCHEMATA_EXTENSIONS` is a nonstandard `INFORMATION_SCHEMA` table.

### 26.3.33 The INFORMATION\_SCHEMA SCHEMA\_PRIVILEGES Table

The `SCHEMA_PRIVILEGES` table provides information about schema (database) privileges. It takes its values from the `mysql.db` system table.

The `SCHEMA_PRIVILEGES` table has these columns:

- `GRANTEE`

The name of the account to which the privilege is granted, in '`user_name`'@'`host_name`' format.

- `TABLE_CATALOG`

The name of the catalog to which the schema belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema.

- `PRIVILEGE_TYPE`

The privilege granted. The value can be any privilege that can be granted at the schema level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per schema privilege held by the grantee.

- `IS_GRANTABLE`

`YES` if the user has the `GRANT OPTION` privilege, `NO` otherwise. The output does not list `GRANT OPTION` as a separate row with `PRIVILEGE_TYPE='GRANT OPTION'`.

## Notes

- `SCHEMA_PRIVILEGES` is a nonstandard `INFORMATION_SCHEMA` table.

The following statements are *not* equivalent:

```
SELECT ... FROM INFORMATION_SCHEMA.SCHEMA_PRIVILEGES
SHOW GRANTS ...
```

### 26.3.34 The INFORMATION\_SCHEMA STATISTICS Table

The `STATISTICS` table provides information about table indexes.

Columns in `STATISTICS` that represent table statistics hold cached values. The `information_schema_stats_expiry` system variable defines the period of time before cached table statistics expire. The default is 86400 seconds (24 hours). If there are no cached statistics or statistics have expired, statistics are retrieved from storage engines when

querying table statistics columns. To update cached values at any time for a given table, use `ANALYZE TABLE`. To always retrieve the latest statistics directly from storage engines, set `information_schema_stats_expiry=0`. For more information, see [Section 8.2.3, “Optimizing INFORMATION\\_SCHEMA Queries”](#).

**Note**

If the `innodb_read_only` system variable is enabled, `ANALYZE TABLE` may fail because it cannot update statistics tables in the data dictionary, which use InnoDB. For `ANALYZE TABLE` operations that update the key distribution, failure may occur even if the operation updates the table itself (for example, if it is a MyISAM table). To obtain the updated distribution statistics, set `information_schema_stats_expiry=0`.

The `STATISTICS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table containing the index belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table containing the index belongs.

- `TABLE_NAME`

The name of the table containing the index.

- `NON_UNIQUE`

0 if the index cannot contain duplicates, 1 if it can.

- `INDEX_SCHEMA`

The name of the schema (database) to which the index belongs.

- `INDEX_NAME`

The name of the index. If the index is the primary key, the name is always `PRIMARY`.

- `SEQ_IN_INDEX`

The column sequence number in the index, starting with 1.

- `COLUMN_NAME`

The column name. See also the description for the `EXPRESSION` column.

- `COLLATION`

How the column is sorted in the index. This can have values `A` (ascending), `D` (descending), or `NULL` (not sorted).

- `CARDINALITY`

An estimate of the number of unique values in the index. To update this number, run `ANALYZE TABLE` or (for MyISAM tables) `myisamchk -a`.

`CARDINALITY` is counted based on statistics stored as integers, so the value is not necessarily exact even for small tables. The higher the cardinality, the greater the chance that MySQL uses the index when doing joins.

- `SUB_PART`

The index prefix. That is, the number of indexed characters if the column is only partly indexed, `NULL` if the entire column is indexed.



### Note

Prefix *limits* are measured in bytes. However, prefix *lengths* for index specifications in `CREATE TABLE`, `ALTER TABLE`, and `CREATE INDEX` statements are interpreted as number of characters for nonbinary string types (`CHAR`, `VARCHAR`, `TEXT`) and number of bytes for binary string types (`BINARY`, `VARBINARY`, `BLOB`). Take this into account when specifying a prefix length for a nonbinary string column that uses a multibyte character set.

For additional information about index prefixes, see [Section 8.3.5, “Column Indexes”](#), and [Section 13.1.15, “CREATE INDEX Statement”](#).

- `PACKED`

Indicates how the key is packed. `NULL` if it is not.

- `NULLABLE`

Contains `YES` if the column may contain `NULL` values and `' '` if not.

- `INDEX_TYPE`

The index method used (`BTREE`, `FULLTEXT`, `HASH`, `RTREE`).

- `COMMENT`

Information about the index not described in its own column, such as `disabled` if the index is disabled.

- `INDEX_COMMENT`

Any comment provided for the index with a `COMMENT` attribute when the index was created.

- `IS_VISIBLE`

Whether the index is visible to the optimizer. See [Section 8.3.12, “Invisible Indexes”](#).

- `EXPRESSION`

MySQL 8.0.13 and higher supports functional key parts (see [Functional Key Parts](#)), which affects both the `COLUMN_NAME` and `EXPRESSION` columns:

- For a nonfunctional key part, `COLUMN_NAME` indicates the column indexed by the key part and `EXPRESSION` is `NULL`.
- For a functional key part, `COLUMN_NAME` column is `NULL` and `EXPRESSION` indicates the expression for the key part.

## Notes

- There is no standard `INFORMATION_SCHEMA` table for indexes. The MySQL column list is similar to what SQL Server 2000 returns for `sp_statistics`, except that `QUALIFIER` and `OWNER` are replaced with `CATALOG` and `SCHEMA`, respectively.

Information about table indexes is also available from the `SHOW INDEX` statement. See [Section 13.7.7.22, “SHOW INDEX Statement”](#). The following statements are equivalent:

```
SELECT * FROM INFORMATION_SCHEMA.STATISTICS
  WHERE table_name = 'tbl_name'
    AND table_schema = 'db_name'
```

```
SHOW INDEX  
  FROM tbl_name  
  FROM db_name
```

In MySQL 8.0.30 and later, information about generated invisible primary key columns is visible in this table by default. You can cause such information to be hidden by setting `show_gipk_in_create_table_and_information_schema = OFF`. For more information, see [Section 13.1.20.11, “Generated Invisible Primary Keys”](#).

### 26.3.35 The INFORMATION\_SCHEMA ST\_GEOMETRY\_COLUMNS Table

The `ST_GEOMETRY_COLUMNS` table provides information about table columns that store spatial data. This table is based on the SQL/MM (ISO/IEC 13249-3) standard, with extensions as noted. MySQL implements `ST_GEOMETRY_COLUMNS` as a view on the `INFORMATION_SCHEMA COLUMNS` table.

The `ST_GEOMETRY_COLUMNS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table containing the column belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table containing the column belongs.

- `TABLE_NAME`

The name of the table containing the column.

- `COLUMN_NAME`

The name of the column.

- `SRS_NAME`

The spatial reference system (SRS) name.

- `SRS_ID`

The spatial reference system ID (SRID).

- `GEOMETRY_TYPE_NAME`

The column data type. Permitted values are: `geometry`, `point`, `linestring`, `polygon`, `multipoint`, `multilinestring`, `multipolygon`, `geometrycollection`. This column is a MySQL extension to the standard.

### 26.3.36 The INFORMATION\_SCHEMA ST\_SPATIAL\_REFERENCE\_SYSTEMS Table

The `ST_SPATIAL_REFERENCE_SYSTEMS` table provides information about available spatial reference systems (SRSs) for spatial data. This table is based on the SQL/MM (ISO/IEC 13249-3) standard.

Entries in the `ST_SPATIAL_REFERENCE_SYSTEMS` table are based on the [European Petroleum Survey Group](#) (EPSG) data set, except for SRID 0, which corresponds to a special SRS used in MySQL that represents an infinite flat Cartesian plane with no units assigned to its axes. For additional information about SRSs, see [Section 11.4.5, “Spatial Reference System Support”](#).

The `ST_SPATIAL_REFERENCE_SYSTEMS` table has these columns:

- `SRS_NAME`

The spatial reference system name. This value is unique.

- [SRS\\_ID](#)

The spatial reference system numeric ID. This value is unique.

[SRS\\_ID](#) values represent the same kind of values as the SRID of geometry values or passed as the SRID argument to spatial functions. SRID 0 (the unitless Cartesian plane) is special. It is always a legal spatial reference system ID and can be used in any computations on spatial data that depend on SRID values.

- [ORGANIZATION](#)

The name of the organization that defined the coordinate system on which the spatial reference system is based.

- [ORGANIZATION\\_COORDSYS\\_ID](#)

The numeric ID given to the spatial reference system by the organization that defined it.

- [DEFINITION](#)

The spatial reference system definition. [DEFINITION](#) values are WKT values, represented as specified in the [Open Geospatial Consortium](#) document [OGC 12-063r5](#).

SRS definition parsing occurs on demand when definitions are needed by GIS functions. Parsed definitions are stored in the data dictionary cache to enable reuse and avoid incurring parsing overhead for every statement that needs SRS information.

- [DESCRIPTION](#)

The spatial reference system description.

## Notes

- The [SRS\\_NAME](#), [ORGANIZATION](#), [ORGANIZATION\\_COORDSYS\\_ID](#), and [DESCRIPTION](#) columns contain information that may be of interest to users, but they are not used by MySQL.

## Example

```
mysql> SELECT * FROM ST_SPATIAL_REFERENCE_SYSTEMS
      WHERE SRS_ID = 4326\G
*****
               1. row ****
      SRS_NAME: WGS 84
      SRS_ID: 4326
      ORGANIZATION: EPSG
      ORGANIZATION_COORDSYS_ID: 4326
      DEFINITION: GEOGCS["WGS 84",DATUM["World Geodetic System 1984",
          SPHEROID["WGS 84",6378137,298.257223563,
          AUTHORITY["EPSG","7030"]],AUTHORITY["EPSG","6326"]],
          PRIMEM["Greenwich",0,AUTHORITY["EPSG","8901"]],
          UNIT["degree",0.017453292519943278,
          AUTHORITY["EPSG","9122"]],
          AXIS["Lat",NORTH],AXIS["Long",EAST],
          AUTHORITY["EPSG","4326"]]
      DESCRIPTION:
```

This entry describes the SRS used for GPS systems. It has a name ([SRS\\_NAME](#)) of WGS 84 and an ID ([SRS\\_ID](#)) of 4326, which is the ID used by the [European Petroleum Survey Group](#) (EPSG).

The [DEFINITION](#) values for projected and geographic SRSs begin with [PROJCS](#) and [GEOGCS](#), respectively. The definition for SRID 0 is special and has an empty [DEFINITION](#) value. The following query determines how many entries in the [ST\\_SPATIAL\\_REFERENCE\\_SYSTEMS](#) table correspond to projected, geographic, and other SRSs, based on [DEFINITION](#) values:

```
mysql> SELECT
      COUNT(*),
      CASE LEFT(DEFINITION, 6)
```

```

WHEN 'PROJCS' THEN 'Projected'
WHEN 'GEOGCS' THEN 'Geographic'
ELSE 'Other'
END AS SRS_TYPE
FROM INFORMATION_SCHEMA.ST_SPATIAL_REFERENCE_SYSTEMS
GROUP BY SRS_TYPE;
+-----+-----+
| COUNT(*) | SRS_TYPE   |
+-----+-----+
|       1 | Other      |
|    4668 | Projected  |
|     483 | Geographic |
+-----+-----+

```

To enable manipulation of SRS entries stored in the data dictionary, MySQL provides these SQL statements:

- [CREATE SPATIAL REFERENCE SYSTEM](#): See [Section 13.1.19, “CREATE SPATIAL REFERENCE SYSTEM Statement”](#). The description for this statement includes additional information about SRS components.
- [DROP SPATIAL REFERENCE SYSTEM](#): See [Section 13.1.31, “DROP SPATIAL REFERENCE SYSTEM Statement”](#).

### 26.3.37 The INFORMATION\_SCHEMA ST\_UNITS\_OF\_MEASURE Table

The `ST_UNITS_OF_MEASURE` table (available as of MySQL 8.0.14) provides information about acceptable units for the `ST_Distance()` function.

The `ST_UNITS_OF_MEASURE` table has these columns:

- `UNIT_NAME`  
The name of the unit.
- `UNIT_TYPE`  
The unit type (for example, `LINEAR`).
- `CONVERSION_FACTOR`  
A conversion factor used for internal calculations.
- `DESCRIPTION`  
A description of the unit.

### 26.3.38 The INFORMATION\_SCHEMA TABLES Table

The `TABLES` table provides information about tables in databases.

Columns in `TABLES` that represent table statistics hold cached values. The `information_schema_stats_expiry` system variable defines the period of time before cached table statistics expire. The default is 86400 seconds (24 hours). If there are no cached statistics or statistics have expired, statistics are retrieved from storage engines when querying table statistics columns. To update cached values at any time for a given table, use `ANALYZE TABLE`. To always retrieve the latest statistics directly from storage engines, set `information_schema_stats_expiry` to 0. For more information, see [Section 8.2.3, “Optimizing INFORMATION\\_SCHEMA Queries”](#).



#### Note

If the `innodb_read_only` system variable is enabled, `ANALYZE TABLE` may fail because it cannot update statistics tables in the data dictionary, which use InnoDB. For `ANALYZE TABLE` operations that update the key distribution, failure may occur even if the operation updates the table itself (for

example, if it is a [MyISAM](#) table). To obtain the updated distribution statistics, set `information_schema_stats_expiry=0`.

The `TABLES` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table belongs.

- `TABLE_NAME`

The name of the table.

- `TABLE_TYPE`

`BASE_TABLE` for a table, `VIEW` for a view, or `SYSTEM VIEW` for an [INFORMATION\\_SCHEMA](#) table.

The `TABLES` table does not list `TEMPORARY` tables.

- `ENGINE`

The storage engine for the table. See [Chapter 15, The InnoDB Storage Engine](#), and [Chapter 16, Alternative Storage Engines](#).

For partitioned tables, `ENGINE` shows the name of the storage engine used by all partitions.

- `VERSION`

This column is unused. With the removal of `.frm` files in MySQL 8.0, this column now reports a hardcoded value of `10`, which is the last `.frm` file version used in MySQL 5.7.

- `ROW_FORMAT`

The row-storage format ([Fixed](#), [Dynamic](#), [Compressed](#), [Redundant](#), [Compact](#)). For [MyISAM](#) tables, [Dynamic](#) corresponds to what `myisamchk -dvv` reports as [Packed](#).

- `TABLE_ROWS`

The number of rows. Some storage engines, such as [MyISAM](#), store the exact count. For other storage engines, such as [InnoDB](#), this value is an approximation, and may vary from the actual value by as much as 40% to 50%. In such cases, use `SELECT COUNT(*)` to obtain an accurate count.

`TABLE_ROWS` is `NULL` for [INFORMATION\\_SCHEMA](#) tables.

For [InnoDB](#) tables, the row count is only a rough estimate used in SQL optimization. (This is also true if the [InnoDB](#) table is partitioned.)

- `AVG_ROW_LENGTH`

The average row length.

- `DATA_LENGTH`

For [MyISAM](#), `DATA_LENGTH` is the length of the data file, in bytes.

For [InnoDB](#), `DATA_LENGTH` is the approximate amount of space allocated for the clustered index, in bytes. Specifically, it is the clustered index size, in pages, multiplied by the [InnoDB](#) page size.

Refer to the notes at the end of this section for information regarding other storage engines.

- [MAX\\_DATA\\_LENGTH](#)

For [MyISAM](#), [MAX\\_DATA\\_LENGTH](#) is maximum length of the data file. This is the total number of bytes of data that can be stored in the table, given the data pointer size used.

Unused for [InnoDB](#).

Refer to the notes at the end of this section for information regarding other storage engines.

- [INDEX\\_LENGTH](#)

For [MyISAM](#), [INDEX\\_LENGTH](#) is the length of the index file, in bytes.

For [InnoDB](#), [INDEX\\_LENGTH](#) is the approximate amount of space allocated for non-clustered indexes, in bytes. Specifically, it is the sum of non-clustered index sizes, in pages, multiplied by the [InnoDB](#) page size.

Refer to the notes at the end of this section for information regarding other storage engines.

- [DATA\\_FREE](#)

The number of allocated but unused bytes.

[InnoDB](#) tables report the free space of the tablespace to which the table belongs. For a table located in the shared tablespace, this is the free space of the shared tablespace. If you are using multiple tablespaces and the table has its own tablespace, the free space is for only that table. Free space means the number of bytes in completely free extents minus a safety margin. Even if free space displays as 0, it may be possible to insert rows as long as new extents need not be allocated.

For NDB Cluster, [DATA\\_FREE](#) shows the space allocated on disk for, but not used by, a Disk Data table or fragment on disk. (In-memory data resource usage is reported by the [DATA\\_LENGTH](#) column.)

For partitioned tables, this value is only an estimate and may not be absolutely correct. A more accurate method of obtaining this information in such cases is to query the [INFORMATION\\_SCHEMA PARTITIONS](#) table, as shown in this example:

```
SELECT SUM(DATA_FREE)
      FROM INFORMATION_SCHEMA.PARTITIONS
     WHERE TABLE_SCHEMA = 'mydb'
       AND TABLE_NAME    = 'mytable';
```

For more information, see [Section 26.3.21, “The INFORMATION\\_SCHEMA PARTITIONS Table”](#).

- [AUTO\\_INCREMENT](#)

The next [AUTO\\_INCREMENT](#) value.

- [CREATE\\_TIME](#)

When the table was created.

- [UPDATE\\_TIME](#)

When the data file was last updated. For some storage engines, this value is [NULL](#). For example, [InnoDB](#) stores multiple tables in its [system tablespace](#) and the data file timestamp does not apply. Even with [file-per-table](#) mode with each [InnoDB](#) table in a separate [.ibd](#) file, [change buffering](#) can delay the write to the data file, so the file modification time is different from the time of the last insert, update, or delete. For [MyISAM](#), the data file timestamp is used; however, on Windows the timestamp is not updated by updates, so the value is inaccurate.

[UPDATE\\_TIME](#) displays a timestamp value for the last [UPDATE](#), [INSERT](#), or [DELETE](#) performed on [InnoDB](#) tables that are not partitioned. For MVCC, the timestamp value reflects the [COMMIT](#) time,

which is considered the last update time. Timestamps are not persisted when the server is restarted or when the table is evicted from the [InnoDB](#) data dictionary cache.

- [CHECK\\_TIME](#)

When the table was last checked. Not all storage engines update this time, in which case, the value is always [NULL](#).

For partitioned [InnoDB](#) tables, [CHECK\\_TIME](#) is always [NULL](#).

- [TABLE\\_COLLATION](#)

The table default collation. The output does not explicitly list the table default character set, but the collation name begins with the character set name.

- [CHECKSUM](#)

The live checksum value, if any.

- [CREATE\\_OPTIONS](#)

Extra options used with [CREATE\\_TABLE](#).

[CREATE\\_OPTIONS](#) shows [partitioned](#) for a partitioned table.

Prior to MySQL 8.0.16, [CREATE\\_OPTIONS](#) shows the [ENCRYPTION](#) clause specified for tables created in file-per-table tablespaces. As of MySQL 8.0.16, it shows the encryption clause for file-per-table tablespaces if the table is encrypted or if the specified encryption differs from the schema encryption. The encryption clause is not shown for tables created in general tablespaces. To identify encrypted file-per-table and general tablespaces, query the [INNODB\\_TABLESPACES\\_ENCRYPTION](#) column.

When creating a table with [strict mode](#) disabled, the storage engine's default row format is used if the specified row format is not supported. The actual row format of the table is reported in the [ROW\\_FORMAT](#) column. [CREATE\\_OPTIONS](#) shows the row format that was specified in the [CREATE\\_TABLE](#) statement.

When altering the storage engine of a table, table options that are not applicable to the new storage engine are retained in the table definition to enable reverting the table with its previously defined options to the original storage engine, if necessary. The [CREATE\\_OPTIONS](#) column may show retained options.

- [TABLE\\_COMMENT](#)

The comment used when creating the table (or information as to why MySQL could not access the table information).

## Notes

- For [NDB](#) tables, the output of this statement shows appropriate values for the [AVG\\_ROW\\_LENGTH](#) and [DATA\\_LENGTH](#) columns, with the exception that [BLOB](#) columns are not taken into account.
- For [NDB](#) tables, [DATA\\_LENGTH](#) includes data stored in main memory only; the [MAX\\_DATA\\_LENGTH](#) and [DATA\\_FREE](#) columns apply to Disk Data.
- For NDB Cluster Disk Data tables, [MAX\\_DATA\\_LENGTH](#) shows the space allocated for the disk part of a Disk Data table or fragment. (In-memory data resource usage is reported by the [DATA\\_LENGTH](#) column.)
- For [MEMORY](#) tables, the [DATA\\_LENGTH](#), [MAX\\_DATA\\_LENGTH](#), and [INDEX\\_LENGTH](#) values approximate the actual amount of allocated memory. The allocation algorithm reserves memory in large amounts to reduce the number of allocation operations.

- For views, most `TABLES` columns are 0 or `NULL` except that `TABLE_NAME` indicates the view name, `CREATE_TIME` indicates the creation time, and `TABLE_COMMENT` says `VIEW`.

Table information is also available from the `SHOW TABLE STATUS` and `SHOW TABLES` statements. See [Section 13.7.7.38, “SHOW TABLE STATUS Statement”](#), and [Section 13.7.7.39, “SHOW TABLES Statement”](#). The following statements are equivalent:

```
SELECT
    TABLE_NAME, ENGINE, VERSION, ROW_FORMAT, TABLE_ROWS, AVG_ROW_LENGTH,
    DATA_LENGTH, MAX_DATA_LENGTH, INDEX_LENGTH, DATA_FREE, AUTO_INCREMENT,
    CREATE_TIME, UPDATE_TIME, CHECK_TIME, TABLE_COLLATION, CHECKSUM,
    CREATE_OPTIONS, TABLE_COMMENT
FROM INFORMATION_SCHEMA.TABLES
WHERE table_schema = 'db_name'
[AND table_name LIKE 'wild']

SHOW TABLE STATUS
FROM db_name
[LIKE 'wild']
```

The following statements are equivalent:

```
SELECT
    TABLE_NAME, TABLE_TYPE
FROM INFORMATION_SCHEMA.TABLES
WHERE table_schema = 'db_name'
[AND table_name LIKE 'wild']

SHOW FULL TABLES
FROM db_name
[LIKE 'wild']
```

### 26.3.39 The INFORMATION\_SCHEMA TABLES\_EXTENSIONS Table

The `TABLES_EXTENSIONS` table (available as of MySQL 8.0.21) provides information about table attributes defined for primary and secondary storage engines.



#### Note

The `TABLES_EXTENSIONS` table is reserved for future use.

The `TABLES_EXTENSIONS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the table belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table belongs.

- `TABLE_NAME`

The name of the table.

- `ENGINE_ATTRIBUTE`

Table attributes defined for the primary storage engine. Reserved for future use.

- `SECONDARY_ENGINE_ATTRIBUTE`

Table attributes defined for the secondary storage engine. Reserved for future use.

### 26.3.40 The INFORMATION\_SCHEMA TABLESPACES Table

This table is unused. It is deprecated; expect it to be removed in a future MySQL release. Other [INFORMATION\\_SCHEMA](#) tables may provide related information:

- For [NDB](#), the [INFORMATION\\_SCHEMA FILES](#) table provides tablespace-related information.
- For [InnoDB](#), the [INFORMATION\\_SCHEMA INNODB\\_TABLESPACES](#) and [INNODB\\_DATAFILES](#) tables provide tablespace metadata.

### 26.3.41 The INFORMATION\_SCHEMA TABLESPACES\_EXTENSIONS Table

The [TABLESPACES\\_EXTENSIONS](#) table (available as of MySQL 8.0.21) provides information about tablespace attributes defined for primary storage engines.



#### Note

The [TABLESPACES\\_EXTENSIONS](#) table is reserved for future use.

The [TABLESPACES\\_EXTENSIONS](#) table has these columns:

- [TABLESPACE\\_NAME](#)

The name of the tablespace.

- [ENGINE\\_ATTRIBUTE](#)

Tablespace attributes defined for the primary storage engine. Reserved for future use.

### 26.3.42 The INFORMATION\_SCHEMA TABLE\_CONSTRAINTS Table

The [TABLE\\_CONSTRAINTS](#) table describes which tables have constraints.

The [TABLE\\_CONSTRAINTS](#) table has these columns:

- [CONSTRAINT\\_CATALOG](#)

The name of the catalog to which the constraint belongs. This value is always [def](#).

- [CONSTRAINT\\_SCHEMA](#)

The name of the schema (database) to which the constraint belongs.

- [TABLE\\_SCHEMA](#)

The name of the schema (database) to which the table belongs.

- [TABLE\\_NAME](#)

The name of the table.

- The [CONSTRAINT\\_TYPE](#)

The type of constraint. The value can be [UNIQUE](#), [PRIMARY KEY](#), [FOREIGN KEY](#), or (as of MySQL 8.0.16) [CHECK](#). This is a [CHAR](#) (not [ENUM](#)) column.

The [UNIQUE](#) and [PRIMARY KEY](#) information is about the same as what you get from the [Key\\_name](#) column in the output from [SHOW INDEX](#) when the [Non\\_unique](#) column is [0](#).

- [ENFORCED](#)

For [CHECK](#) constraints, the value is [YES](#) or [NO](#) to indicate whether the constraint is enforced. For other constraints, the value is always [YES](#).

This column was added in MySQL 8.0.16.

## 26.3.43 The INFORMATION\_SCHEMA TABLE\_CONSTRAINTS\_EXTENSIONS Table

The `TABLE_CONSTRAINTS_EXTENSIONS` table (available as of MySQL 8.0.21) provides information about table constraint attributes defined for primary and secondary storage engines.



### Note

The `TABLE_CONSTRAINTS_EXTENSIONS` table is reserved for future use.

The `TABLE_CONSTRAINTS_EXTENSIONS` table has these columns:

- `CONSTRAINT_CATALOG`

The name of the catalog to which the table belongs.

- `CONSTRAINT_SCHEMA`

The name of the schema (database) to which the table belongs.

- `CONSTRAINT_NAME`

The name of the constraint.

- `TABLE_NAME`

The name of the table.

- `ENGINE_ATTRIBUTE`

Constraint attributes defined for the primary storage engine. Reserved for future use.

- `SECONDARY_ENGINE_ATTRIBUTE`

Constraint attributes defined for the secondary storage engine. Reserved for future use.

## 26.3.44 The INFORMATION\_SCHEMA TABLE\_PRIVILEGES Table

The `TABLE_PRIVILEGES` table provides information about table privileges. It takes its values from the `mysql.tables_priv` system table.

The `TABLE_PRIVILEGES` table has these columns:

- `GRANTEE`

The name of the account to which the privilege is granted, in '`user_name`'@'`host_name`' format.

- `TABLE_CATALOG`

The name of the catalog to which the table belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the table belongs.

- `TABLE_NAME`

The name of the table.

- `PRIVILEGE_TYPE`

The privilege granted. The value can be any privilege that can be granted at the table level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per table privilege held by the grantee.

- `IS_GRANTABLE`

`YES` if the user has the `GRANT OPTION` privilege, `NO` otherwise. The output does not list `GRANT OPTION` as a separate row with `PRIVILEGE_TYPE='GRANT OPTION'`.

## Notes

- `TABLE_PRIVILEGES` is a nonstandard `INFORMATION_SCHEMA` table.

The following statements are *not* equivalent:

```
SELECT ... FROM INFORMATION_SCHEMA.TABLE_PRIVILEGES  
SHOW GRANTS ...
```

### 26.3.45 The INFORMATION\_SCHEMA TRIGGERS Table

The `TRIGGERS` table provides information about triggers. To see information about a table's triggers, you must have the `TRIGGER` privilege for the table.

The `TRIGGERS` table has these columns:

- `TRIGGER_CATALOG`

The name of the catalog to which the trigger belongs. This value is always `def`.

- `TRIGGER_SCHEMA`

The name of the schema (database) to which the trigger belongs.

- `TRIGGER_NAME`

The name of the trigger.

- `EVENT_MANIPULATION`

The trigger event. This is the type of operation on the associated table for which the trigger activates. The value is `INSERT` (a row was inserted), `DELETE` (a row was deleted), or `UPDATE` (a row was modified).

- `EVENT_OBJECT_CATALOG`, `EVENT_OBJECT_SCHEMA`, and `EVENT_OBJECT_TABLE`

As noted in [Section 25.3, “Using Triggers”](#), every trigger is associated with exactly one table. These columns indicate the catalog and schema (database) in which this table occurs, and the table name, respectively. The `EVENT_OBJECT_CATALOG` value is always `def`.

- `ACTION_ORDER`

The ordinal position of the trigger's action within the list of triggers on the same table with the same `EVENT_MANIPULATION` and `ACTION_TIMING` values.

- `ACTION_CONDITION`

This value is always `NULL`.

- `ACTION_STATEMENT`

The trigger body; that is, the statement executed when the trigger activates. This text uses UTF-8 encoding.

- `ACTION_ORIENTATION`

This value is always `ROW`.

- `ACTION_TIMING`

Whether the trigger activates before or after the triggering event. The value is `BEFORE` or `AFTER`.

- `ACTION_REFERENCE_OLD_TABLE`

This value is always `NULL`.

- `ACTION_REFERENCE_NEW_TABLE`

This value is always `NULL`.

- `ACTION_REFERENCE_OLD_ROW` and `ACTION_REFERENCE_NEW_ROW`

The old and new column identifiers, respectively. The `ACTION_REFERENCE_OLD_ROW` value is always `OLD` and the `ACTION_REFERENCE_NEW_ROW` value is always `NEW`.

- `CREATED`

The date and time when the trigger was created. This is a `TIMESTAMP(2)` value (with a fractional part in hundredths of seconds) for triggers.

- `SQL_MODE`

The SQL mode in effect when the trigger was created, and under which the trigger executes. For the permitted values, see [Section 5.1.11, “Server SQL Modes”](#).

- `DEFINER`

The account named in the `DEFINER` clause (often the user who created the trigger), in `'user_name'@'host_name'` format.

- `CHARACTER_SET_CLIENT`

The session value of the `character_set_client` system variable when the trigger was created.

- `COLLATION_CONNECTION`

The session value of the `collation_connection` system variable when the trigger was created.

- `DATABASE_COLLATION`

The collation of the database with which the trigger is associated.

## Example

The following example uses the `ins_sum` trigger defined in [Section 25.3, “Using Triggers”](#):

```
mysql> SELECT * FROM INFORMATION_SCHEMA.TRIGGERS
      WHERE TRIGGER_SCHEMA='test' AND TRIGGER_NAME='ins_sum'\G
***** 1. row *****
    TRIGGER_CATALOG: def
    TRIGGER_SCHEMA: test
    TRIGGER_NAME: ins_sum
    EVENT_MANIPULATION: INSERT
    EVENT_OBJECT_CATALOG: def
    EVENT_OBJECT_SCHEMA: test
    EVENT_OBJECT_TABLE: account
    ACTION_ORDER: 1
    ACTION_CONDITION: NULL
    ACTION_STATEMENT: SET @sum = @sum + NEW.amount
    ACTION_ORIENTATION: ROW
    ACTION_TIMING: BEFORE
ACTION_REFERENCE_OLD_TABLE: NULL
ACTION_REFERENCE_NEW_TABLE: NULL
    ACTION_REFERENCE_OLD_ROW: OLD
    ACTION_REFERENCE_NEW_ROW: NEW
    CREATED: 2018-08-08 10:10:12.61
```

```

SQL_MODE: ONLY_FULL_GROUP_BY,STRICT_TRANS_TABLES,
          NO_ZERO_IN_DATE,NO_ZERO_DATE,
          ERROR_FOR_DIVISION_BY_ZERO,
          NO_ENGINE_SUBSTITUTION
DEFINER: me@localhost
CHARACTER_SET_CLIENT: utf8mb4
COLLATION_CONNECTION: utf8mb4_0900_ai_ci
DATABASE_COLLATION: utf8mb4_0900_ai_ci

```

Trigger information is also available from the `SHOW TRIGGERS` statement. See [Section 13.7.7.40, “SHOW TRIGGERS Statement”](#).

### 26.3.46 The INFORMATION\_SCHEMA USER\_ATTRIBUTES Table

The `USER_ATTRIBUTES` table (available as of MySQL 8.0.21) provides information about user comments and user attributes. It takes its values from the `mysql.user` system table.

The `USER_ATTRIBUTES` table has these columns:

- `USER`

The user name portion of the account to which the `ATTRIBUTE` column value applies.

- `HOST`

The host name portion of the account to which the `ATTRIBUTE` column value applies.

- `ATTRIBUTE`

The user comment, user attribute, or both belonging to the account specified by the `USER` and `HOST` columns. The value is in JSON object notation. Attributes are shown exactly as set using `CREATE USER` and `ALTER USER` statements with `ATTRIBUTE` or `COMMENT` options. A comment is shown as a key-value pair having `comment` as the key. For additional information and examples, see [CREATE USER Comment and Attribute Options](#).

### Notes

- `USER_ATTRIBUTES` is a nonstandard `INFORMATION_SCHEMA` table.
- To obtain only the user comment for a given user as an unquoted string, you can employ a query such as this one:

```

mysql> SELECT ATTRIBUTE->>"$.comment" AS Comment
      ->      FROM INFORMATION_SCHEMA.USER_ATTRIBUTES
      ->      WHERE USER='bill' AND HOST='localhost';
+-----+
| Comment |
+-----+
| A comment |
+-----+

```

Similarly, you can obtain the unquoted value for a given user attribute using its key.

- Prior to MySQL 8.0.22, `USER_ATTRIBUTES` contents are accessible by anyone. As of MySQL 8.0.22, `USER_ATTRIBUTES` contents are accessible as follows:
  - All rows are accessible if:
    - The current thread is a replica thread.
    - The access control system has not been initialized (for example, the server was started with the `--skip-grant-tables` option).
    - The currently authenticated account has the `UPDATE` or `SELECT` privilege for the `mysql.user` system table.

- The currently authenticated account has the `CREATE USER` and `SYSTEM_USER` privileges.
- Otherwise, the currently authenticated account can see the row for that account. Additionally, if the account has the `CREATE USER` privilege but not the `SYSTEM_USER` privilege, it can see rows for all other accounts that do not have the `SYSTEM_USER` privilege.

For more information about specifying account comments and attributes, see [Section 13.7.1.3, “CREATE USER Statement”](#).

## 26.3.47 The INFORMATION\_SCHEMA USER\_PRIVILEGES Table

The `USER_PRIVILEGES` table provides information about global privileges. It takes its values from the `mysql.user` system table.

The `USER_PRIVILEGES` table has these columns:

- `GRANTEE`

The name of the account to which the privilege is granted, in '`user_name`'@'`host_name`' format.

- `TABLE_CATALOG`

The name of the catalog. This value is always `def`.

- `PRIVILEGE_TYPE`

The privilege granted. The value can be any privilege that can be granted at the global level; see [Section 13.7.1.6, “GRANT Statement”](#). Each row lists a single privilege, so there is one row per global privilege held by the grantee.

- `IS_GRANTABLE`

`YES` if the user has the `GRANT OPTION` privilege, `NO` otherwise. The output does not list `GRANT OPTION` as a separate row with `PRIVILEGE_TYPE='GRANT OPTION'`.

### Notes

- `USER_PRIVILEGES` is a nonstandard `INFORMATION_SCHEMA` table.

The following statements are *not* equivalent:

```
SELECT ... FROM INFORMATION_SCHEMA.USER_PRIVILEGES  
SHOW GRANTS ...
```

## 26.3.48 The INFORMATION\_SCHEMA VIEWS Table

The `VIEWS` table provides information about views in databases. You must have the `SHOW VIEW` privilege to access this table.

The `VIEWS` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the view belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the view belongs.

- `TABLE_NAME`

The name of the view.

- [VIEW\\_DEFINITION](#)

The `SELECT` statement that provides the definition of the view. This column has most of what you see in the `Create Table` column that `SHOW CREATE VIEW` produces. Skip the words before `SELECT` and skip the words `WITH CHECK OPTION`. Suppose that the original statement was:

```
CREATE VIEW v AS
  SELECT s2,s1 FROM t
  WHERE s1 > 5
  ORDER BY s1
  WITH CHECK OPTION;
```

Then the view definition looks like this:

```
SELECT s2,s1 FROM t WHERE s1 > 5 ORDER BY s1
```

- [CHECK\\_OPTION](#)

The value of the `CHECK_OPTION` attribute. The value is one of `NONE`, `CASCADE`, or `LOCAL`.

- [IS\\_UPDATABLE](#)

MySQL sets a flag, called the view updatability flag, at `CREATE VIEW` time. The flag is set to `YES` (true) if `UPDATE` and `DELETE` (and similar operations) are legal for the view. Otherwise, the flag is set to `NO` (false). The `IS_UPDATABLE` column in the `VIEWS` table displays the status of this flag. It means that the server always knows whether a view is updatable.

If a view is not updatable, statements such `UPDATE`, `DELETE`, and `INSERT` are illegal and are rejected. (Even if a view is updatable, it might not be possible to insert into it; for details, refer to [Section 25.5.3, “Updatable and Insertable Views”](#).)

- [DEFINER](#)

The account of the user who created the view, in '`'user_name'@'host_name'`' format.

- [SECURITY\\_TYPE](#)

The view `SQL SECURITY` characteristic. The value is one of `DEFINER` or `INVOKER`.

- [CHARACTER\\_SET\\_CLIENT](#)

The session value of the `character_set_client` system variable when the view was created.

- [COLLATION\\_CONNECTION](#)

The session value of the `collation_connection` system variable when the view was created.

## Notes

MySQL permits different `sql_mode` settings to tell the server the type of SQL syntax to support. For example, you might use the `ANSI` SQL mode to ensure MySQL correctly interprets the standard SQL concatenation operator, the double bar (`||`), in your queries. If you then create a view that concatenates items, you might worry that changing the `sql_mode` setting to a value different from `ANSI` could cause the view to become invalid. But this is not the case. No matter how you write out a view definition, MySQL always stores it the same way, in a canonical form. Here is an example that shows how the server changes a double bar concatenation operator to a `CONCAT()` function:

```
mysql> SET sql_mode = 'ANSI';
Query OK, 0 rows affected (0.00 sec)

mysql> CREATE VIEW test.v AS SELECT 'a' || 'b' as coll;
Query OK, 0 rows affected (0.00 sec)

mysql> SELECT VIEW_DEFINITION FROM INFORMATION_SCHEMA.VIEWS
```

```
WHERE TABLE_SCHEMA = 'test' AND TABLE_NAME = 'v';
+-----+
| VIEW_DEFINITION          |
+-----+
| select concat('a','b') AS `coll` |
+-----+
1 row in set (0.00 sec)
```

The advantage of storing a view definition in canonical form is that changes made later to the value of `sql_mode` do not affect the results from the view. However, an additional consequence is that comments prior to `SELECT` are stripped from the definition by the server.

### 26.3.49 The INFORMATION\_SCHEMA VIEW\_ROUTINE\_USAGE Table

The `VIEW_ROUTINE_USAGE` table (available as of MySQL 8.0.13) provides access to information about stored functions used in view definitions. The table does not list information about built-in (native) functions or loadable functions used in the definitions.

You can see information only for views for which you have some privilege, and only for functions for which you have some privilege.

The `VIEW_ROUTINE_USAGE` table has these columns:

- `TABLE_CATALOG`

The name of the catalog to which the view belongs. This value is always `def`.

- `TABLE_SCHEMA`

The name of the schema (database) to which the view belongs.

- `TABLE_NAME`

The name of the view.

- `SPECIFIC_CATALOG`

The name of the catalog to which the function used in the view definition belongs. This value is always `def`.

- `SPECIFIC_SCHEMA`

The name of the schema (database) to which the function used in the view definition belongs.

- `SPECIFIC_NAME`

The name of the function used in the view definition.

### 26.3.50 The INFORMATION\_SCHEMA VIEW\_TABLE\_USAGE Table

The `VIEW_TABLE_USAGE` table (available as of MySQL 8.0.13) provides access to information about tables and views used in view definitions.

You can see information only for views for which you have some privilege, and only for tables for which you have some privilege.

The `VIEW_TABLE_USAGE` table has these columns:

- `VIEW_CATALOG`

The name of the catalog to which the view belongs. This value is always `def`.

- `VIEW_SCHEMA`

The name of the schema (database) to which the view belongs.

- [VIEW\\_NAME](#)  
The name of the view.
- [TABLE\\_CATALOG](#)  
The name of the catalog to which the table or view used in the view definition belongs. This value is always `def`.
- [TABLE\\_SCHEMA](#)  
The name of the schema (database) to which the table or view used in the view definition belongs.
- [TABLE\\_NAME](#)  
The name of the table or view used in the view definition.

## 26.4 INFORMATION\_SCHEMA InnoDB Tables

This section provides table definitions for [INFORMATION\\_SCHEMA InnoDB](#) tables. For related information and examples, see [Section 15.15, “InnoDB INFORMATION\\_SCHEMA Tables”](#).

[INFORMATION\\_SCHEMA InnoDB](#) tables can be used to monitor ongoing [InnoDB](#) activity, to detect inefficiencies before they turn into issues, or to troubleshoot performance and capacity issues. As your database becomes bigger and busier, running up against the limits of your hardware capacity, you monitor and tune these aspects to keep the database running smoothly.

### 26.4.1 INFORMATION\_SCHEMA InnoDB Table Reference

The following table summarizes [INFORMATION\\_SCHEMA InnoDB](#) tables. For greater detail, see the individual table descriptions.

**Table 26.3 INFORMATION\_SCHEMA InnoDB Tables**

Table Name	Description	Introduced
<a href="#">INNODB_BUFFER_PAGE</a>	Pages in InnoDB buffer pool	
<a href="#">INNODB_BUFFER_PAGE_LRU</a>	LRU ordering of pages in InnoDB buffer pool	
<a href="#">INNODB_BUFFER_POOL_STATS</a>	InnoDB buffer pool statistics	
<a href="#">INNODB_CACHED_INDEXES</a>	Number of index pages cached per index in InnoDB buffer pool	
<a href="#">INNODB_CMP</a>	Status for operations related to compressed InnoDB tables	
<a href="#">INNODB_CMP_PER_INDEX</a>	Status for operations related to compressed InnoDB tables and indexes	
<a href="#">INNODB_CMP_PER_INDEX_RESET</a>	Status for operations related to compressed InnoDB tables and indexes	
<a href="#">INNODB_CMP_RESET</a>	Status for operations related to compressed InnoDB tables	
<a href="#">INNODB_CMPMEM</a>	Status for compressed pages within InnoDB buffer pool	
<a href="#">INNODB_CMPMEM_RESET</a>	Status for compressed pages within InnoDB buffer pool	

## The INFORMATION\_SCHEMA INNODB\_BUFFER\_PAGE Table

---

Table Name	Description	Introduced
INNODB_COLUMNS	Columns in each InnoDB table	
INNODB_DATAFILES	Data file path information for InnoDB file-per-table and general tablespaces	
INNODB_FIELDS	Key columns of InnoDB indexes	
INNODB_FOREIGN	InnoDB foreign-key metadata	
INNODB_FOREIGN_COLS	InnoDB foreign-key column status information	
INNODB_FT_BEING_DELETED	Snapshot of INNODB_FT_DELETED table	
INNODB_FT_CONFIG	Metadata for InnoDB table FULLTEXT index and associated processing	
INNODB_FT_DEFAULT_STOPWORD	Default list of stopwords for InnoDB FULLTEXT indexes	
INNODB_FT_DELETED	Rows deleted from InnoDB table FULLTEXT index	
INNODB_FT_INDEX_CACHE	Token information for newly inserted rows in InnoDB FULLTEXT index	
INNODB_FT_INDEX_TABLE	Inverted index information for processing text searches against InnoDB table FULLTEXT index	
INNODB_INDEXES	InnoDB index metadata	
INNODB_METRICS	InnoDB performance information	
INNODB_SESSION_TEMP_TABLES	Session temporary-tablespace metadata	8.0.13
INNODB_TABLES	InnoDB table metadata	
INNODB_TABLESPACES	InnoDB file-per-table, general, and undo tablespace metadata	
INNODB_TABLESPACES_BRIEF	Brief file-per-table, general, undo, and system tablespace metadata	
INNODB_TABLESTATS	InnoDB table low-level status information	
INNODB_TEMP_TABLE_INFO	Information about active user-created InnoDB temporary tables	
INNODB_TRX	Active InnoDB transaction information	
INNODB_VIRTUAL	InnoDB virtual generated column metadata	

### 26.4.2 The INFORMATION\_SCHEMA INNODB\_BUFFER\_PAGE Table

The `INNODB_BUFFER_PAGE` table provides information about each page in the InnoDB buffer pool.

For related usage information and examples, see [Section 15.15.5, “InnoDB INFORMATION\\_SCHEMA Buffer Pool Tables”](#).

**Warning**

Querying the `INNODB_BUFFER_PAGE` table can affect performance. Do not query this table on a production system unless you are aware of the performance impact and have determined it to be acceptable. To avoid impacting performance on a production system, reproduce the issue you want to investigate and query buffer pool statistics on a test instance.

The `INNODB_BUFFER_PAGE` table has these columns:

- `POOL_ID`

The buffer pool ID. This is an identifier to distinguish between multiple buffer pool instances.

- `BLOCK_ID`

The buffer pool block ID.

- `SPACE`

The tablespace ID; the same value as `INNODB_TABLES.SPACE`.

- `PAGE_NUMBER`

The page number.

- `PAGE_TYPE`

The page type. The following table shows the permitted values.

**Table 26.4 INNODB\_BUFFER\_PAGE.PAGE\_TYPE Values**

Page Type	Description
<code>ALLOCATED</code>	Freshly allocated page
<code>BLOB</code>	Uncompressed BLOB page
<code>COMPRESSED_BLOB2</code>	Subsequent comp BLOB page
<code>COMPRESSED_BLOB</code>	First compressed BLOB page
<code>ENCRYPTED_RTREE</code>	Encrypted R-tree
<code>EXTENT_DESCRIPTOR</code>	Extent descriptor page
<code>FILE_SPACE_HEADER</code>	File space header
<code>FIL_PAGE_TYPE_UNUSED</code>	Unused
<code>IBUF_BITMAP</code>	Insert buffer bitmap
<code>IBUF_FREE_LIST</code>	Insert buffer free list
<code>IBUF_INDEX</code>	Insert buffer index
<code>INDEX</code>	B-tree node
<code>INODE</code>	Index node
<code>LOB_DATA</code>	Uncompressed LOB data
<code>LOB_FIRST</code>	First page of uncompressed LOB
<code>LOB_INDEX</code>	Uncompressed LOB index
<code>PAGE_IO_COMPRESSED</code>	Compressed page
<code>PAGE_IO_COMPRESSED_ENCRYPTED</code>	Compressed and encrypted page
<code>PAGE_IO_ENCRYPTED</code>	Encrypted page
<code>RSEG_ARRAY</code>	Rollback segment array

Page Type	Description
RTREE_INDEX	R-tree index
SDI_BLOB	Uncompressed SDI BLOB
SDI_COMPRESSED_BLOB	Compressed SDI BLOB
SDI_INDEX	SDI index
SYSTEM	System page
TRX_SYSTEM	Transaction system data
UNDO_LOG	Undo log page
UNKNOWN	Unknown
ZLOB_DATA	Compressed LOB data
ZLOB_FIRST	First page of compressed LOB
ZLOB_FRAG	Compressed LOB fragment
ZLOB_FRAG_ENTRY	Compressed LOB fragment index
ZLOB_INDEX	Compressed LOB index

- [FLUSH\\_TYPE](#)

The flush type.

- [FIX\\_COUNT](#)

The number of threads using this block within the buffer pool. When zero, the block is eligible to be evicted.

- [IS\\_HASHED](#)

Whether a hash index has been built on this page.

- [NEWEST\\_MODIFICATION](#)

The Log Sequence Number of the youngest modification.

- [OLDEST\\_MODIFICATION](#)

The Log Sequence Number of the oldest modification.

- [ACCESS\\_TIME](#)

An abstract number used to judge the first access time of the page.

- [TABLE\\_NAME](#)

The name of the table the page belongs to. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#). The column is [NULL](#) if the server has not yet accessed the table.

- [INDEX\\_NAME](#)

The name of the index the page belongs to. This can be the name of a clustered index or a secondary index. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#).

- [NUMBER\\_RECORDS](#)

The number of records within the page.

- [DATA\\_SIZE](#)

The sum of the sizes of the records. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#).

- [COMPRESSED\\_SIZE](#)

The compressed page size. [NULL](#) for pages that are not compressed.

- [PAGE\\_STATE](#)

The page state. The following table shows the permitted values.

**Table 26.5 INNODB\_BUFFER\_PAGE.PAGE\_STATE Values**

Page State	Description
<a href="#">FILE_PAGE</a>	A buffered file page
<a href="#">MEMORY</a>	Contains a main memory object
<a href="#">NOT_USED</a>	In the free list
<a href="#">NULL</a>	Clean compressed pages, compressed pages in the flush list, pages used as buffer pool watch sentinels
<a href="#">READY_FOR_USE</a>	A free page
<a href="#">REMOVE_HASH</a>	Hash index should be removed before placing in the free list

- [IO\\_FIX](#)

Whether any I/O is pending for this page: [IO\\_NONE](#) = no pending I/O, [IO\\_READ](#) = read pending, [IO\\_WRITE](#) = write pending, [IO\\_PIN](#) = relocation and removal from the flush not permitted.

- [IS\\_OLD](#)

Whether the block is in the sublist of old blocks in the LRU list.

- [FREE\\_PAGE\\_CLOCK](#)

The value of the [freed\\_page\\_clock](#) counter when the block was the last placed at the head of the LRU list. The [freed\\_page\\_clock](#) counter tracks the number of blocks removed from the end of the LRU list.

- [IS\\_STALE](#)

Whether the page is stale. Added in MySQL 8.0.24.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_BUFFER_PAGE LIMIT 1\G
***** 1. row *****
      POOL_ID: 0
      BLOCK_ID: 0
        SPACE: 97
    PAGE_NUMBER: 2473
    PAGE_TYPE: INDEX
    FLUSH_TYPE: 1
    FIX_COUNT: 0
    IS_HASHED: YES
NEWEST_MODIFICATION: 733855581
OLDEST_MODIFICATION: 0
    ACCESS_TIME: 3378385672
    TABLE_NAME: `employees`.`salaries`
    INDEX_NAME: PRIMARY
  NUMBER_RECORDS: 468
```

```
DATA_SIZE: 14976
COMPRESSED_SIZE: 0
PAGE_STATE: FILE_PAGE
IO_FIX: IO_NONE
IS_OLD: YES
FREE_PAGE_CLOCK: 66
IS_STALE: NO
```

## Notes

- This table is useful primarily for expert-level performance monitoring, or when developing performance-related extensions for MySQL.
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- When tables, table rows, partitions, or indexes are deleted, associated pages remain in the buffer pool until space is required for other data. The `INNODB_BUFFER_PAGE` table reports information about these pages until they are evicted from the buffer pool. For more information about how the `InnoDB` manages buffer pool data, see [Section 15.5.1, “Buffer Pool”](#).

### 26.4.3 The INFORMATION\_SCHEMA INNODB\_BUFFER\_PAGE\_LRU Table

The `INNODB_BUFFER_PAGE_LRU` table provides information about the pages in the `InnoDB` buffer pool; in particular, how they are ordered in the LRU list that determines which pages to `evict` from the buffer pool when it becomes full.

The `INNODB_BUFFER_PAGE_LRU` table has the same columns as the `INNODB_BUFFER_PAGE` table with a few exceptions. It has `LRU_POSITION` and `COMPRESSED` columns instead of `BLOCK_ID` and `PAGE_STATE` columns, and it does not include and `IS_STALE` column.

For related usage information and examples, see [Section 15.15.5, “InnoDB INFORMATION\\_SCHEMA Buffer Pool Tables”](#).



#### Warning

Querying the `INNODB_BUFFER_PAGE_LRU` table can affect performance. Do not query this table on a production system unless you are aware of the performance impact and have determined it to be acceptable. To avoid impacting performance on a production system, reproduce the issue you want to investigate and query buffer pool statistics on a test instance.

The `INNODB_BUFFER_PAGE_LRU` table has these columns:

- `POOL_ID`  
The buffer pool ID. This is an identifier to distinguish between multiple buffer pool instances.
- `LRU_POSITION`  
The position of the page in the LRU list.
- `SPACE`  
The tablespace ID; the same value as `INNODB_TABLES . SPACE`.
- `PAGE_NUMBER`  
The page number.
- `PAGE_TYPE`  
The page type. The following table shows the permitted values.

**Table 26.6 INNODB\_BUFFER\_PAGE\_LRU.PAGE\_TYPE Values**

<b>Page Type</b>	<b>Description</b>
ALLOCATED	Freshly allocated page
BLOB	Uncompressed BLOB page
COMPRESSED_BLOB2	Subsequent comp BLOB page
COMPRESSED_BLOB	First compressed BLOB page
ENCRYPTED_RTREE	Encrypted R-tree
EXTENT_DESCRIPTOR	Extent descriptor page
FILE_SPACE_HEADER	File space header
FIL_PAGE_TYPE_UNUSED	Unused
IBUF_BITMAP	Insert buffer bitmap
IBUF_FREE_LIST	Insert buffer free list
IBUF_INDEX	Insert buffer index
INDEX	B-tree node
INODE	Index node
LOB_DATA	Uncompressed LOB data
LOB_FIRST	First page of uncompressed LOB
LOB_INDEX	Uncompressed LOB index
PAGE_IO_COMPRESSED	Compressed page
PAGE_IO_COMPRESSED_ENCRYPTED	Compressed and encrypted page
PAGE_IO_ENCRYPTED	Encrypted page
RSEG_ARRAY	Rollback segment array
RTREE_INDEX	R-tree index
SDI_BLOB	Uncompressed SDI BLOB
SDI_COMPRESSED_BLOB	Compressed SDI BLOB
SDI_INDEX	SDI index
SYSTEM	System page
TRX_SYSTEM	Transaction system data
UNDO_LOG	Undo log page
UNKNOWN	Unknown
ZLOB_DATA	Compressed LOB data
ZLOB_FIRST	First page of compressed LOB
ZLOB_FRAG	Compressed LOB fragment
ZLOB_FRAG_ENTRY	Compressed LOB fragment index
ZLOB_INDEX	Compressed LOB index

- **FLUSH\_TYPE**

The flush type.

- **FIX\_COUNT**

The number of threads using this block within the buffer pool. When zero, the block is eligible to be evicted.

- [IS\\_HASHED](#)

Whether a hash index has been built on this page.

- [NEWEST\\_MODIFICATION](#)

The Log Sequence Number of the youngest modification.

- [OLDEST\\_MODIFICATION](#)

The Log Sequence Number of the oldest modification.

- [ACCESS\\_TIME](#)

An abstract number used to judge the first access time of the page.

- [TABLE\\_NAME](#)

The name of the table the page belongs to. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#). The column is [NULL](#) if the server has not yet accessed the table.

- [INDEX\\_NAME](#)

The name of the index the page belongs to. This can be the name of a clustered index or a secondary index. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#).

- [NUMBER\\_RECORDS](#)

The number of records within the page.

- [DATA\\_SIZE](#)

The sum of the sizes of the records. This column is applicable only to pages with a [PAGE\\_TYPE](#) value of [INDEX](#).

- [COMPRESSED\\_SIZE](#)

The compressed page size. [NULL](#) for pages that are not compressed.

- [COMPRESSED](#)

Whether the page is compressed.

- [IO\\_FIX](#)

Whether any I/O is pending for this page: [IO\\_NONE](#) = no pending I/O, [IO\\_READ](#) = read pending, [IO\\_WRITE](#) = write pending.

- [IS\\_OLD](#)

Whether the block is in the sublist of old blocks in the LRU list.

- [FREE\\_PAGE\\_CLOCK](#)

The value of the [freed\\_page\\_clock](#) counter when the block was the last placed at the head of the LRU list. The [freed\\_page\\_clock](#) counter tracks the number of blocks removed from the end of the LRU list.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_BUFFER_PAGE_LRU LIMIT 1\G
***** 1. row *****
    POOL_ID: 0
    LRU_POSITION: 0
```

```

SPACE: 97
PAGE_NUMBER: 1984
PAGE_TYPE: INDEX
FLUSH_TYPE: 1
FIX_COUNT: 0
IS_HASHED: YES
NEWEST_MODIFICATION: 719490396
OLDEST_MODIFICATION: 0
ACCESS_TIME: 3378383796
TABLE_NAME: `employees`.`salaries`
INDEX_NAME: PRIMARY
NUMBER_RECORDS: 468
DATA_SIZE: 14976
COMPRESSED_SIZE: 0
COMPRESSED: NO
IO_FIX: IO_NONE
IS_OLD: YES
FREE_PAGE_CLOCK: 0

```

## Notes

- This table is useful primarily for expert-level performance monitoring, or when developing performance-related extensions for MySQL.
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- Querying this table can require MySQL to allocate a large block of contiguous memory, more than 64 bytes times the number of active pages in the buffer pool. This allocation could potentially cause an out-of-memory error, especially for systems with multi-gigabyte buffer pools.
- Querying this table requires MySQL to lock the data structure representing the buffer pool while traversing the LRU list, which can reduce concurrency, especially for systems with multi-gigabyte buffer pools.
- When tables, table rows, partitions, or indexes are deleted, associated pages remain in the buffer pool until space is required for other data. The `INNODB_BUFFER_PAGE_LRU` table reports information about these pages until they are evicted from the buffer pool. For more information about how the InnoDB manages buffer pool data, see [Section 15.5.1, “Buffer Pool”](#).

## 26.4.4 The INFORMATION\_SCHEMA INNODB\_BUFFER\_POOL\_STATS Table

The `INNODB_BUFFER_POOL_STATS` table provides much of the same buffer pool information provided in `SHOW ENGINE INNODB STATUS` output. Much of the same information may also be obtained using InnoDB buffer pool [server status variables](#).

The idea of making pages in the buffer pool “young” or “not young” refers to transferring them between the [sublists](#) at the head and tail of the buffer pool data structure. Pages made “young” take longer to age out of the buffer pool, while pages made “not young” are moved much closer to the point of [eviction](#).

For related usage information and examples, see [Section 15.15.5, “InnoDB INFORMATION\\_SCHEMA Buffer Pool Tables”](#).

The `INNODB_BUFFER_POOL_STATS` table has these columns:

- `POOL_ID`

The buffer pool ID. This is an identifier to distinguish between multiple buffer pool instances.

- `POOL_SIZE`

The InnoDB buffer pool size in pages.

- **FREE\_BUFFERS**

The number of free pages in the [InnoDB](#) buffer pool.

- **DATABASE\_PAGES**

The number of pages in the [InnoDB](#) buffer pool containing data. This number includes both dirty and clean pages.

- **OLD\_DATABASE\_PAGES**

The number of pages in the [old](#) buffer pool sublist.

- **MODIFIED\_DATABASE\_PAGES**

The number of modified (dirty) database pages.

- **PENDING\_DECOMPRESS**

The number of pages pending decompression.

- **PENDING\_READS**

The number of pending reads.

- **PENDING\_FLUSH\_LRU**

The number of pages pending flush in the LRU.

- **PENDING\_FLUSH\_LIST**

The number of pages pending flush in the flush list.

- **PAGES\_MADE\_YOUNG**

The number of pages made young.

- **PAGES\_NOT\_MADE\_YOUNG**

The number of pages not made young.

- **PAGES\_MADE\_YOUNG\_RATE**

The number of pages made young per second (pages made young since the last printout / time elapsed).

- **PAGES\_MADE\_NOT\_YOUNG\_RATE**

The number of pages not made per second (pages not made young since the last printout / time elapsed).

- **NUMBER\_PAGES\_READ**

The number of pages read.

- **NUMBER\_PAGES\_CREATED**

The number of pages created.

- **NUMBER\_PAGES\_WRITTEN**

The number of pages written.

- **PAGES\_READ\_RATE**

The number of pages read per second (pages read since the last printout / time elapsed).

- [PAGES\\_CREATE\\_RATE](#)

The number of pages created per second (pages created since the last printout / time elapsed).

- [PAGES\\_WRITTEN\\_RATE](#)

The number of pages written per second (pages written since the last printout / time elapsed).

- [NUMBER\\_PAGES\\_GET](#)

The number of logical read requests.

- [HIT\\_RATE](#)

The buffer pool hit rate.

- [YOUNG\\_MAKE\\_PER\\_THOUSAND\\_GETS](#)

The number of pages made young per thousand gets.

- [NOT\\_YOUNG\\_MAKE\\_PER\\_THOUSAND\\_GETS](#)

The number of pages not made young per thousand gets.

- [NUMBER\\_PAGES\\_READ\\_AHEAD](#)

The number of pages read ahead.

- [NUMBER\\_READ\\_AHEAD\\_EVICTED](#)

The number of pages read into the [InnoDB](#) buffer pool by the read-ahead background thread that were subsequently evicted without having been accessed by queries.

- [READ\\_AHEAD\\_RATE](#)

The read-ahead rate per second (pages read ahead since the last printout / time elapsed).

- [READ\\_AHEAD\\_EVICTED\\_RATE](#)

The number of read-ahead pages evicted without access per second (read-ahead pages not accessed since the last printout / time elapsed).

- [LRU\\_IO\\_TOTAL](#)

Total LRU I/O.

- [LRU\\_IO\\_CURRENT](#)

LRU I/O for the current interval.

- [UNCOMPRESS\\_TOTAL](#)

The total number of pages decompressed.

- [UNCOMPRESS\\_CURRENT](#)

The number of pages decompressed in the current interval.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_BUFFER_POOL_STATS\G
***** 1. row *****
```

```
POOL_ID: 0
POOL_SIZE: 8192
FREE_BUFFERS: 1
DATABASE_PAGES: 8085
OLD_DATABASE_PAGES: 2964
MODIFIED_DATABASE_PAGES: 0
PENDING_DECOMPRESS: 0
PENDING_READS: 0
PENDING_FLUSH_LRU: 0
PENDING_FLUSH_LIST: 0
PAGES_MADE_YOUNG: 22821
PAGES_NOT_MADE_YOUNG: 3544303
PAGES_MADE_YOUNG_RATE: 357.62602199870594
PAGES_MADE_NOT_YOUNG_RATE: 0
NUMBER_PAGES_READ: 2389
NUMBER_PAGES_CREATED: 12385
NUMBER_PAGES_WRITTEN: 13111
PAGES_READ_RATE: 0
PAGES_CREATE_RATE: 0
PAGES_WRITTEN_RATE: 0
NUMBER_PAGES_GET: 33322210
HIT_RATE: 1000
YOUNG_MAKE_PER_THOUSAND_GETS: 18
NOT_YOUNG_MAKE_PER_THOUSAND_GETS: 0
NUMBER_PAGES_READ_AHEAD: 2024
NUMBER_READ_AHEAD_EVICTED: 0
READ_AHEAD_RATE: 0
READ_AHEAD_EVICTED_RATE: 0
LRU_IO_TOTAL: 0
LRU_IO_CURRENT: 0
UNCOMPRESS_TOTAL: 0
UNCOMPRESS_CURRENT: 0
```

## Notes

- This table is useful primarily for expert-level performance monitoring, or when developing performance-related extensions for MySQL.
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

### 26.4.5 The INFORMATION\_SCHEMA INNODB\_CACHED\_INDEXES Table

The `INNODB_CACHED_INDEXES` table reports the number of index pages cached in the `InnoDB` buffer pool for each index.

For related usage information and examples, see [Section 15.15.5, “InnoDB INFORMATION\\_SCHEMA Buffer Pool Tables”](#).

The `INNODB_CACHED_INDEXES` table has these columns:

- `SPACE_ID`

The tablespace ID.

- `INDEX_ID`

An identifier for the index. Index identifiers are unique across all the databases in an instance.

- `N_CACHED_PAGES`

The number of index pages cached in the `InnoDB` buffer pool.

## Examples

This query returns the number of index pages cached in the `InnoDB` buffer pool for a specific index:

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_CACHED_INDEXES WHERE INDEX_ID=65\G
***** 1. row *****
    SPACE_ID: 4294967294
    INDEX_ID: 65
N_CACHED_PAGES: 45
```

This query returns the number of index pages cached in the InnoDB buffer pool for each index, using the `INNODB_INDEXES` and `INNODB_TABLES` tables to resolve the table name and index name for each `INDEX_ID` value.

```
SELECT
    tables.NAME AS table_name,
    indexes.NAME AS index_name,
    cached.N_CACHED_PAGES AS n_cached_pages
FROM
    INFORMATION_SCHEMA.INNODB_CACHED_INDEXES AS cached,
    INFORMATION_SCHEMA.INNODB_INDEXES AS indexes,
    INFORMATION_SCHEMA.INNODB_TABLES AS tables
WHERE
    cached.INDEX_ID = indexes.INDEX_ID
    AND indexes.TABLE_ID = tables.TABLE_ID;
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.6 The INFORMATION\_SCHEMA INNODB\_CMP and INNODB\_CMP\_RESET Tables

The `INNODB_CMP` and `INNODB_CMP_RESET` tables provide status information on operations related to compressed InnoDB tables.

The `INNODB_CMP` and `INNODB_CMP_RESET` tables have these columns:

- `PAGE_SIZE`

The compressed page size in bytes.

- `COMPRESS_OPS`

The number of times a B-tree page of size `PAGE_SIZE` has been compressed. Pages are compressed whenever an empty page is created or the space for the uncompressed modification log runs out.

- `COMPRESS_OPS_OK`

The number of times a B-tree page of size `PAGE_SIZE` has been successfully compressed. This count should never exceed `COMPRESS_OPS`.

- `COMPRESS_TIME`

The total time in seconds used for attempts to compress B-tree pages of size `PAGE_SIZE`.

- `UNCOMPRESS_OPS`

The number of times a B-tree page of size `PAGE_SIZE` has been uncompressed. B-tree pages are uncompressed whenever compression fails or at first access when the uncompressed page does not exist in the buffer pool.

- `UNCOMPRESS_TIME`

The total time in seconds used for uncompressing B-tree pages of the size `PAGE_SIZE`.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_CMP\G
***** 1. row *****
    page_size: 1024
    compress_ops: 0
compress_ops_ok: 0
    compress_time: 0
    uncompress_ops: 0
uncompress_time: 0
***** 2. row *****
    page_size: 2048
    compress_ops: 0
compress_ops_ok: 0
    compress_time: 0
    uncompress_ops: 0
uncompress_time: 0
***** 3. row *****
    page_size: 4096
    compress_ops: 0
compress_ops_ok: 0
    compress_time: 0
    uncompress_ops: 0
uncompress_time: 0
***** 4. row *****
    page_size: 8192
    compress_ops: 86955
compress_ops_ok: 81182
    compress_time: 27
    uncompress_ops: 26828
uncompress_time: 5
***** 5. row *****
    page_size: 16384
    compress_ops: 0
compress_ops_ok: 0
    compress_time: 0
    uncompress_ops: 0
uncompress_time: 0
```

## Notes

- Use these tables to measure the effectiveness of [InnoDB](#) table [compression](#) in your database.
- You must have the [PROCESS](#) privilege to query this table.
- Use the [INFORMATION\\_SCHEMA COLUMNS](#) table or the [SHOW COLUMNS](#) statement to view additional information about the columns of this table, including data types and default values.
- For usage information, see [Section 15.9.1.4, “Monitoring InnoDB Table Compression at Runtime”](#) and [Section 15.15.1.3, “Using the Compression Information Schema Tables”](#). For general information about [InnoDB](#) table compression, see [Section 15.9, “InnoDB Table and Page Compression”](#).

## 26.4.7 The INFORMATION\_SCHEMA INNODB\_CMPMEM and INNODB\_CMPMEM\_RESET Tables

The [INNODB\\_CMPMEM](#) and [INNODB\\_CMPMEM\\_RESET](#) tables provide status information on compressed pages within the [InnoDB](#) buffer pool.

The [INNODB\\_CMPMEM](#) and [INNODB\\_CMPMEM\\_RESET](#) tables have these columns:

- [PAGE\\_SIZE](#)

The block size in bytes. Each record of this table describes blocks of this size.

- [BUFFER\\_POOL\\_INSTANCE](#)

A unique identifier for the buffer pool instance.

- [PAGES\\_USED](#)

The number of blocks of size [PAGE\\_SIZE](#) that are currently in use.

- [PAGES\\_FREE](#)

The number of blocks of size [PAGE\\_SIZE](#) that are currently available for allocation. This column shows the external fragmentation in the memory pool. Ideally, these numbers should be at most 1.

- [RELOCATION\\_OPS](#)

The number of times a block of size [PAGE\\_SIZE](#) has been relocated. The buddy system can relocate the allocated “buddy neighbor” of a freed block when it tries to form a bigger freed block. Reading from the [INNODB\\_CMPMEM\\_RESET](#) table resets this count.

- [RELOCATION\\_TIME](#)

The total time in microseconds used for relocating blocks of size [PAGE\\_SIZE](#). Reading from the table [INNODB\\_CMPMEM\\_RESET](#) resets this count.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_CMPMEM\G
***** 1. row *****
    page_size: 1024
buffer_pool_instance: 0
    pages_used: 0
    pages_free: 0
    relocation_ops: 0
    relocation_time: 0
***** 2. row *****
    page_size: 2048
buffer_pool_instance: 0
    pages_used: 0
    pages_free: 0
    relocation_ops: 0
    relocation_time: 0
***** 3. row *****
    page_size: 4096
buffer_pool_instance: 0
    pages_used: 0
    pages_free: 0
    relocation_ops: 0
    relocation_time: 0
***** 4. row *****
    page_size: 8192
buffer_pool_instance: 0
    pages_used: 7673
    pages_free: 15
    relocation_ops: 4638
    relocation_time: 0
***** 5. row *****
    page_size: 16384
buffer_pool_instance: 0
    pages_used: 0
    pages_free: 0
    relocation_ops: 0
    relocation_time: 0
```

## Notes

- Use these tables to measure the effectiveness of [InnoDB](#) table [compression](#) in your database.
- You must have the [PROCESS](#) privilege to query this table.
- Use the [INFORMATION\\_SCHEMA COLUMNS](#) table or the [SHOW COLUMNS](#) statement to view additional information about the columns of this table, including data types and default values.

- For usage information, see [Section 15.9.1.4, “Monitoring InnoDB Table Compression at Runtime”](#) and [Section 15.15.1.3, “Using the Compression Information Schema Tables”](#). For general information about InnoDB table compression, see [Section 15.9, “InnoDB Table and Page Compression”](#).

## 26.4.8 The INFORMATION\_SCHEMA INNODB\_CMP\_PER\_INDEX and INNODB\_CMP\_PER\_INDEX\_RESET Tables

The `INNODB_CMP_PER_INDEX` and `INNODB_CMP_PER_INDEX_RESET` tables provide status information on operations related to compressed InnoDB tables and indexes, with separate statistics for each combination of database, table, and index, to help you evaluate the performance and usefulness of compression for specific tables.

For a compressed InnoDB table, both the table data and all the [secondary indexes](#) are compressed. In this context, the table data is treated as just another index, one that happens to contain all the columns: the [clustered index](#).

The `INNODB_CMP_PER_INDEX` and `INNODB_CMP_PER_INDEX_RESET` tables have these columns:

- `DATABASE_NAME`

The schema (database) containing the applicable table.

- `TABLE_NAME`

The table to monitor for compression statistics.

- `INDEX_NAME`

The index to monitor for compression statistics.

- `COMPRESS_OPS`

The number of compression operations attempted. [Pages](#) are compressed whenever an empty page is created or the space for the uncompressed modification log runs out.

- `COMPRESS_OPS_OK`

The number of successful compression operations. Subtract from the `COMPRESS_OPS` value to get the number of [compression failures](#). Divide by the `COMPRESS_OPS` value to get the percentage of compression failures.

- `COMPRESS_TIME`

The total time in seconds used for compressing data in this index.

- `UNCOMPRESS_OPS`

The number of uncompression operations performed. Compressed InnoDB pages are uncompressed whenever compression [fails](#), or the first time a compressed page is accessed in the [buffer pool](#) and the uncompressed page does not exist.

- `UNCOMPRESS_TIME`

The total time in seconds used for uncompressing data in this index.

### Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_CMP_PER_INDEX\G
***** 1. row *****
  database_name: employees
    table_name: salaries
```

```

index_name: PRIMARY
compress_ops: 0
compress_ops_ok: 0
compress_time: 0
uncompress_ops: 23451
uncompress_time: 4
***** 2. row *****
database_name: employees
table_name: salaries
index_name: emp_no
compress_ops: 0
compress_ops_ok: 0
compress_time: 0
uncompress_ops: 1597
uncompress_time: 0

```

## Notes

- Use these tables to measure the effectiveness of [InnoDB](#) table [compression](#) for specific tables, indexes, or both.
- You must have the [PROCESS](#) privilege to query these tables.
- Use the [INFORMATION\\_SCHEMA COLUMNS](#) table or the [SHOW COLUMNS](#) statement to view additional information about the columns of these tables, including data types and default values.
- Because collecting separate measurements for every index imposes substantial performance overhead, [INNODB\\_CMP\\_PER\\_INDEX](#) and [INNODB\\_CMP\\_PER\\_INDEX\\_RESET](#) statistics are not gathered by default. You must enable the [innodb\\_cmp\\_per\\_index\\_enabled](#) system variable before performing the operations on compressed tables that you want to monitor.
- For usage information, see [Section 15.9.1.4, “Monitoring InnoDB Table Compression at Runtime”](#) and [Section 15.15.1.3, “Using the Compression Information Schema Tables”](#). For general information about [InnoDB](#) table compression, see [Section 15.9, “InnoDB Table and Page Compression”](#).

### 26.4.9 The INFORMATION\_SCHEMA INNODB\_COLUMNS Table

The [INNODB\\_COLUMNS](#) table provides metadata about [InnoDB](#) table columns.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The [INNODB\\_COLUMNS](#) table has these columns:

- [TABLE\\_ID](#)

An identifier representing the table associated with the column; the same value as [INNODB\\_TABLES.TABLE\\_ID](#).

- [NAME](#)

The name of the column. These names can be uppercase or lowercase depending on the [lower\\_case\\_table\\_names](#) setting. There are no special system-reserved names for columns.

- [POS](#)

The ordinal position of the column within the table, starting from 0 and incrementing sequentially. When a column is dropped, the remaining columns are reordered so that the sequence has no gaps. The [POS](#) value for a virtual generated column encodes the column sequence number and ordinal position of the column. For more information, see the [POS](#) column description in [Section 26.4.29, “The INFORMATION\\_SCHEMA INNODB\\_VIRTUAL Table”](#).

- [MTYPE](#)

Stands for “main type”. A numeric identifier for the column type. 1 = `VARCHAR`, 2 = `CHAR`, 3 = `FIXBINARY`, 4 = `BINARY`, 5 = `BLOB`, 6 = `INT`, 7 = `SYS_CHILD`, 8 = `SYS`, 9 = `FLOAT`, 10 = `DOUBLE`, 11 = `DECIMAL`, 12 = `VARMYSQL`, 13 = `MYSQL`, 14 = `GEOMETRY`.

- `PRTYPE`

The `InnoDB` “precise type”, a binary value with bits representing MySQL data type, character set code, and nullability.

- `LEN`

The column length, for example 4 for `INT` and 8 for `BIGINT`. For character columns in multibyte character sets, this length value is the maximum length in bytes needed to represent a definition such as `VARCHAR(N)`; that is, it might be  $2^*N$ ,  $3^*N$ , and so on depending on the character encoding.

- `HAS_DEFAULT`

A boolean value indicating whether a column that was added instantly using `ALTER TABLE ... ADD COLUMN` with `ALGORITHM=INSTANT` has a default value. All columns added instantly have a default value, which makes this column an indicator of whether the column was added instantly.

- `DEFAULT_VALUE`

The initial default value of a column that was added instantly using `ALTER TABLE ... ADD COLUMN` with `ALGORITHM=INSTANT`. If the default value is `NULL` or was not specified, this column reports `NULL`. An explicitly specified non-`NULL` default value is shown in an internal binary format. Subsequent modifications of the column default value do not change the value reported by this column.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_COLUMNS where TABLE_ID = 71\G
*****1. row *****
    TABLE_ID: 71
      NAME: col1
      POS: 0
     MTYPE: 6
    PRTYPE: 1027
      LEN: 4
   HAS_DEFAULT: 0
DEFAULT_VALUE: NULL
*****2. row *****
    TABLE_ID: 71
      NAME: col2
      POS: 1
     MTYPE: 2
    PRTYPE: 524542
      LEN: 10
   HAS_DEFAULT: 0
DEFAULT_VALUE: NULL
*****3. row *****
    TABLE_ID: 71
      NAME: col3
      POS: 2
     MTYPE: 1
    PRTYPE: 524303
      LEN: 10
   HAS_DEFAULT: 0
DEFAULT_VALUE: NULL
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.10 The INFORMATION\_SCHEMA INNODB\_DATAFILES Table

The `INNODB_DATAFILES` table provides data file path information for InnoDB file-per-table and general tablespaces.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).



### Note

The `INFORMATION_SCHEMA FILES` table reports metadata for InnoDB tablespace types including file-per-table tablespaces, general tablespaces, the system tablespace, the global temporary tablespace, and undo tablespaces.

The `INNODB_DATAFILES` table has these columns:

- `SPACE`

The tablespace ID.

- `PATH`

The tablespace data file path. If a `file-per-table` tablespace is created in a location outside the MySQL data directory, the path value is a fully qualified directory path. Otherwise, the path is relative to the data directory.

### Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_DATAFILES WHERE SPACE = 57\G
***** 1. row *****
SPACE: 57
PATH: ./test/t1.ibd
```

### Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.11 The INFORMATION\_SCHEMA INNODB\_FIELDS Table

The `INNODB_FIELDS` table provides metadata about the key columns (fields) of InnoDB indexes.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The `INNODB_FIELDS` table has these columns:

- `INDEX_ID`

An identifier for the index associated with this key field; the same value as `INNODB_INDEXES.INDEX_ID`.

- `NAME`

The name of the original column from the table; the same value as `INNODB_COLUMNS.NAME`.

- `POS`

The ordinal position of the key field within the index, starting from 0 and incrementing sequentially. When a column is dropped, the remaining columns are reordered so that the sequence has no gaps.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FIELDS WHERE INDEX_ID = 117\G
***** 1. row ****
INDEX_ID: 117
  NAME: col1
  POS: 0
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.12 The INFORMATION\_SCHEMA INNODB\_FOREIGN Table

The `INNODB_FOREIGN` table provides metadata about InnoDB foreign keys.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The `INNODB_FOREIGN` table has these columns:

- `ID`

The name (not a numeric value) of the foreign key index, preceded by the schema (database) name (for example, `test/products_fk`).

- `FOR_NAME`

The name of the `child table` in this foreign key relationship.

- `REF_NAME`

The name of the `parent table` in this foreign key relationship.

- `N_COLS`

The number of columns in the foreign key index.

- `TYPE`

A collection of bit flags with information about the foreign key column, ORed together. 0 = `ON DELETE/UPDATE RESTRICT`, 1 = `ON DELETE CASCADE`, 2 = `ON DELETE SET NULL`, 4 = `ON UPDATE CASCADE`, 8 = `ON UPDATE SET NULL`, 16 = `ON DELETE NO ACTION`, 32 = `ON UPDATE NO ACTION`.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FOREIGN\G
***** 1. row ****
      ID: test/fk1
FOR_NAME: test/child
REF_NAME: test/parent
  N_COLS: 1
    TYPE: 1
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.13 The INFORMATION\_SCHEMA INNODB\_FOREIGN\_COLS Table

The `INNODB_FOREIGN_COLS` table provides status information about InnoDB foreign key columns.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The `INNODB_FOREIGN_COLS` table has these columns:

- `ID`

The foreign key index associated with this index key field; the same value as `INNODB_FOREIGN.ID`.

- `FOR_COL_NAME`

The name of the associated column in the child table.

- `REF_COL_NAME`

The name of the associated column in the parent table.

- `POS`

The ordinal position of this key field within the foreign key index, starting from 0.

### Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FOREIGN_COLS WHERE ID = 'test/fk1'\G
***** 1. row *****
ID: test/fk1
FOR_COL_NAME: parent_id
REF_COL_NAME: id
    POS: 0
```

### Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.14 The INFORMATION\_SCHEMA INNODB\_FT\_BEING\_DELETED Table

The `INNODB_FT_BEING_DELETED` table provides a snapshot of the `INNODB_FT_DELETED` table; it is used only during an `OPTIMIZE TABLE` maintenance operation. When `OPTIMIZE TABLE` is run, the `INNODB_FT_BEING_DELETED` table is emptied, and `DOC_ID` values are removed from the `INNODB_FT_DELETED` table. Because the contents of `INNODB_FT_BEING_DELETED` typically have a short lifetime, this table has limited utility for monitoring or debugging. For information about running `OPTIMIZE TABLE` on tables with `FULLTEXT` indexes, see [Section 12.10.6, “Fine-Tuning MySQL Full-Text Search”](#).

This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT` index (for example, `test/articles`). The output appears similar to the example provided for the `INNODB_FT_DELETED` table.

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The `INNODB_FT_BEING_DELETED` table has these columns:

- `DOC_ID`

The document ID of the row that is in the process of being deleted. This value might reflect the value of an ID column that you defined for the underlying table, or it can be a sequence value generated by [InnoDB](#) when the table contains no suitable column. This value is used when you perform text searches, to skip rows in the [INNODB\\_FT\\_INDEX\\_TABLE](#) table before data for deleted rows is physically removed from the [FULLTEXT](#) index by an [OPTIMIZE TABLE](#) statement. For more information, see [Optimizing InnoDB Full-Text Indexes](#).

## Notes

- Use the [INFORMATION\\_SCHEMA COLUMNS](#) table or the [SHOW COLUMNS](#) statement to view additional information about the columns of this table, including data types and default values.
- You must have the [PROCESS](#) privilege to query this table.
- For more information about [InnoDB FULLTEXT](#) search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

## 26.4.15 The INFORMATION\_SCHEMA INNODB\_FT\_CONFIG Table

The [INNODB\\_FT\\_CONFIG](#) table provides metadata about the [FULLTEXT](#) index and associated processing for an [InnoDB](#) table.

This table is empty initially. Before querying it, set the value of the [innodb\\_ft\\_aux\\_table](#) system variable to the name (including the database name) of the table that contains the [FULLTEXT](#) index (for example, `test/articles`).

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The [INNODB\\_FT\\_CONFIG](#) table has these columns:

- **KEY**

The name designating an item of metadata for an [InnoDB](#) table containing a [FULLTEXT](#) index.

The values for this column might change, depending on the needs for performance tuning and debugging for [InnoDB](#) full-text processing. The key names and their meanings include:

- [optimize\\_checkpoint\\_limit](#): The number of seconds after which an [OPTIMIZE TABLE](#) run stops.
- [synced\\_doc\\_id](#): The next [DOC\\_ID](#) to be issued.
- [stopword\\_table\\_name](#): The [database/table](#) name for a user-defined stopword table. The [VALUE](#) column is empty if there is no user-defined stopword table.
- [use\\_stopword](#): Indicates whether a stopword table is used, which is defined when the [FULLTEXT](#) index is created.
- **VALUE**

The value associated with the corresponding **KEY** column, reflecting some limit or current value for an aspect of a [FULLTEXT](#) index for an [InnoDB](#) table.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FT_CONFIG;
+-----+-----+
| KEY | VALUE |
+-----+-----+
```

optimize_checkpoint_limit	180
synced_doc_id	0
stopword_table_name	test/my_stopwords
use_stopword	1

## Notes

- This table is intended only for internal configuration. It is not intended for statistical information purposes.
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- For more information about `InnoDB FULLTEXT` search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

## 26.4.16 The INFORMATION\_SCHEMA INNODB\_FT\_DEFAULT\_STOPWORD Table

The `INNODB_FT_DEFAULT_STOPWORD` table holds a list of [stopwords](#) that are used by default when creating a `FULLTEXT` index on `InnoDB` tables. For information about the default `InnoDB` stopword list and how to define your own stopword lists, see [Section 12.10.4, “Full-Text Stopwords”](#).

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The `INNODB_FT_DEFAULT_STOPWORD` table has these columns:

- `value`

A word that is used by default as a stopword for `FULLTEXT` indexes on `InnoDB` tables. This is not used if you override the default stopword processing with either the `innodb_ft_server_stopword_table` or the `innodb_ft_user_stopword_table` system variable.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FT_DEFAULT_STOPWORD;
+-----+
| value |
+-----+
| a      |
| about |
| an     |
| are    |
| as     |
| at     |
| be     |
| by     |
| com    |
| de     |
| en     |
| for    |
| from   |
| how    |
| i      |
| in     |
| is     |
| it     |
| la     |
| of     |
| on     |
+-----+
```

```
| or      |
| that    |
| the     |
| this    |
| to      |
| was     |
| what    |
| when    |
| where   |
| who     |
| will    |
| with    |
| und     |
| the     |
| www     |
+-----+
36 rows in set (0.00 sec)
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- For more information about InnoDB FULLTEXT search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

### 26.4.17 The INFORMATION\_SCHEMA INNODB\_FT\_DELETED Table

The `INNODB_FT_DELETED` table stores rows that are deleted from the `FULLTEXT` index for an InnoDB table. To avoid expensive index reorganization during DML operations for an InnoDB `FULLTEXT` index, the information about newly deleted words is stored separately, filtered out of search results when you do a text search, and removed from the main search index only when you issue an `OPTIMIZE TABLE` statement for the InnoDB table. For more information, see [Optimizing InnoDB Full-Text Indexes](#).

This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT` index (for example, `test/articles`).

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The `INNODB_FT_DELETED` table has these columns:

- `DOC_ID`

The document ID of the newly deleted row. This value might reflect the value of an ID column that you defined for the underlying table, or it can be a sequence value generated by InnoDB when the table contains no suitable column. This value is used when you perform text searches, to skip rows in the `INNODB_FT_INDEX_TABLE` table before data for deleted rows is physically removed from the `FULLTEXT` index by an `OPTIMIZE TABLE` statement. For more information, see [Optimizing InnoDB Full-Text Indexes](#).

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_FT_DELETED;
+-----+
| DOC_ID |
+-----+
|      6 |
|      7 |
|      8 |
```

```
+-----+
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- For more information about `InnoDB FULLTEXT` search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

## 26.4.18 The INFORMATION\_SCHEMA INNODB\_FT\_INDEX\_CACHE Table

The `INNODB_FT_INDEX_CACHE` table provides token information about newly inserted rows in a `FULLTEXT` index. To avoid expensive index reorganization during DML operations, the information about newly indexed words is stored separately, and combined with the main search index only when `OPTIMIZE TABLE` is run, when the server is shut down, or when the cache size exceeds a limit defined by the `innodb_ft_cache_size` or `innodb_ft_total_cache_size` system variable.

This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT` index (for example, `test/articles`).

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The `INNODB_FT_INDEX_CACHE` table has these columns:

- `WORD`

A word extracted from the text of a newly inserted row.

- `FIRST_DOC_ID`

The first document ID in which this word appears in the `FULLTEXT` index.

- `LAST_DOC_ID`

The last document ID in which this word appears in the `FULLTEXT` index.

- `DOC_COUNT`

The number of rows in which this word appears in the `FULLTEXT` index. The same word can occur several times within the cache table, once for each combination of `DOC_ID` and `POSITION` values.

- `DOC_ID`

The document ID of the newly inserted row. This value might reflect the value of an ID column that you defined for the underlying table, or it can be a sequence value generated by `InnoDB` when the table contains no suitable column.

- `POSITION`

The position of this particular instance of the word within the relevant document identified by the `DOC_ID` value. The value does not represent an absolute position; it is an offset added to the `POSITION` of the previous instance of that word.

## Notes

- This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT`

index (for example `test/articles`). The following example demonstrates how to use the `innodb_ft_aux_table` system variable to show information about a `FULLTEXT` index for a specified table.

```
mysql> USE test;

mysql> CREATE TABLE articles (
    id INT UNSIGNED AUTO_INCREMENT NOT NULL PRIMARY KEY,
    title VARCHAR(200),
    body TEXT,
    FULLTEXT (title,body)
) ENGINE=InnoDB;

mysql> INSERT INTO articles (title,body) VALUES
    ('MySQL Tutorial','DBMS stands for DataBase ...'),
    ('How To Use MySQL Well','After you went through a ...'),
    ('Optimizing MySQL','In this tutorial we show ...'),
    ('1001 MySQL Tricks','1. Never run mysqld as root. 2. ...'),
    ('MySQL vs. YourSQL','In the following database comparison ...'),
    ('MySQL Security','When configured properly, MySQL ...');

mysql> SET GLOBAL innodb_ft_aux_table = 'test/articles';

mysql> SELECT WORD, DOC_COUNT, DOC_ID, POSITION
    FROM INFORMATION_SCHEMA.INNODB_FT_INDEX_CACHE LIMIT 5;
+-----+-----+-----+-----+
| WORD | DOC_COUNT | DOC_ID | POSITION |
+-----+-----+-----+-----+
| 1001 | 1 | 4 | 0 |
| after | 1 | 2 | 22 |
| comparison | 1 | 5 | 44 |
| configured | 1 | 6 | 20 |
| database | 2 | 1 | 31 |
+-----+-----+-----+-----+
```

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- For more information about `InnoDB FULLTEXT` search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

## 26.4.19 The INFORMATION\_SCHEMA INNODB\_FT\_INDEX\_TABLE Table

The `INNODB_FT_INDEX_TABLE` table provides information about the inverted index used to process text searches against the `FULLTEXT` index of an `InnoDB` table.

This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT` index (for example, `test/articles`).

For related usage information and examples, see [Section 15.15.4, “InnoDB INFORMATION\\_SCHEMA FULLTEXT Index Tables”](#).

The `INNODB_FT_INDEX_TABLE` table has these columns:

- `WORD`

A word extracted from the text of the columns that are part of a `FULLTEXT`.

- `FIRST_DOC_ID`

The first document ID in which this word appears in the `FULLTEXT` index.

- `LAST_DOC_ID`

The last document ID in which this word appears in the `FULLTEXT` index.

- `DOC_COUNT`

The number of rows in which this word appears in the `FULLTEXT` index. The same word can occur several times within the cache table, once for each combination of `DOC_ID` and `POSITION` values.

- `DOC_ID`

The document ID of the row containing the word. This value might reflect the value of an ID column that you defined for the underlying table, or it can be a sequence value generated by `InnoDB` when the table contains no suitable column.

- `POSITION`

The position of this particular instance of the word within the relevant document identified by the `DOC_ID` value.

## Notes

- This table is empty initially. Before querying it, set the value of the `innodb_ft_aux_table` system variable to the name (including the database name) of the table that contains the `FULLTEXT` index (for example, `test/articles`). The following example demonstrates how to use the `innodb_ft_aux_table` system variable to show information about a `FULLTEXT` index for a specified table. Before information for newly inserted rows appears in `INNODB_FT_INDEX_TABLE`, the `FULLTEXT` index cache must be flushed to disk. This is accomplished by running an `OPTIMIZE TABLE` operation on the indexed table with the `innodb_optimize_fulltext_only` system variable enabled. (The example disables that variable again at the end because it is intended to be enabled only temporarily.)

```
mysql> USE test;
mysql> CREATE TABLE articles (
    id INT UNSIGNED AUTO_INCREMENT NOT NULL PRIMARY KEY,
    title VARCHAR(200),
    body TEXT,
    FULLTEXT (title,body)
) ENGINE=InnoDB;

mysql> INSERT INTO articles (title,body) VALUES
    ('MySQL Tutorial','DBMS stands for DataBase ...'),
    ('How To Use MySQL Well','After you went through a ...'),
    ('Optimizing MySQL','In this tutorial we show ...'),
    ('1001 MySQL Tricks','1. Never run mysqld as root. 2. ....'),
    ('MySQL vs. YourSQL','In the following database comparison ...'),
    ('MySQL Security','When configured properly, MySQL ...');

mysql> SET GLOBAL innodb_optimize_fulltext_only=ON;

mysql> OPTIMIZE TABLE articles;
+-----+-----+-----+
| Table | Op   | Msg_type | Msg_text |
+-----+-----+-----+
| test.articles | optimize | status   | OK      |
+-----+-----+-----+

mysql> SET GLOBAL innodb_ft_aux_table = 'test/articles';

mysql> SELECT WORD, DOC_COUNT, DOC_ID, POSITION
    FROM INFORMATION_SCHEMA.INNODB_FT_INDEX_TABLE LIMIT 5;
+-----+-----+-----+-----+
| WORD | DOC_COUNT | DOC_ID | POSITION |
+-----+-----+-----+-----+
| 1001 | 1          | 4      | 0        |
| after | 1          | 2      | 22       |
| comparison | 1          | 5      | 44       |
+-----+-----+-----+-----+
```

```

+-----+-----+-----+-----+
| configured |      1 |      6 |     20 |
| database   |      2 |      1 |     31 |
+-----+-----+-----+-----+
mysql> SET GLOBAL innodb_optimize_fulltext_only=OFF;

```

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- For more information about `InnoDB FULLTEXT` search, see [Section 15.6.2.4, “InnoDB Full-Text Indexes”](#), and [Section 12.10, “Full-Text Search Functions”](#).

## 26.4.20 The INFORMATION\_SCHEMA INNODB\_INDEXES Table

The `INNODB_INDEXES` table provides metadata about `InnoDB` indexes.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The `INNODB_INDEXES` table has these columns:

- `INDEX_ID`

An identifier for the index. Index identifiers are unique across all the databases in an instance.

- `NAME`

The name of the index. Most indexes created implicitly by `InnoDB` have consistent names but the index names are not necessarily unique. Examples: `PRIMARY` for a primary key index, `GEN_CLUST_INDEX` for the index representing a primary key when one is not specified, and `ID_IND`, `FOR_IND`, and `REF_IND` for foreign key constraints.

- `TABLE_ID`

An identifier representing the table associated with the index; the same value as `INNODB_TABLES.TABLE_ID`.

- `TYPE`

A numeric value derived from bit-level information that identifies the index type. 0 = nonunique secondary index; 1 = automatically generated clustered index (`GEN_CLUST_INDEX`); 2 = unique nonclustered index; 3 = clustered index; 32 = full-text index; 64 = spatial index; 128 = secondary index on a [virtual generated column](#).

- `N_FIELDS`

The number of columns in the index key. For `GEN_CLUST_INDEX` indexes, this value is 0 because the index is created using an artificial value rather than a real table column.

- `PAGE_NO`

The root page number of the index B-tree. For full-text indexes, the `PAGE_NO` column is unused and set to -1 (`FIL_NULL`) because the full-text index is laid out in several B-trees (auxiliary tables).

- `SPACE`

An identifier for the tablespace where the index resides. 0 means the `InnoDB` system tablespace. Any other number represents a table created with a separate `.ibd` file in file-per-table mode. This identifier stays the same after a `TRUNCATE TABLE` statement. Because all indexes for a table reside in the same tablespace as the table, this value is not necessarily unique.

- **MERGE\_THRESHOLD**

The merge threshold value for index pages. If the amount of data in an index page falls below the `MERGE_THRESHOLD` value when a row is deleted or when a row is shortened by an update operation, InnoDB attempts to merge the index page with the neighboring index page. The default threshold value is 50%. For more information, see [Section 15.8.11, “Configuring the Merge Threshold for Index Pages”](#).

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_INDEXES WHERE TABLE_ID = 34\G
***** 1. row *****
  INDEX_ID: 39
    NAME: GEN_CLUST_INDEX
  TABLE_ID: 34
     TYPE: 1
  N_FIELDS: 0
    PAGE_NO: 3
      SPACE: 23
MERGE_THRESHOLD: 50
***** 2. row *****
  INDEX_ID: 40
    NAME: i1
  TABLE_ID: 34
     TYPE: 0
  N_FIELDS: 1
    PAGE_NO: 4
      SPACE: 23
MERGE_THRESHOLD: 50
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.21 The INFORMATION\_SCHEMA INNODB\_METRICS Table

The `INNODB_METRICS` table provides a wide variety of InnoDB performance information, complementing the specific focus areas of the Performance Schema tables for InnoDB. With simple queries, you can check the overall health of the system. With more detailed queries, you can diagnose issues such as performance bottlenecks, resource shortages, and application issues.

Each monitor represents a point within the InnoDB source code that is instrumented to gather counter information. Each counter can be started, stopped, and reset. You can also perform these actions for a group of counters using their common module name.

By default, relatively little data is collected. To start, stop, and reset counters, set one of the system variables `innodb_monitor_enable`, `innodb_monitor_disable`, `innodb_monitor_reset`, or `innodb_monitor_reset_all`, using the name of the counter, the name of the module, a wildcard match for such a name using the “%” character, or the special keyword `all`.

For usage information, see [Section 15.15.6, “InnoDB INFORMATION\\_SCHEMA Metrics Table”](#).

The `INNODB_METRICS` table has these columns:

- **NAME**

A unique name for the counter.

- **SUBSYSTEM**

The aspect of InnoDB that the metric applies to.

- [COUNT](#)

The value since the counter was enabled.

- [MAX\\_COUNT](#)

The maximum value since the counter was enabled.

- [MIN\\_COUNT](#)

The minimum value since the counter was enabled.

- [AVG\\_COUNT](#)

The average value since the counter was enabled.

- [COUNT\\_RESET](#)

The counter value since it was last reset. (The [\\_RESET](#) columns act like the lap counter on a stopwatch: you can measure the activity during some time interval, while the cumulative figures are still available in [COUNT](#), [MAX\\_COUNT](#), and so on.)

- [MAX\\_COUNT\\_RESET](#)

The maximum counter value since it was last reset.

- [MIN\\_COUNT\\_RESET](#)

The minimum counter value since it was last reset.

- [AVG\\_COUNT\\_RESET](#)

The average counter value since it was last reset.

- [TIME\\_ENABLED](#)

The timestamp of the last start.

- [TIME\\_DISABLED](#)

The timestamp of the last stop.

- [TIME\\_ELAPSED](#)

The elapsed time in seconds since the counter started.

- [TIME\\_RESET](#)

The timestamp of the last reset.

- [STATUS](#)

Whether the counter is still running ([enabled](#)) or stopped ([disabled](#)).

- [TYPE](#)

Whether the item is a cumulative counter, or measures the current value of some resource.

- [COMMENT](#)

The counter description.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_METRICS WHERE NAME='dml_inserts' \G
```

```
***** 1. row *****
  NAME: dml_inserts
  SUBSYSTEM: dml
  COUNT: 3
  MAX_COUNT: 3
  MIN_COUNT: NULL
  AVG_COUNT: 0.046153846153846156
  COUNT_RESET: 3
MAX_COUNT_RESET: 3
MIN_COUNT_RESET: NULL
AVG_COUNT_RESET: NULL
  TIME_ENABLED: 2014-12-04 14:18:28
  TIME_DISABLED: NULL
  TIME_ELAPSED: 65
  TIME_RESET: NULL
  STATUS: enabled
  TYPE: status_counter
  COMMENT: Number of rows inserted
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.
- Transaction counter `COUNT` values may differ from the number of transaction events reported in Performance Schema `EVENTS_TRANSACTIONS_SUMMARY` tables. InnoDB counts only those transactions that it executes, whereas Performance Schema collects events for all non-aborted transactions initiated by the server, including empty transactions.

### 26.4.22 The INFORMATION\_SCHEMA INNODB\_SESSION\_TEMP\_TABLESPACES Table

The `INNODB_SESSION_TEMP_TABLESPACES` table provides metadata about session temporary tablespaces used for internal and user-created temporary tables. This table was added in MySQL 8.0.13.

The `INNODB_SESSION_TEMP_TABLESPACES` table has these columns:

- `ID`

The process or session ID.

- `SPACE`

The tablespace ID. A range of 400 thousand space IDs is reserved for session temporary tablespaces. Session temporary tablespaces are recreated each time the server is started. Space IDs are not persisted when the server is shut down and may be reused.

- `PATH`

The tablespace data file path. A session temporary tablespace has an `.ibt` file extension.

- `SIZE`

The size of the tablespace, in bytes.

- `STATE`

The state of the tablespace. `ACTIVE` indicates that the tablespace is currently used by a session. `INACTIVE` indicates that the tablespace is in the pool of available session temporary tablespaces.

- `PURPOSE`

The purpose of the tablespace. [INTRINSIC](#) indicates that the tablespace is used for optimized internal temporary tables use by the optimizer. [SLAVE](#) indicates that the tablespace is allocated for storing user-created temporary tables on a replication slave. [USER](#) indicates that the tablespace is used for user-created temporary tables. [NONE](#) indicates that the tablespace is not in use.

## Example

ID	SPACE	PATH	SIZE	STATE	PURPOSE
8	4294566162	./#innodb_temp/temp_10.ibt	81920	ACTIVE	INTRINSIC
8	4294566161	./#innodb_temp/temp_9.ibt	98304	ACTIVE	USER
0	4294566153	./#innodb_temp/temp_1.ibt	81920	INACTIVE	NONE
0	4294566154	./#innodb_temp/temp_2.ibt	81920	INACTIVE	NONE
0	4294566155	./#innodb_temp/temp_3.ibt	81920	INACTIVE	NONE
0	4294566156	./#innodb_temp/temp_4.ibt	81920	INACTIVE	NONE
0	4294566157	./#innodb_temp/temp_5.ibt	81920	INACTIVE	NONE
0	4294566158	./#innodb_temp/temp_6.ibt	81920	INACTIVE	NONE
0	4294566159	./#innodb_temp/temp_7.ibt	81920	INACTIVE	NONE
0	4294566160	./#innodb_temp/temp_8.ibt	81920	INACTIVE	NONE

## Notes

- You must have the [PROCESS](#) privilege to query this table.
- Use the [INFORMATION\\_SCHEMA COLUMNS](#) table or the [SHOW COLUMNS](#) statement to view additional information about the columns of this table, including data types and default values.

## 26.4.23 The INFORMATION\_SCHEMA INNODB\_TABLES Table

The [INNODB\\_TABLES](#) table provides metadata about [InnoDB](#) tables.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The [INNODB\\_TABLES](#) table has these columns:

- [TABLE\\_ID](#)

An identifier for the [InnoDB](#) table. This value is unique across all databases in the instance.

- [NAME](#)

The name of the table, preceded by the schema (database) name where appropriate (for example, `test/t1`). Names of databases and user tables are in the same case as they were originally defined, possibly influenced by the [lower\\_case\\_table\\_names](#) setting.

- [FLAG](#)

A numeric value that represents bit-level information about table format and storage characteristics.

- [N\\_COLS](#)

The number of columns in the table. The number reported includes three hidden columns that are created by [InnoDB](#) ([DB\\_ROW\\_ID](#), [DB\\_TRX\\_ID](#), and [DB\\_ROLL\\_PTR](#)). The number reported also includes [virtual generated columns](#), if present.

- [SPACE](#)

An identifier for the tablespace where the table resides. 0 means the [InnoDB system tablespace](#). Any other number represents either a [file-per-table](#) tablespace or a general tablespace. This

identifier stays the same after a `TRUNCATE TABLE` statement. For file-per-table tablespaces, this identifier is unique for tables across all databases in the instance.

- `ROW_FORMAT`

The table's row format (`Compact`, `Redundant`, `Dynamic`, or `Compressed`).

- `ZIP_PAGE_SIZE`

The zip page size. Applies only to tables with a row format of `Compressed`.

- `SPACE_TYPE`

The type of tablespace to which the table belongs. Possible values include `System` for the system tablespace, `General` for general tablespaces, and `Single` for file-per-table tablespaces. Tables assigned to the system tablespace using `CREATE TABLE` or `ALTER TABLE TABLESPACE=innodb_system` have a `SPACE_TYPE` of `General`. For more information, see `CREATE TABLESPACE`.

- `INSTANT_COLS`

The number of columns that existed before the first instant column was added using `ALTER TABLE ... ADD COLUMN` with `ALGORITHM=INSTANT`. This column is no longer used as of MySQL 8.0.29 but continues to show information for tables with columns that were added instantly prior to MySQL 8.0.29.

- `TOTAL_ROW_VERSIONS`

The number of row versions for the table. The initial value is 0. The value is incremented by `ALTER TABLE ... ALGORITHM=INSTANT` operations that add or remove columns. When a table with instantly added or dropped columns is rebuilt due to a table-rebuilding `ALTER TABLE` or `OPTIMIZE TABLE` operation, the value is reset to 0. For more information, see [Column Operations](#).

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TABLES WHERE TABLE_ID = 214\G
***** 1. row *****
    TABLE_ID: 1064
      NAME: test/t1
      FLAG: 33
     N_COLS: 6
      SPACE: 3
    ROW_FORMAT: Dynamic
 ZIP_PAGE_SIZE: 0
   SPACE_TYPE: Single
 INSTANT_COLS: 0
TOTAL_ROW_VERSIONS: 3
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.24 The INFORMATION\_SCHEMA INNODB\_TABLESPACES Table

The `INNODB_TABLESPACES` table provides metadata about InnoDB file-per-table, general, and undo tablespaces.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

**Note**

The [INFORMATION\\_SCHEMA FILES](#) table reports metadata for [InnoDB](#) tablespace types including file-per-table tablespaces, general tablespaces, the system tablespace, the global temporary tablespace, and undo tablespaces.

The [INNODB\\_TABLESPACES](#) table has these columns:

- [SPACE](#)

The tablespace ID.

- [NAME](#)

The schema (database) and table name.

- [FLAG](#)

A numeric value that represents bit-level information about tablespace format and storage characteristics.

- [ROW\\_FORMAT](#)

The tablespace row format ([Compact or Redundant](#), [Dynamic or Compressed](#), or [Undo](#)). The data in this column is interpreted from the tablespace flag information that resides in the data file.

There is no way to determine from this flag information if the tablespace row format is [Redundant](#) or [Compact](#), which is why one of the possible [ROW\\_FORMAT](#) values is [Compact or Redundant](#).

- [PAGE\\_SIZE](#)

The tablespace page size. The data in this column is interpreted from the tablespace flags information that resides in the [.ibd file](#).

- [ZIP\\_PAGE\\_SIZE](#)

The tablespace zip page size. The data in this column is interpreted from the tablespace flags information that resides in the [.ibd file](#).

- [SPACE\\_TYPE](#)

The type of tablespace. Possible values include [General](#) for general tablespaces, [Single](#) for file-per-table tablespaces, [System](#) for the system tablespace, and [Undo](#) for undo tablespaces.

- [FS\\_BLOCK\\_SIZE](#)

The file system block size, which is the unit size used for hole punching. This column pertains to the [InnoDB transparent page compression](#) feature.

- [FILE\\_SIZE](#)

The apparent size of the file, which represents the maximum size of the file, uncompressed. This column pertains to the [InnoDB transparent page compression](#) feature.

- [ALLOCATED\\_SIZE](#)

The actual size of the file, which is the amount of space allocated on disk. This column pertains to the [InnoDB transparent page compression](#) feature.

- [AUTOEXTEND\\_SIZE](#)

The auto-extend size of the tablespace. This column was added in MySQL 8.0.23.

- [SERVER\\_VERSION](#)

The MySQL version that created the tablespace, or the MySQL version into which the tablespace was imported, or the version of the last major MySQL version upgrade. The value is unchanged by a release series upgrade, such as an upgrade from MySQL 8.0.x to 8.0.y. The value can be considered a “creation” marker or “certified” marker for the tablespace.

- **SPACE\_VERSION**

The tablespace version, used to track changes to the tablespace format.

- **ENCRYPTION**

Whether the tablespace is encrypted. This column was added in MySQL 8.0.13.

- **STATE**

The tablespace state. This column was added in MySQL 8.0.14.

For file-per-table and general tablespaces, states include:

- **normal**: The tablespace is normal and active.
- **discarded**: The tablespace was discarded by an `ALTER TABLE ... DISCARD TABLESPACE` statement.
- **corrupted**: The tablespace is identified by `InnoDB` as corrupted.

For undo tablespaces, states include:

- **active**: Rollback segments in the undo tablespace can be allocated to new transactions.
- **inactive**: Rollback segments in the undo tablespace are no longer used by new transactions. The truncate process is in progress. The undo tablespace was either selected by the purge thread implicitly or was made inactive by an `ALTER UNDO TABLESPACE ... SET INACTIVE` statement.
- **empty**: The undo tablespace was truncated and is no longer active. It is ready to be dropped or made active again by an `ALTER UNDO TABLESPACE ... SET INACTIVE` statement.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TABLESPACES WHERE SPACE = 26\G
***** 1. row *****
    SPACE: 26
    NAME: test/t1
    FLAG: 0
    ROW_FORMAT: Compact or Redundant
    PAGE_SIZE: 16384
    ZIP_PAGE_SIZE: 0
    SPACE_TYPE: Single
    FS_BLOCK_SIZE: 4096
    FILE_SIZE: 98304
    ALLOCATED_SIZE: 65536
    AUTOEXTEND_SIZE: 0
    SERVER_VERSION: 8.0.23
    SPACE_VERSION: 1
    ENCRYPTION: N
    STATE: normal
```

## Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.25 The INFORMATION\_SCHEMA INNODB\_TABLESPACES\_BRIEF Table

The `INNODB_TABLESPACES_BRIEF` table provides space ID, name, path, flag, and space type metadata for file-per-table, general, undo, and system tablespaces.

`INNODB_TABLESPACES` provides the same metadata but loads more slowly because other metadata provided by the table, such as `FS_BLOCK_SIZE`, `FILE_SIZE`, and `ALLOCATED_SIZE`, must be loaded dynamically.

Space and path metadata is also provided by the `INNODB_DATAFILES` table.

The `INNODB_TABLESPACES_BRIEF` table has these columns:

- `SPACE`

The tablespace ID.

- `NAME`

The tablespace name. For file-per-table tablespaces, the name is in the form of `schema/table_name`.

- `PATH`

The tablespace data file path. If a `file-per-table` tablespace is created in a location outside the MySQL data directory, the path value is a fully qualified directory path. Otherwise, the path is relative to the data directory.

- `FLAG`

A numeric value that represents bit-level information about tablespace format and storage characteristics.

- `SPACE_TYPE`

The type of tablespace. Possible values include `General` for InnoDB general tablespaces, `Single` for InnoDB file-per-table tablespaces, and `System` for the InnoDB system tablespace.

### Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TABLESPACES_BRIEF WHERE SPACE = 7;
+-----+-----+-----+-----+
| SPACE | NAME | PATH          | FLAG | SPACE_TYPE |
+-----+-----+-----+-----+
| 7     | test/t1 | ./test/t1.ibd | 16417 | Single      |
+-----+-----+-----+-----+
```

### Notes

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.26 The INFORMATION\_SCHEMA INNODB\_TABLESTATS View

The `INNODB_TABLESTATS` table provides a view of low-level status information about InnoDB tables. This data is used by the MySQL optimizer to calculate which index to use when querying an InnoDB table. This information is derived from in-memory data structures rather than data stored on disk. There is no corresponding internal InnoDB system table.

InnoDB tables are represented in this view if they have been opened since the last server restart and have not aged out of the table cache. Tables for which persistent stats are available are always represented in this view.

Table statistics are updated only for `DELETE` or `UPDATE` operations that modify indexed columns. Statistics are not updated by operations that modify only nonindexed columns.

`ANALYZE TABLE` clears table statistics and sets the `STATS_INITIALIZED` column to `Uninitialized`. Statistics are collected again the next time the table is accessed.

For related usage information and examples, see [Section 15.15.3, “InnoDB INFORMATION\\_SCHEMA Schema Object Tables”](#).

The `INNODB_TABLESTATS` table has these columns:

- `TABLE_ID`

An identifier representing the table for which statistics are available; the same value as `INNODB_TABLES.TABLE_ID`.

- `NAME`

The name of the table; the same value as `INNODB_TABLES.NAME`.

- `STATS_INITIALIZED`

The value is `Initialized` if the statistics are already collected, `Uninitialized` if not.

- `NUM_ROWS`

The current estimated number of rows in the table. Updated after each DML operation. The value could be imprecise if uncommitted transactions are inserting into or deleting from the table.

- `CLUST_INDEX_SIZE`

The number of pages on disk that store the clustered index, which holds the InnoDB table data in primary key order. This value might be null if no statistics are collected yet for the table.

- `OTHER_INDEX_SIZE`

The number of pages on disk that store all secondary indexes for the table. This value might be null if no statistics are collected yet for the table.

- `MODIFIED_COUNTER`

The number of rows modified by DML operations, such as `INSERT`, `UPDATE`, `DELETE`, and also cascade operations from foreign keys. This column is reset each time table statistics are recalculated

- `AUTOINC`

The next number to be issued for any auto-increment-based operation. The rate at which the `AUTOINC` value changes depends on how many times auto-increment numbers have been requested and how many numbers are granted per request.

- `REF_COUNT`

When this counter reaches zero, the table metadata can be evicted from the table cache.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TABLESTATS where TABLE_ID = 71\G
***** 1. row *****
```

```
TABLE_ID: 71
  NAME: test/t1
STATS_INITIALIZED: Initialized
  NUM_ROWS: 1
CLUST_INDEX_SIZE: 1
OTHER_INDEX_SIZE: 0
MODIFIED_COUNTER: 1
  AUTOINC: 0
  REF_COUNT: 1
```

## Notes

- This table is useful primarily for expert-level performance monitoring, or when developing performance-related extensions for MySQL.
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

### 26.4.27 The INFORMATION\_SCHEMA INNODB\_TEMP\_TABLE\_INFO Table

The `INNODB_TEMP_TABLE_INFO` table provides information about user-created InnoDB temporary tables that are active in an InnoDB instance. It does not provide information about internal InnoDB temporary tables used by the optimizer. The `INNODB_TEMP_TABLE_INFO` table is created when first queried, exists only in memory, and is not persisted to disk.

For usage information and examples, see [Section 15.15.7, “InnoDB INFORMATION\\_SCHEMA Temporary Table Info Table”](#).

The `INNODB_TEMP_TABLE_INFO` table has these columns:

- `TABLE_ID`

The table ID of the temporary table.

- `NAME`

The name of the temporary table.

- `N_COLS`

The number of columns in the temporary table. The number includes three hidden columns created by InnoDB (`DB_ROW_ID`, `DB_TRX_ID`, and `DB_ROLL_PTR`).

- `SPACE`

The ID of the temporary tablespace where the temporary table resides.

## Example

```
mysql> CREATE TEMPORARY TABLE t1 (c1 INT PRIMARY KEY) ENGINE=INNODB;
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TEMP_TABLE_INFO\G
***** 1. row *****
TABLE_ID: 97
  NAME: #sql18c88_43_0
  N_COLS: 4
  SPACE: 76
```

## Notes

- This table is useful primarily for expert-level monitoring.

- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.28 The INFORMATION\_SCHEMA INNODB\_TRX Table

The `INNODB_TRX` table provides information about every transaction currently executing inside `InnoDB`, including whether the transaction is waiting for a lock, when the transaction started, and the SQL statement the transaction is executing, if any.

For usage information, see [Section 15.15.2.1, “Using InnoDB Transaction and Locking Information”](#).

The `INNODB_TRX` table has these columns:

- `TRX_ID`

A unique transaction ID number, internal to `InnoDB`. These IDs are not created for transactions that are read only and nonlocking. For details, see [Section 8.5.3, “Optimizing InnoDB Read-Only Transactions”](#).

- `TRX_WEIGHT`

The weight of a transaction, reflecting (but not necessarily the exact count of) the number of rows altered and the number of rows locked by the transaction. To resolve a deadlock, `InnoDB` selects the transaction with the smallest weight as the “victim” to roll back. Transactions that have changed nontransactional tables are considered heavier than others, regardless of the number of altered and locked rows.

- `TRX_STATE`

The transaction execution state. Permitted values are `RUNNING`, `LOCK_WAIT`, `ROLLING BACK`, and `COMMITTING`.

- `TRX_STARTED`

The transaction start time.

- `TRX_REQUESTED_LOCK_ID`

The ID of the lock the transaction is currently waiting for, if `TRX_STATE` is `LOCK_WAIT`; otherwise `NULL`. To obtain details about the lock, join this column with the `ENGINE_LOCK_ID` column of the Performance Schema `data_locks` table.

- `TRX_WAIT_STARTED`

The time when the transaction started waiting on the lock, if `TRX_STATE` is `LOCK_WAIT`; otherwise `NULL`.

- `TRX_MYSQL_THREAD_ID`

The MySQL thread ID. To obtain details about the thread, join this column with the `ID` column of the `INFORMATION_SCHEMA PROCESSLIST` table, but see [Section 15.15.2.3, “Persistence and Consistency of InnoDB Transaction and Locking Information”](#).

- `TRX_QUERY`

The SQL statement that is being executed by the transaction.

- `TRX_OPERATION_STATE`

The transaction's current operation, if any; otherwise `NULL`.

- [TRX\\_TABLES\\_IN\\_USE](#)

The number of [InnoDB](#) tables used while processing the current SQL statement of this transaction.

- [TRX\\_TABLES\\_LOCKED](#)

The number of [InnoDB](#) tables that the current SQL statement has row locks on. (Because these are row locks, not table locks, the tables can usually still be read from and written to by multiple transactions, despite some rows being locked.)

- [TRX\\_LOCK\\_STRUCTS](#)

The number of locks reserved by the transaction.

- [TRX\\_LOCK\\_MEMORY\\_BYTES](#)

The total size taken up by the lock structures of this transaction in memory.

- [TRX\\_ROWS\\_LOCKED](#)

The approximate number of rows locked by this transaction. The value might include delete-marked rows that are physically present but not visible to the transaction.

- [TRX\\_ROWS\\_MODIFIED](#)

The number of modified and inserted rows in this transaction.

- [TRX\\_CONCURRENCY\\_TICKETS](#)

A value indicating how much work the current transaction can do before being swapped out, as specified by the [innodb\\_concurrency\\_tickets](#) system variable.

- [TRX\\_ISOLATION\\_LEVEL](#)

The isolation level of the current transaction.

- [TRX\\_UNIQUE\\_CHECKS](#)

Whether unique checks are turned on or off for the current transaction. For example, they might be turned off during a bulk data load.

- [TRX\\_FOREIGN\\_KEY\\_CHECKS](#)

Whether foreign key checks are turned on or off for the current transaction. For example, they might be turned off during a bulk data load.

- [TRX\\_LAST\\_FOREIGN\\_KEY\\_ERROR](#)

The detailed error message for the last foreign key error, if any; otherwise [NULL](#).

- [TRX\\_ADAPTIVE\\_HASH\\_LATCHED](#)

Whether the adaptive hash index is locked by the current transaction. When the adaptive hash index search system is partitioned, a single transaction does not lock the entire adaptive hash index. Adaptive hash index partitioning is controlled by [innodb\\_adaptive\\_hash\\_index\\_parts](#), which is set to 8 by default.

- [TRX\\_ADAPTIVE\\_HASH\\_TIMEOUT](#)

Whether to relinquish the search latch immediately for the adaptive hash index, or reserve it across calls from MySQL. When there is no adaptive hash index contention, this value remains zero and statements reserve the latch until they finish. During times of contention, it counts down to zero, and statements release the latch immediately after each row lookup. When the adaptive hash index

search system is partitioned (controlled by `innodb_adaptive_hash_index_parts`), the value remains 0.

- `TRX_IS_READ_ONLY`

A value of 1 indicates the transaction is read only.

- `TRX_AUTOCOMMIT_NON_LOCKING`

A value of 1 indicates the transaction is a `SELECT` statement that does not use the `FOR UPDATE` or `LOCK IN SHARED MODE` clauses, and is executing with `autocommit` enabled so that the transaction contains only this one statement. When this column and `TRX_IS_READ_ONLY` are both 1, InnoDB optimizes the transaction to reduce the overhead associated with transactions that change table data.

- `TRX_SCHEDULE_WEIGHT`

The transaction schedule weight assigned by the Contention-Aware Transaction Scheduling (CATS) algorithm to transactions waiting for a lock. The value is relative to the values of other transactions. A higher value has a greater weight. A value is computed only for transactions in a `LOCK_WAIT` state, as reported by the `TRX_STATE` column. A NULL value is reported for transactions that are not waiting for a lock. The `TRX_SCHEDULE_WEIGHT` value is different from the `TRX_WEIGHT` value, which is computed by a different algorithm for a different purpose.

## Example

```
mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_TRX\G
***** 1. row *****
      trx_id: 1510
      trx_state: RUNNING
      trx_started: 2014-11-19 13:24:40
      trx_requested_lock_id: NULL
      trx_wait_started: NULL
      trx_weight: 586739
      trx_mysql_thread_id: 2
      trx_query: DELETE FROM employees.salaries WHERE salary > 65000
      trx_operation_state: updating or deleting
      trx_tables_in_use: 1
      trx_tables_locked: 1
      trx_lock_structs: 3003
      trx_lock_memory_bytes: 450768
      trx_rows_locked: 1407513
      trx_rows_modified: 583736
      trx_concurrency_tickets: 0
      trx_isolation_level: REPEATABLE READ
      trx_unique_checks: 1
      trx_foreign_key_checks: 1
      trx_last_foreign_key_error: NULL
      trx_adaptive_hash_latched: 0
      trx_adaptive_hash_timeout: 10000
      trx_is_read_only: 0
      trx_autocommit_non_locking: 0
      trx_schedule_weight: NULL
```

## Notes

- Use this table to help diagnose performance problems that occur during times of heavy concurrent load. Its contents are updated as described in [Section 15.15.2.3, “Persistence and Consistency of InnoDB Transaction and Locking Information”](#).
- You must have the `PROCESS` privilege to query this table.
- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.4.29 The INFORMATION\_SCHEMA INNODB\_VIRTUAL Table

The [INNODB\\_VIRTUAL](#) table provides metadata about [InnoDB virtual generated columns](#) and columns upon which virtual generated columns are based.

A row appears in the [INNODB\\_VIRTUAL](#) table for each column upon which a virtual generated column is based.

The [INNODB\\_VIRTUAL](#) table has these columns:

- [TABLE\\_ID](#)

An identifier representing the table associated with the virtual column; the same value as [INNODB\\_TABLES.TABLE\\_ID](#).

- [POS](#)

The position value of the [virtual generated column](#). The value is large because it encodes the column sequence number and ordinal position. The formula used to calculate the value uses a bitwise operation:

```
((nth virtual generated column for the InnoDB instance + 1) << 16)
+ the ordinal position of the virtual generated column
```

For example, if the first virtual generated column in the [InnoDB](#) instance is the third column of the table, the formula is `(0 + 1) << 16) + 2`. The first virtual generated column in the [InnoDB](#) instance is always number 0. As the third column in the table, the ordinal position of the virtual generated column is 2. Ordinal positions are counted from 0.

- [BASE\\_POS](#)

The ordinal position of the columns upon which a virtual generated column is based.

## Example

```
mysql> CREATE TABLE `t1` (
    `a` int(11) DEFAULT NULL,
    `b` int(11) DEFAULT NULL,
    `c` int(11) GENERATED ALWAYS AS (a+b) VIRTUAL,
    `h` varchar(10) DEFAULT NULL
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4;

mysql> SELECT * FROM INFORMATION_SCHEMA.INNODB_VIRTUAL
WHERE TABLE_ID IN
    (SELECT TABLE_ID FROM INFORMATION_SCHEMA.INNODB_TABLES
     WHERE NAME LIKE "test/t1");
+-----+-----+-----+
| TABLE_ID | POS   | BASE_POS |
+-----+-----+-----+
|      98  | 65538 |       0  |
|      98  | 65538 |       1  |
+-----+-----+-----+
```

## Notes

- If a constant value is assigned to a [virtual generated column](#), as in the following table, an entry for the column does not appear in the [INNODB\\_VIRTUAL](#) table. For an entry to appear, a virtual generated column must have a base column.

```
CREATE TABLE `t1` (
    `a` int(11) DEFAULT NULL,
    `b` int(11) DEFAULT NULL,
    `c` int(11) GENERATED ALWAYS AS (5) VIRTUAL
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4;
```

However, metadata for such a column does appear in the [INNODB\\_COLUMNS](#) table.

- You must have the [PROCESS](#) privilege to query this table.

- Use the `INFORMATION_SCHEMA COLUMNS` table or the `SHOW COLUMNS` statement to view additional information about the columns of this table, including data types and default values.

## 26.5 INFORMATION\_SCHEMA Thread Pool Tables



### Note

As of MySQL 8.0.14, the `INFORMATION_SCHEMA` thread pool tables are also available as Performance Schema tables. (See [Section 27.12.16, “Performance Schema Thread Pool Tables”](#).) The `INFORMATION_SCHEMA` tables are deprecated; expect them to be removed in a future version of MySQL. Applications should transition away from the old tables to the new tables. For example, if an application uses this query:

```
SELECT * FROM INFORMATION_SCHEMA.TP_THREAD_STATE;
```

The application should use this query instead:

```
SELECT * FROM performance_schema.tp_thread_state;
```

The following sections describe the `INFORMATION_SCHEMA` tables associated with the thread pool plugin (see [Section 5.6.3, “MySQL Enterprise Thread Pool”](#)). They provide information about thread pool operation:

- `TP_THREAD_GROUP_STATE`: Information about thread pool thread group states
- `TP_THREAD_GROUP_STATS`: Thread group statistics
- `TP_THREAD_STATE`: Information about thread pool thread states

Rows in these tables represent snapshots in time. In the case of `TP_THREAD_STATE`, all rows for a thread group comprise a snapshot in time. Thus, the MySQL server holds the mutex of the thread group while producing the snapshot. But it does not hold mutexes on all thread groups at the same time, to prevent a statement against `TP_THREAD_STATE` from blocking the entire MySQL server.

The `INFORMATION_SCHEMA` thread pool tables are implemented by individual plugins and the decision whether to load one can be made independently of the others (see [Section 5.6.3.2, “Thread Pool Installation”](#)). However, the content of all the tables depends on the thread pool plugin being enabled. If a table plugin is enabled but the thread pool plugin is not, the table becomes visible and can be accessed but is empty.

### 26.5.1 INFORMATION\_SCHEMA Thread Pool Table Reference

The following table summarizes `INFORMATION_SCHEMA` thread pool tables. For greater detail, see the individual table descriptions.

**Table 26.7 INFORMATION\_SCHEMA Thread Pool Tables**

Table Name	Description
<code>TP_THREAD_GROUP_STATE</code>	Thread pool thread group states
<code>TP_THREAD_GROUP_STATS</code>	Thread pool thread group statistics
<code>TP_THREAD_STATE</code>	Thread pool thread information

### 26.5.2 The INFORMATION\_SCHEMA TP\_THREAD\_GROUP\_STATE Table



### Note

As of MySQL 8.0.14, the thread pool `INFORMATION_SCHEMA` tables are also available as Performance Schema tables. (See [Section 27.12.16, “Performance Schema Thread Pool Tables”](#).) The `INFORMATION_SCHEMA`

tables are deprecated; expect them to be removed in a future version of MySQL. Applications should transition away from the old tables to the new tables. For example, if an application uses this query:

```
SELECT * FROM INFORMATION_SCHEMA.TP_THREAD_GROUP_STATE;
```

The application should use this query instead:

```
SELECT * FROM performance_schema.tp_thread_group_state;
```

The `TP_THREAD_GROUP_STATE` table has one row per thread group in the thread pool. Each row provides information about the current state of a group.

For descriptions of the columns in the `INFORMATION_SCHEMA TP_THREAD_GROUP_STATE` table, see [Section 27.12.16.1, “The tp\\_thread\\_group\\_state Table”](#). The Performance Schema `tp_thread_group_state` table has equivalent columns.

### 26.5.3 The INFORMATION\_SCHEMA TP\_THREAD\_GROUP\_STATS Table



#### Note

As of MySQL 8.0.14, the thread pool `INFORMATION_SCHEMA` tables are also available as Performance Schema tables. (See [Section 27.12.16, “Performance Schema Thread Pool Tables”](#).) The `INFORMATION_SCHEMA` tables are deprecated; expect them to be removed in a future version of MySQL. Applications should transition away from the old tables to the new tables. For example, if an application uses this query:

```
SELECT * FROM INFORMATION_SCHEMA.TP_THREAD_GROUP_STATS;
```

The application should use this query instead:

```
SELECT * FROM performance_schema.tp_thread_group_stats;
```

The `TP_THREAD_GROUP_STATS` table reports statistics per thread group. There is one row per group.

For descriptions of the columns in the `INFORMATION_SCHEMA TP_THREAD_GROUP_STATS` table, see [Section 27.12.16.2, “The tp\\_thread\\_group\\_stats Table”](#). The Performance Schema `tp_thread_group_stats` table has equivalent columns.

### 26.5.4 The INFORMATION\_SCHEMA TP\_THREAD\_STATE Table



#### Note

As of MySQL 8.0.14, the thread pool `INFORMATION_SCHEMA` tables are also available as Performance Schema tables. (See [Section 27.12.16, “Performance Schema Thread Pool Tables”](#).) The `INFORMATION_SCHEMA` tables are deprecated; expect them to be removed in a future version of MySQL. Applications should transition away from the old tables to the new tables. For example, if an application uses this query:

```
SELECT * FROM INFORMATION_SCHEMA.TP_THREAD_STATE;
```

The application should use this query instead:

```
SELECT * FROM performance_schema.tp_thread_state;
```

The `TP_THREAD_STATE` table has one row per thread created by the thread pool to handle connections.

For descriptions of the columns in the `INFORMATION_SCHEMA TP_THREAD_STATE` table, see [Section 27.12.16.3, “The tp\\_thread\\_state Table”](#). The Performance Schema `tp_thread_state` table has equivalent columns.

## 26.6 INFORMATION\_SCHEMA Connection-Control Tables

The following sections describe the [INFORMATION\\_SCHEMA](#) tables associated with the [CONNECTION\\_CONTROL](#) plugin.

### 26.6.1 INFORMATION\_SCHEMA Connection-Control Table Reference

The following table summarizes [INFORMATION\\_SCHEMA](#) connection-control tables. For greater detail, see the individual table descriptions.

**Table 26.8 INFORMATION\_SCHEMA Connection-Control Tables**

Table Name	Description
<a href="#">CONNECTION_CONTROL_FAILED_LOGIN_ATTEMPTS</a>	Current number of consecutive failed connection attempts per account

### 26.6.2 The INFORMATION\_SCHEMA CONNECTION\_CONTROL\_FAILED\_LOGIN\_ATTEMPTS Table

This table provides information about the current number of consecutive failed connection attempts per account (user/host combination).

[CONNECTION\\_CONTROL\\_FAILED\\_LOGIN\\_ATTEMPTS](#) has these columns:

- [USERHOST](#)

The user/host combination indicating an account that has failed connection attempts, in '`user_name '@ host_name'`' format.

- [FAILED\\_ATTEMPTS](#)

The current number of consecutive failed connection attempts for the [USERHOST](#) value. This counts all failed attempts, regardless of whether they were delayed. The number of attempts for which the server added a delay to its response is the difference between the [FAILED\\_ATTEMPTS](#) value and the [connection\\_control\\_failed\\_connections\\_threshold](#) system variable value.

#### Notes

- The [CONNECTION\\_CONTROL\\_FAILED\\_LOGIN\\_ATTEMPTS](#) plugin must be activated for this table to be available, and the [CONNECTION\\_CONTROL](#) plugin must be activated or the table contents are always empty. See [Section 6.4.2, “The Connection-Control Plugins”](#).
- The table contains rows only for accounts that have had one or more consecutive failed connection attempts without a subsequent successful attempt. When an account connects successfully, its failed-connection count is reset to zero and the server removes any row corresponding to the account.
- Assigning a value to the [connection\\_control\\_failed\\_connections\\_threshold](#) system variable at runtime resets all accumulated failed-connection counters to zero, which causes the table to become empty.

## 26.7 INFORMATION\_SCHEMA MySQL Enterprise Firewall Tables

The following sections describe the [INFORMATION\\_SCHEMA](#) tables associated with MySQL Enterprise Firewall (see [Section 6.4.7, “MySQL Enterprise Firewall”](#)). They provide views into the firewall in-memory data cache. These tables are available only if the appropriate firewall plugins are enabled.

### 26.7.1 INFORMATION\_SCHEMA Firewall Table Reference

The following table summarizes [INFORMATION\\_SCHEMA](#) firewall tables. For greater detail, see the individual table descriptions.

**Table 26.9 INFORMATION\_SCHEMA Firewall Tables**

Table Name	Description	Deprecated
MYSQL_FIREWALL_USERS	Firewall in-memory data for account profiles	8.0.26
MYSQL_FIREWALL_WHITELIST	Firewall in-memory data for account profile allowlists	8.0.26

## 26.7.2 The INFORMATION\_SCHEMA MYSQL\_FIREWALL\_USERS Table

The `MYSQL_FIREWALL_USERS` table provides a view into the in-memory data cache for MySQL Enterprise Firewall. It lists names and operational modes of registered firewall account profiles. It is used in conjunction with the `mysql.firewall_users` system table that provides persistent storage of firewall data; see [MySQL Enterprise Firewall Tables](#).

The `MYSQL_FIREWALL_USERS` table has these columns:

- `USERHOST`

The account profile name. Each account name has the format `user_name@host_name`.

- `MODE`

The current operational mode for the profile. Permitted mode values are `OFF`, `DETECTING`, `PROTECTING`, `RECORDING`, and `RESET`. For details about their meanings, see [Firewall Concepts](#).

As of MySQL 8.0.26, this table is deprecated and subject to removal in a future MySQL version. See [Migrating Account Profiles to Group Profiles](#).

## 26.7.3 The INFORMATION\_SCHEMA MYSQL\_FIREWALL\_WHITELIST Table

The `MYSQL_FIREWALL_WHITELIST` table provides a view into the in-memory data cache for MySQL Enterprise Firewall. It lists allowlist rules of registered firewall account profiles. It is used in conjunction with the `mysql.firewall_whitelist` system table that provides persistent storage of firewall data; see [MySQL Enterprise Firewall Tables](#).

The `MYSQL_FIREWALL_WHITELIST` table has these columns:

- `USERHOST`

The account profile name. Each account name has the format `user_name@host_name`.

- `RULE`

A normalized statement indicating an acceptable statement pattern for the profile. A profile allowlist is the union of its rules.

As of MySQL 8.0.26, this table is deprecated and subject to removal in a future MySQL version. See [Migrating Account Profiles to Group Profiles](#).

## 26.8 Extensions to SHOW Statements

Some extensions to `SHOW` statements accompany the implementation of [INFORMATION\\_SCHEMA](#):

- `SHOW` can be used to get information about the structure of [INFORMATION\\_SCHEMA](#) itself.

- Several `SHOW` statements accept a `WHERE` clause that provides more flexibility in specifying which rows to display.

`INFORMATION_SCHEMA` is an information database, so its name is included in the output from `SHOW DATABASES`. Similarly, `SHOW TABLES` can be used with `INFORMATION_SCHEMA` to obtain a list of its tables:

```
mysql> SHOW TABLES FROM INFORMATION_SCHEMA;
+-----+
| Tables_in_INFORMATION_SCHEMA |
+-----+
| CHARACTER_SETS
| COLLATIONS
| COLLATION_CHARACTER_SET_APPLICABILITY
| COLUMNS
| COLUMN_PRIVILEGES
| ENGINES
| EVENTS
| FILES
| KEY_COLUMN_USAGE
| PARTITIONS
| PLUGINS
| PROCESSLIST
| REFERENTIAL_CONSTRAINTS
| ROUTINES
| SCHEMATA
| SCHEMA_PRIVILEGES
| STATISTICS
| TABLES
| TABLE_CONSTRAINTS
| TABLE_PRIVILEGES
| TRIGGERS
| USER_PRIVILEGES
| VIEWS
+-----+
```

`SHOW COLUMNS` and `DESCRIBE` can display information about the columns in individual `INFORMATION_SCHEMA` tables.

`SHOW` statements that accept a `LIKE` clause to limit the rows displayed also permit a `WHERE` clause that specifies more general conditions that selected rows must satisfy:

```
SHOW CHARACTER SET
SHOW COLLATION
SHOW COLUMNS
SHOW DATABASES
SHOW FUNCTION STATUS
SHOW INDEX
SHOW OPEN TABLES
SHOW PROCEDURE STATUS
SHOW STATUS
SHOW TABLE STATUS
SHOW TABLES
SHOW TRIGGERS
SHOW VARIABLES
```

The `WHERE` clause, if present, is evaluated against the column names displayed by the `SHOW` statement. For example, the `SHOW CHARACTER SET` statement produces these output columns:

```
mysql> SHOW CHARACTER SET;
+-----+-----+-----+-----+
| Charset | Description          | Default collation | Maxlen |
+-----+-----+-----+-----+
| big5   | Big5 Traditional Chinese | big5_chinese_ci    | 2      |
| dec8   | DEC West European     | dec8_swedish_ci    | 1      |
| cp850  | DOS West European     | cp850_general_ci  | 1      |
| hp8    | HP West European      | hp8_english_ci    | 1      |
| koi8r  | KOI8-R Relcom Russian | koi8r_general_ci  | 1      |
| latin1 | cp1252 West European  | latin1_swedish_ci | 1      |
| latin2 | ISO 8859-2 Central European | latin2_general_ci | 1      |
```

...

To use a `WHERE` clause with `SHOW CHARACTER SET`, you would refer to those column names. As an example, the following statement displays information about character sets for which the default collation contains the string '`japanese`':

```
mysql> SHOW CHARACTER SET WHERE `Default collation` LIKE '%japanese%';
+-----+-----+-----+
| Charset | Description           | Default collation | Maxlen |
+-----+-----+-----+
| ujis   | EUC-JP Japanese        | ujis_japanese_ci | 3      |
| sjis   | Shift-JIS Japanese     | sjis_japanese_ci | 2      |
| cp932  | SJIS for Windows Japanese | cp932_japanese_ci | 2      |
| eucjpms | UJIS for Windows Japanese | eucjpms_japanese_ci | 3      |
+-----+-----+-----+
```

This statement displays the multibyte character sets:

```
mysql> SHOW CHARACTER SET WHERE Maxlen > 1;
+-----+-----+-----+
| Charset | Description           | Default collation | Maxlen |
+-----+-----+-----+
| big5   | Big5 Traditional Chinese | big5_chinese_ci   | 2      |
| cp932  | SJIS for Windows Japanese | cp932_japanese_ci | 2      |
| eucjpms | UJIS for Windows Japanese | eucjpms_japanese_ci | 3      |
| euckr  | EUC-KR Korean          | euckr_korean_ci  | 2      |
| gbk    | GB18030                 | gbk_chinese_ci   | 2      |
| gb18030 | China National Standard GB18030 | gb18030_chinese_ci | 4      |
| gb2312 | GB2312 Simplified Chinese | gb2312_chinese_ci | 2      |
| gbk    | GBK Simplified Chinese | gbk_chinese_ci   | 2      |
| sjis   | Shift-JIS Japanese     | sjis_japanese_ci | 2      |
| ucs2   | UCS-2 Unicode           | ucs2_general_ci | 2      |
| ujis   | EUC-JP Japanese        | ujis_japanese_ci | 3      |
| utf16  | UTF-16 Unicode          | utf16_general_ci | 4      |
| utf16le | UTF-16LE Unicode       | utf16le_general_ci | 4      |
| utf32  | UTF-32 Unicode          | utf32_general_ci | 4      |
| utf8mb3 | UTF-8 Unicode           | utf8mb3_general_ci | 3      |
| utf8mb4 | UTF-8 Unicode           | utf8mb4_0900_ai_ci | 4      |
+-----+-----+-----+
```

---

# Chapter 27 MySQL Performance Schema

## Table of Contents

27.1 Performance Schema Quick Start .....	4955
27.2 Performance Schema Build Configuration .....	4961
27.3 Performance Schema Startup Configuration .....	4961
27.4 Performance Schema Runtime Configuration .....	4963
27.4.1 Performance Schema Event Timing .....	4964
27.4.2 Performance Schema Event Filtering .....	4966
27.4.3 Event Pre-Filtering .....	4968
27.4.4 Pre-Filtering by Instrument .....	4968
27.4.5 Pre-Filtering by Object .....	4970
27.4.6 Pre-Filtering by Thread .....	4971
27.4.7 Pre-Filtering by Consumer .....	4973
27.4.8 Example Consumer Configurations .....	4976
27.4.9 Naming Instruments or Consumers for Filtering Operations .....	4981
27.4.10 Determining What Is Instrumented .....	4981
27.5 Performance Schema Queries .....	4982
27.6 Performance Schema Instrument Naming Conventions .....	4982
27.7 Performance Schema Status Monitoring .....	4986
27.8 Performance Schema Atom and Molecule Events .....	4989
27.9 Performance Schema Tables for Current and Historical Events .....	4989
27.10 Performance Schema Statement Digests and Sampling .....	4991
27.11 Performance Schema General Table Characteristics .....	4995
27.12 Performance Schema Table Descriptions .....	4996
27.12.1 Performance Schema Table Reference .....	4996
27.12.2 Performance Schema Setup Tables .....	5001
27.12.3 Performance Schema Instance Tables .....	5010
27.12.4 Performance Schema Wait Event Tables .....	5015
27.12.5 Performance Schema Stage Event Tables .....	5020
27.12.6 Performance Schema Statement Event Tables .....	5026
27.12.7 Performance Schema Transaction Tables .....	5037
27.12.8 Performance Schema Connection Tables .....	5044
27.12.9 Performance Schema Connection Attribute Tables .....	5048
27.12.10 Performance Schema User-Defined Variable Tables .....	5053
27.12.11 Performance Schema Replication Tables .....	5053
27.12.12 Performance Schema NDB Cluster Tables .....	5076
27.12.13 Performance Schema Lock Tables .....	5079
27.12.14 Performance Schema System Variable Tables .....	5088
27.12.15 Performance Schema Status Variable Tables .....	5093
27.12.16 Performance Schema Thread Pool Tables .....	5094
27.12.17 Performance Schema Firewall Tables .....	5099
27.12.18 Performance Schema Keyring Tables .....	5101
27.12.19 Performance Schema Clone Tables .....	5102
27.12.20 Performance Schema Summary Tables .....	5105
27.12.21 Performance Schema Miscellaneous Tables .....	5133
27.13 Performance Schema Option and Variable Reference .....	5151
27.14 Performance Schema Command Options .....	5155
27.15 Performance Schema System Variables .....	5156
27.16 Performance Schema Status Variables .....	5175
27.17 The Performance Schema Memory-Allocation Model .....	5178
27.18 Performance Schema and Plugins .....	5179
27.19 Using the Performance Schema to Diagnose Problems .....	5179
27.19.1 Query Profiling Using Performance Schema .....	5180
27.19.2 Obtaining Parent Event Information .....	5182

The MySQL Performance Schema is a feature for monitoring MySQL Server execution at a low level. The Performance Schema has these characteristics:

- The Performance Schema provides a way to inspect internal execution of the server at runtime. It is implemented using the `PERFORMANCE_SCHEMA` storage engine and the `performance_schema` database. The Performance Schema focuses primarily on performance data. This differs from `INFORMATION_SCHEMA`, which serves for inspection of metadata.
- The Performance Schema monitors server events. An “event” is anything the server does that takes time and has been instrumented so that timing information can be collected. In general, an event could be a function call, a wait for the operating system, a stage of an SQL statement execution such as parsing or sorting, or an entire statement or group of statements. Event collection provides access to information about synchronization calls (such as for mutexes) file and table I/O, table locks, and so forth for the server and for several storage engines.
- Performance Schema events are distinct from events written to the server's binary log (which describe data modifications) and Event Scheduler events (which are a type of stored program).
- Performance Schema events are specific to a given instance of the MySQL Server. Performance Schema tables are considered local to the server, and changes to them are not replicated or written to the binary log.
- Current events are available, as well as event histories and summaries. This enables you to determine how many times instrumented activities were performed and how much time they took. Event information is available to show the activities of specific threads, or activity associated with particular objects such as a mutex or file.
- The `PERFORMANCE_SCHEMA` storage engine collects event data using “instrumentation points” in server source code.
- Collected events are stored in tables in the `performance_schema` database. These tables can be queried using `SELECT` statements like other tables.
- Performance Schema configuration can be modified dynamically by updating tables in the `performance_schema` database through SQL statements. Configuration changes affect data collection immediately.
- Tables in the Performance Schema are in-memory tables that use no persistent on-disk storage. The contents are repopulated beginning at server startup and discarded at server shutdown.
- Monitoring is available on all platforms supported by MySQL.

Some limitations might apply: The types of timers might vary per platform. Instruments that apply to storage engines might not be implemented for all storage engines. Instrumentation of each third-party engine is the responsibility of the engine maintainer. See also [Section 27.20, “Restrictions on Performance Schema”](#).

- Data collection is implemented by modifying the server source code to add instrumentation. There are no separate threads associated with the Performance Schema, unlike other features such as replication or the Event Scheduler.

The Performance Schema is intended to provide access to useful information about server execution while having minimal impact on server performance. The implementation follows these design goals:

- Activating the Performance Schema causes no changes in server behavior. For example, it does not cause thread scheduling to change, and it does not cause query execution plans (as shown by `EXPLAIN`) to change.
- Server monitoring occurs continuously and unobtrusively with very little overhead. Activating the Performance Schema does not make the server unusable.

- The parser is unchanged. There are no new keywords or statements.
- Execution of server code proceeds normally even if the Performance Schema fails internally.
- When there is a choice between performing processing during event collection initially or during event retrieval later, priority is given to making collection faster. This is because collection is ongoing whereas retrieval is on demand and might never happen at all.
- Most Performance Schema tables have indexes, which gives the optimizer access to execution plans other than full table scans. For more information, see [Section 8.2.4, “Optimizing Performance Schema Queries”](#).
- It is easy to add new instrumentation points.
- Instrumentation is versioned. If the instrumentation implementation changes, previously instrumented code continues to work. This benefits developers of third-party plugins because it is not necessary to upgrade each plugin to stay synchronized with the latest Performance Schema changes.



#### Note

The MySQL `sys` schema is a set of objects that provides convenient access to data collected by the Performance Schema. The `sys` schema is installed by default. For usage instructions, see [Chapter 28, MySQL sys Schema](#).

## 27.1 Performance Schema Quick Start

This section briefly introduces the Performance Schema with examples that show how to use it. For additional examples, see [Section 27.19, “Using the Performance Schema to Diagnose Problems”](#).

The Performance Schema is enabled by default. To enable or disable it explicitly, start the server with the `performance_schema` variable set to an appropriate value. For example, use these lines in the server `my.cnf` file:

```
[mysqld]
performance_schema=ON
```

When the server starts, it sees `performance_schema` and attempts to initialize the Performance Schema. To verify successful initialization, use this statement:

```
mysql> SHOW VARIABLES LIKE 'performance_schema';
+-----+-----+
| Variable_name      | Value   |
+-----+-----+
| performance_schema | ON      |
+-----+-----+
```

A value of `ON` means that the Performance Schema initialized successfully and is ready for use. A value of `OFF` means that some error occurred. Check the server error log for information about what went wrong.

The Performance Schema is implemented as a storage engine, so you can see it listed in the output from the Information Schema `ENGINES` table or the `SHOW ENGINES` statement:

```
mysql> SELECT * FROM INFORMATION_SCHEMA.ENGINES
      WHERE ENGINE='PERFORMANCE_SCHEMA'\G
***** 1. row *****
      ENGINE: PERFORMANCE_SCHEMA
      SUPPORT: YES
      COMMENT: Performance Schema
TRANSACTIONS: NO
          XA: NO
    SAVEPOINTS: NO

mysql> SHOW ENGINES\G
...
```

```

Engine: PERFORMANCE_SCHEMA
Support: YES
Comment: Performance Schema
Transactions: NO
XA: NO
Savepoints: NO
...

```

The `PERFORMANCE_SCHEMA` storage engine operates on tables in the `performance_schema` database. You can make `performance_schema` the default database so that references to its tables need not be qualified with the database name:

```
mysql> USE performance_schema;
```

Performance Schema tables are stored in the `performance_schema` database. Information about the structure of this database and its tables can be obtained, as for any other database, by selecting from the `INFORMATION_SCHEMA` database or by using `SHOW` statements. For example, use either of these statements to see what Performance Schema tables exist:

```
mysql> SELECT TABLE_NAME FROM INFORMATION_SCHEMA.TABLES
      WHERE TABLE_SCHEMA = 'performance_schema';
+-----+
| TABLE_NAME
+-----+
| accounts
| cond_instances
...
| events_stages_current
| events_stages_history
| events_stages_history_long
| events_stages_summary_by_account_by_event_name
| events_stages_summary_by_host_by_event_name
| events_stages_summary_by_thread_by_event_name
| events_stages_summary_by_user_by_event_name
| events_stages_summary_global_by_event_name
| events_statements_current
| events_statements_history
| events_statements_history_long
...
| file_instances
| file_summary_by_event_name
| file_summary_by_instance
| host_cache
| hosts
| memory_summary_by_account_by_event_name
| memory_summary_by_host_by_event_name
| memory_summary_by_thread_by_event_name
| memory_summary_by_user_by_event_name
| memory_summary_global_by_event_name
| metadata_locks
| mutex_instances
| objects_summary_global_by_type
| performance_timers
| replication_connection_configuration
| replication_connection_status
| replication_applier_configuration
| replication_applier_status
| replication_applier_status_by_coordinator
| replication_applier_status_by_worker
| rwlock_instances
| session_account_connect_attrs
| session_connect_attrs
| setup_actors
| setup_consumers
| setup_instruments
| setup_objects
| socket_instances
| socket_summary_by_event_name
| socket_summary_by_instance
| table_handles
| table_io_waits_summary_by_index_usage

```

```

| table_io_waits_summary_by_table
| table_lock_waits_summary_by_table
| threads
| users
+-----+
mysql> SHOW TABLES FROM performance_schema;
+-----+
| Tables_in_performance_schema
+-----+
| accounts
| cond_instances
| events_stages_current
| events_stages_history
| events_stages_history_long
+-----+
...
```

The number of Performance Schema tables increases over time as implementation of additional instrumentation proceeds.

The name of the `performance_schema` database is lowercase, as are the names of tables within it. Queries should specify the names in lowercase.

To see the structure of individual tables, use `SHOW CREATE TABLE`:

```

mysql> SHOW CREATE TABLE performance_schema.setup_consumers\G
***** 1. row *****
      Table: setup_consumers
Create Table: CREATE TABLE `setup_consumers` (
  `NAME` varchar(64) NOT NULL,
  `ENABLED` enum('YES','NO') NOT NULL,
  PRIMARY KEY (`NAME`)
) ENGINE=PERFORMANCE_SCHEMA DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
```

Table structure is also available by selecting from tables such as `INFORMATION_SCHEMA.COLUMNS` or by using statements such as `SHOW COLUMNS`.

Tables in the `performance_schema` database can be grouped according to the type of information in them: Current events, event histories and summaries, object instances, and setup (configuration) information. The following examples illustrate a few uses for these tables. For detailed information about the tables in each group, see [Section 27.12, “Performance Schema Table Descriptions”](#).

Initially, not all instruments and consumers are enabled, so the performance schema does not collect all events. To turn all of these on and enable event timing, execute two statements (the row counts may differ depending on MySQL version):

```

mysql> UPDATE performance_schema.setup_instruments
      SET ENABLED = 'YES', TIMED = 'YES';
Query OK, 560 rows affected (0.04 sec)
mysql> UPDATE performance_schema.setup_consumers
      SET ENABLED = 'YES';
Query OK, 10 rows affected (0.00 sec)
```

To see what the server is doing at the moment, examine the `events_waits_current` table. It contains one row per thread showing each thread's most recent monitored event:

```

mysql> SELECT *
      FROM performance_schema.events_waits_current\G
***** 1. row *****
      THREAD_ID: 0
      EVENT_ID: 5523
    END_EVENT_ID: 5523
      EVENT_NAME: wait/synch/mutex/mysys/THR_LOCK::mutex
      SOURCE: thr_lock.c:525
    TIMER_START: 201660494489586
    TIMER_END: 201660494576112
    TIMER_WAIT: 86526
      SPINS: NULL
```

```

OBJECT_SCHEMA: NULL
OBJECT_NAME: NULL
INDEX_NAME: NULL
OBJECT_TYPE: NULL
OBJECT_INSTANCE_BEGIN: 142270668
NESTING_EVENT_ID: NULL
NESTING_EVENT_TYPE: NULL
OPERATION: lock
NUMBER_OF_BYTES: NULL
FLAGS: 0
...

```

This event indicates that thread 0 was waiting for 86,526 picoseconds to acquire a lock on `THR_LOCK::mutex`, a mutex in the `mysys` subsystem. The first few columns provide the following information:

- The ID columns indicate which thread the event comes from and the event number.
- `EVENT_NAME` indicates what was instrumented and `SOURCE` indicates which source file contains the instrumented code.
- The timer columns show when the event started and stopped and how long it took. If an event is still in progress, the `TIMER_END` and `TIMER_WAIT` values are `NULL`. Timer values are approximate and expressed in picoseconds. For information about timers and event time collection, see [Section 27.4.1, “Performance Schema Event Timing”](#).

The history tables contain the same kind of rows as the current-events table but have more rows and show what the server has been doing “recently” rather than “currently.” The `events_waits_history` and `events_waits_history_long` tables contain the most recent 10 events per thread and most recent 10,000 events, respectively. For example, to see information for recent events produced by thread 13, do this:

```

mysql> SELECT EVENT_ID, EVENT_NAME, TIMER_WAIT
      FROM performance_schema.events_waits_history
     WHERE THREAD_ID = 13
       ORDER BY EVENT_ID;
+-----+-----+-----+
| EVENT_ID | EVENT_NAME          | TIMER_WAIT |
+-----+-----+-----+
|    86    | wait/synch/mutex/my... | 686322    |
|    87    | wait/synch/mutex/my... | 320535    |
|    88    | wait/synch/mutex/my... | 339390    |
|    89    | wait/synch/mutex/my... | 377100    |
|    90    | wait/synch/mutex/sql/... | 614673    |
|    91    | wait/synch/mutex/sql/... | 659925    |
|    92    | wait/synch/mutex/sql/THD:... | 494001    |
|    93    | wait/synch/mutex/my... | 222489    |
|    94    | wait/synch/mutex/my... | 214947    |
|    95    | wait/synch/mutex/my... | 312993    |
+-----+-----+-----+

```

As new events are added to a history table, older events are discarded if the table is full.

Summary tables provide aggregated information for all events over time. The tables in this group summarize event data in different ways. To see which instruments have been executed the most times or have taken the most wait time, sort the `events_waits_summary_global_by_event_name` table on the `COUNT_STAR` or `SUM_TIMER_WAIT` column, which correspond to a `COUNT(*)` or `SUM(TIMER_WAIT)` value, respectively, calculated over all events:

```

mysql> SELECT EVENT_NAME, COUNT_STAR
      FROM performance_schema.events_waits_summary_global_by_event_name
       ORDER BY COUNT_STAR DESC LIMIT 10;
+-----+-----+
| EVENT_NAME          | COUNT_STAR |
+-----+-----+
| wait/synch/mutex/my... | 6419      |
| wait/io/file/sql/FRM | 452       |
+-----+-----+

```

```

| wait/synch/mutex/sql/LOCK_plugin          |      337 |
| wait/synch/mutex/mysys/THR_LOCK_open     |      187 |
| wait/synch/mutex/mysys/LOCK_alarm        |      147 |
| wait/synch/mutex/sql/THD::LOCK_thd_data |      115 |
| wait/io/file/myisam/kfile                |      102 |
| wait/synch/mutex/sql/LOCK_global_system_variables |      89 |
| wait/synch/mutex/mysys/THR_LOCK::mutex   |      89 |
| wait/synch/mutex/sql/LOCK_open           |      88 |
+-----+-----+-----+-----+-----+-----+
mysql> SELECT EVENT_NAME, SUM_TIMER_WAIT
      FROM performance_schema.events_waits_summary_global_by_event_name
     ORDER BY SUM_TIMER_WAIT DESC LIMIT 10;
+-----+-----+
| EVENT_NAME          | SUM_TIMER_WAIT |
+-----+-----+
| wait/io/file/sql/MYSQL_LOG               | 1599816582 |
| wait/synch/mutex/mysys/THR_LOCK_malloc   | 1530083250 |
| wait/io/file/sql/binlog_index            | 1385291934 |
| wait/io/file/sql/FRM                     | 1292823243 |
| wait/io/file/myisam/kfile                | 411193611  |
| wait/io/file/myisam/dfile                | 322401645  |
| wait/synch/mutex/mysys/LOCK_alarm       | 145126935  |
| wait/io/file/sql/casetest               | 104324715  |
| wait/synch/mutex/sql/LOCK_plugin        | 86027823   |
| wait/io/file/sql/pid                   | 72591750   |
+-----+-----+

```

These results show that the `THR_LOCK_malloc` mutex is “hot,” both in terms of how often it is used and amount of time that threads wait attempting to acquire it.



### Note

The `THR_LOCK_malloc` mutex is used only in debug builds. In production builds it is not hot because it is nonexistent.

Instance tables document what types of objects are instrumented. An instrumented object, when used by the server, produces an event. These tables provide event names and explanatory notes or status information. For example, the `file_instances` table lists instances of instruments for file I/O operations and their associated files:

```

mysql> SELECT *
      FROM performance_schema.file_instances\G
***** 1. row *****
FILE_NAME: /opt/mysql-log/60500/binlog.000007
EVENT_NAME: wait/io/file/sql/binlog
OPEN_COUNT: 0
***** 2. row *****
FILE_NAME: /opt/mysql/60500/data/mysql/tables_priv.MYI
EVENT_NAME: wait/io/file/myisam/kfile
OPEN_COUNT: 1
***** 3. row *****
FILE_NAME: /opt/mysql/60500/data/mysql/columns_priv.MYI
EVENT_NAME: wait/io/file/myisam/kfile
OPEN_COUNT: 1
...

```

Setup tables are used to configure and display monitoring characteristics. For example, `setup_instruments` lists the set of instruments for which events can be collected and shows which of them are enabled:

```

mysql> SELECT NAME, ENABLED, TIMED
      FROM performance_schema.setup_instruments;
+-----+-----+-----+
| NAME          | ENABLED | TIMED |
+-----+-----+-----+
| stage/sql/end | NO      | NO    |
| stage/sql/executing | NO      | NO    |
+-----+-----+-----+

```

stage/sql/init	NO	NO	
stage/sql/insert	NO	NO	
...			
statement/sql/load	YES	YES	
statement/sql/grant	YES	YES	
statement/sql/check	YES	YES	
statement/sql/flush	YES	YES	
...			
wait/synch/mutex/sql/LOCK_global_read_lock	YES	YES	
wait/synch/mutex/sql/LOCK_global_system_variables	YES	YES	
wait/synch/mutex/sql/LOCK_lock_db	YES	YES	
wait/synch/mutex/sql/LOCK_manager	YES	YES	
...			
wait/synch/rwlock/sql/LOCK_grant	YES	YES	
wait/synch/rwlock/sql/LOGGER::LOCK_logger	YES	YES	
wait/synch/rwlock/sql/LOCK_sys_init_connect	YES	YES	
wait/synch/rwlock/sql/LOCK_sys_init_slave	YES	YES	
...			
wait/io/file/sql/binlog	YES	YES	
wait/io/file/sql/binlog_index	YES	YES	
wait/io/file/sql/casetest	YES	YES	
wait/io/file/sql/dbopt	YES	YES	
...			

To understand how to interpret instrument names, see [Section 27.6, “Performance Schema Instrument Naming Conventions”](#).

To control whether events are collected for an instrument, set its `ENABLED` value to `YES` or `NO`. For example:

```
mysql> UPDATE performance_schema.setup_instruments
    SET ENABLED = 'NO'
    WHERE NAME = 'wait/synch/mutex/sql/LOCK_mysql_create_db';
```

The Performance Schema uses collected events to update tables in the `performance_schema` database, which act as “consumers” of event information. The `setup_consumers` table lists the available consumers and which are enabled:

```
mysql> SELECT * FROM performance_schema.setup_consumers;
+-----+-----+
| NAME           | ENABLED |
+-----+-----+
| events_stages_current | NO      |
| events_stages_history | NO      |
| events_stages_history_long | NO      |
| events_statements_cpu | NO      |
| events_statements_current | YES     |
| events_statements_history | YES     |
| events_statements_history_long | NO      |
| events_transactions_current | YES     |
| events_transactions_history | YES     |
| events_transactions_history_long | NO      |
| events_waits_current | NO      |
| events_waits_history | NO      |
| events_waits_history_long | NO      |
| global_instrumentation | YES     |
| thread_instrumentation | YES     |
| statements_digest | YES     |
+-----+-----+
```

To control whether the Performance Schema maintains a consumer as a destination for event information, set its `ENABLED` value.

For more information about the setup tables and how to use them to control event collection, see [Section 27.4.2, “Performance Schema Event Filtering”](#).

There are some miscellaneous tables that do not fall into any of the previous groups. For example, `performance_timers` lists the available event timers and their characteristics. For information about timers, see [Section 27.4.1, “Performance Schema Event Timing”](#).

## 27.2 Performance Schema Build Configuration

The Performance Schema is mandatory and always compiled in. It is possible to exclude certain parts of the Performance Schema instrumentation. For example, to exclude stage and statement instrumentation, do this:

```
$> cmake . \
  -DDISABLE_PSI_STAGE=1 \
  -DDISABLE_PSI_STATEMENT=1
```

For more information, see the descriptions of the `DISABLE_PSI_XXX` CMake options in [Section 2.8.7, “MySQL Source-Configuration Options”](#).

If you install MySQL over a previous installation that was configured without the Performance Schema (or with an older version of the Performance Schema that has missing or out-of-date tables). One indication of this issue is the presence of messages such as the following in the error log:

```
[ERROR] Native table 'performance_schema'.events_waits_history
has the wrong structure
[ERROR] Native table 'performance_schema'.events_waits_history_long'
has the wrong structure
...
```

To correct that problem, perform the MySQL upgrade procedure. See [Section 2.10, “Upgrading MySQL”](#).

Because the Performance Schema is configured into the server at build time, a row for `PERFORMANCE_SCHEMA` appears in the output from `SHOW ENGINES`. This means that the Performance Schema is available, not that it is enabled. To enable it, you must do so at server startup, as described in the next section.

## 27.3 Performance Schema Startup Configuration

To use the MySQL Performance Schema, it must be enabled at server startup to enable event collection to occur.

The Performance Schema is enabled by default. To enable or disable it explicitly, start the server with the `performance_schema` variable set to an appropriate value. For example, use these lines in the server `my.cnf` file:

```
[mysqld]
performance_schema=ON
```

If the server is unable to allocate any internal buffer during Performance Schema initialization, the Performance Schema disables itself and sets `performance_schema` to `OFF`, and the server runs without instrumentation.

The Performance Schema also permits instrument and consumer configuration at server startup.

To control an instrument at server startup, use an option of this form:

```
--performance-schema-instrument='instrument_name=value'
```

Here, `instrument_name` is an instrument name such as `wait/synch/mutex/sql/LOCK_open`, and `value` is one of these values:

- `OFF`, `FALSE`, or `0`: Disable the instrument
- `ON`, `TRUE`, or `1`: Enable and time the instrument
- `COUNTED`: Enable and count (rather than time) the instrument

Each `--performance-schema-instrument` option can specify only one instrument name, but multiple instances of the option can be given to configure multiple instruments. In addition, patterns

are permitted in instrument names to configure instruments that match the pattern. To configure all condition synchronization instruments as enabled and counted, use this option:

```
--performance-schema-instrument='wait/synch/cond/=%=COUNTED'
```

To disable all instruments, use this option:

```
--performance-schema-instrument='=%=OFF'
```

Exception: The `memory/performance_schema/%` instruments are built in and cannot be disabled at startup.

Longer instrument name strings take precedence over shorter pattern names, regardless of order. For information about specifying patterns to select instruments, see [Section 27.4.9, “Naming Instruments or Consumers for Filtering Operations”](#).

An unrecognized instrument name is ignored. It is possible that a plugin installed later may create the instrument, at which time the name is recognized and configured.

To control a consumer at server startup, use an option of this form:

```
--performance-schema-consumer-consumer_name=value
```

Here, `consumer_name` is a consumer name such as `events_waits_history`, and `value` is one of these values:

- `OFF`, `FALSE`, or `0`: Do not collect events for the consumer
- `ON`, `TRUE`, or `1`: Collect events for the consumer

For example, to enable the `events_waits_history` consumer, use this option:

```
--performance-schema-consumer-events-waits-history=ON
```

The permitted consumer names can be found by examining the `setup_consumers` table. Patterns are not permitted. Consumer names in the `setup_consumers` table use underscores, but for consumers set at startup, dashes and underscores within the name are equivalent.

The Performance Schema includes several system variables that provide configuration information:

```
mysql> SHOW VARIABLES LIKE 'perf%';
+-----+-----+
| Variable_name          | Value   |
+-----+-----+
| performance_schema      | ON      |
| performance_schema_accounts_size | 100    |
| performance_schema_digests_size | 200    |
| performance_schema_events_stages_history_long_size | 10000 |
| performance_schema_events_stages_history_size       | 10     |
| performance_schema_events_statements_history_long_size | 10000 |
| performance_schema_events_statements_history_size | 10     |
| performance_schema_events_waits_history_long_size | 10000 |
| performance_schema_events_waits_history_size       | 10     |
| performance_schema_hosts_size                      | 100    |
| performance_schema_max_cond_classes               | 80     |
| performance_schema_max_cond_instances             | 1000   |
...

```

The `performance_schema` variable is `ON` or `OFF` to indicate whether the Performance Schema is enabled or disabled. The other variables indicate table sizes (number of rows) or memory allocation values.



#### Note

With the Performance Schema enabled, the number of Performance Schema instances affects the server memory footprint, perhaps to a large extent. The Performance Schema autoscales many parameters to use memory only as

required; see [Section 27.17, “The Performance Schema Memory-Allocation Model”](#).

To change the value of Performance Schema system variables, set them at server startup. For example, put the following lines in a `my.cnf` file to change the sizes of the history tables for wait events:

```
[mysqld]
performance_schema
performance_schema_events_waits_history_size=20
performance_schema_events_waits_history_long_size=15000
```

The Performance Schema automatically sizes the values of several of its parameters at server startup if they are not set explicitly. For example, it sizes the parameters that control the sizes of the events waits tables this way. The Performance Schema allocates memory incrementally, scaling its memory use to actual server load, instead of allocating all the memory it needs during server startup. Consequently, many sizing parameters need not be set at all. To see which parameters are autosized or autoscaled, use `mysqld --verbose --help` and examine the option descriptions, or see [Section 27.15, “Performance Schema System Variables”](#).

For each autosized parameter that is not set at server startup, the Performance Schema determines how to set its value based on the value of the following system values, which are considered as “hints” about how you have configured your MySQL server:

```
max_connections
open_files_limit
table_definition_cache
table_open_cache
```

To override autosizing or autoscaling for a given parameter, set it to a value other than `-1` at startup. In this case, the Performance Schema assigns it the specified value.

At runtime, `SHOW VARIABLES` displays the actual values that autosized parameters were set to. Autoscaled parameters display with a value of `-1`.

If the Performance Schema is disabled, its autosized and autoscaled parameters remain set to `-1` and `SHOW VARIABLES` displays `-1`.

## 27.4 Performance Schema Runtime Configuration

Specific Performance Schema features can be enabled at runtime to control which types of event collection occur.

Performance Schema setup tables contain information about monitoring configuration:

```
mysql> SELECT TABLE_NAME FROM INFORMATION_SCHEMA.TABLES
      WHERE TABLE_SCHEMA = 'performance_schema'
        AND TABLE_NAME LIKE 'setup%';
+-----+
| TABLE_NAME |
+-----+
| setup_actors |
| setup_consumers |
| setup_instruments |
| setup_objects |
| setup_threads |
+-----+
```

You can examine the contents of these tables to obtain information about Performance Schema monitoring characteristics. If you have the `UPDATE` privilege, you can change Performance Schema operation by modifying setup tables to affect how monitoring occurs. For additional details about these tables, see [Section 27.12.2, “Performance Schema Setup Tables”](#).

The `setup_instruments` and `setup_consumers` tables list the instruments for which events can be collected and the types of consumers for which event information actually is collected, respectively.

Other setup tables enable further modification of the monitoring configuration. [Section 27.4.2, “Performance Schema Event Filtering”](#), discusses how you can modify these tables to affect event collection.

If there are Performance Schema configuration changes that must be made at runtime using SQL statements and you would like these changes to take effect each time the server starts, put the statements in a file and start the server with the `init_file` system variable set to name the file. This strategy can also be useful if you have multiple monitoring configurations, each tailored to produce a different kind of monitoring, such as casual server health monitoring, incident investigation, application behavior troubleshooting, and so forth. Put the statements for each monitoring configuration into their own file and specify the appropriate file as the `init_file` value when you start the server.

## 27.4.1 Performance Schema Event Timing

Events are collected by means of instrumentation added to the server source code. Instruments time events, which is how the Performance Schema provides an idea of how long events take. It is also possible to configure instruments not to collect timing information. This section discusses the available timers and their characteristics, and how timing values are represented in events.

### Performance Schema Timers

Performance Schema timers vary in precision and amount of overhead. To see what timers are available and their characteristics, check the `performance_timers` table:

TIMER_NAME	TIMER_FREQUENCY	TIMER_RESOLUTION	TIMER_OVERHEAD
CYCLE	2389029850	1	72
NANOSECOND	1000000000	1	112
MICROSECOND	1000000	1	136
MILLISECOND	1036	1	168
THREAD_CPU	339101694	1	798

If the values associated with a given timer name are `NULL`, that timer is not supported on your platform.

The columns have these meanings:

- The `TIMER_NAME` column shows the names of the available timers. `CYCLE` refers to the timer that is based on the CPU (processor) cycle counter.
- `TIMER_FREQUENCY` indicates the number of timer units per second. For a cycle timer, the frequency is generally related to the CPU speed. The value shown was obtained on a system with a 2.4GHz processor. The other timers are based on fixed fractions of seconds.
- `TIMER_RESOLUTION` indicates the number of timer units by which timer values increase at a time. If a timer has a resolution of 10, its value increases by 10 each time.
- `TIMER_OVERHEAD` is the minimal number of cycles of overhead to obtain one timing with the given timer. The overhead per event is twice the value displayed because the timer is invoked at the beginning and end of the event.

The Performance Schema assigns timers as follows:

- The wait timer uses `CYCLE`.
- The idle, stage, statement, and transaction timers use `NANOSECOND` on platforms where the `NANOSECOND` timer is available, `MICROSECOND` otherwise.

At server startup, the Performance Schema verifies that assumptions made at build time about timer assignments are correct, and displays a warning if a timer is not available.

To time wait events, the most important criterion is to reduce overhead, at the possible expense of the timer accuracy, so using the `CYCLE` timer is the best.

The time a statement (or stage) takes to execute is in general orders of magnitude larger than the time it takes to execute a single wait. To time statements, the most important criterion is to have an accurate measure, which is not affected by changes in processor frequency, so using a timer which is not based on cycles is the best. The default timer for statements is `NANOSECOND`. The extra “overhead” compared to the `CYCLE` timer is not significant, because the overhead caused by calling a timer twice (once when the statement starts, once when it ends) is orders of magnitude less compared to the CPU time used to execute the statement itself. Using the `CYCLE` timer has no benefit here, only drawbacks.

The precision offered by the cycle counter depends on processor speed. If the processor runs at 1 GHz (one billion cycles/second) or higher, the cycle counter delivers sub-nanosecond precision. Using the cycle counter is much cheaper than getting the actual time of day. For example, the standard `gettimeofday()` function can take hundreds of cycles, which is an unacceptable overhead for data gathering that may occur thousands or millions of times per second.

Cycle counters also have disadvantages:

- End users expect to see timings in wall-clock units, such as fractions of a second. Converting from cycles to fractions of seconds can be expensive. For this reason, the conversion is a quick and fairly rough multiplication operation.
- Processor cycle rate might change, such as when a laptop goes into power-saving mode or when a CPU slows down to reduce heat generation. If a processor's cycle rate fluctuates, conversion from cycles to real-time units is subject to error.
- Cycle counters might be unreliable or unavailable depending on the processor or the operating system. For example, on Pentiums, the instruction is `RDTSC` (an assembly-language rather than a C instruction) and it is theoretically possible for the operating system to prevent user-mode programs from using it.
- Some processor details related to out-of-order execution or multiprocessor synchronization might cause the counter to seem fast or slow by up to 1000 cycles.

MySQL works with cycle counters on x386 (Windows, macOS, Linux, Solaris, and other Unix flavors), PowerPC, and IA-64.

## Performance Schema Timer Representation in Events

Rows in Performance Schema tables that store current events and historical events have three columns to represent timing information: `TIMER_START` and `TIMER_END` indicate when an event started and finished, and `TIMER_WAIT` indicates event duration.

The `setup_instruments` table has an `ENABLED` column to indicate the instruments for which to collect events. The table also has a `TIMED` column to indicate which instruments are timed. If an instrument is not enabled, it produces no events. If an enabled instrument is not timed, events produced by the instrument have `NULL` for the `TIMER_START`, `TIMER_END`, and `TIMER_WAIT` timer values. This in turn causes those values to be ignored when calculating aggregate time values in summary tables (sum, minimum, maximum, and average).

Internally, times within events are stored in units given by the timer in effect when event timing begins. For display when events are retrieved from Performance Schema tables, times are shown in picoseconds (trillionths of a second) to normalize them to a standard unit, regardless of which timer is selected.

The timer baseline (“time zero”) occurs at Performance Schema initialization during server startup. `TIMER_START` and `TIMER_END` values in events represent picoseconds since the baseline. `TIMER_WAIT` values are durations in picoseconds.

Picosecond values in events are approximate. Their accuracy is subject to the usual forms of error associated with conversion from one unit to another. If the `CYCLE` timer is used and the processor rate varies, there might be drift. For these reasons, it is not reasonable to look at the `TIMER_START` value for an event as an accurate measure of time elapsed since server startup. On the other hand, it is reasonable to use `TIMER_START` or `TIMER_WAIT` values in `ORDER BY` clauses to order events by start time or duration.

The choice of picoseconds in events rather than a value such as microseconds has a performance basis. One implementation goal was to show results in a uniform time unit, regardless of the timer. In an ideal world this time unit would look like a wall-clock unit and be reasonably precise; in other words, microseconds. But to convert cycles or nanoseconds to microseconds, it would be necessary to perform a division for every instrumentation. Division is expensive on many platforms. Multiplication is not expensive, so that is what is used. Therefore, the time unit is an integer multiple of the highest possible `TIMER_FREQUENCY` value, using a multiplier large enough to ensure that there is no major precision loss. The result is that the time unit is “picoseconds.” This precision is spurious, but the decision enables overhead to be minimized.

While a wait, stage, statement, or transaction event is executing, the respective current-event tables display current-event timing information:

```
events_waits_current
events_stages_current
events_statements_current
events_transactions_current
```

To make it possible to determine how long a not-yet-completed event has been running, the timer columns are set as follows:

- `TIMER_START` is populated.
- `TIMER_END` is populated with the current timer value.
- `TIMER_WAIT` is populated with the time elapsed so far (`TIMER_END - TIMER_START`).

Events that have not yet completed have an `END_EVENT_ID` value of `NULL`. To assess time elapsed so far for an event, use the `TIMER_WAIT` column. Therefore, to identify events that have not yet completed and have taken longer than `N` picoseconds thus far, monitoring applications can use this expression in queries:

```
WHERE END_EVENT_ID IS NULL AND TIMER_WAIT > N
```

Event identification as just described assumes that the corresponding instruments have `ENABLED` and `TIMED` set to `YES` and that the relevant consumers are enabled.

## 27.4.2 Performance Schema Event Filtering

Events are processed in a producer/consumer fashion:

- Instrumented code is the source for events and produces events to be collected. The `setup_instruments` table lists the instruments for which events can be collected, whether they are enabled, and (for enabled instruments) whether to collect timing information:

```
mysql> SELECT NAME, ENABLED, TIMED
    FROM performance_schema.setup_instruments;
+-----+-----+-----+
| NAME           | ENABLED | TIMED |
+-----+-----+-----+
| ...            |         |       |
| wait/synch/mutex/sql/LOCK_global_read_lock | YES    | YES   |
| wait/synch/mutex/sql/LOCK_global_system_variables | YES    | YES   |
| wait/synch/mutex/sql/LOCK_lock_db      | YES    | YES   |
| wait/synch/mutex/sql/LOCK_manager     | YES    | YES   |
```

...

The `setup_instruments` table provides the most basic form of control over event production. To further refine event production based on the type of object or thread being monitored, other tables may be used as described in [Section 27.4.3, “Event Pre-Filtering”](#).

- Performance Schema tables are the destinations for events and consume events. The `setup_consumers` table lists the types of consumers to which event information can be sent and whether they are enabled:

NAME	ENABLED
events_stages_current	NO
events_stages_history	NO
events_stages_history_long	NO
events_statements_cpu	NO
events_statements_current	YES
events_statements_history	YES
events_statements_history_long	NO
events_transactions_current	YES
events_transactions_history	YES
events_transactions_history_long	NO
events_waits_current	NO
events_waits_history	NO
events_waits_history_long	NO
global_instrumentation	YES
thread_instrumentation	YES
statements_digest	YES

Filtering can be done at different stages of performance monitoring:

- **Pre-filtering.** This is done by modifying Performance Schema configuration so that only certain types of events are collected from producers, and collected events update only certain consumers. To do this, enable or disable instruments or consumers. Pre-filtering is done by the Performance Schema and has a global effect that applies to all users.

Reasons to use pre-filtering:

- To reduce overhead. Performance Schema overhead should be minimal even with all instruments enabled, but perhaps you want to reduce it further. Or you do not care about timing events and want to disable the timing code to eliminate timing overhead.
- To avoid filling the current-events or history tables with events in which you have no interest. Pre-filtering leaves more “room” in these tables for instances of rows for enabled instrument types. If you enable only file instruments with pre-filtering, no rows are collected for nonfile instruments. With post-filtering, nonfile events are collected, leaving fewer rows for file events.
- To avoid maintaining some kinds of event tables. If you disable a consumer, the server does not spend time maintaining destinations for that consumer. For example, if you do not care about event histories, you can disable the history table consumers to improve performance.

- **Post-filtering.** This involves the use of `WHERE` clauses in queries that select information from Performance Schema tables, to specify which of the available events you want to see. Post-filtering is performed on a per-user basis because individual users select which of the available events are of interest.

Reasons to use post-filtering:

- To avoid making decisions for individual users about which event information is of interest.
- To use the Performance Schema to investigate a performance issue when the restrictions to impose using pre-filtering are not known in advance.

The following sections provide more detail about pre-filtering and provide guidelines for naming instruments or consumers in filtering operations. For information about writing queries to retrieve information (post-filtering), see [Section 27.5, “Performance Schema Queries”](#).

### 27.4.3 Event Pre-Filtering

Pre-filtering is done by the Performance Schema and has a global effect that applies to all users. Pre-filtering can be applied to either the producer or consumer stage of event processing:

- To configure pre-filtering at the producer stage, several tables can be used:
  - `setup_instruments` indicates which instruments are available. An instrument disabled in this table produces no events regardless of the contents of the other production-related setup tables. An instrument enabled in this table is permitted to produce events, subject to the contents of the other tables.
  - `setup_objects` controls whether the Performance Schema monitors particular table and stored program objects.
  - `threads` indicates whether monitoring is enabled for each server thread.
  - `setup_actors` determines the initial monitoring state for new foreground threads.
- To configure pre-filtering at the consumer stage, modify the `setup_consumers` table. This determines the destinations to which events are sent. `setup_consumers` also implicitly affects event production. If a given event is not sent to any destination (that is, it is never consumed), the Performance Schema does not produce it.

Modifications to any of these tables affect monitoring immediately, with the exception that modifications to the `setup_actors` table affect only foreground threads created subsequent to the modification, not existing threads.

When you change the monitoring configuration, the Performance Schema does not flush the history tables. Events already collected remain in the current-events and history tables until displaced by newer events. If you disable instruments, you might need to wait a while before events for them are displaced by newer events of interest. Alternatively, use `TRUNCATE TABLE` to empty the history tables.

After making instrumentation changes, you might want to truncate the summary tables. Generally, the effect is to reset the summary columns to 0 or `NULL`, not to remove rows. This enables you to clear collected values and restart aggregation. That might be useful, for example, after you have made a runtime configuration change. Exceptions to this truncation behavior are noted in individual summary table sections.

The following sections describe how to use specific tables to control Performance Schema pre-filtering.

### 27.4.4 Pre-Filtering by Instrument

The `setup_instruments` table lists the available instruments:

NAME	ENABLED	TIMED
stage/sql/end	NO	NO
stage/sql/executing	NO	NO
stage/sql/init	NO	NO
stage/sql/insert	NO	NO
statement/sql/load	YES	YES
statement/sql/grant	YES	YES
statement/sql/check	YES	YES
statement/sql/flush	YES	YES

...			
wait/synch/mutex/sql/LOCK_global_read_lock	YES	YES	
wait/synch/mutex/sql/LOCK_global_system_variables	YES	YES	
wait/synch/mutex/sql/LOCK_lock_db	YES	YES	
wait/synch/mutex/sql/LOCK_manager	YES	YES	
...			
wait/synch/rwlock/sql/LOCK_grant	YES	YES	
wait/synch/rwlock/sql/LOGGER::LOCK_logger	YES	YES	
wait/synch/rwlock/sql/LOCK_sys_init_connect	YES	YES	
wait/synch/rwlock/sql/LOCK_sys_init_slave	YES	YES	
...			
wait/io/file/sql/binlog	YES	YES	
wait/io/file/sql/binlog_index	YES	YES	
wait/io/file/sql/casetest	YES	YES	
wait/io/file/sql/dbopt	YES	YES	
...			

To control whether an instrument is enabled, set its `ENABLED` column to `YES` or `NO`. To configure whether to collect timing information for an enabled instrument, set its `TIMED` value to `YES` or `NO`. Setting the `TIMED` column affects Performance Schema table contents as described in [Section 27.4.1, “Performance Schema Event Timing”](#).

Modifications to most `setup_instruments` rows affect monitoring immediately. For some instruments, modifications are effective only at server startup; changing them at runtime has no effect. This affects primarily mutexes, conditions, and rwlocks in the server, although there may be other instruments for which this is true.

The `setup_instruments` table provides the most basic form of control over event production. To further refine event production based on the type of object or thread being monitored, other tables may be used as described in [Section 27.4.3, “Event Pre-Filtering”](#).

The following examples demonstrate possible operations on the `setup_instruments` table. These changes, like other pre-filtering operations, affect all users. Some of these queries use the `LIKE` operator and a pattern match instrument names. For additional information about specifying patterns to select instruments, see [Section 27.4.9, “Naming Instruments or Consumers for Filtering Operations”](#).

- Disable all instruments:

```
UPDATE performance_schema.setup_instruments
SET ENABLED = 'NO';
```

Now no events are collected.

- Disable all file instruments, adding them to the current set of disabled instruments:

```
UPDATE performance_schema.setup_instruments
SET ENABLED = 'NO'
WHERE NAME LIKE 'wait/io/file/%';
```

- Disable only file instruments, enable all other instruments:

```
UPDATE performance_schema.setup_instruments
SET ENABLED = IF(NAME LIKE 'wait/io/file/%', 'NO', 'YES');
```

- Enable all but those instruments in the `mysys` library:

```
UPDATE performance_schema.setup_instruments
SET ENABLED = CASE WHEN NAME LIKE '%/mysys/%' THEN 'YES' ELSE 'NO' END;
```

- Disable a specific instrument:

```
UPDATE performance_schema.setup_instruments
SET ENABLED = 'NO'
WHERE NAME = 'wait/synch/mutex/mysys/TMPDIR_mutex';
```

- To toggle the state of an instrument, “flip” its `ENABLED` value:

```
UPDATE performance_schema.setup_instruments
```

```
SET ENABLED = IF(ENABLED = 'YES', 'NO', 'YES')
WHERE NAME = 'wait/synch/mutex/mysys/TMPDIR_mutex';
```

- Disable timing for all events:

```
UPDATE performance_schema.setup_instruments
SET TIMED = 'NO';
```

## 27.4.5 Pre-Filtering by Object

The `setup_objects` table controls whether the Performance Schema monitors particular table and stored program objects. The initial `setup_objects` contents look like this:

OBJECT_TYPE	OBJECT_SCHEMA	OBJECT_NAME	ENABLED	TIMED
EVENT	mysql	%	NO	NO
EVENT	performance_schema	%	NO	NO
EVENT	information_schema	%	NO	NO
EVENT	%	%	YES	YES
FUNCTION	mysql	%	NO	NO
FUNCTION	performance_schema	%	NO	NO
FUNCTION	information_schema	%	NO	NO
FUNCTION	%	%	YES	YES
PROCEDURE	mysql	%	NO	NO
PROCEDURE	performance_schema	%	NO	NO
PROCEDURE	information_schema	%	NO	NO
PROCEDURE	%	%	YES	YES
TABLE	mysql	%	NO	NO
TABLE	performance_schema	%	NO	NO
TABLE	information_schema	%	NO	NO
TABLE	%	%	YES	YES
TRIGGER	mysql	%	NO	NO
TRIGGER	performance_schema	%	NO	NO
TRIGGER	information_schema	%	NO	NO
TRIGGER	%	%	YES	YES

Modifications to the `setup_objects` table affect object monitoring immediately.

The `OBJECT_TYPE` column indicates the type of object to which a row applies. `TABLE` filtering affects table I/O events (`wait/io/table/sql/handler` instrument) and table lock events (`wait/lock/table/sql/handler` instrument).

The `OBJECT_SCHEMA` and `OBJECT_NAME` columns should contain a literal schema or object name, or '`%`' to match any name.

The `ENABLED` column indicates whether matching objects are monitored, and `TIMED` indicates whether to collect timing information. Setting the `TIMED` column affects Performance Schema table contents as described in [Section 27.4.1, “Performance Schema Event Timing”](#).

The effect of the default object configuration is to instrument all objects except those in the `mysql`, `INFORMATION_SCHEMA`, and `performance_schema` databases. (Tables in the `INFORMATION_SCHEMA` database are not instrumented regardless of the contents of `setup_objects`; the row for `information_schema.%` simply makes this default explicit.)

When the Performance Schema checks for a match in `setup_objects`, it tries to find more specific matches first. For rows that match a given `OBJECT_TYPE`, the Performance Schema checks rows in this order:

- Rows with `OBJECT_SCHEMA='literal'` and `OBJECT_NAME='literal'`.
- Rows with `OBJECT_SCHEMA='literal'` and `OBJECT_NAME='%'`.
- Rows with `OBJECT_SCHEMA='%'` and `OBJECT_NAME='%'`.

For example, with a table `db1.t1`, the Performance Schema looks in `TABLE` rows for a match for '`db1`' and '`t1`', then for '`db1`' and '`%`', then for '`%`' and '`%`'. The order in which matching occurs matters because different matching `setup_objects` rows can have different `ENABLED` and `TIMED` values.

For table-related events, the Performance Schema combines the contents of `setup_objects` with `setup_instruments` to determine whether to enable instruments and whether to time enabled instruments:

- For tables that match a row in `setup_objects`, table instruments produce events only if `ENABLED` is `YES` in both `setup_instruments` and `setup_objects`.
- The `TIMED` values in the two tables are combined, so that timing information is collected only when both values are `YES`.

For stored program objects, the Performance Schema takes the `ENABLED` and `TIMED` columns directly from the `setup_objects` row. There is no combining of values with `setup_instruments`.

Suppose that `setup_objects` contains the following `TABLE` rows that apply to `db1`, `db2`, and `db3`:

OBJECT_TYPE	OBJECT_SCHEMA	OBJECT_NAME	ENABLED	TIMED
TABLE	db1	t1	YES	YES
TABLE	db1	t2	NO	NO
TABLE	db2	%	YES	YES
TABLE	db3	%	NO	NO
TABLE	%	%	YES	YES

If an object-related instrument in `setup_instruments` has an `ENABLED` value of `NO`, events for the object are not monitored. If the `ENABLED` value is `YES`, event monitoring occurs according to the `ENABLED` value in the relevant `setup_objects` row:

- `db1.t1` events are monitored
- `db1.t2` events are not monitored
- `db2.t3` events are monitored
- `db3.t4` events are not monitored
- `db4.t5` events are monitored

Similar logic applies for combining the `TIMED` columns from the `setup_instruments` and `setup_objects` tables to determine whether to collect event timing information.

If a persistent table and a temporary table have the same name, matching against `setup_objects` rows occurs the same way for both. It is not possible to enable monitoring for one table but not the other. However, each table is instrumented separately.

## 27.4.6 Pre-Filtering by Thread

The `threads` table contains a row for each server thread. Each row contains information about a thread and indicates whether monitoring is enabled for it. For the Performance Schema to monitor a thread, these things must be true:

- The `thread_instrumentation` consumer in the `setup_consumers` table must be `YES`.
- The `threads.INSTRUMENTED` column must be `YES`.
- Monitoring occurs only for those thread events produced from instruments that are enabled in the `setup_instruments` table.

The `threads` table also indicates for each server thread whether to perform historical event logging. This includes wait, stage, statement, and transaction events and affects logging to these tables:

```
events_waits_history
events_waits_history_long
events_stages_history
events_stages_history_long
events_statements_history
events_statements_history_long
events_transactions_history
events_transactions_history_long
```

For historical event logging to occur, these things must be true:

- The appropriate history-related consumers in the `setup_consumers` table must be enabled. For example, wait event logging in the `events_waits_history` and `events_waits_history_long` tables requires the corresponding `events_waits_history` and `events_waits_history_long` consumers to be `YES`.
- The `threads.HISTORY` column must be `YES`.
- Logging occurs only for those thread events produced from instruments that are enabled in the `setup_instruments` table.

For foreground threads (resulting from client connections), the initial values of the `INSTRUMENTED` and `HISTORY` columns in `threads` table rows are determined by whether the user account associated with a thread matches any row in the `setup_actors` table. The values come from the `ENABLED` and `HISTORY` columns of the matching `setup_actors` table row.

For background threads, there is no associated user. `INSTRUMENTED` and `HISTORY` are `YES` by default and `setup_actors` is not consulted.

The initial `setup_actors` contents look like this:

```
mysql> SELECT * FROM performance_schema.setup_actors;
+-----+-----+-----+-----+-----+
| HOST | USER | ROLE | ENABLED | HISTORY |
+-----+-----+-----+-----+-----+
| %    | %    | %    | YES   | YES   |
+-----+-----+-----+-----+-----+
```

The `HOST` and `USER` columns should contain a literal host or user name, or `'%'` to match any name.

The `ENABLED` and `HISTORY` columns indicate whether to enable instrumentation and historical event logging for matching threads, subject to the other conditions described previously.

When the Performance Schema checks for a match for each new foreground thread in `setup_actors`, it tries to find more specific matches first, using the `USER` and `HOST` columns (`ROLE` is unused):

- Rows with `USER='literal'` and `HOST='literal'`.
- Rows with `USER='literal'` and `HOST='%'`.
- Rows with `USER='%'` and `HOST='literal'`.
- Rows with `USER='%'` and `HOST='%'`.

The order in which matching occurs matters because different matching `setup_actors` rows can have different `USER` and `HOST` values. This enables instrumenting and historical event logging to be applied selectively per host, user, or account (user and host combination), based on the `ENABLED` and `HISTORY` column values:

- When the best match is a row with `ENABLED=YES`, the `INSTRUMENTED` value for the thread becomes `YES`. When the best match is a row with `HISTORY=YES`, the `HISTORY` value for the thread becomes `YES`.
- When the best match is a row with `ENABLED=NO`, the `INSTRUMENTED` value for the thread becomes `NO`. When the best match is a row with `HISTORY=NO`, the `HISTORY` value for the thread becomes `NO`.